
Collective Semantic Memory: Merge, Trust, and Staleness for Peer Memory in Multi-Agent Navigation

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 **Problem.** Collective memory is often treated as monotone: more shared information should help. In cooperative navigation with scarce semantic targets, this
2 intuition fails. Fast global sharing improves target materialization but causes agents
3 to race for the same targets; slow sharing avoids contention but leaves agents with
4 stale or incomplete memories. **Approach.** We instantiate four memory topologies
5 over the same symbolic-semantic substrate: independent private memories, a
6 shared event bus, a central consensus aggregator, and peer-to-peer broadcast with
7 cadence K . We then introduce a minimal *collective semantic memory* (CSM): peer
8 broadcast plus explicit *merge*, *trust*, and *staleness* rules. Q/R/M/C stage rates are
9 used only as measurement instrumentation, not as the contribution of this paper.
10 **Findings.** On a 35,640-run multi-agent sweep, the fixed-cadence peer family
11 exhibits a bottleneck shift: fast broadcast saturates materialization but collapses
12 completion, while slow broadcast has the opposite profile. Mean time-to-success
13 has an interior optimum at $K=8$. On a 3,240-run scarcity benchmark, the minimal
14 CSM strictly dominates five fixed-cadence peer baselines on success rate: CSM
15 reaches 0.617 with 95% CI [0.610, 0.624], above every peer-K upper CI. **Scope.**
16 Experiments are grid-world substrates. The claim is not that this is a complete
17 theory of collective memory; it is a reproducible architecture split showing that
18 merge/trust/staleness are the right explicit control surfaces after cadence alone is
19 exhausted.
20

21 1 Introduction

22 Many multi-agent systems share information either by writing to a single global memory, by sending
23 learned messages, or by periodically exchanging local state. These mechanisms are usually evaluated
24 end-to-end. If success improves, communication is judged useful; if it does not, the reason is opaque.
25 This paper studies a narrower architectural question: when multiple agents maintain semantic memory
26 about places, what should be shared, at what rate, and with what weight?

27 The answer is not “share everything as soon as possible.” In a scarcity setting, where N agents
28 compete for $T < N$ semantic targets, fresh global memory can make all agents select the same
29 target. The memory is then accurate but unhelpful. Conversely, no communication prevents herding
30 but starves agents of useful discoveries made by peers. The design problem is therefore a trade-off
31 between semantic freshness and occupancy contention.

32 We make three contributions.

- 33 1. **Architecture split.** We implement and compare four memory topologies over the same symbolic
34 substrate: independent, shared-bus, central-aggregator, and peer-to-peer broadcast.




Figure 1: Four memory topologies. **(a)** Independent: per-agent private memory, no exchange. **(b)** Shared bus: one global memory, every agent reads and writes each tick. **(c)** Central aggregator: private memories plus a hub that merges snapshots into a consensus report. **(d)** Peer-to-peer broadcast: private memories with snapshot exchange every K ticks and per-agent merged views.

- 35 **2. Cadence result.** We show that peer broadcast exposes a real coordination parameter K : fast
 36 broadcast improves $P(M^*|R^*)$ but hurts $P(C^*|M^*)$; the efficiency metric has an interior
 37 optimum.
- 38 **3. Minimal CSM.** We instantiate collective semantic memory with three explicit rules: merge,
 39 trust, and staleness. The minimal CSM reaches a new point on the $M \times C$ plane and strictly
 40 dominates fixed-cadence peer baselines on success rate in the scarcity benchmark.
- 41 **Relation to the companion Q/R/M/C paper.** The companion paper treats Q/R/M/C as a diagnostic
 42 framework. Here those stage rates are used as measuring instruments. The main object is different:
 43 the design of distributed symbolic memory itself.

44 2 Memory substrate and task

45 Each agent observes a grid environment and emits semantic place evidence such as `water_source`,
 46 `hazard_edge`, and `open`. A local memory stores place records, tag confidences, visit counts,
 47 freshness, and transition statistics. Planner queries retrieve candidate places by matching required,
 48 preferred, and penalty tags. A selected candidate is materialized into a concrete cell and executed by
 49 a local controller.

50 The multi-agent task uses a 12×10 grid. Each agent must reach a cell tagged `water_source`. Targets
 51 are scarce: $N \in \{3, 5, 8\}$ agents, $T \leq N$ targets, three layouts, three hazard densities, and 40 seeds for
 52 the main cadence sweep. Every run logs Q/R/M/C events per agent:

$$Q^* \rightarrow R^* \rightarrow M^* \rightarrow C^*,$$

53 where M^* means a semantic candidate was locked as an actionable target and C^* means the target
 54 was reached. These logs are not the architectural contribution; they let us see where each topology
 55 fails.

56 3 Four collective-memory topologies

57 The four topologies share the same local memory and planner. They differ only in where merge
 58 happens.

59 **Independent.** Each agent owns a private event store and concept graph. There is no communication
 60 path.

61 **Shared bus.** All observations are appended to a single global event bus. One shared concept graph
 62 is rebuilt from the union of all events.

63 **Central aggregator.** Each agent owns a private concept graph; a central consensus layer pulls
 64 snapshots and computes one trust-weighted report for everyone.




Figure 2: Conditional rates and mean time-to-success from the 35,640-run sweep. Fast sharing improves target materialization but harms completion through contention; slow sharing has the opposite profile. Mean time-to-success has an interior optimum at $K=8$.

65 **Peer broadcast.** Each agent owns a private concept graph and a private merged view. Every K ticks,
 66 agents broadcast snapshots; each agent merges its own memory with the latest snapshot from each
 67 peer.

68 The key distinction is not the clustering primitive. It is the level at which that primitive is instantiated:

independent : N graphs, 0 mergers
 shared bus : 1 graph, 1 implicit global merger
 central : N graphs + 1 central merger
 peer : N graphs + N private mergers.

69 4 Cadence exposes the failure mode

70 The peer topology has one explicit parameter: broadcast cadence K . The main sweep crosses
 71 four architectures and eight peer cadences $K \in \{1, 2, 4, 8, 16, 32, 48, 64\}$ for 35,640 runs. Figure 2
 72 summarizes the result.

73 Fast peer broadcast and centralized sharing live in the M-saturation/C-collapse regime: agents can
 74 lock targets reliably, but many locks point to targets that peers also pursue. Independent and slow
 75 peer broadcast live in the M-deficit/C-saturation regime: fewer agents lock onto good targets, but
 76 those that do are less likely to collide with peer occupancy.

77 Across all peer runs, Spearman of $P(M^* | R^*)$ vs K is -0.608 and of $P(C^* | M^*)$ vs K is $+0.420$
 78 (both $p < 10^{-4}$). The stronger signature is visible at fixed cadence: the within-cadence Pearson
 79 correlation between the M and C links peaks at -0.61 for $K=4$ and decays toward noise as K
 80 grows. Cadence therefore does not simply tune “amount of communication.” It moves the bottleneck
 81 between semantic commitment and shared-resource completion.

82 5 Minimal collective semantic memory

83 The cadence sweep shows that fixed-rate snapshot exchange exposes only one control surface. A
 84 collective semantic memory should expose at least three:

85 **Merge.** How peer evidence about the same semantic place is combined.

86 **Trust.** How evidence from different peers is weighted when it conflicts or becomes unreliable.




Figure 3: Bottleneck shift across peer cadences. $P(M^*|R^*)$ falls and $P(C^*|M^*)$ rises as K grows. The practical efficiency metric, mean time-to-success, is minimized at $K=8$.

87 **Staleness.** How old evidence decays before it influences retrieval.

88 The minimal CSM uses peer broadcast at $K=8$ and overlays three rules:

89 **Merge.** Per-place score equals own evidence at weight 1 plus peer evidence at weight $\text{trust}_i \cdot e^{-\alpha \cdot \text{age}}$,
 90 with $\alpha = 0.05$. Inclusion threshold $\tau = 0.30$.

91 **Trust.** Per-peer EMA on retrieval-success consistency. Peers whose latest snapshot supports the
 92 locally locked target gain trust; others decay toward 0.4. Update rate $\beta = 0.10$.

93 **Staleness.** Each broadcast snapshot is timestamped; merge weight decays as $e^{-\alpha \cdot \text{age}}$. Snapshots
 94 older than $10K$ ticks are pruned.

95 The implementation is a drop-in replacement for fixed-cadence peer memory:
 96 `experiments/collective_semantic_memory/csm_memory.py`.

97 5.1 Results

98 We compare CSM against fixed-cadence peer broadcast at $K \in \{1, 4, 8, 16, 64\}$ on scarcity cells
 99 $(N, T) \in \{(3, 2), (5, 3), (8, 5)\}$, three layouts, three hazard densities, and 20 seeds: 3,240 runs total.

100 CSM occupies a third regime. Fast broadcasts have excellent materialization but poor completion;
 101 slow broadcasts have the reverse. CSM keeps both conditional links above 0.78. The absolute
 102 success gain is modest, +1.7 to +4.0 percentage points, but the stage decomposition shows why it
 103 is meaningful: the architecture changes the joint $M \times C$ operating point rather than merely moving
 104 along the fixed-cadence curve.

105 **Why CSM wins: the gate, not the sharing.** Provenance annotation of the same runs gives the
 106 mechanism. CSM does not share *better* than fixed-cadence peers — it *refuses to act on stale or*
 107 *untrusted foreign evidence*. The inclusion rule of Section 5 implies a closed-form horizon beyond

Table 1: Minimal CSM vs fixed-cadence peer broadcast on the scarcity benchmark. CSM’s lower 95% bootstrap CI bound exceeds every peer-K upper CI.

Architecture	$P(M^* R^*)$	$P(C^* M^*)$	Success rate (95% CI)
peer ($K=1$)	0.991	0.586	0.577 [0.572, 0.586]
peer ($K=4$)	0.955	0.623	0.581 [0.572, 0.590]
peer ($K=8$)	0.730	0.871	0.599 [0.591, 0.607]
peer ($K=16$)	0.697	0.892	0.597 [0.589, 0.606]
peer ($K=64$)	0.684	0.900	0.601 [0.594, 0.609]
CSM (minimal)	0.782	0.825	0.617 [0.610, 0.624]

which a peer’s evidence can no longer form an intention, $\text{age}_{\max}(\text{trust}_i) = \alpha^{-1} \ln(\text{trust}_i \cdot c/\tau)$, and this horizon is empirically exact: measured intention-formation cutoffs match the formula tick-for-tick across trust levels (23/23.05, 18/18.59, 12/12.84 at $\text{trust} = 1.0/0.8/0.6$), including a registered-prediction hold-out on six never-run cells (6/6 exact integer breakpoints) and a discrete trust-flip: $\text{trust} \leq \tau/c$ zeroes foreign-evidence commitments entirely. Fixed-cadence peers have no such gate — their foreign evidence never expires — which is precisely why their completion collapses when sharing is fast. The full provenance-conditioned analysis appears in a companion manuscript; here it serves as the mechanistic reading of Table 1.

6 Place identity from meaning: removing the shared frame

Everything above — and the categorical reading below — assumes a shared spatial basis X : ‘the same place’ means ‘the same (x, y) in a frame the environment hands to every agent’. This section removes that assumption and reports what it was hiding. Each agent now holds a *private* frame (the true grid under a secret translation, and, in the final experiment, a secret 90° -multiple rotation); broadcast snapshots carry the sender’s private coordinates, which are meaningless to the receiver until it decides what the sender’s places *mean*.

Mechanism. A *place fingerprint* is the local semantic constellation around a visited cell: which content tags (hazards, walls, water) sit at which relative offsets — translation-invariant and coordinate-free. Two agents match places by *mutually unique* fingerprint correspondence (each side is the other’s sole candidate; one-directional uniqueness aliases distant look-alikes into the receiver’s neighbourhood, and border-only patterns are excluded as descriptions of the world’s shape rather than of a place). The inter-agent frame is the modal translation of ≥ 3 agreeing matched pairs; under rotations, one hypothesis per 90° reuses the same aligner and the winner must dominate the runner-up’s consensus by $2\times$. Foreign evidence — including places the receiver has never seen — is then transported through the recovered frame.

Results. Registered before execution; full protocol in the companion manuscript.

- Exact recovery.** Witness–traveler: the secret offset is recovered exactly in 100% of landmark-world runs, and under full $SE(2)$ frames in 80/80 rotation $\times$ offset $\times$ seed combinations, with 100% completion on foreign-committed locks.
- Failure asymmetry.** Where identification is impossible — a featureless region, or a 180° -symmetric layout in which the frame is unrecoverable *in principle* — semantic identity makes *zero* commitments (fails closed), while coordinate identity keeps confidently locking wrong cells (100% wrong locks; fails open).
- The full stack.** With $N=4$ agents in distinct secret frames, alignment develops online (onset in 100% of episodes, mean tick ≈ 15) and recovers 88% of the shared-frame performance gap (success 0.444 vs. 0.456 oracle vs. 0.350 coordinate); under random $SE(2)$ frames every semantic commitment points at real water (8/8) while the coordinate stack produces 992 phantom locks (0 real).
- The gate is orthogonal.** Re-running the $\text{age}_{\max}$ analysis above with the CSM inclusion rule ON TOP of semantic identity, under a frame offset and under a 90° rotation, reproduces the SAME integer breakpoints (20/14/9; bp = 20 rotated) — the trust $\times$ staleness gate and frame recovery are independent layers.

The stronger claim promised in earlier drafts’ limitations is therefore discharged constructively: place identity can be carried by meaning alone, it composes with merge/trust/staleness unchanged, and — unlike the coordinate convention it replaces — it knows when it does not know.

7 Categorical BDI reading of collective semantic memory

This section gives a compact theoretical reading of the water-search example in terms of belief, desire, and intention. The goal is not to introduce a new planner, but to make explicit the mathematical structure already used by the collective semantic memory architecture.

We assume a finite set of agents $A = \{1, \dots, N\}$, a shared spatial basis X , and a shared semantic alphabet Σ . In the present grid-world implementation, X is the common coordinate frame and Σ contains tags such as `water_source`, `water_candidate`, `hazard_edge`, and `open`. This is an explicit modelling assumption. The agents may have different observations, different memories, and different confidence values, but they interpret places in a common coordinate framework and express semantic evidence in a common tag alphabet. Without the first assumption, peer merge requires a place-alignment problem to be solved first — Section 6 solves it constructively (fingerprint consensus recovers the inter-agent frame), and Section 7.1 states what that construction is, categorically. Symbol alignment (a shared Σ) remains assumed.

At time t , agent i maintains a local semantic context $\mathcal{K}_t^i = (G_t^i, \Sigma, I_t^i)$, where G_t^i is the set of place records stored in the agent’s private memory, and

$$I_t^i \subseteq G_t^i \times \Sigma$$

is the thresholded incidence relation saying that place $g \in G_t^i$ carries tag $\sigma \in \Sigma$ with sufficient confidence. The graded implementation uses tag confidences and weighted scores; the crisp relation I_t^i is the thresholded skeleton needed for the theoretical description.

This formal context induces a Galois connection between sets of tags and sets of places:

$$\text{ext}_t^i(B) = \{g \in G_t^i : \forall \sigma \in B, (g, \sigma) \in I_t^i\}, \quad \text{int}_t^i(S) = \{\sigma \in \Sigma : \forall g \in S, (g, \sigma) \in I_t^i\}.$$

For a water-search task, the desired semantic predicate is $B_{\text{water}} = \{\text{water_source}\}$, or, in the implemented weighted form, a predicate combining positive evidence for `water_source`, `water_candidate`, and `water_nearby` with penalties for `hazard_edge` or nearby risk. Retrieval succeeds for agent i at time t exactly when

$$\text{ext}_t^i(B_{\text{water}}) \neq \emptyset.$$

Thus the retrieval stage is not a coordinate lookup. It is the test that the queried semantic meaning already has a non-empty extent in the agent’s realized memory.

We can now define the BDI components.

Belief. The belief state of agent i is the structured semantic memory induced by $\mathcal{K}_t^i$. Equivalently, it is the concept lattice $L_t^i = \text{Concepts}(\mathcal{K}_t^i)$, whose nodes pair extents of places with intensions of tags. A water belief is not the proposition “there is water at coordinate x .” It is the region of L_t^i whose intension contains the water predicate. If this region is empty, the agent does not merely choose badly; the relevant semantic object is absent from its current memory.

Desire. Desire is a state-conditioned semantic predicate over the belief lattice. In the water example, the agent’s internal state contains a thirst variable. The active desire can be written as

$$D_t^i = \begin{cases} B_{\text{water}}, & \text{if } \text{thirst}_t^i \geq \theta_{\text{on}}, \\ B_{\text{goal}}, & \text{otherwise.} \end{cases}$$

Thus desire is not an additional metaphysical object. It is the rule that selects which semantic predicate should be queried from the current belief structure.

Intention. Intention is the commitment to one materialized candidate. Once desire has selected a predicate and retrieval has produced a non-empty extent, the planner applies a satisficing choice function

$$\chi_t^i : \mathcal{P}(G_t^i) \rightarrow G_t^i \cup \{\perp\}, \quad g_t^{i,*} = \chi_t^i(\text{ext}_t^i(D_t^i)),$$

where $\perp$ denotes failure to materialize a target. In Q/R/M/C terms, desire determines the query Q^* , belief supports retrieval R^* , and intention corresponds to materialization M^* . Completion C^* is downstream of this semantic structure: it depends on whether the physical controller reaches the locked target and whether the target remains available under multi-agent occupancy.

This gives a compact categorical reading. Let $\mathbf{Ctx}_\Sigma$ be the category whose objects are semantic contexts (G, Σ, I) over the fixed alphabet Σ , and whose morphisms preserve semantic incidence. A local memory trajectory for agent i is a time-indexed diagram

$$\mathcal{K}_0^i \rightarrow \mathcal{K}_1^i \rightarrow \dots \rightarrow \mathcal{K}_t^i \rightarrow \mathcal{K}_{t+1}^i \rightarrow \dots$$

in $\mathbf{Ctx}_\Sigma$. Each arrow is induced by observation and memory update. When new evidence is consolidated, the formal context changes; therefore its concept lattice changes. This is the order-theoretic content of memory growth.

In Q/R/M/C terms, an empty extent $\text{ext}_t^i(D_t^i) = \emptyset$ is a belief/retrieval failure, a non-empty extent with no stable lock is a materialization failure, and a successful lock that cannot be reached is a completion failure — the semantic and the physical parts of the agent’s state fail in different, separately measurable ways.

The collective case lifts this by one construction. At broadcast times the receiver i holds a diagram of local memory objects $\mathcal{D}_t^i = (\mathcal{K}_t^i, \mathcal{K}_t^{j_1}, \dots, \mathcal{K}_t^{j_m})$ together with alignment maps, and queries the merged consensus object $\tilde{\mathcal{K}}_t^i = \text{Merge}_{\text{trust, stale}}(\mathcal{D}_t^i)$, with peer evidence weighted exactly by the rules of Section 5 ($\text{trust}_{j \rightarrow i} \cdot e^{-\alpha \text{age}_j}$). The same BDI definitions apply to the merged context: intention is now selected from $\text{ext}_{\tilde{\mathcal{K}}_t^i}(D_t^i)$ rather than from the private extent alone.

This explains why communication cadence produces a bottleneck shift. Faster broadcast increases the freshness of peer evidence. This tends to enlarge or improve the water extent $\text{ext}_{\tilde{\mathcal{K}}_t^i}(B_{\text{water}})$, so agents can materialize targets more reliably; in Q/R/M/C terms, this raises $P(M^* \mid R^*)$. But the same fresh shared belief also synchronizes intentions. Several agents may lock onto the same scarce water source. The memory is semantically accurate, but the physical target is now contested. This lowers $P(C^* \mid M^*)$. The bottleneck is therefore physically grounded. It is not a formal artifact of the decomposition: it arises because semantic belief is shared through memory, while completion is constrained by occupancy and scarce resources. Collective semantic memory must therefore regulate not only how much information is shared, but also how it is merged, whose evidence is trusted, and when stale evidence should stop influencing intention formation.

In this sense the “collective” is not a central mind and not a global map but a stabilised family of peer-indexed processes; the next subsection makes that statement precise.

The dual, process-level reading — agents as coalgebras (codata), the collective as a final coalgebra, communication as depth-one corecursion, and the bottleneck shift as a competition between belief- and intention-bisimilarity — is developed in Appendix A. It is a lens, not a mechanism: nothing downstream depends on it, which is why it lives in the appendix.

7.1 Colimit merge and its invariants

The merge object above can be described more precisely. Working in a timestamp- and provenance-enriched version of $\mathbf{Ctx}_\Sigma$, a snapshot merge from agents $1, \dots, N$ is a diagram of local memory objects plus alignment maps induced by the place-identity relation. The consensus object $\tilde{\mathcal{K}}_t^i$ is the colimit of this diagram, implemented concretely by the union-find merge over aligned concepts. This reading gives three useful constraints.

1. **Order independence.** The merged object is unique up to isomorphism, so the result is not an artefact of processing agent snapshots in a particular order.
2. **Chaining is expected.** If concept A aligns with B and B aligns with C, union-find performs the transitive closure. Over-merging is therefore not a bug in the data structure; it is the cost of the chosen equivalence relation.
3. **Decay is not a morphism.** Staleness is kept as a weight field outside the alignment morphisms. Otherwise age-then-merge and merge-then-age need not commute, breaking the clean interpretation of broadcast snapshots as one diagram.

In earlier drafts this section ended by conceding that the model does not answer which perceptual features should identify two places — the alignment maps were assumed, not constructed. That gap is now closed by Section 6, and the closure gives the colimit reading its missing scientific function:

1. **The alignment morphisms are constructed, not assumed.** Fingerprint consensus IS an algorithm that produces the alignment maps of the diagram — mutually-unique landmark matches determine the frame transform through which every local memory object embeds.
2. **Existence becomes an empirical predicate.** The colimit exists iff the consensus does. In a 180° -symmetric world the alignment morphism does not exist *in principle* (two incomparable diagrams support equal consensus), and the implementation detects exactly this and refuses to merge — the fail-closed behaviour of Section 6 is the categorical non-existence made operational.
3. **Decay stays outside the morphisms** (invariant 3), and empirically the trust $\times$ staleness gate is unchanged by frame recovery — the same integer breakpoints with and without a shared frame — confirming that the two layers commute as the diagram reading requires.

8 Limitations and next experiments

Grid substrate. The current result is intentionally controlled: all experiments are in grid worlds. A continuous-state substrate such as VMAS, Habitat, ProcTHOR, or MiniWorld is the next portability step.

Coordinate anchor. Removed in Section 6: local fingerprint signatures recover the inter-agent frame (translation and 90° rotations) exactly, and the merge/trust/staleness results are unchanged on top of the recovered frames. What remains open is the harder end of the same axis: continuous $SE(2)$, appearance noise (the current matcher uses exact constellation equality because the substrate is noise-free), and symbol alignment — agents with *different tag alphabets* still cannot merge.

Minimal CSM. The CSM rules are deliberately simple. Adaptive cadence K_t , intent-conditioned trust, and selective consolidation of under-represented tags are the next architecture variants.

Learned baselines. CommNet-style learned communication is a useful comparison, but the more relevant future baseline is an explicit memory LLM or vector-memory multi-agent system measured on the same task.

9 Conclusion

Collective memory is not monotone. In scarce multi-agent navigation, instant information sharing can reduce completion by causing agents to converge on the same targets. Peer cadence exposes the trade-off, but cadence alone is a one-dimensional control surface. A collective semantic memory should make merge, trust, and staleness explicit. The minimal CSM does exactly that and moves the system to a new operating point: neither fast-broadcast M-saturation nor slow-broadcast C-saturation, but a more balanced M $\times$ C profile with higher success than every fixed peer cadence in the scarcity benchmark.

References

- P. Aczel. *Non-Well-Founded Sets*. CSLI Lecture Notes 14, 1988.
- M. Bettini, R. Kortvelesy, J. Blumenkamp, and A. Prorok. VMAS: A vectorized multi-agent simulator for collective robot learning. In *16th International Symposium on Distributed Autonomous Robotic Systems (DARS)*, 2022. arXiv:2207.03530.
- S. Boyd, A. Ghosh, B. Prabhakar, and D. Shah. Randomized gossip algorithms. *IEEE Trans. Information Theory*, 52(6):2508–2530, 2006.
- M. Chevalier-Boisvert, B. Lahlou, S. Willems, and P. S. Castro. MiniGrid and MiniWorld: Modular reinforcement learning environments. *arXiv:2306.13831*, 2023.
- M. Deitke, E. V. Wijmans, D. Schwenk, O. Salvador, K. M. Ehsani, et al. ProcTHOR: Large-scale embodied AI using procedural generation. In *NeurIPS Datasets and Benchmarks*, 2022.

- 285 J. N. Foerster, G. Farquhar, T. Afouras, N. Nardelli, and S. Whiteson. Counterfactual multi-agent policy gradients.
286 In *AAAI*, 2018.
- 287 B. Jacobs. *Introduction to Coalgebra: Towards Mathematics of States and Observation*. Cambridge University
288 Press, 2016.
- 289 S. Mac Lane. *Categories for the Working Mathematician*. Springer, 2nd edition, 1998.
- 290 A. S. Rao and M. P. Georgeff. BDI agents: From theory to practice. In *ICMAS*, 1995.
- 291 J. J. M. M. Rutten. Universal coalgebra: a theory of systems. *Theoretical Computer Science*, 249(1):3–80, 2000.
- 292 D. Sangiorgi. *Introduction to Bisimulation and Coinduction*. Cambridge University Press, 2011.
- 293 M. Savva, A. Kadian, O. Maksymets, Y. Zhao, E. Wijmans, B. Jain, J. Straub, J. Liu, V. Koltun, J. Malik, et al.
294 Habitat: A platform for embodied AI research. In *ICCV*, 2019.
- 295 D. Sperber and D. Wilson. *Relevance: Communication and Cognition*. Blackwell, 2nd edition, 1995.
- 296 S. Sukhbaatar, A. Szlam, and R. Fergus. Learning multiagent communication with backpropagation. In *NeurIPS*,
297 2016.
- 298 L. Xiao, S. Boyd, and S. Lall. A scheme for robust distributed sensor fusion based on average consensus. *IPSN*,
299 2007.

300 A The coalgebraic layer: individuals as codata, the collective as a final 301 coalgebra

302 The reading of §7 fixes belief *content*: at a state, the concept lattice $L_t^i = \text{Concepts}(\mathcal{K}_t^i)$ says what
303 is retrievable, and the merge of §7.1 says how peer contents combine into $\tilde{\mathcal{K}}_t^i$. Both are *algebraic*:
304 they build a belief object. Neither speaks to the agent as a non-terminating *process* that must emit an
305 action and a signal at every tick under a moving environment, nor to the sense in which a *population*
306 comes to share a convention. That process side is *coalgebraic* (Rutten, 2000; Jacobs, 2016). We
307 make it precise on the water-search example and show the two sides are dual, meeting inside a single
308 state transition. (The belief/desire/intention vocabulary follows the BDI agent tradition (Rao and
309 Georgeff, 1995), here given coalgebraic content.)

310 **Signature.** Fix an action set Act (move, dig, noop), a per-agent observation set O_i (agent i ’s local
311 percept — its *unique view*), and an *ostension* set

$$M = 1 + (X \times \Sigma),$$

312 whose values are either *silence* (the summand 1) or a pointer “place x carries tag σ ” (whistle-and-point
313 at a water-marked cell). That silence is a first-class value of M matters: the ostension channel *can be*
314 *informatively empty*, and it is separate from the state update below. This is exactly the observability
315 that a learned real-valued message vector destroys (its channel is never empty and never separates
316 signalling from stepping), and it is why the symbolic stages remain readable.

317 **State.** We take the coalgebra carrier to be the H-MDP high-level state of §7,

$$\mathcal{S}_i = \{ h = (\mathcal{K}^i, D^i, g^{i,*}, s^i) \},$$

318 carrying the same information as $(L^i, D^i, g^{i,*}, s^i)$ since $L^i = \text{Concepts}(\mathcal{K}^i)$. Thus *belief lives in*
319 *the state*; the lattice is a projection of it.

320 **Definition (codata: the agent’s cotype).** The observable behaviour of agent i is a coalgebra for the
321 endofunctor

$$B_i(S) = \underbrace{\text{Act}}_{\text{act now}} \times \underbrace{M}_{\text{signal now}} \times \underbrace{S^{O_i \times M^{A \setminus \{i\}}}}_{\text{continue given (my next percept, peers’ signals)}},$$

322 i.e. a map $c_i = \langle \text{observe_act}_i, \text{observe_ostension}_i, \text{step_intension}_i \rangle : \mathcal{S}_i \rightarrow B_i(\mathcal{S}_i)$
323 with three destructors

$$\text{observe_act}_i : \mathcal{S}_i \rightarrow \text{Act}, \quad \text{observe_ostension}_i : \mathcal{S}_i \rightarrow M, \quad \text{step_intension}_i : \mathcal{S}_i \rightarrow S_i^{O_i \times M^{A \setminus \{i\}}}.$$

Here $M^{A \setminus \{i\}} = \prod_{j \neq i} M$ is the tuple of peers' ostensions delivered this tick. Operationally, `step_intensioni(h)(o, (mj)j≠i)` (i) folds the percept o and the peer pointers (m_j) into the context by the trust/staleness merge, $\mathcal{K}_t^i = \text{Merge}_{\text{trust, stale}}(\dots)$ of §7.1; (ii) recomputes the lattice; (iii) re-evaluates the desire predicate D^i ; and (iv) re-locks the intention $g^{i,*} = \chi_t^i(\text{ext}(D^i))$ of §7. So Merge and χ are *components of one destructor*. The agent is the codata generated by c_i : it has no finite terminal plan, only the standing ability to emit an action and an ostension and to continue. This is the “infinite survival generator” of the motivating example.

Definition (indexed family). The system is the A -indexed family $\{(\mathcal{S}_i, c_i)\}_{i \in A}$, a coalgebra in Set^A . The functor B_i depends on the index i through O_i (and through the peer set $A \setminus \{i\}$): the same survival cotype is instantiated at each agent through that agent's local viewpoint. Agent 1 on the dune generates from “I see the horizon”; agent 2 in the hollow generates from “I feel the soil.” The index is the viewpoint.

Definition (codependence: the collective coalgebra). Let the environment deliver a joint percept $o = (o_i)_{i \in A} \in \prod_i O_i$. The collective one-tick map on $\prod_i \mathcal{S}_i$ wires each agent's ostension output into the others' step inputs,

$$\Phi((h_i)_i)(o) = \left(\text{step_intension}_i(h_i)(o_i, (\text{observe_ostension}_j(h_j))_{j \neq i}) \right)_{i \in A}.$$

Because i 's next state is a function of the *destructor outputs* of the others, the family is not N independent processes but one coalgebra for the product functor $\mathbf{B}(S) = \left(\prod_i (\text{Act} \times M) \right) \times S^{\prod_i O_i}$ on $\prod_i \mathcal{S}_i$. This wiring is the *codependence*: agent 1 cannot step without passing agent 2's ostension through its own `step_intension`. Broadcast cadence enters here as a gate: at non-broadcast ticks the delivered tuple (m_j) is the last received (stale) tuple, or silence; K is the delivery period.

Definition (behaviour and bisimulation). B_i is a container (polynomial) functor, so it has a final coalgebra $(\nu B_i, \text{beh}_i)$ (Aczel, 1988; Rutten, 2000) whose elements are the $(O_i \times M^{A \setminus \{i\}})$ -branching, $(\text{Act} \times M)$ -labelled behaviour trees, and there is a unique coalgebra morphism $!_i : \mathcal{S}_i \rightarrow \nu B_i$ sending a state to its whole future behaviour. A B_i -bisimulation is a relation $R \subseteq \mathcal{S}_i \times \mathcal{S}_i$ under which related states emit equal action and equal ostension and, on every input, step to R -related states. Container functors preserve weak pullbacks, so the *greatest bisimulation* $\sim_i$ exists, is an equivalence, and equals $\ker(!_i)$: two states are bisimilar iff they have the same behaviour (Sangiorgi, 2011). The same holds for the collective coalgebra $(\prod_i \mathcal{S}_i, \langle (\text{observe_act}_i)_i, (\text{observe_ostension}_i)_i, \Phi \rangle)$, with greatest bisimulation $\approx$.

Definition (society / culture). The *collective* is the image of the states in the final coalgebra $\nu \mathbf{B}$, equivalently the quotient $(\prod_i \mathcal{S}_i) / \approx$. It is neither a central server nor a shared map: it is a bisimilarity class of joint behaviour. A *water-search convention* is this invariant read on the water sub-signature. Let $M_{\text{water}} \subseteq M$ be the pointers whose tag lies in the water family, and let $\nu B_i \rightarrow \nu B^w$ be the forgetful map to the behaviour trees restricted to receiving M_{water} and emitting the corresponding intention lock. The population *holds the convention* when all agents' reachable behaviours agree under this map: hearing “point-and-dig at a water-marked place” drives every agent's intention destructor to lock that place (up to occupancy). Governance is then automatic in both directions. A new agent (the example's fourth agent) is a member of the society exactly when its behaviour lands in the class $[\cdot]_{\approx}$ on M_{water} — otherwise its ostensions are not understood and its own step never converges (top-down); and the class itself was reached by the local, coupled steps of the first agents at one green shrub (bottom-up).

Remark (communication = lazy, depth-one corecursion). The behaviour $!_i(h)$ is defined by unfolding c_i indefinitely. On receiving $m_j = \text{observe_ostension}_j(h_j)$, agent i uses only this *depth-one* destructor output inside `step_intensioni`; it does *not* compute $!_j(h_j) \in \nu B_j$, the peer's full behaviour, which is the “I think that you think...” regress. Communication is therefore corecursion truncated to the least depth that resolves the receiver's own next step — a finite approximation of the final-coalgebra map — with a relevance-theoretic stopping rule (Sperber and Wilson, 1995) deciding that depth. The regress collapses into one effective action.

Table 2: The motivating example’s terms as formal objects of the coalgebraic reading. Each row is the water-search meaning of a construction already used in §A.

Informal term	Formal object	Water-search meaning
Cotype (codata)	B_i -coalgebra (S_i, c_i)	the agent’s non-terminating action/signal generator
Indexed type	family over A ; dependence of B_i on O_i	the agent’s unique local viewpoint (dune vs. hollow)
Codependent type	wiring $\text{observe_ostension}_j \rightarrow \text{step_intension}_i$ in Φ	“I step only through your ostension”
Individual (instance)	a state $h \in S_i$ and its behaviour $!_i(h)$	the running agent digging and whistling now
Destructors (cogenerators)	observe_act , observe_ostension , step_intension	step-and-dig; whistle-and-point; fold neighbour’s signal into next state
Lazy evaluation	depth-one corecursion + relevance stop	unfold the peer’s coalgebra one step, not its whole behaviour
Final coinductive core / society	νB ; greatest bisimulation $\approx$ (on M_{water})	the stable, distributed water-search culture

Remark (algebra–coalgebra bridge to §7–7.1). Belief is algebraic: the FCA lattice L_t^i built by the colimit merge is the *content at a state*. The agent is coalgebraic: the codata observed by the three destructors is the *process over states*. They are dual and meet in one place — the colimit merge and the satisfying choice χ are the internals of step_intension . The memory-growth diagram $\mathcal{K}_0^i \rightarrow \mathcal{K}_1^i \rightarrow \dots$ of §7 is precisely the belief-projection of the coalgebra trajectory $h_0^i \mapsto h_1^i \mapsto \dots$.

Coalgebraic reading of the bottleneck shift. The result of §4 and Table 1 has a one-line reading here. Two kinds of bisimilarity compete. Bisimilarity on the *ostension / belief* sub-signature is the wanted invariant: it is the shared water convention. Bisimilarity on the raw *intention* output over a *scarce* target is the unwanted one: several states mapping to the same g^* is a collision. Occupancy $\mathcal{O}_{-i}$ is the extra-coalgebraic constraint that penalises the second. Cadence K tunes how hard the coupling Φ drives the collective toward the *full* greatest bisimulation $\approx$: small K delivers fresh ostensions often, enlarging the merged water extent so intention locks reliably ($P(M^* \mid R^*) \uparrow$), but over-synchronises the intention destructor so agents collide ($P(C^* \mid M^*) \downarrow$). The interior optimum $K=8$ is where belief converges without intention collapsing onto one cell; the minimal CSM adds trust and staleness precisely so that step_intension reaches belief-bisimilarity (fresh, trusted, aligned evidence) while damping intention-bisimilarity (stale evidence decays before it can herd every agent onto the same lock). We state this as a reading consistent with the data, not as a theorem: the container functors, their final coalgebras, and bisimulation are standard constructions used here as a lens, and the relevance-theoretic stopping rule is an interpretive gloss rather than a fitted mechanism.

B Reproducibility commands

- Main cadence sweep: `pythonexperiments/big_experiment/exp_cadence_phase.py --modefull --workers8 --out_dirtmp/big_experiment_qrmc_40`
- Extended cadences: `pythonexperiments/big_experiment/exp_extra_K.py --workers8 --out_dirtmp/big_experiment_extraK`
- Q/R/M/C analysis: `pythonexperiments/big_experiment/analyze_qrmc.py --runs_csvtmp/big_experiment_qrmc_40/runs.csv --out_dirtmp/big_experiment_qrmc_40`
- Minimal CSM benchmark: `python-mexperiments.collective_semantic_memory.run_csm_benchmark --out_dirtmp/csm_benchmark`
- MiniGrid portability sweep: `python-mexperiments.minigrid_multiagent_wrapper.run_portability_sweep --out_dirtmp/minigrid_portability`

403 **C Asset and license note**

404 The experiments use MiniGrid and MiniWorld (Chevalier-Boisvert et al., 2023). Locally implemented
405 benchmark and memory code will be released under the MIT License as part of the artifact package.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and introduction present the paper primarily as a four-stage diagnostic framework with a reference symbolic-semantic instantiation, and they explicitly scope the evidence to benchmarked navigation environments with a limited portability check (Abstract; Sections 1 and 6).

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: The paper includes an explicit “Limitations and future work” section that discusses environment difficulty, the observable-query and observable-occupancy assumptions, the slope-level scope of Empirical Claim 1, the learned-baseline scope, and agent-count scaling (Section 6).

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: The paper states the conditional stage-Markov property and positivity assumptions in Section 2.2, gives a proof sketch in the main text, and includes a full formal argument for Definition 1 (single-agent stage estimability) in Appendix A together with the multi-agent extension (Definition 2) and the slope-level argument for Empirical Claim 1.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The paper specifies the task families, environments, seeds, episode counts, assist modes, baselines, statistical protocol, and benchmark semantics; the single-agent setup is in Section 3.1 and the multi-agent sweep protocol in Section 4.4; the operational stage definitions used to extract Q^* , R^* , M^* , C^* from runtime logs are in Appendix B.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The paper describes a unified benchmark runner and source-of-truth artifacts. An anonymized public release of the codebase — single-agent benchmark, multi-agent grid environment, 35,640-run sweep driver, Q/R/M/C analyser, oracle-intervention scripts, and exact commands to reproduce every table and figure — accompanies the submission as supplemental material (anonymized GitHub repository; single-command reproduction of the main multi-agent sweep on an 8-core machine in ≈ 22 minutes, matching Appendix H).

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: The benchmark protocol, task definitions, assist settings, seeds, episode counts, baseline families, and the learned baselines are described in the main text together with the measurement and statistical protocols: single-agent setup in Section 3.1, multi-agent sweep in Section 4.4, and per-baseline hyperparameters in Appendix G.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: The paper reports 95% bootstrap confidence intervals over seed-level aggregates ($B=5,000$ resamples; Section 3.1, “Protocol” paragraph) and uses paired Wilcoxon signed-rank tests with Holm-Bonferroni correction for assist comparisons. CIs are shown in Tables 3 and 5 (Best semantic per-task rates; the 35,640-run MOP column) and in Figures 4–6 (bootstrap shaded bands).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: The paper discloses the hardware platform (10-core ARM CPU, 32 GB RAM; CPU-only), the scale of each benchmark, and wall-clock costs — single-agent oracle sweep < 1 minute (Section 3.1); main multi-agent sweep (35,640 runs) ≈ 22 minutes on 8 cores (Section 4.4); $N=12$ scale-up (8,640 runs) ≈ 12 minutes (Appendix D); PPO CommNet baseline (200 updates $\times$ 64 rollouts) ≈ 96 minutes on a single CPU core (Appendix G); MiniGrid portability sweep (100 runs) ≈ 2 minutes on a single core (§8).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn’t make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

Answer: [Yes]

Justification: The work is a simulation-only study in standard benchmark environments, uses anonymized authorship for review, and does not involve human subjects, private data, or deployment claims that conflict with the NeurIPS Code of Ethics.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: The paper includes a dedicated “Broader impacts” section that discusses the primary methodological benefits of the work and the indirect risks of improving navigation systems without adequate safeguards (Section 7).

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper does not release high-risk models or scraped datasets; it studies navigation policies in simulated environments.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.

- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [\[Yes\]](#)

Justification: The paper cites the upstream benchmark platforms (MiniGrid, BabyAI) in the main references. Appendix I states that both are MIT-licensed and that the benchmark scripts introduced here will be released under the MIT License at camera-ready.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[No\]](#)

Justification: The submission introduces benchmark scripts and analysis artifacts. A formal anonymized release package with documentation is planned for camera-ready; Appendix H lists all scripts needed to reproduce each reported result, and Appendix I states the intended license for the new assets.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [\[N/A\]](#)

Justification: The paper does not involve crowdsourcing or human-subject experiments.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: The paper does not involve human subjects or human-participant data collection.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: LLMs are not an important, original, or non-standard component of the core methods or experiments reported in the paper.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.