

---

# Q/R/M/C: A Stage-Decomposition Measurement Framework for Memory-Based Navigation

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

**Setting.** Success-only metrics collapse qualitatively different memory-based navigation failures into a single scalar, hiding *which stage* of the retrieval pipeline is responsible. This paper contributes an *evaluation protocol*: the Q/R/M/C stage decomposition (**Q**uery formation, **R**etrieval satisfaction, target **M**aterialization, **C**ompletion after retrieval), operationalised so that the four conditional rates are computable directly from standard trajectory logs (Definition 1). The arithmetic is the chain rule, deliberately so — just as decomposing accuracy into precision and recall added no mathematics but changed what the field could *see*, the value here is identifiability from logs and what the decomposition reveals, not a theorem.

**Single-agent measurement.** On a 10-seed MiniGrid benchmark (FourRooms, LavaGapS7) and a BabyAI portability check, the framework localises bottlenecks: hazard recovery is retrieval-limited (oracle retrieval lifts success  $0.39 \rightarrow 0.60$ ), goal-region rejoin is downstream-limited. Log inspection refines the attribution: candidate ranking is near-optimal (96.9%) while candidate generation fails at 91% of decision points — a memory-accumulation, not a ranking, limit.

**Multi-agent measurement.** On a 35,640-run multi-agent sweep over four communication architectures (independent, shared-bus, central-aggregator, peer-broadcast at eight cadences  $K \in \{1, 2, 4, 8, 16, 32, 48, 64\}$ , 40 seeds), the bottleneck-shift between the M and C links appears directly as a negative within-cadence correlation peaking at  $K=4$  (Pearson  $-0.61$ ); mean time-to-success has an interior minimum at  $K=8$  (robust in location but shallow in margin — not resolved from the adjacent fast cadences). Both slope signs (Empirical Claim 1) hold as moderate, cluster-robust effect sizes (Spearman  $-0.53$  and  $+0.43$  vs.  $K$ , 95% CIs  $[-0.59, -0.47]$  and  $[+0.37, +0.50]$  over 81 configuration cells; Appendix F.3). Non-claims (notably on the conditional product  $\Phi$ ) are consolidated in one place (§1, Scope and non-claims).

**Diagnostic use.** The same four rates also expose why learned policies with implicit communication can be competitive on success while collapsing Q/R/M/C observability: their communication vectors are never informatively empty and they lack an explicit target-lock interface. On two local LLM agents (qwen3:4b, llama3.1 on TextNav), the same logger localizes a completion-limited failure: under a tight step budget both agents form the query, retrieve, and lock the target ( $Q^*=R^*=M^*=1$ ) yet none complete ( $C^*=0$ ), so an otherwise opaque success collapse is resolved into a specific failing stage rather than a leaderboard number. Thus the contribution is a measurement protocol for making memory failures comparable across symbolic, learned, LLM, single-agent, and multi-agent systems.

**Scope.** The primary experiments are in grid-world substrates (MiniGrid, BabyAI, a custom  $12 \times 10$  grid). The contribution is the diagnostic framework and its validation; a preliminary continuous-state check on VMAS reproduces both slope signs directionally (Appendix F.7), and a full continuous/3D validation with cluster-robust CIs (Habitat, ProcTHOR) is future work.




Figure 1: Q/R/M/C decomposes one semantic-navigation episode into four observable stages. **Top:** the four stages and their event-level definitions. **(a)** A successful water-search episode anchors at one event per stage, producing logged starred events  $Q^*$ ,  $R^*$ ,  $M^*$ ,  $C^*$ . **(b)** The same scalar “end-to-end success” summary masks qualitatively different failures on the four single-agent task families: water and rest are strong; goal-region rejoin (FourRooms) is downstream-limited; hazard recovery (LavaGapS7) is retrieval-limited (oracle interventions, §3.4).

## 1 Introduction

Classical navigation benchmarks often assume that the agent is given an explicit destination coordinate or a dense task reward. Many practical navigation problems are different: the agent must reach a place of a certain *type*, such as a safe rest area, a water-like resource, a post-hazard recovery exit, or a goal-associated region, specified by a semantic pattern rather than a fixed  $(x, y)$  waypoint. The place must therefore be *remembered and retrieved by meaning*. This setting opens a methodological question: can we measure *where* memory-based semantic navigation breaks down, not only *whether* it succeeds?

We answer this question with a four-stage diagnostic framework that applies to both single agents and multi-agent collectives. The decomposition is Q/R/M/C: **Q**uery formation, **R**etrieval satisfaction, **M**aterialization, and **C**ompletion after retrieval. The decomposition factorizes per-episode semantic-path completion and, under a conditional stage-Markov property, its factors are estimable from trajectory logs (Definition 1). The rest of the paper develops a single argument in three parts.

**Single agent.** Stage attribution is non-trivial. Success-only metrics collapse qualitatively different failures into a single scalar. In a 10-seed MiniGrid benchmark, water and rest tasks are strong end-to-end (0.95, 0.93 success), but goal-region rejoin and hazard recovery are substantially harder, and for *different reasons*. Oracle-stage interventions show that hazard recovery is retrieval-limited (oracle retrieval lifts success  $0.39 \rightarrow 0.60$ ), while goal-region rejoin is downstream-limited (oracle retrieval rescues semantic-path completion without changing end-to-end success). Inspection of the framework’s own logs then sharpens the attribution: the retrieval bottleneck is not candidate *ranking* (precision 96.9% when candidates exist) but candidate *generation*: the memory contains no eligible symbols at 91% of decision points. The single-agent bottleneck is a *memory-accumulation* limit: evidence is observed but never consolidated into queryable structure.

**Multi-agent.** Consolidation acquires a rate. When  $N$  agents share memory through a broadcast protocol with cadence  $K$ , the factorization extends with enlarged conditioning sets (Definition 2), and  $K$  becomes a structural axis: it is, mechanically, a knob on *how fast private memories consolidate into a collective one*. A slope-level claim (Empirical Claim 1) predicts a bottleneck shift: as  $K$  decreases,  $P(M^*|R^*)$  rises (fresher peer memory enables reliable commitment) and  $P(C^*|M^*)$  falls (peers race for the same targets). Both slope signs are confirmed on a 35,640-run sweep as moderate, cluster-robust effect sizes (Appendix F.3), the shift is directly visible as a negative within-cadence

71 correlation between the two links, and the resulting trade-off produces an interior optimum of mean  
72 time-to-success at  $K=8$ : consolidation has a non-trivial optimal rate.

73 **Diagnostic synthesis.** Both regimes point at the same measurement object. The single-agent diagnosis  
74 says that evidence consolidation into symbols can be the binding constraint; the multi-agent diagnosis  
75 says that the rate and cost of sharing consolidated memory determine the optimum.  $Q/R/M/C$  is the  
76 common instrument: candidate systems can be compared by how they move  $P(R^*|Q^*)$ ,  $P(M^*|R^*)$ ,  
77 and  $P(C^*|M^*)$ , even when their end-to-end success rates look similar.

78 **Scope and non-claims (consolidated).** All caveats of this paper in one place; the running text  
no longer repeats them. (i) The decomposition’s consistency argument is the chain rule plus  
frequency estimation — no mathematical novelty is claimed; the contribution is the logging  
discipline and identifiability (Definition 1–2). (ii) Empirical Claim 1 is a verified slope-level  
observation on the conditional rates, not a theorem. (iii) We make no claim about the shape  
of the conditional product  $\Phi$ ; in our data it is monotone in  $K$ , and the interior optimum  
lives on mean time-to-success (§4.5). (iv) Primary substrates are grid worlds; VMAS is a  
directional continuous-state check, and the LLM studies are small and controlled (two local  
models on TextNav; an  $N=3$  collective; a three-framework external validation at  $n=8$  per  
cell, §5). (v) Multi-agent place identity assumes a shared coordinate frame (§6). A stage-level  
operationalisation is testable on specific summaries and can be *wrong* on others — (iii) is an  
instance we report — which is what separates a measurement protocol from a stage-agnostic  
guess.

79 **Why this framing is timely.** Embodied and goal-conditioned agents increasingly retrieve places,  
80 objects, or tools by meaning (Andrychowicz et al., 2017; Ghosh et al., 2021; Savva et al., 2019).  
81 LLM-agent systems lean on explicit memory layers but evaluate mostly end-to-end (Wang et al., 2023;  
82 Shinn et al., 2023; Packer et al., 2023). Multi-agent reinforcement learning likewise increasingly  
83 involves communicating populations (Lowe et al., 2017; Foerster et al., 2018), but the question  
84 “which agent should know what, and when?” lacks a unified diagnostic vocabulary. Our framework  
85 supplies one.

86 **Contributions.** This is a methodology paper. We are explicit about what is bookkeeping vs. empirical  
87 claim vs. instantiated artefact.

- 88 1. **Measurement protocol (single-agent).** The  $Q/R/M/C$  decomposition is operationalised so  
89 that the conditional rates are estimable from trajectory logs (Definition 1); the contribution  
90 is the *logging discipline* that makes the decomposition usable on standard navigation traces  
91 (non-claims: see box above).
- 92 2. **Measurement protocol (multi-agent).** An extension to  $N$  communicating agents with two  
93 coupling channels (memory  $\mathcal{M}_{-i}^{(K)}$ , occupancy  $\mathcal{O}_{-i}$ ) under a peer-broadcast protocol with  
94 cadence  $K$  (Definition 2).
- 95 3. **Empirical Claim 1 (bottleneck shift).** On a 35,640-run multi-agent sweep, the two stylised  
96 slope conditions **(F)** and **(O)** are both confirmed as moderate, cluster-robust effect sizes (Spear-  
97 man  $-0.53$  and  $+0.43$  vs.  $K$ , 95% CIs excluding zero over 81 configuration cells; Appendix F.3;  
98 full benchmark in §4.5), with a provenance-level mechanism (§4.5).
- 99 4. **System and benchmarks.** A symbolic-semantic memory stack with four communication  
100 architectures (independent, shared-bus, central-aggregator, peer-to-peer broadcast); single-agent  
101 oracle benchmark (10-seed MiniGrid, BabyAI portability); 35,640-run multi-agent sweep at  
102  $N \in \{3, 5, 8\} + 8,640$ -run scale-up at  $N=12$ ; 100-run portability sweep on a standard-MiniGrid  
103 wrapper.
- 104 5. **Diagnostic comparison across architectures.** The same logging contract separates symbolic  
105 memory systems from learned policies with implicit communication: the latter can be competi-  
106 tive on success while collapsing  $Q/R/M/C$  stage observability. This is an evaluation result, not a  
107 claim that one architecture class should always dominate.
- 108 6. **Scope and honesty.** The primary experiments are in grid-world substrates. We treat that  
109 as a methodology-paper choice and state it explicitly (§6); a preliminary VMAS check on a  
110 continuous substrate is included (Appendix F.7), and full continuous/3D validation (Habitat,  
111 ProcTHOR, MineDojo) is future work.

## 2 The Q/R/M/C decomposition

### 2.1 Problem formulation and notation

The agent operates in a partially observable navigation environment. At each step it receives an observation  $o_t$ , maintains a memory state  $\mathcal{M}_t$ , and forms a query  $q_t$  encoding the current task intent. A retrieval step returns a candidate set of places matching  $q_t$  from  $\mathcal{M}_t$ ; a chosen candidate is materialized into an actionable target; a local controller executes actions until either the target is reached or the episode times out. The target is not a coordinate but a semantic pattern, so success is a function of the chosen candidate’s semantic profile, not of any pre-supplied coordinate.

Habitat, AI2-THOR, and ProcTHOR made visually rich navigation evaluation standard (Savva et al., 2019; Kolve et al., 2017; Deitke et al., 2022). We deliberately focus on MiniGrid-class environments to obtain clean and reproducible stage measurements, and use BabyAI only as a portability check. We do not claim to solve LLM-agent memory here; rather, we argue that the same Q/R/M/C decomposition is a useful template for evaluating memory-grounded agent decisions beyond navigation.

The central empirical question is therefore not just whether the agent succeeds, but *which stage fails first* when it does not. Two related measurement layers are used. **Operational rates** count, per query attempt, whether the query yielded a non-empty candidate set, whether the chosen candidate satisfied the semantic predicate, and whether it was materialized into an actionable target; *completion after retrieval* is an episode-level indicator. **Nested episode-level events**  $Q^* \Rightarrow R^* \Rightarrow M^* \Rightarrow C^*$  are used in the factorization below and audited in §3.2. Full operational definitions are given in Appendix C.

**Notation.** The four stages are  $Q, R, M, C$  with episode-level starred events  $Q^*, R^*, M^*, C^* \in \{0, 1\}$ , nested by construction.  $P(\cdot)$  denotes the population probability;  $\hat{P}(\cdot)$  its empirical-frequency estimator. **The letter  $M$  always refers to the Materialization stage; the number of targets in the multi-agent environment is written  $T$  (not  $M$ , constraint  $T \leq N$ , “scarcity” =  $N > T$ ).** Other symbols:  $\mathcal{C}_K$  peer-broadcast protocol with cadence  $K$ ,  $\mathcal{M}_{-i}^{(K)}$  peer memory snapshots,  $\mathcal{O}_{-i}$  peer occupancy,  $\nu(\cdot)$  freshness functional,  $\Phi(K) = P(M^* | R^*; K) \cdot P(C^* | M^*; K)$ ,  $t_{\text{succ}}$  mean time-to-success. Full symbol table in Appendix D.1.

### 2.2 Single-agent factorization

Let  $Y$  denote overall episode success, and let  $Q^*, R^*, M^*, C^*$  denote the nested episode-level semantic-path events. The framework factorizes *semantic-path completion* rather than raw task success:

$$Q^* \rightarrow R^* \rightarrow M^* \rightarrow C^*.$$

For any query  $q$  and memory state  $\mathcal{M}$ ,

$$\begin{aligned} P(C^* = 1 \mid q, \mathcal{M}) &= P(Q^* = 1 \mid q, \mathcal{M}) P(R^* = 1 \mid Q^* = 1, q, \mathcal{M}) \\ &\quad \cdot P(M^* = 1 \mid R^* = 1, q, \mathcal{M}) P(C^* = 1 \mid M^* = 1, q, \mathcal{M}). \end{aligned}$$

Overall task success  $Y$  can exceed  $C^*$  because some episodes succeed without traversing the full semantic retrieval path.

**Assumption 1 (conditional stage-Markov property).** Conditioned on the current stage being successful, the probability of the next stage depends only on the current successful stage, the current query, and the current memory state. Formally,

$$P(R^* \mid Q^*, q, \mathcal{M}, \text{history}) = P(R^* \mid Q^*, q, \mathcal{M}),$$

and analogously for  $M^*$  and  $C^*$ .

**Assumption 2 (positivity).** Each conditioning event used above has non-zero probability on the support of the benchmark.

**Definition 1 (single-agent stage estimability).** Under Assumptions 1–2, the four conditional rates in the factorization above are well-defined population quantities and the empirical-frequency estimators are consistent estimators of them; their product converges to semantic-path completion.




Figure 2: Stage-structured interpretation of the Q/R/M/C framework. Assist on/off ablations target intermediate mechanisms; in the multi-agent extension the cadence protocol  $\mathcal{C}_K$  acts on the  $R \rightarrow M$  and  $M \rightarrow C$  edges.

The argument is elementary — the chain-rule factorization plus consistency of frequency estimators on Bernoulli events — and the value of stating it is operational, not mathematical: it pins down the *logging discipline* (which events, conditioned on what) under which the decomposition is readable from standard trajectory logs, and it is the contract that every experiment in this paper audits against (§3.2). Full statement and elementary argument in Appendix A.

**When Assumption 1 may fail.** Assumption 1 holds tightly here because  $q$  and  $\mathcal{M}$  are explicitly logged at every stage event. Two practically relevant violations remain: adaptive in-episode query rewrites not surfaced to the logger, and hidden between-stage memory updates. In those cases the factorization should be read as a projection onto observable traces rather than as a full structural identification claim. Appendix F.6 quantifies the estimator bias under a controlled violation of this kind: it is negligible for mild unlogged couplings and one-signed (an upward bias on  $P(C^*|M^*)$ ), which bounds the framework’s exposure when it is applied to learned or LLM agents.

**Mechanism-level reading.** The same decomposition admits a stage-structured mechanism interpretation. Figure 2 treats stage outcomes as a directed acyclic graph. Under this interpretation, assist on/off ablations are mechanism-targeted perturbations of materialization and post-retrieval execution rather than undifferentiated sensitivity analyses; in the multi-agent setting the same graph is replicated per agent and connected through the cadence protocol  $\mathcal{C}_K$ , which acts on  $\mathcal{M}_{-i}^{(K)}$  at the  $R \rightarrow M$  edge and on  $\mathcal{O}_{-i}$  at the  $M \rightarrow C$  edge — exactly the two edges where Empirical Claim 1 will predict the slope shift.

**Relation to prior work.** Closest are (i) cognitive-map and semantic-map frameworks (Tolman, 1948; O’Keefe and Nadel, 1978; Kuipers, 2000; Gupta et al., 2017; Savinov et al., 2018), (ii) goal-conditioned and successor-representation RL (Andrychowicz et al., 2017; Ghosh et al., 2021; Dayan, 1993; Stachenfeld et al., 2017; Momennejad et al., 2017), (iii) multi-agent communication (Lowe et al., 2017; Foerster et al., 2018; Sukhbaatar et al., 2016; Foerster et al., 2016) and gossip/consensus (Boyd et al., 2006; Xiao et al., 2007), and (iv) LLM-agent memory layers evaluated end-to-end (Wang et al., 2023; Shinn et al., 2023; Packer et al., 2023). Our framework is orthogonal to all four: it *diagnoses* where in the Q/R/M/C chain a given architecture succeeds or fails and lifts the cadence axis from network protocols to a cognitive-architecture parameter (full discussion in Appendix F.9).

### 2.3 Event semantics, bounded rationality, and an order-theoretic reading

The four starred events anchor at concrete logged triggers. We state the triggers once, because every empirical claim in the paper is a claim about them (per-regime formal definitions in Appendix C), and we give each stage its reading in terms of the resource it bounds.

**$Q^*$  (intent issued).** The planner emits a semantic intent  $q = (q^{\text{req}}, q^{\text{pref}}, q^{\text{pen}})$ : small sets of required, preferred, and penalty tags. The episode-level event fires on emission, so it is saturated whenever the task is semantic ( $\hat{P}(Q^*)=1.00$  in Table 1); the multi-agent per-tick logger additionally requires a non-empty candidate set (Appendix C.1), a bookkeeping difference we record rather than hide. Resource bound: attention — the intent is a few tags, not the full state.

**$R^*$  (memory answers).** Retrieval evaluates  $q$  against the memory and returns a candidate set; the event fires when the set contains a candidate satisfying the task predicate (within  $\epsilon$  of a true target in the multi-agent logger). The 91% empty-candidate-set failures of §3.5 are  $R$ -side events: no

195 stored place clears the required-tag threshold. Resource bound: memory — retrieval can only  
 196 surface what consolidation has stored.

197  **$M^*$  (target lock).** The planner materializes one candidate into a concrete actionable cell. Candidacy  
 198 itself is threshold-gated (a place enters the candidate set only if its score clears the planner’s gate,  
 199 Appendix D), and the lock commits to a single candidate once per query rather than re-optimizing  
 200 against the full memory at every tick. The event fires on lock acquisition (within  $\epsilon$  of a true target  
 201 in the multi-agent logger). Resource bound: deliberation.

202  **$C^*$  (closed-loop completion).** The local controller executes until the locked cell is reached or the  
 203 episode times out; the event fires on arrival.  $C$  is the only stage whose failure modes are extra-  
 204 semantic: controller competence and, in the multi-agent regime, occupancy  $\mathcal{O}_{-i}$ . Resource bound:  
 205 control, and in a collective, shared targets.

206 **Why a deterministic strategy: bounded and ecological rationality.** A natural objection is that the  
 207 semantic stack is a fixed, deterministic strategy rather than a learned policy. We treat the determinism  
 208 as principled and name it: the pipeline is a satisficing agent in Simon’s sense (Simon, 1955, 1956).  
 209 Candidacy is threshold-gated instead of exhaustively scored, the target commitment is made once  
 210 instead of re-optimized at every tick, and the intent is a few tags instead of the full state. Bounded  
 211 rationality is also what makes the four failure loci observable: each stage is a distinct resource bound  
 212 (attention, memory, deliberation, control), and Q/R/M/C is the audit trail of where the bound binds.  
 213 The learned baselines make the converse point empirically: policies without an explicit query/lock  
 214 interface collapse the stages even at competitive success (BC in §3.3; CommNet PPO in §4.5), so  
 215 observability is traded away exactly where end-to-end optimization is traded in. That the *same*  
 216 heuristic is strong on water and rest, retrieval-starved on hazard recovery, and downstream-limited on  
 217 goal-region (Figure 1b) is the ecological-rationality point that heuristic success is a property of the  
 218 heuristic–environment pair (Simon, 1956; Gigerenzer and Goldstein, 1996); the framework is the  
 219 instrument that locates the mismatch. The same lens reads the  $\Phi$ -versus- $t_{\text{succ}}$  split of §4.5: the chain  
 220 probability  $\Phi$  ignores the cost of traversing the chain, so under a resource-rational reading (Lieder  
 221 and Griffiths, 2020) an optimum should be sought on a cost-bearing metric; that is where it appears  
 222 ( $K=8$  on mean time-to-success) and not on  $\Phi$ .

223 **An order-theoretic reading (deferred to Appendix A.4).** The query–retrieval interface carries a  
 224 standard order-theoretic structure: places and tags form a formal context whose Galois connection  
 225 realizes the memory’s concept lattice, so retrieval fails empty exactly when the queried intent’s extent  
 226 is empty — “absence of evidence,” the 91%/96.9% regime of §3.5, and the regime where growing  
 227 the context can only grow extents. We use this as a naming device, not a contribution; the statement,  
 228 its four-line argument, and its scope limits are in Appendix A.4.

229 **A dynamic (options / H-MDP) reading (deferred to Appendix B).** The stages also admit a  
 230 temporally-extended-action reading: each query instantiates an option  $o = \langle \mathcal{I}, \pi, \beta \rangle$  (Sutton et al.,  
 231 1999) whose internal structure Q/R/M/C audits edge-for-edge (initiation  $\rightarrow$  high-level policy over  
 232 options  $\rightarrow$  execution-to-termination), and the semi-MDP execution duration  $\tau$  explains structurally  
 233 why the interior optimum appears on  $t_{\text{succ}}$  but not on the  $\tau$ -blind chain probability  $\Phi$ . Like the  
 234 order-theoretic reading, this is an interpretive lens, not a contribution; the full development, the  
 235 belief–desire–intention gloss (shared with the collective-memory companion), and the multi-agent  
 236 option-collision account of the bottleneck shift are in Appendix B.

### 237 3 Single agent: stage diagnostics localize failure

#### 238 3.1 System, tasks, and protocol

239 **Memory stack.** The system uses a graph substrate with a semantic retrieval interface on top:  
 240 observations  $o_t$  are tokenized into symbolic scene tokens and semantic tags that update a memory  
 241 of place records, transitions, episodic traces, and concept/relation records. Planner queries return  
 242 candidates by a weighted log-score over required, preferred, and penalty tags. The retrieval machinery  
 243 is reused across all tasks and architectures; full memory schema, update rules, and scoring formula  
 244 are in Appendix D.

Table 1: Stage-factor consistency audit for the best semantic method in each task family. The product of nested conditional estimators matches semantic-path completion; overall success can be higher because some episodes succeed without traversing the full semantic retrieval path.

| Task            | Method    | Overall | SP-comp | $\hat{P}(Q^*)$ | $\hat{P}(R^* Q^*)$ | $\hat{P}(M^* R^*)$ | $\hat{P}(C^* M^*)$ | Product |
|-----------------|-----------|---------|---------|----------------|--------------------|--------------------|--------------------|---------|
| Water           | water-cp  | 0.95    | 0.70    | 1.00           | 0.98               | 0.74               | 0.97               | 0.70    |
| Rest            | rest-s2   | 0.93    | 0.73    | 1.00           | 1.00               | 0.80               | 0.91               | 0.73    |
| Goal region     | goal-n1   | 0.54    | 0.00    | 1.00           | 0.54               | 0.01               | 0.00               | 0.00    |
| Hazard recovery | hazard-s7 | 0.40    | 0.16    | 1.00           | 0.56               | 0.58               | 0.50               | 0.16    |

**Tasks, baselines, protocol.** Four single-agent task families are evaluated (water, safe-rest, goal-region on Empty-6x6 / FourRooms (Sutton et al., 1999), hazard recovery on LavaGapS7), 10 seeds  $\times$  8 episodes each,  $\times$  2 assist modes. Semantic stack (node, concept, state-conditioned variants) is compared against four baselines: random, graph-only, SPTM-like, learned BC. Reported: end-to-end success, steps, four operational stage rates, assist deltas, weakest stage per task and per method. All means are seed-aggregated; 95% bootstrap CIs use  $B=5,000$  resamples (Efron, 1979); assist comparisons use paired Wilcoxon signed-rank tests (Wilcoxon, 1945) with Holm-Bonferroni correction (Holm, 1979) (Appendix G). For the two hardest tasks (FourRooms, LavaGapS7), four oracle regimes (*base*, *oracle R/M/C*) are run — each replaces exactly one stage while leaving the rest of the pipeline unchanged. BabyAI portability check in Appendix E. All single-agent experiments ran on a 10-core ARM CPU, 32 GB RAM (CPU-only); oracle sweep  $< 1$  min wall-clock.

**LLM diagnostic block.** To check that Q/R/M/C is not tied to the symbolic planner implementation, we also run a fixed TextNav substrate with an LLM-driven observe/query/retrieve/decide loop. The TextNav task, prompts, parser contract, and event definitions are held fixed while only the local Ollama model changes; the backend uses model-compatible completion settings when an Ollama chat template otherwise hides short structured answers in a separate thinking field. This block is a controlled step-budget study (Table 3): sweeping the completion budget drives both LLM agents into a failure that Q/R/M/C localizes to the  $C$  stage while  $Q/R/M$  stay saturated, exhibiting the framework’s diagnostic property on real language agents. Reproducing full ReAct/Reflexion/Voyager/MemGPT-style agents is a separate systems contribution; the comparison to those systems is a coverage comparison in Appendix F.10.

### 3.2 Stage-estimator validation

Before the main benchmark, the factorization is validated. On a synthetic four-stage process with known probabilities the nested-estimator product tracks empirical semantic-path completion across sample sizes. On the real benchmark, the same product matches semantic-path completion rather than overall success (Table 1) — exactly the contract of Definition 1.

### 3.3 Success-only summary and what it hides

Water and rest are strong end-to-end tasks (0.95 and 0.93 success; Figure 1b). Goal-region reaches 0.54 success at the task-family level — an average that pools a solved environment (Empty-6x6, 1.00) with a hard one (FourRooms, 0.075; §3.4), so the family number is bimodal rather than a uniform mid-difficulty — but the bottleneck sits downstream of candidate selection (§3.4). Hazard recovery is the hardest retrieval-sensitive task at 0.40 success, with retrieval satisfaction as the weakest stage. Under success-only evaluation the two hard families look like the same kind of weakness; the stage view will show they are opposites. Per-task operational rates with 95% CIs are in Table 5 (Appendix D.3).

**Baseline comparison.** Table 2 compares the semantic stack against random, graph-only, SPTM-like, and learned behavior-cloned baselines. The learned baseline dominates the easiest resource tasks (1.00 on water and rest) and hazard recovery (0.74), while the semantic stack remains strongest on the relation-heavy goal-region family (0.54). The relevant contrast for this paper, however, is not the success column: the semantic stack is the only method family whose decisions expose interpretable Q/R/M/C structure. On the BC baseline the stage events are degenerate for the same structural reason

Table 2: Baseline comparison. The learned behavior-cloned baseline dominates the easiest resource tasks and hazard recovery, while the semantic stack remains strongest on the relation-heavy goal-region family and remains the only family with interpretable Q/R/M/C structure.

| Task            | Random | Graph-only | SPTM-like | Learned | Best semantic |
|-----------------|--------|------------|-----------|---------|---------------|
| Water           | 0.85   | 0.90       | 0.90      | 1.00    | 0.95          |
| Rest            | 0.85   | 0.90       | 0.90      | 1.00    | 0.93          |
| Goal region     | 0.41   | 0.49       | 0.52      | 0.50    | 0.54          |
| Hazard recovery | 0.08   | 0.39       | 0.00      | 0.74    | 0.40          |

as on the multi-agent learned baseline analysed in §4.5 — a policy without an explicit query/lock interface cannot emit separable stage events, so the diagnostic decomposition has nothing to attach to. End-to-end learning and stage diagnostics answer different questions; this paper is about the second.

### 3.4 Oracle-stage interventions localize the bottleneck

On goal-region the semantic variants solve Empty-6x6 (1.00 success) but collapse on FourRooms (0.075). For hazard recovery (LavaGapS7), base success 0.39 rises only to 0.43 under oracle controller, but jumps to 0.59 under oracle materialization and 0.60 under oracle retrieval; goal-region rejoin is downstream-limited (oracle R lifts semantic-path completion to 0.08 without moving end-to-end success). The per-task attribution as a stacked waterfall is in Figure 5 (Appendix D.3). Two tasks identical under success-only evaluation fail at opposite ends of the Q/R/M/C chain.

### 3.5 The logs refine the diagnosis: a memory-accumulation limit

Inspecting the framework’s own logs sharpens the retrieval attribution. On the 10-seed hazard-recovery run, 349 retrieval calls were logged: only 32 (9.2%) returned at least one candidate, and of these 31 (96.9% precision) selected a semantically satisfied candidate. The bottleneck is therefore not candidate *ranking* but candidate *generation* — the symbolic memory does not contain hazard-recovery-tagged nodes at 91% of decision points. A retrained 2-layer candidate-generation MLP trained on per-state features (test accuracy 0.59 vs majority-class 0.81; Appendix H) confirms that instantaneous state features cannot substitute for the missing structure. The framework’s “retrieval-limited” attribution thus refines to a *memory-accumulation* limit: candidate generation cannot be solved from state features alone, and the binding constraint is how trajectory evidence is consolidated into persistent symbolic structure. This attribution is not an artefact of the required-tag threshold  $\theta$ : sweeping  $\theta$  across the full range leaves the empty-candidate-set rate invariant (Appendix F.5), because the empties are absence of any stored candidate, not sub-threshold filtering. This sets up the multi-agent section: when consolidation becomes a shared process across  $N$  agents, it acquires a measurable *rate*.

## 4 Multi-agent: memory as a shared resource

### 4.1 Four memory architectures and the cadence axis

Four communication architectures are instantiated; they share the single-agent retrieval substrate but differ in how memory is distributed across agents (Figure 3).

**Independent.** Each agent maintains a private event store and concept graph. No inter-agent communication mechanism exists. Lower-bound baseline.

**Shared bus.** All agents publish observations to a single *collective memory* event bus. One shared concept graph is rebuilt from the union of events. Maximally centralized.

**Central aggregator (ConsensusLayer).** Each agent has its own private concept graph; a central layer pulls snapshots and computes a trust-weighted merge, producing a single *consensus report* consumed by all agents.

**Peer-to-peer broadcast (cadence  $K$ ).** Each agent has its own concept graph and its own merged view. Every  $K$  ticks each agent broadcasts a snapshot to a passive bus; each agent independently






Figure 3: Four multi-agent memory architectures. **(a) Independent:** per-agent private memory, no exchange. **(b) Shared bus:** one global memory, every agent reads/writes each tick. **(c) Central aggregator:** per-agent memories plus a hub that merges them into a consensus report. **(d) Peer-to-peer broadcast:** per-agent memories with pairwise snapshot exchange every  $K$  ticks (the cadence axis of Empirical Claim 1).

merges its own snapshot with the latest snapshot received from each peer. Different agents may hold different merged views; there is no central report.

The first three correspond to the three regimes in the multi-agent communication literature: no exchange, CTDE-style aggregation, and decentralized peer communication. The fourth is the architecture for which Definition 2 and Empirical Claim 1 are directly applicable, since the cadence  $K$  is a tunable parameter rather than a fixed protocol. Mechanically,  $K$  is the rate at which private memories consolidate into a collective view, which is why it is the structural axis of this section.

## 4.2 Multi-agent factorization

Extend the factorization to a population of  $N$  agents that share information through a communication protocol  $\mathcal{C}_K$  parameterized by broadcast cadence  $K$ . Each agent  $i$  has its own memory state  $\mathcal{M}_i$  and issues its own query  $q_i$ . Let  $\mathcal{M}_i^{(K)}$  denote the projection of agent  $i$ 's memory at the moment it queries, after absorbing every peer message received under cadence  $K$ . Two coupling channels are separated: (i) memory coupling  $\mathcal{M}_{-i}^{(K)}$ , the latest peer memory snapshots accessible to agent  $i$ ; (ii) occupancy coupling  $\mathcal{O}_{-i}$ , the set of targets that other agents have already claimed by the time  $i$  acts.

**Definition 2 (multi-agent stage estimability).** For each agent  $i$ ,

$$P(C_i^* \mid q_i, \mathcal{M}_i, \mathcal{C}_K) = P(Q_i^* \mid q_i, \mathcal{M}_i) P(R_i^* \mid Q_i^*, q_i, \mathcal{M}_i, \mathcal{M}_{-i}^{(K)}) \\ \cdot P(M_i^* \mid R_i^*, q_i, \mathcal{M}_i, \mathcal{M}_{-i}^{(K)}) P(C_i^* \mid M_i^*, q_i, \mathcal{O}_{-i}).$$

Under a multi-agent stage-Markov assumption (Appendix C), the four factors are estimable from per-agent trajectory logs and the empirical-frequency estimators are consistent.

**When Definition 2 may fail.** The single-agent stage-Markov property absorbs trajectory history into  $(q, \mathcal{M})$ . In the multi-agent setting, the analogous absorption requires that dependence on past peer decisions enters only through  $\mathcal{M}_{-i}^{(K)}$  at the M stage and through  $\mathcal{O}_{-i}$  at the C stage. This step is non-trivial:  $\mathcal{O}_{-i}$  summarises target occupancy at decision time but loses information about *how* peers arrived there. The factorization is therefore a projection onto observable per-agent traces under that summary. Two specific failure modes arise: (i) inter-agent communication side-channels not logged through  $\mathcal{C}_K$ ; (ii) coordinated peer policies whose decision at the C stage depends on hidden peer state not summarised by  $\mathcal{O}_{-i}$ .

## 4.3 Bottleneck shift

**Empirical Claim 1 (bottleneck shift  $M \leftrightarrow C$ , slope-level).** Assume scarcity ( $N > T$  targets, see §2.1) and a peer-broadcast architecture with cadence  $K \in \mathbb{N}$ . Two stylized assumptions are posited: **(F)** memory freshness  $\nu(\mathcal{M}_{-i}^{(K)})$  is non-decreasing pointwise as  $K$  decreases, and the probability

that agent  $i$ 's planner locks onto a reachable target is non-decreasing in  $\nu$ ; **(O)** the probability that an arbitrary committed target is already in  $\mathcal{O}_{-i}$  is non-increasing as  $K$  increases. Write  $P(M^*|R^*; K)$  and  $P(C^*|M^*; K)$  for the population-level conditional probabilities (averaged over agents and the seed distribution of the benchmark);  $\hat{P}(\cdot)$  denotes the corresponding empirical-frequency estimator computed from logs. Under (F) and (O),

$$\frac{\partial P(M^*|R^*; K)}{\partial K^{-1}} > 0, \quad \frac{\partial P(C^*|M^*; K)}{\partial K^{-1}} < 0.$$

Both slope conditions are direct empirical predictions that are testable on logs through  $\hat{P}$ .

**What this claim is and is not.** Empirical Claim 1 is a *slope-level statement on conditional rates*: each of the two slopes is testable and is confirmed on the 35,640-run sweep (§4.5) and independently on a 100-run standard-MiniGrid portability sweep (§6). We report the slopes as *effect sizes with cluster-robust confidence intervals* (resampled over the 81 independent configuration cells rather than the correlated run count) in Appendix F.3. The claim's boundary with respect to the conditional product  $\Phi$  is stated once, in the Scope-and-non-claims box (§1) and in §4.5.

**Chronology.** Empirical Claim 1 is post-hoc. It was formulated *after* an initial 12,960-run sweep conducted under four stage-agnostic hypotheses (Appendix F); the two slope conditions were then confirmed on the full 35,640-run sweep, on the  $N = 12$  scale-up (Appendix E), and on the MiniGrid portability wrapper (§6). We state this explicitly: the verified slope signs are predictions on enlarged data, not a pre-registered prediction. A fully formal version is left for future work.

#### 4.4 Sweep protocol

The four architectures are evaluated on a custom multi-agent grid ( $12 \times 10$  cells); each agent must reach a cell tagged `water_source`, choosing actions via a shared memory-driven planner so only the memory layer differs. The sweep crosses  $N \in \{3, 5, 8\}$ ,  $T \leq N$ , three layouts (symmetric/asymmetric/random), three hazard densities, four architectures, eight peer cadences  $K \in \{1, 2, 4, 8, 16, 32, 48, 64\}$ , 40 seeds — 35,640 runs in  $\approx 22$  min on 8 cores. Per-agent per-tick Q/R/M/C indicators are logged (Appendix C); episode-level starred events are the disjunction over ticks. A separate oracle-intervention block on scarcity cases  $(N, T) \in \{(3, 2), (5, 3), (8, 5)\}$  at five cadences evaluates oracle R (memory returns ground-truth water positions), oracle M (planner locks onto nearest real water), and oracle C (completion of a real materialised lock is guaranteed; Appendix F.12).

#### 4.5 Results: stage decomposition and bottleneck shift

The 35,640-run sweep is reported through four observations, structured to mirror Empirical Claim 1: (i) the conditional decomposition isolates the M and C stages as the differentiating links; (ii) the bottleneck shift appears directly as a negative within-cadence correlation between  $P(M^*|R^*)$  and  $P(C^*|M^*)$ , peaking at  $K = 4$  and decaying to noise at  $K = 64$ ; (iii) on mean time-to-success, the resulting trade-off produces an interior optimum at  $K = 8$  — robust in location but shallow in margin, with the advantage over the adjacent fast cadences not resolved at this sample size (Appendix F.3); (iv) on the conditional product  $\Phi$ , the same trade-off produces a monotone trend rather than an interior peak — a summary statistic on which we make no formal claim. A learned-communication baseline and a multi-agent oracle-intervention loop close the subsection.

##### 4.5.1 Conditional decomposition reveals where each architecture fails.

Figure 4 shows per-architecture conditional Q/R/M/C rates; full numerics with bootstrap CIs are in Appendix F.1, Table 8. Retrieval conditional  $P(R^*|Q^*)$  is saturated; the differentiating signal lives at M and C. Shared-bus, central-aggregator, and fast-broadcast peer ( $K \leq 4$ ) have high  $P(M^*|R^*)$  but low  $P(C^*|M^*)$ ; independent and slow-broadcast peer ( $K \geq 16$ ) show the reverse — the  $M \leftrightarrow C$  trade-off predicted by Empirical Claim 1.

##### 4.5.2 Bottleneck shift and interior efficiency optimum.

Across all peer runs, Spearman of  $P(M^*|R^*)$  vs  $K$  is  $-0.53$  (95% cluster-robust CI  $[-0.59, -0.47]$ ) and of  $P(C^*|M^*)$  vs  $K$  is  $+0.43$  (CI  $[+0.37, +0.50]$ ); both are moderate effect sizes whose




Figure 4: Per-architecture conditional rates (top) and mean time-to-success (bottom) from the 35,640-run sweep. Shared bus, central aggregator, and fast peer broadcast ( $K \leq 4$ ) live in the M-saturation / C-collapse regime; independent and slow peer ( $K \geq 16$ ) live in the M-deficit / C-saturation regime. Mean time-to-success has an interior minimum at  $K=8$  (7.47 ticks), below the independent-agent baseline (8.02); the location is robust but shallow — the margin over the adjacent fast cadences is not resolved (Appendix F.3).

401 CIs exclude zero under resampling of the 81 configuration cells (Appendix F.3), match-  
 402 ing the slope conditions of Empirical Claim 1. We report these rather than  $p$ -values,  
 403 which are uninformative at this correlated sample size. The stronger empirical signature  
 404 is at the per-configuration level: Pearson of  $P(M^* | R^*)$  vs  $P(C^* | M^*)$  within a fixed  $K$  is  
 405  $-0.43, -0.54, -0.61, -0.19, -0.05, -0.04, -0.02, -0.02$  at  $K=1, 2, 4, 8, 16, 32, 48, 64$ , peaking  
 406 at  $K=4$  and decaying to noise by  $K=64$  (Figure 8, Appendix F.1). Configurations where  $M$  is  
 407 high systematically have low  $C$ : the bottleneck migrates between the two stages within the same  
 408 ( $N, T$ , layout, hazard, seed) cell as  $K$  varies. On the efficiency metric, mean  $t_{\text{succ}}$  has a U-shape with  
 409 interior minimum at  $K=8$  (7.47 ticks vs 8.02 for the independent baseline); the minimum’s location  
 410 is robust but shallow — its margin over the adjacent fast cadences ( $K=1, 4$ ) is not resolved under  
 411 cell-level resampling (paired-bootstrap CIs in Appendices F.1 and F.3). Scaling to  $N=12$  sharpens  
 412 the signature: the peak Pearson rises to  $-0.70$  and migrates outward to  $K=8$ , consistent with  
 413 stronger occupancy coupling at higher  $N$  delaying the cost-regime transition (Appendix E, Figure 6).

414 **Provenance stratification: the mechanism behind the shift.** Annotating each  $M^*$  lock with the  
 415 *provenance* of the evidence behind it makes the bottleneck shift mechanistic rather than correlational.  
 416 Let  $\varphi$  be the foreign share of the evidence mass supporting the locked candidate, computed from  
 417 the same trust/support weights the merge itself used and frozen at lock time, and call a lock a *warp*  
 418 ( $W^*$ ) when  $\varphi \geq 0.8$ : the agent is committing to a target it knows essentially only from peers. On  
 419 a dedicated peer sweep over the scarcity cells (random layouts, 20 seeds,  $K \in \{1, 2, 4, 8, 16\}$ ) the  
 420 warp share  $P(W^* | M^*)$  falls from 0.55 at  $K=1$  to 0.11 at  $K=8$  while the own-evidence stratum  
 421 completes at a flat 0.26–0.34 and the warp stratum completes below 0.05 throughout: *the C-collapse*  
 422 *of fast sharing is concentrated in the warp stratum*. What cadence controls, mechanically, is how  
 423 much of the collective’s materialisation runs on foreign evidence — and that is precisely the stratum  
 424 that loses the occupancy race. The same stratification reproduces on the continuous VMAS bridge  
 425 ( $0.72 \rightarrow 0.00$  over  $K=1/4/16$ ) and on an LLM text substrate ( $0.42 \rightarrow 0.08$  over  $K=2/8$ ), with the  
 426 curves nearly coinciding across substrates. A full provenance-conditioned analysis (counterfactual  
 427 attribution by masked replay, a closed-form staleness horizon with a registered 6/6-exact hold-out,  
 428 and a decentralized repair protocol) appears in a companion manuscript; here the stratification serves  
 429 as the mechanistic reading of Empirical Claim 1.

430 **What we do not claim about  $\Phi$ .** The conditional product  $\Phi = P(M^*|R^*) \cdot P(C^*|M^*)$  does  
 431 *not* have an interior peak in the tested range (mean-of-products estimator increases monotonically  
 432 from 0.569 at  $K=4$  to 0.603 at  $K=64$ , converging to the independent-agent boundary 0.605). The  
 433 slope-level statement of Empirical Claim 1 is on the conditional rates themselves, not on  $\Phi$ ; we  
 434 make no formal claim about the shape of  $\Phi(K)$ . Jensen-gap detail and POM-MOP comparison in  
 435 Appendix F.1.

436 **Learned-communication baseline: stage events are structurally degenerate even at competi-**  
 437 **itive success.** A CommNet-style policy (Sukhbaatar et al., 2016) trained with PPO (Schulman  
 438 et al., 2017) on the same scenario ( $N=3$ ,  $T=2$ ) reaches 0.667 held-out success ( $n=100$  seeds)  
 439 — matching the symbolic peer architecture at  $K=8$  (0.65). Despite this, the Q/R/M/C profile re-  
 440 mains structurally degenerate:  $Q^*=1.00$  saturated by construction;  $R^*=0.997$ ,  $M^*=0.997$  collapsed  
 441 together ( $P(M^*|R^*)=1.00$ ). This is the verified structural claim: a real-valued comm vector is  
 442 never informatively empty (Q saturates) and lacks an explicit target-lock interface (R/M merge),  
 443 regardless of training budget. The multi-agent counterpart of the single-agent learned-BC analysis  
 444 (§3.3). Full architecture, PPO hyperparameters, and side-by-side REINFORCE vs PPO comparison  
 445 in Appendix H.

446 **LLM model sweep: the logger localizes a real failure in language agents.** Table 3 runs the  
 447 same Q/R/M/C event contract on two local LLM agents (qwen3:4b and llama3.1 via Ollama) on  
 448 TextNav, sweeping the step budget over  $\{6, 12, 25\}$  (8 episodes each; the task needs  $\sim 10$  low-level  
 449 steps). At a generous budget (12, 25) both agents succeed and all four stages fire. At a budget  
 450 below the requirement (6), success drops to 0/8 for *both* models — and the decomposition shows  
 451 *where*:  $Q^*=R^*=M^*=1.00$  (the agents form the query, retrieve, and lock the correct target every  
 452 episode) while  $C^*=0.00$  (none complete in time). The framework localizes the failure precisely to  
 453 a budget-limited completion ( $C$ ) rather than to query formation, retrieval, or materialization; that  
 454 both models fail identically shows the failure is structural, not a model artifact. This is the diagnostic  
 455 property the framework claims, now exhibited on *real language agents*: an otherwise opaque success  
 456 collapse ( $1.00 \rightarrow 0.00$ ) is resolved into a specific failing stage. (Instrumentation note: a naive  
 457 qwen3:4b/Ollama chat-template call initially returned empty response strings because the model  
 458 spent the short structured-output budget in a separate thinking field; we use raw structured completion  
 459 for qwen3:4b and a first-line parser, recorded for reproducibility.)

460 **The multi-agent phenomenology reproduces on a live LLM collective.** Beyond the single-agent  
 461 budget study, the same event contract instruments a collective of  $N=3$  LLM agents (llama3.1)  
 462 sharing symbolic place memory by peer broadcast under scarcity ( $M=2$  takeable targets, four layouts,  
 463  $K \in \{2, 8\}$ ), with all predictions registered before the runs. The cadence phenomenology of §4.5  
 464 carries over to the language substrate: a paired comparison against the no-communication baseline on  
 465 identical layouts shows raw sharing *lowering* success from 2/3 to 1/3 per episode — the occupancy-  
 466 racing C-collapse, now produced by agents whose queries, tag extraction, and commitments are  
 467 LLM-made — and target collisions appear in 7/8 shared-memory episodes. Attaching the reservation-  
 468 and-rollback repair protocol (the constructive answer to open problem 3 of Appendix F.8) recovers  
 469 part of the loss ( $0.333 \rightarrow 0.417$  at both cadences, rollback latency again  $\approx K$ ). Two interface findings  
 470 transfer to any LLM-over-symbolic-memory design: LLM query formers free-generate retrieval  
 471 tags outside the memory’s vocabulary (requiring tolerant retrieval gates), and small local models are  
 472 reliable at semantic commitment but not at coordinate arithmetic or systematic exploration (requiring  
 473 layered execution). Full protocol and provenance analysis in the companion manuscript.

474 Scope: these are controlled studies on local models; scaling the same contract to full  
 475 Voyager/MemGPT-style agents on richer tasks with multiple failure modes is developed next (external-  
 476 framework validation) and in Appendix F.10.

477 **Multi-agent oracle interventions close the diagnose→intervene→verify loop.** On the scarcity  
 478 cases, mean  $t_{\text{succ}}$  at small/intermediate cadence ( $K \in \{1, 2, 4\}$ ) drops from  $\sim 10$ –11 in baseline to  
 479  $\sim 7.2$  under either oracle R or oracle M, and further to  $\sim 6.1$  under oracle C, which additionally  
 480 removes post-lock travel; at slow cadence ( $K=16$ ) the baseline is already within  $\sim 1.2$  of the R/M  
 481 oracle floor (Table 15 in Appendix F.11). Crucially, oracle C gives the largest time reduction but  
 482 the smallest success-rate gain (Appendix F.12): in scarcity, completion costs time while retrieval

Table 3: LLM TextNav step-budget sweep: Q/R/M/C localizes a completion-limited failure in two local Ollama agents. Eight episodes per (model, step limit), temperature 0.0; the task needs  $\sim 10$  low-level steps. At a budget below that requirement (step limit 6) success drops to 0.00 for both models, yet  $Q^*=R^*=M^*=1.00$  (the agents form the query, retrieve, and lock the correct target every episode) while  $C^*=0.00$ : the framework localizes the failure to the C stage, not to Q/R/M. Both models behave identically, so the failure is structural (budget), not model-specific. Artifacts: tmp/llm\_steplimit\_{6,12,25}.

| Model    | Step limit | Success | $Q^*$ | $R^*$ | $M^*$ | $C^*$       |
|----------|------------|---------|-------|-------|-------|-------------|
| qwen3:4b | 6          | 0.00    | 1.00  | 1.00  | 1.00  | <b>0.00</b> |
| llama3.1 | 6          | 0.00    | 1.00  | 1.00  | 1.00  | <b>0.00</b> |
| qwen3:4b | 12         | 1.00    | 1.00  | 1.00  | 1.00  | 1.00        |
| llama3.1 | 12         | 1.00    | 1.00  | 1.00  | 1.00  | 1.00        |
| qwen3:4b | 25         | 1.00    | 1.00  | 1.00  | 1.00  | 1.00        |
| llama3.1 | 25         | 1.00    | 1.00  | 1.00  | 1.00  | 1.00        |

483 and materialisation decide success. The cadence-induced excess time at small/intermediate  $K$  is  
484 attributable to the  $M \leftrightarrow C$  transition; at large  $K$  the bottleneck has moved into solo exploration.

485 The multi-agent block thus delivers its half of the argument. The cadence  $K$  is, mechanically, the rate  
486 at which private memories consolidate into a collective view; the data show that this rate has a non-  
487 trivial interior optimum ( $K=8$  on  $t_{\text{succ}}$ ), that its cost channel is occupancy coupling, and that both  
488 effects are visible stage-by-stage on the same Q/R/M/C chain used for single agents. Consolidation is  
489 not only the single-agent limit; in a collective it is a tunable, measurable design axis.

## 490 5 The contract on third-party agent stacks

491 The strongest external-validity question a reader can ask of a diagnostic protocol is whether it  
492 works on systems its authors did not build. We therefore instrument three third-party agent stacks  
493 — a function-calling loop on the OpenAI SDK, LangGraph’s prebuilt ReAct agent, and AutoGen’s  
494 AssistantAgent/UserProxy pair, all driving the same local model (llama3.1, temperature 0) — with  
495 the identical Q/R/M/C contract, defined once over *tool traces*: an episodic household memory is  
496 exposed through `recall/goto/take` tools, and  $Q^*$  fires on querying memory,  $R^*$  on *attribution* (a  
497 returned record names the item at its true place),  $M^*$  on an episode-level lock (a `goto` to the true  
498 place, faithful to Definition 1; first-commitment discipline  $M_1$  is logged as an auxiliary lock-quality  
499 diagnostic), and  $C^*$  on task completion within a tool budget. Three variants engineer one failure each  
500 — a consolidation gap (the location fact was never written to memory), a stale-distractor ambiguity,  
501 a tool budget below the task’s requirement — plus a clean control; all predictions were registered  
502 before the runs, and the two revisions of the event semantics (an attribution leak; an over-tightened  
503  $M^*$ ) are recorded in the registration file.

504 Three findings (Table 4). **(1) Structural localizations are exact and framework-independent.**  
505 When the memory lacks the fact,  $R^*=0.00$  on every stack; when the budget cannot fit completion,  
506  $C^*=0.00$  at  $Q^*=R^*=1$  on every stack. The contract needs no per-framework adaptation. **(2) The**  
507 **clean control separates the frameworks, and the decomposition says how.** End-to-end success  
508 on the SAME clean episodes ranges 0.38 (AutoGen) to 0.75 (OpenAI SDK). A leaderboard would  
509 report “different quality”; the stages show a shared pathology plus different recovery loops: first-  
510 commitment discipline is weak everywhere ( $M_1=0.12\text{--}0.50$  — the agent’s first `goto` frequently  
511 ignores its own just-retrieved evidence), and the frameworks differ in whether they *recover* from  
512 the bad first lock within budget (OpenAI SDK re-queries and recovers,  $M^*=0.75 \rightarrow C^*=0.75$ ;  
513 AutoGen locks the true place eventually but burns its turns,  $M^*=0.88$  vs.  $C^*=0.38$ ). **(3) The stale**  
514 **distractor pulls, measurably and unevenly.** The outdated record captures the first commitment in  
515 0.38–0.62 of ambiguous episodes — most strongly for AutoGen — the LLM-memory analogue of the  
516 staleness gating that the collective experiments showed to be decisive. Honestly failed registrations:  
517 the behavioral predictions that assumed a clean control (all-stages  $\geq 6/8$ ; an ambiguity-induced  
518 completion drop below control) broke *because the frameworks themselves fail a clean memory task*  
519 *25–62% of the time* — which is finding (2), not noise; both are recorded as registered failures with

Table 4: Q/R/M/C on three third-party agent frameworks (llama3.1, 8 episodes per cell, identical episodes/tools/budgets). Structural failures localize identically everywhere: the consolidation gap is  $R=0.00$  and the budget failure is  $C=0.00$  at  $Q=R=1$ , for all three stacks. On the *clean* control the frameworks spread from 0.38 to 0.75 end-to-end — and the decomposition attributes the spread (see text).  $M_1$ : first goto targets the true place;  $stale_1$ : first goto targets the stale distractor.

| Framework  | Variant           | $Q^*$ | $R^*$       | $M^*$ | $C^*$       | $M_1$ | $stale_1$ |
|------------|-------------------|-------|-------------|-------|-------------|-------|-----------|
| OpenAI SDK | control           | 1.00  | 1.00        | 0.75  | 0.75        | 0.12  | —         |
|            | consolidation gap | 1.00  | <b>0.00</b> | 0.38  | 0.38        | 0.12  | —         |
|            | stale ambiguity   | 1.00  | 1.00        | 0.88  | 0.12        | 0.12  | 0.38      |
|            | tight budget      | 1.00  | 1.00        | 0.12  | <b>0.00</b> | 0.12  | —         |
| LangGraph  | control           | 1.00  | 1.00        | 0.75  | 0.50        | 0.50  | —         |
|            | consolidation gap | 1.00  | <b>0.00</b> | 0.38  | 0.25        | 0.12  | —         |
|            | stale ambiguity   | 1.00  | 1.00        | 0.62  | 0.62        | 0.38  | 0.38      |
|            | tight budget      | 1.00  | 1.00        | 0.12  | <b>0.00</b> | 0.12  | —         |
| AutoGen    | control           | 1.00  | 1.00        | 0.88  | 0.38        | 0.38  | —         |
|            | consolidation gap | 1.00  | <b>0.00</b> | 0.50  | 0.25        | 0.12  | —         |
|            | stale ambiguity   | 1.00  | 1.00        | 0.38  | 0.25        | 0.12  | 0.62      |
|            | tight budget      | 1.00  | 1.00        | 0.12  | <b>0.00</b> | 0.12  | —         |

mechanisms. Scope: one model,  $n=8$  per cell, three stacks; Letta/CrewAI (Python  $\geq 3.10$ ) and API-scale models are the natural extension.

## 6 Limitations and future work

**Portability across substrates: slope conditions hold on MiniGrid.** The slope conditions of Empirical Claim 1 also hold on a standard-MiniGrid substrate. A FourRooms-based multi-agent wrapper (built from `minigrid.core.grid.Grid + Wall/Floor/Lava` primitives, identical operational  $Q^*, R^*, M^*, C^*$  definitions; `experiments/minigrid_multiagent_wrapper/`) was used to run a 100-run portability sweep ( $K \in \{1, 4, 8, 16, 64\}$ , 20 seeds,  $N=3$ ,  $T=2$ ): Spearman  $P(M^*|R^*)$  vs  $K$  is  $-0.50$  and Spearman  $P(C^*|M^*)$  vs  $K$  is  $+0.50$  (moderate rank correlations; this is a small 100-run sweep, so we report the effect sizes and do not lean on significance). The interior  $t_{succ}$  minimum does not separate at this minimal sweep size (5 cadences, 20 seeds, one hazard density); a denser sweep is the natural next step. The two-substrate seam is therefore closed at the slope-level claim.

**Empirical Claim 1 is slope-level on conditional rates only.** The two slope conditions hold as moderate, cluster-robust effect sizes (Appendix F.3); we make no formal claim on the shape of the conditional product  $\Phi$  (§4.3). A future statement directly in terms of the time-cost surface, or a derivation under a generative model (Poisson-renewal memory, birth-death occupancy), would convert the empirical claim into a theorem.

**Estimability under hidden state.** With unlogged in-episode query rewrites or peer side-channels, Definitions 1 and 2 are projections onto logged traces, not full identification.

**Shared coordinate frame.** All multi-agent results assume agents interpret places in a common coordinate frame: cross-agent place identity is coordinate proximity, with semantic profiles as a tie-breaker inside a spatial radius. The place-correspondence problem — whether two agents’ records denote the same place at all — is therefore bracketed here, not solved. A companion manuscript removes the assumption constructively (fingerprint-based place identity with consensus frame recovery up to translation and  $90^\circ$  rotation) and shows the Q/R/M/C measurements, including the cadence slopes and the staleness horizon, unchanged on top of the recovered frames; the harder variants (continuous  $SE(2)$ , appearance noise, symbol alignment) remain open.

**Learned baselines.** The CommNet PPO baseline reaches 0.667 success on the scarcity scenario — competitive with symbolic peer at  $K=8$  (0.65) — and confirms Q-saturation ( $Q^*=1.00$ ) and R/M collapse ( $P(M^*|R^*)=1.00$ ) at competitive success. The structural prediction now holds verified, not asserted; ablating on architectures with explicit symbolic target-lock channels is left for future work.

552 **Scale.** Multi-agent sweep bounded at  $N=12$  (Appendix E; signature sharpens, peak Pearson  $-0.70$  at  
553  $K=8$ , interior  $t_{\text{succ}}$  minimum 4.77). Denser scaling at  $N \geq 16$  is future work. **Continuous substrate.**  
554 A preliminary VMAS (Bettini et al., 2022) check reproduces *both* slope signs of Empirical Claim 1  
555 in the coupling regime (Spearman  $-1.00$  and  $+1.00$  over  $K \leq 16$ , 40 seeds; Appendix F.7), reusing  
556 the symbolic memory through a discretisation bridge. This is a directional signal only — a perfect  
557 rank correlation over five monotone points, one configuration, discretised memory — so it indicates  
558 the shift is not an artefact of grid cells, but a full continuous-state validation with cluster-robust CIs  
559 remains future work.

## 560 7 Broader impacts

561 The paper proposes diagnostic evaluation for memory-based agents in simulated navigation. Its posi-  
562 tive impact is methodological: making memory-based failures more measurable across architectures,  
563 and making the cadence axis of distributed multi-agent memory tunable on an empirically grounded  
564 criterion. Potential negative impacts are indirect and would arise if similar diagnostics were used to  
565 improve agents in safety-critical settings (drone swarms, autonomous vehicles, distributed sensor  
566 networks) without adequate safeguards. The current work is benchmark-level infrastructure, not a  
567 deployment-ready system.

## 568 8 Conclusion

569 We proposed Q/R/M/C, a four-stage decomposition of memory-based navigation whose factors are  
570 estimable from standard trajectory logs, and used it as a diagnostic instrument across symbolic,  
571 learned, LLM-driven, single-agent, and multi-agent traces. The main result is not a new agent  
572 architecture: it is a reproducible logging and estimation protocol that separates query formation,  
573 retrieval satisfaction, target materialization, and post-retrieval completion. Across the experiments,  
574 this protocol localizes failures that success-only evaluation cannot, exposes degenerate observability  
575 in implicit learned policies, and makes memory-sharing cadence measurable as a stage-level design  
576 variable.

## 577 References

- 578 M. Andrychowicz, F. Wolski, A. Ray, J. Schneider, R. Fong, P. Welinder, B. McGrew, J. Tobin, P. Abbeel, and  
579 W. Zaremba. Hindsight experience replay. In *NeurIPS*, 2017.
- 580 R. Bělohávek. Fuzzy Galois connections. *Mathematical Logic Quarterly*, 45(4):497–504, 1999.
- 581 M. Bettini, R. Kortvelesy, J. Blumenkamp, and A. Prorok. VMAS: A vectorized multi-agent simulator for  
582 collective robot learning. In *16th International Symposium on Distributed Autonomous Robotic Systems*  
583 (*DARS*), 2022. arXiv:2207.03530.
- 584 S. Boyd, A. Ghosh, B. Prabhakar, and D. Shah. Randomized gossip algorithms. *IEEE Trans. Information Theory*,  
585 52(6):2508–2530, 2006.
- 586 M. Chevalier-Boisvert, D. Bhoopchand, H. Abdolmaleki, and A. Mordatch. BabyAI: A platform to study the  
587 sample efficiency of grounded language learning. In *ICLR*, 2019.
- 588 M. Chevalier-Boisvert, B. Lahlou, S. Willems, and P. S. Castro. MiniGrid and MiniWorld: Modular reinforcement  
589 learning environments. *arXiv:2306.13831*, 2023.
- 590 P. Chhikara, D. Khant, S. Aryan, T. Singh, and D. Yadav. Mem0: Building production-ready AI agents with  
591 scalable long-term memory. *arXiv:2504.19413*, 2025.
- 592 P. Dayan. Improving generalization for temporal difference learning: The successor representation. *Neural*  
593 *Computation*, 5(4):613–624, 1993.
- 594 M. Deitke, E. V. Wijnmans, D. Schwenk, O. Salvador, K. M. Ehsani, et al. ProcTHOR: Large-scale embodied AI  
595 using procedural generation. In *NeurIPS Datasets and Benchmarks*, 2022.
- 596 B. Efron. Bootstrap methods: Another look at the jackknife. *Annals of Statistics*, 7(1):1–26, 1979.
- 597 J. N. Foerster, Y. M. Assael, N. de Freitas, and S. Whiteson. Learning to communicate with deep multi-agent  
598 reinforcement learning. In *NeurIPS*, 2016.

599 J. N. Foerster, G. Farquhar, T. Afouras, N. Nardelli, and S. Whiteson. Counterfactual multi-agent policy gradients.  
600 In *AAAI*, 2018.

601 B. Ganter and R. Wille. *Formal Concept Analysis: Mathematical Foundations*. Springer, 1999.

602 D. Ghosh, A. Gupta, A. Reddy, J. Fu, S. Levine, and C. Finn. Learning to reach goals via iterated supervised  
603 learning. In *ICLR*, 2021.

604 G. Gigerenzer and D. G. Goldstein. Reasoning the fast and frugal way: Models of bounded rationality.  
605 *Psychological Review*, 103(4):650–669, 1996.

606 S. Gupta, J. Davidson, S. Levine, R. Sukthankar, and J. Malik. Cognitive mapping and planning for visual  
607 navigation. In *CVPR*, 2017.

608 S. Holm. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70,  
609 1979.

610 E. Kolve, R. Mottaghi, D. Gordon, Y. Zhu, A. Gupta, A. Farhadi, and D. Fox. AI2-THOR: An interactive 3D  
611 environment for visual AI. *arXiv:1712.05474*, 2017.

612 B. Kuipers. The spatial semantic hierarchy. *Artificial Intelligence*, 119(1–2):191–233, 2000.

613 F. Lieder and T. L. Griffiths. Resource-rational analysis: Understanding human cognition as the optimal use of  
614 limited computational resources. *Behavioral and Brain Sciences*, 43:e1, 2020.

615 R. Lowe, Y. Wu, A. Tamar, J. Harb, P. Abbeel, and I. Mordatch. Multi-agent actor-critic for mixed cooperative-  
616 competitive environments. In *NeurIPS*, 2017.

617 S. Mac Lane. *Categories for the Working Mathematician*. Springer, 2nd edition, 1998.

618 I. Momennejad, E. M. Russek, J. H. Cheong, M. M. Botvinick, N. Daw, and S. J. Gershman. The successor  
619 representation in human reinforcement learning. *Nature Human Behaviour*, 1:680–692, 2017.

620 J. O’Keefe and L. Nadel. *The Hippocampus as a Cognitive Map*. Oxford University Press, 1978.

621 C. Packer, V. Fang, S. Patil, K. Lin, S. Wooders, and J. Gonzalez. MemGPT: Towards LLMs as operating  
622 systems. *arXiv:2310.08560*, 2023.

623 J. S. Park, J. C. O’Brien, C. J. Cai, M. R. Morris, P. Liang, and M. S. Bernstein. Generative agents: Interactive  
624 simulacra of human behavior. *arXiv:2304.03442*, 2023.

625 T. Rashid, M. Samvelyan, C. Schroeder de Witt, G. Farquhar, J. Foerster, and S. Whiteson. QMIX: Monotonic  
626 value function factorisation for deep multi-agent reinforcement learning. In *ICML*, 2018.

627 N. Savinov, A. Dosovitskiy, and V. Koltun. Semi-parametric topological memory for navigation. In *ICLR*, 2018.

628 M. Savva, A. Kadian, O. Maksymets, Y. Zhao, E. Wijmans, B. Jain, J. Straub, J. Liu, V. Koltun, J. Malik, et al.  
629 Habitat: A platform for embodied AI research. In *ICCV*, 2019.

630 J. Schulman, P. Moritz, S. Levine, M. Jordan, and P. Abbeel. High-dimensional continuous control using  
631 generalized advantage estimation. In *ICLR*, 2016. *arXiv:1506.02438*.

632 J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov. Proximal policy optimization algorithms.  
633 *arXiv:1707.06347*, 2017.

634 N. Shinn, F. Cassano, A. Berman, and A. G. Labash. Reflexion: Language agents with verbal reinforcement  
635 learning. In *NeurIPS*, 2023.

636 H. A. Simon. A behavioral model of rational choice. *Quarterly Journal of Economics*, 69(1):99–118, 1955.

637 H. A. Simon. Rational choice and the structure of the environment. *Psychological Review*, 63(2):129–138, 1956.

638 K. L. Stachenfeld, M. M. Botvinick, and S. J. Gershman. The hippocampus as a predictive map. *Nature  
639 Neuroscience*, 20(11):1643–1653, 2017.

640 S. Sukhbaatar, A. Szlam, and R. Fergus. Learning multiagent communication with backpropagation. In *NeurIPS*,  
641 2016.

642 R. S. Sutton, D. Precup, and S. Singh. Between MDPs and semi-MDPs: A framework for temporal abstraction  
643 in reinforcement learning. *Artificial Intelligence*, 112(1–2):181–211, 1999.



- 644 E. C. Tolman. Cognitive maps in rats and men. *Psychological Review*, 55(4):189–208, 1948.
- 645 G. Wang, Y. Xie, Z. Jiang, A. Mandelkar, C. Xie, A. Yu, and P. Liang. Voyager: An open-ended embodied agent  
646 with large language models. *arXiv:2305.16291*, 2023.
- 647 Q. Wu, G. Bansal, J. Zhang, Y. Wu, B. Li, E. Zhu, L. Jiang, X. Zhang, S. Zhang, J. Liu, A. H. Awadallah, R. W.  
648 White, D. Burger, and C. Wang. AutoGen: Enabling next-gen LLM applications via multi-agent conversation.  
649 *arXiv:2308.08155*, 2023.
- 650 F. Wilcoxon. Individual comparisons by ranking methods. *Biometrics Bulletin*, 1(6):80–83, 1945.
- 651 R. J. Williams. Simple statistical gradient-following algorithms for connectionist reinforcement learning.  
652 *Machine Learning*, 8(3–4):229–256, 1992.
- 653 S. Yao, J. Zhao, D. Yu, N. Du, I. Shafran, K. Narasimhan, and Y. Cao. ReAct: Synergizing reasoning and acting  
654 in language models. In *ICLR*, 2023.
- 655 L. Xiao, S. Boyd, and S. Lall. A scheme for robust distributed sensor fusion based on average consensus. *IPSN*,  
656 2007.
- 657 C. Yu, A. Velu, E. Vinitisky, J. Gao, Y. Wang, A. Bayen, and Y. Wu. The surprising effectiveness of PPO in  
658 cooperative multi-agent games. In *NeurIPS Datasets and Benchmarks Track*, 2022.

## 659 A Formal proofs

### 660 A.1 Argument for Definition 1 (single-agent estimability)

661 For each episode  $i \in \{1, \dots, n\}$  let  $Q_i^*, R_i^*, M_i^*, C_i^*$  denote the nested episode-level indicators. By  
662 construction these are measurable functions of the runtime trajectory log and satisfy  $C_i^* \Rightarrow M_i^* \Rightarrow$   
663  $R_i^* \Rightarrow Q_i^*$ . The empirical estimators

$$\hat{p}_Q = \frac{1}{n} \sum_i Q_i^*, \quad \hat{p}_{R|Q} = \frac{\sum_i R_i^*}{\sum_i Q_i^*}, \quad \hat{p}_{M|R} = \frac{\sum_i M_i^*}{\sum_i R_i^*}, \quad \hat{p}_{C|M} = \frac{\sum_i C_i^*}{\sum_i M_i^*}$$

664 are consistent by the chain rule plus the conditional stage-Markov property, and the empirical product  
665 converges to semantic-path completion.  $\square$

### 666 A.2 Argument for Definition 2 (multi-agent estimability)

667 Conditioning sets are enlarged by  $\mathcal{M}_{-i}^{(K)}$  and  $\mathcal{O}_{-i}$ . Under the enlarged stage-Markov property, the  
668 proof of Definition 1 transfers verbatim and each enlarged conditional is a Bernoulli probability on  
669 a measurable event. The qualifications stated in §4.2 (“When Definition 2 may fail”) describe two  
670 specific situations in which the enlarged stage-Markov property is violated.  $\square$

### 671 A.3 Argument for Empirical Claim 1 (slope-level statement)

672 Empirical Claim 1 is an empirical claim, not a theorem. The argument below is the qualitative  
673 justification that motivates the two slope signs; the data are what verify them. We deliberately do not  
674 call this a proof.

675 *Slope (a), under (F).* As  $K$  decreases,  $\nu(\mathcal{M}_{-i}^{(K)})$  is non-decreasing pointwise. Conditional on  
676 a successful retrieval, the probability of locking onto a reachable target is non-decreasing in  $\nu$ .  
677 Therefore  $\partial P(M^*|R^*; K)/\partial K^{-1} > 0$ .

678 *Slope (b), under (O).* As  $K$  decreases, peers receive updated memory simultaneously and con-  
679 verge on the same closest target; the marginal occupancy  $\Pr[\text{target} \in \mathcal{O}_{-i}]$  rises. Therefore  
680  $\partial P(C^*|M^*; K)/\partial K^{-1} < 0$ .

681 *Note on the conditional product  $\Phi$ .* Empirical Claim 1 makes no statement about the shape of  
682  $\Phi(K) = P(M^*|R^*; K) \cdot P(C^*|M^*; K)$ . Under (F) and (O) the slope of  $\Phi$  in  $K^{-1}$  is the sum of two  
683 opposing terms, and whether it changes sign depends on quantitative magnitudes that are not pinned  
684 down by (F) and (O) alone. In our data  $\Phi$  is monotone in  $K$  over  $\{4, 8, 16, 32, 48, 64\}$  and converges  
685 to the independent-agent boundary  $\Phi_\infty = 0.605$ . The practically relevant efficiency surface — mean  
686 time-to-success — does have an interior minimum at  $K = 8$ ; see §4.5.  $\square$

#### 687 A.4 The query–retrieval Galois connection: statement, argument, scope

688 **Fact (standard).** Let  $G$  and  $\Sigma$  be sets and  $I \subseteq G \times \Sigma$  a relation. For  $B \subseteq \Sigma$  define  $\text{ext}(B) =$   
689  $\{p \in G : \forall \tau \in B, (p, \tau) \in I\}$ , and for  $A \subseteq G$  define  $\text{int}(A) = \{\tau \in \Sigma : \forall p \in A, (p, \tau) \in I\}$ .  
690 Then for all  $A \subseteq G, B \subseteq \Sigma$ :

$$A \subseteq \text{ext}(B) \iff B \subseteq \text{int}(A);$$

691 both maps are antitone, both composites  $\text{ext} \circ \text{int}$  and  $\text{int} \circ \text{ext}$  are closure operators, and the closed  
692 pairs  $(A, B)$  with  $A = \text{ext}(B)$  and  $B = \text{int}(A)$ , ordered by extent inclusion, form a complete lattice  
693 — the concept lattice of the context  $(G, \Sigma, I)$  (Ganter and Wille, 1999).

694 **Argument.**  $A \subseteq \text{ext}(B)$  unfolds to  $\forall p \in A \forall \tau \in B : (p, \tau) \in I$ , a condition symmetric in  $A$   
695 and  $B$ , which is exactly  $B \subseteq \text{int}(A)$ . Antitonicity and the closure properties follow by the standard  
696 one-line arguments (Ganter and Wille, 1999, Ch. 1). Viewing the power sets as thin categories (one  
697 with reversed order), the displayed equivalence says precisely that  $\text{int}$  and  $\text{ext}$  form an adjoint pair; a  
698 Galois connection is an adjunction between posets (Mac Lane, 1998, Ch. IV).  $\square$

699 **Instantiation and scope.** In our memory, take  $G$  to be the place records,  $\Sigma$  the tag alphabet,  
700 and  $I_\theta = \{(p, \tau) : \text{the per-place tag table assigns } \tau \text{ to } p \text{ with confidence } \geq \theta\}$  (Appendix D); the  
701 intent of a query is  $B = q^{\text{req}}$  and R-availability is  $\text{ext}(B) \neq \emptyset$ . Three deliberate scope restrictions.  
702 (i) *Graded scores.* The deployed retrieval score is graded (the weighted log-score over required,  
703 preferred, and penalty tags of Appendix D), so the crisp connection above formalizes the thresholded  
704 required-tag skeleton of candidacy; the graded analogue is the theory of fuzzy Galois connections  
705 (Bělohlávek, 1999) and we do not use it here. (ii) *Non-monotone updates.* Dirichlet–multinomial  
706 updates are not literally monotone in  $I_\theta$ : contrary evidence can remove an incidence. On the  
707 reported benchmarks the binding failure is absence of evidence rather than contradiction (91% empty  
708 extents, 96.9% precision when non-empty, §3.5), so the monotone-growth reading is the empirically  
709 relevant regime, and we flag the exception rather than assume it away. (iii) *No structure for  $M$  or  $C$ .*  
710 Materialization is a satisficing choice function on non-empty extents and Completion is a control  
711 problem under occupancy; neither carries the order structure, and we attach no categorical claims to  
712 them. The reading earns its place by naming the failure mode exactly (a queried intent whose extent  
713 is empty lies outside the realized concept lattice), by giving consolidation a precise object (growth of  
714 the context, hence of every extent), and by locating the cadence axis structurally: memory coupling  
715  $\mathcal{M}_{-i}^{(K)}$  merges peer contexts at rate  $K^{-1}$ , while occupancy  $\mathcal{O}_{-i}$  is extra-contextual — which is the  
716 structural content behind the two slope signs of Empirical Claim 1.

## 717 B An options / hierarchical-MDP reading of the stages

718 This appendix develops the dynamic reading pointed to from §2.3. It introduces no new algorithm: it  
719 names the object the pipeline already computes and makes the  $\Phi$ -versus- $t_{\text{succ}}$  split a property of that  
720 object. Where the order-theoretic reading (Appendix A.4) fixes the *static* content of the stages, this  
721 one addresses time and cost.

722 **Q/R/M/C as an audit of an option.** Each query instantiates one temporally extended action—an  
723 option  $o = \langle \mathcal{I}, \pi, \beta \rangle$  in the sense of Sutton et al. (1999): an initiation set  $\mathcal{I}$ , an internal policy  $\pi$ , and  
724 a termination condition  $\beta$ . Standard trajectory-level analysis treats an option as a black box scored by  
725 its return. The Q/R/M/C decomposition is instead an audit of the option’s *internal* structure, and it  
726 lines up edge-for-edge with the factorization  $Q^* \rightarrow R^* \rightarrow M^* \rightarrow C^*$ :

727  $Q^* \rightarrow R^*$  (**initiation**). The option “reach a place satisfying  $q$ ” is well-initiated only when its initiation  
728 set is non-empty. In the vocabulary of §2.3 this is exactly  $\text{ext}(q^{\text{req}}) \neq \emptyset$ : an empty extent is not  
729 a bad choice of option but the absence of any admissible one, so  $R$  fails before a target is ever  
730 locked.

731  $R^* \rightarrow M^*$  (**policy over options**). Materialization is the high-level policy that commits to a single  
732 option out of the admissible set—the target lock  $g^* = \chi(\text{ext}(q^{\text{req}}))$  under the satisficing choice  
733 function  $\chi$  of §2.3.

734  $M^* \rightarrow C^*$  (**execution to termination**). Completion is the roll-out of the internal policy  $\pi$  until the  
735 termination condition  $\beta$  fires (arrival at  $g^*$ , or timeout). This is the only edge whose failure modes  
736 are extra-semantic: controller competence and, in a collective, occupancy  $\mathcal{O}_{-i}$ .

737 So Q/R/M/C is, on this reading, a stage-resolved instrument for the internal structure of a hierarchical-  
738 MDP option generator, rather than a bespoke logger for a single environment.

739 **Why the optimum lives on  $t_{\text{succ}}$  and not on  $\Phi$ : a semi-MDP reading.** In the options / semi-MDP  
740 formalism, an option is evaluated by a return accumulated over a *random* execution duration  $\tau$ ,

$$R(s, o) = \mathbb{E} \left[ \sum_{t=0}^{\tau-1} \gamma^t r_{t+1} \mid s_0 = s, o \right],$$

741 where  $\tau$  is the number of low-level steps until  $\beta$  fires (Sutton et al., 1999). The chain probability  
742  $\Phi = P(M^*|R^*) \cdot P(C^*|M^*)$  records only *whether* the option-audit succeeds; it is by construction  
743 blind to  $\tau$ , which is the cost of traversing it. The cost-bearing metric  $t_{\text{succ}}$  is  $\tau$ . This is the structural  
744 reason—and not merely an empirical coincidence—that the resource-rational reading of §2.3 (Lieder  
745 and Griffiths, 2020) locates an interior optimum on  $t_{\text{succ}}$  ( $K=8$ ) while  $\Phi$  stays monotone: the two  
746 metrics measure different faces of the same option, success versus duration.

747 **The bottleneck shift as multi-agent option collision.** The dynamic reading makes the  $M \leftrightarrow C$   
748 trade-off of Empirical Claim 1 physically concrete. Faster broadcast (smaller  $K$ ) merges peer contexts  
749  $\mathcal{M}_{-i}^{(K)}$  more often, which can only enlarge each agent’s extent and therefore raises the probability  
750 that a target is admissible and locked:  $P(M^*|R^*)$  increases. But because agents share a spatial basis  
751 and a tag alphabet (§2.1), fresher shared memory also *synchronizes* the high-level policy: several  
752 agents commit to the same option over the same scarce target. Under occupancy  $\mathcal{O}_{-i}$  this inflates  
753 the execution duration  $\tau$  and depresses  $P(C^*|M^*)$ . Cadence  $K$  thus regulates the dispersion of the  
754 agents’ option distribution: rare exchange desynchronizes option selection and shortens  $\tau$ , giving the  
755 minimum of  $t_{\text{succ}}$  at an interior  $K$ . The bottleneck is therefore grounded in a competition between  
756 initiation-set growth (semantic, contextual) and option-collision cost (extra-semantic, occupancy),  
757 exactly the two edges the cadence protocol  $\mathcal{C}_K$  acts on in Figure 2.

758 **Belief–desire–intention gloss and scope.** The high-level state audited above can be named in  
759 belief–desire–intention terms, which is the form used by the collective-memory companion: *belief* is  
760 the realized concept lattice of §2.3 (what is retrievable), *desire* is the state-conditioned predicate that  
761 selects the query  $q$  (e.g. a thirst variable switching the required tags to `water_source`), and *intention*  
762 is the locked option  $g^*$ . The hierarchical state is then  $h = (\text{lattice}, \text{desire}, g^*, s)$  with  $s$  the low-level  
763 physical state, and the three failure loci separate as before: an empty extent is a belief/ $R$  failure, a  
764 non-empty extent with no stable lock is an  $M$  failure, and a lock unreachable under occupancy is a  $C$   
765 failure. Two scope caveats. We claim no options-*learning* result: the pipeline is a fixed satisficing  
766 option generator (Simon, 1955) and Q/R/M/C audits it. And we attach the option structure only  
767 to the semantic path— $\mathcal{I}$  to  $Q \rightarrow R$ ,  $(\pi, \beta)$  to  $M \rightarrow C$ —while the semi-MDP return is used as an  
768 explanatory lens for the  $\Phi$ -versus- $t_{\text{succ}}$  split, not fitted to the data. Neither this options reading nor the  
769 order-theoretic reading of Appendix A.4 is claimed as a contribution: both are standard constructions  
770 used as interpretive lenses on the measurement protocol, which is the paper’s contribution. Their  
771 value is that each makes a testable statement on the logs—an empty extent localises the  $R$  failure,  
772 and the semi-MDP duration  $\tau$  predicts that the optimum appears on  $t_{\text{succ}}$  and not on  $\Phi$ .

## 773 C Operational vs. nested stage estimators

774 **Operational diagnostic rates.** The main benchmark logs per-query and per-episode diagnos-  
775 tic quantities directly from runtime traces (query formed, retrieval satisfied, target materialized,  
776 completion after retrieval). These are operational rates, not direct multiplicative factors.

777 **Nested stage estimators for factorization.** For the factorization, episode-level nested events  
778 from the first failure stage recorded in the runtime taxonomy are used:  $Q^*, R^*, M^*, C^*$ , nested by  
779 construction. The consistency audit in §3.2 reports  $\hat{P}(Q^*)$ ,  $\hat{P}(R^*|Q^*)$ ,  $\hat{P}(M^*|R^*)$ ,  $\hat{P}(C^*|M^*)$ ,  
780 their product, and its gap to semantic-path completion.

### 781 C.1 Multi-agent operational definitions

782 For the multi-agent sweep, every metric is defined per agent and per tick, then aggregated. Let  
783  $\text{Targets}_i(t)$  be the candidate set returned by the memory query of agent  $i$  at tick  $t$ ,  $\text{Locked}_i(t)$  be the  
784 agent’s planner-internal locked target, and  $\mathcal{W}$  the ground-truth target set.

785 **Per-tick events.**  $Q_i(t) = \mathbf{1}[\text{Targets}_i(t) \neq \emptyset]$ ;  $R_i(t) = \mathbf{1}[\exists w \in \mathcal{W} : \min_c \|c - w\| \leq \epsilon]$  ( $\epsilon = 0.6$ );  
 786  $M_i(t) = \mathbf{1}[\text{Locked}_i(t) \neq \emptyset \wedge \min_w \|\text{Locked}_i(t) - w\| \leq \epsilon]$ .

787 **Episode-level starred events.**  $Q_i^* = \bigvee_t Q_i(t)$ ; analogously for  $R^*, M^*$ ;  $C_i^* =$   
 788  $\mathbf{1}[\text{agent } i \text{ reached a target}]$ .

789 **Aggregation.** Per-run rates are averaged over the  $N$  agents; per-configuration rates are averaged  
 790 over the 40 seeds. Conditional rates  $P(R^*|Q^*)$ ,  $P(M^*|R^*)$ ,  $P(C^*|M^*)$  are computed as ratios of  
 791 episode-level counts within each configuration before averaging across configurations.

## 792 D Memory and task schema

### 793 D.1 Full notation table

794 Symbols are fixed throughout the paper.

795 **Stages.**  $Q, R, M, C$  — the four stages: **Q**uery formation, **R**etrieval, **M**aterialization, **C**ompletion  
 796 after retrieval. They name both the stage and (in expressions like “the M link”) the corresponding  
 797 edge of the chain.

798 **Events.**  $Q^*, R^*, M^*, C^* \in \{0, 1\}$  — episode-level starred events recording whether each stage  
 799 succeeded (Appendix C). Nested by construction:  $Q^*=1 \Leftarrow R^*=1 \Leftarrow M^*=1 \Leftarrow C^*=1$ .

800 **Probabilities.**  $P(\cdot)$  — population-level probability (averaged over the seed distribution and, in the  
 801 multi-agent case, over agents).  $\hat{P}(\cdot)$  — empirical-frequency estimator computed from logs.

802 **Single-agent.**  $q$  — query;  $\mathcal{M}$  — memory state;  $Y$  — end-to-end task success (can exceed  $C^*$ ).

803 **Multi-agent.**  $N$  — number of agents in the population.  $T$  — number of targets in the environment  
 804 (constraint  $T \leq N$ ; “scarcity” means  $N > T$ ).  $i$  — agent index.  $\mathcal{M}_i$  — agent  $i$ ’s private memory.  $q_i$   
 805 — agent  $i$ ’s query.

806 **Coupling.**  $\mathcal{C}_K$  — the peer-broadcast communication protocol with cadence  $K \in \mathbb{N}$  (every  $K$   
 807 ticks each agent broadcasts a memory snapshot, §4.1).  $\mathcal{M}_{-i}^{(K)}$  — the latest peer memory snapshots  
 808 accessible to agent  $i$  under  $\mathcal{C}_K$  (the *memory coupling* channel).  $\mathcal{O}_{-i}$  — the set of targets already  
 809 claimed by other agents by the time  $i$  acts (the *occupancy coupling* channel).

810 **Slope variables.**  $\nu(\mathcal{M}_{-i}^{(K)})$  — a freshness functional on peer memory; non-decreasing pointwise as  
 811  $K$  decreases (assumption (F) of Empirical Claim 1). The occupancy slope is assumption (O).

812 **Summary statistics.**  $\Phi(K) := P(M^*|R^*; K) \cdot P(C^*|M^*; K)$  — the conditional product, the chain  
 813 probability  $P(M^* \wedge C^* | R^*)$ . POM (product-of-means) and MOP (mean-of-products) are two  
 814 estimators of  $\Phi$ ; they differ by  $\text{Cov}(P(M^*|R^*), P(C^*|M^*))$ , which is the Jensen gap (§4.5).  $t_{\text{succ}}$   
 815 — mean time-to-success in environment ticks.

### 816 D.2 Memory stack and task patterns

817 The memory stack stores symbolic place records (per-place tag confidences, counts, visitation,  
 818 freshness, geometry), transition records (success, risk, cost statistics with fast/slow estimates),  
 819 episodic traces (intent + state at each visit), and concept/relation records. At each step, the tokenizer  
 820 emits  $(z_t, \tau_t, e_t)$  from  $o_t$ . Per-place tag tables use a Dirichlet-multinomial update; transitions use  
 821 exponentially smoothed statistics; episodic records attach the active intent and agent state. A planner  
 822 query is answered by retrieving candidate places whose semantic profiles match the current task intent  
 823 and state, via a weighted log-score  $s(p | q_t, s_t) = \sum_{\tau \in q^{\text{req}}} \log P(\tau|p) + \alpha \sum_{\tau \in q^{\text{pref}}} \log P(\tau|p) -$   
 824  $\beta \sum_{\tau \in q^{\text{pen}}} \log P(\tau|p) + \gamma u_{\text{epi}}(p, s_t)$  over required, preferred, and penalty tags. The four task  
 825 families are defined by semantic place patterns: water search (target {water}), safe rest ({rest} with  
 826 state-conditioning), goal-region rejoin ({goal-region}), hazard recovery ({post-hazard-exit}).

### 827 D.3 Single-agent supporting tables and figures

Table 5: Best semantic result per task family from the final 10-seed benchmark, with 95% bootstrap CIs. Hazard-recovery is bottlenecked by retrieval; goal-region remains downstream-bottlenecked.

| Task            | Best method | Assist | Success (95% CI)  | Steps | Query (op.) | Retrieval (op.) | Mat. (op.) | Post (op.) |
|-----------------|-------------|--------|-------------------|-------|-------------|-----------------|------------|------------|
| Water           | water-cp    | on     | 0.95 [0.88, 1.00] | 17.99 | 0.71        | 0.71            | 0.71       | 0.70       |
| Rest            | rest-s2     | on     | 0.93 [0.85, 0.99] | 23.35 | 0.61        | 0.61            | 0.61       | 0.73       |
| Goal region     | goal-n1     | off    | 0.54 [0.52, 0.56] | 49.34 | 0.00        | 0.00            | 0.00       | 0.00       |
| Hazard recovery | hazard-s7   | on     | 0.40 [0.35, 0.45] | 10.18 | 0.11        | 0.11            | 0.11       | 0.16       |




Figure 5: Oracle stage interventions on the two hardest single-agent task families (10 seeds  $\times$  8 episodes). **(a)** FourRooms goal-rejoin: replacing any single stage with an oracle leaves end-to-end success at 0.07; the bottleneck is downstream of the stages in the chain. **(b)** LavaGapS7 hazard recovery: oracle controller adds only +0.04, but oracle materialization adds +0.16 and oracle retrieval adds a further +0.01 to reach 0.60. The framework localises the bottleneck stage-by-stage.

## 828 E Portability and scale-up

### 829 E.1 BabyAI portability transfer (10 seeds, full benchmark)

830 A 10-seed portability benchmark on BabyAI-GoToObjMaze-v0 (Chevalier-Boisvert et al., 2018)  
831 (8 episodes per seed, max-steps 200) measured all four single-agent methods used in the main  
832 benchmark (random, greedy, babyai\_semantic\_node\_v1, babyai\_semantic\_plan\_v1) with  
833 the same Q/R/M/C logger applied to BabyAI traces. The numbers in Table 6 confirm that (a) the  
834 Q/R/M/C events emit non-trivially on the BabyAI substrate, (b) the two semantic variants expose  
835 stage-separable structure where the non-semantic baselines have  $Q^*=R^*=M^*=0$  (no semantic  
836 queries are issued), and (c) the stage-rate ordering matches the MiniGrid benchmark (Table 1) within  
837 bootstrap CIs. GoToObjMaze is a harder task substrate than the LavaGapS7/FourRooms benchmark  
838 used in the main text — end-to-end success caps at 0.15 across all methods — which makes the stage  
839 decomposition especially informative: the semantic stack is not winning on success but it is the only  
840 method family producing measurable per-stage events.

Table 6: BabyAI portability benchmark: 10 seeds  $\times$  8 episodes on BabyAI-GoToObjMaze-v0, max-steps 200. The semantic-stack variants succeed at the same end-to-end rate as the greedy baseline (0.15) but are the only methods that produce non-zero Q/R/M/C events — the framework’s protocol transfers to the BabyAI substrate without modification. Q/R/M/C operational rates here are evaluated on the same denominator (number of decision points), so  $Q=R=M$  for both semantic variants because GoToObjMaze emits one semantic target type per episode.

| Method        | Success          | Steps | $Q$ (op.) | $R$ (op.) | $M$ (op.) | Post (op.) |
|---------------|------------------|-------|-----------|-----------|-----------|------------|
| random        | $0.125 \pm 0.06$ | 182.7 | 0.000     | 0.000     | 0.000     | 0.000      |
| greedy        | $0.150 \pm 0.05$ | 171.3 | 0.000     | 0.000     | 0.000     | 0.000      |
| semantic node | $0.150 \pm 0.05$ | 171.3 | 0.032     | 0.032     | 0.032     | 0.025      |
| semantic plan | $0.150 \pm 0.05$ | 171.3 | 0.032     | 0.032     | 0.032     | 0.025      |

**Takeaway.** The framework’s measurement protocol — not its task-specific results — transfers to a second substrate. End-to-end success on BabyAI is dominated by the much lower episode-level success ceiling of GoToObjMaze, not by the stage structure; conversely, the stage structure that the semantic stack exposes is identical in shape to what the same stack exposes on MiniGrid. This is exactly the substrate-independence property the methodology is intended to provide. Raw data: tmp/babyai\_deeper\_10seed/.

## E.2 Multi-agent scale-up at $N = 12$

To check that the multi-agent results of §4.5 are not an artefact of the  $N \leq 8$  range, an additional 8,640-run sweep was run at  $N=12$ ,  $T \in \{6, 8, 10\}$ , the same three layouts, three hazard densities, eight peer cadences, and 40 seeds. Wall-clock  $\approx 12$  minutes on 8 cores. The conditional-rate signature of §4.5 persists and intensifies (Table 7).




Figure 6:  $N = 12$  scale-up bottleneck signature ( $n = 8,640$  runs, 40 seeds). Top: the per-cadence mean rates move in the predicted directions in  $K$ ; the within-cadence Pearson correlation (red dashed, right axis) peaks at  $-0.70$  at  $K = 8$  (this within-cadence peak is what sharpens at  $N=12$ ; the M-side  $K$ -slope itself is not cluster-robust at this agent count, Table 7). Bottom: mean time-to-success retains an interior minimum at  $K = 8$  (4.77 ticks); as on the main sweep, the location is robust but the margin is shallow — the gap to the adjacent  $K = 4$  (4.82 ticks) is  $\approx 0.05$  tick and is not resolved at this sample size. The bottleneck-shift peak migrates outward from  $K = 4$  at  $N \leq 8$  to  $K = 8$  at  $N = 12$ , consistent with stronger occupancy coupling at higher agent count delaying the cost-regime transition.

## F Extended diagnostics

### F.1 Conditional-product detail and the Jensen gap

The conditional product  $\Phi = P(M^*|R^*) \cdot P(C^*|M^*)$  in the 35,640-run sweep is monotonically increasing in  $K$  over  $K \in \{4, 8, 16, 32, 48, 64\}$  and converges to the independent-agent boundary

Table 7:  $N=12$  scale-up sweep numerical detail (peer architecture; 27 configuration cells). Effect sizes with cluster-robust 95% CIs (block bootstrap over cells): Spearman of  $P(C^*|M^*)$  with  $K$  is  $+0.31$  (CI  $[+0.21, +0.41]$ , excludes zero), while Spearman of  $P(M^*|R^*)$  with  $K$  is only  $-0.14$  (CI  $[-0.28, +0.02]$ , *not* resolved from zero at  $N=12$ ). The C-slope of Empirical Claim 1 holds; the M-slope is directionally consistent but not cluster-robust at this agent count. What sharpens at  $N=12$  is the within-cadence correlation and the  $t_{\text{succ}}$  minimum (Figure 6), not the M-side  $K$ -slope.

| $K$ | $P(M^* R^*)$ | $P(C^* M^*)$ | MOP   | mean $t_{\text{succ}}$ | corr         |
|-----|--------------|--------------|-------|------------------------|--------------|
| 1   | 0.758        | 0.679        | 0.475 | 5.49                   | -0.59        |
| 2   | 0.745        | 0.692        | 0.474 | 5.29                   | -0.61        |
| 4   | 0.679        | 0.768        | 0.482 | 4.82                   | -0.64        |
| 8   | 0.731        | 0.767        | 0.539 | <b>4.77</b>            | <b>-0.70</b> |
| 16  | 0.687        | 0.803        | 0.538 | 5.04                   | -0.62        |
| 32  | 0.676        | 0.814        | 0.538 | 5.75                   | -0.51        |
| 48  | 0.677        | 0.814        | 0.539 | 5.96                   | -0.51        |
| 64  | 0.677        | 0.816        | 0.541 | 6.12                   | -0.51        |

Table 8: Per-architecture conditional Q/R/M/C rates from the 35,640-run sweep (40 seeds).  $P(R^*|Q^*)$  is saturated. POM = product-of-means; MOP = mean-of-products (the proper per-configuration expected product). The two summaries disagree systematically at small/intermediate  $K$  because the within-cadence correlation between  $P(M^*|R^*)$  and  $P(C^*|M^*)$  is strongly negative there (last column) — the bottleneck-shift signature. **Bold:** MOP-best of the peer cadences. Mean  $t_{\text{succ}}$  (last column) is the practically relevant efficiency measure; its interior minimum at  $K=8$  is the predicted-shape result, robust in location but shallow in margin (adjacent fast cadences not resolved; Appendix F.3).

| Architecture       | $P(R^* Q^*)$ | $P(M^* R^*)$ | $P(C^* M^*)$ | POM   | MOP (95% CI)                | corr  | mean $t_{\text{succ}}$ |
|--------------------|--------------|--------------|--------------|-------|-----------------------------|-------|------------------------|
| independent        | 0.99         | 0.68         | 0.88         | 0.598 | 0.605 [0.597, 0.614]        | +0.00 | 8.02                   |
| shared bus         | 1.00         | 1.00         | 0.60         | 0.600 | 0.595 [0.586, 0.604]        | -0.10 | 8.33                   |
| central aggregator | 1.00         | 0.97         | 0.61         | 0.592 | 0.572 [0.564, 0.580]        | -0.48 | 7.64                   |
| peer ( $K=1$ )     | 1.00         | 0.97         | 0.60         | 0.585 | 0.573 [0.564, 0.581]        | -0.43 | 7.80                   |
| peer ( $K=2$ )     | 1.00         | 0.97         | 0.62         | 0.594 | 0.576 [0.568, 0.584]        | -0.54 | 7.59                   |
| peer ( $K=4$ )     | 1.00         | 0.92         | 0.66         | 0.599 | 0.569 [0.561, 0.577]        | -0.61 | 8.00                   |
| peer ( $K=8$ )     | 0.99         | 0.69         | 0.87         | 0.597 | 0.586 [0.578, 0.595]        | -0.19 | <b>7.47</b>            |
| peer ( $K=16$ )    | 0.99         | 0.67         | 0.89         | 0.592 | 0.590 [0.581, 0.598]        | -0.05 | 7.87                   |
| peer ( $K=32$ )    | 0.99         | 0.67         | 0.89         | 0.595 | 0.598 [0.589, 0.606]        | -0.04 | 8.28                   |
| peer ( $K=48$ )    | 0.99         | 0.68         | 0.88         | 0.598 | 0.602 [0.593, 0.611]        | -0.02 | 8.79                   |
| peer ( $K=64$ )    | 0.99         | 0.68         | 0.88         | 0.599 | <b>0.603</b> [0.594, 0.611] | -0.02 | 9.19                   |

856  $\Phi(K \rightarrow \infty) = 0.605$ . Paired bootstrap:  $\Delta(K=32-K=16) = +0.008$  (CI  $[+0.006, +0.010]$ ),  
857  $\Delta(K=48-K=16) = +0.011$ ,  $\Delta(K=64-K=16) = +0.012$ ; the independent-agent level sits  
858 above every peer  $K$  except  $K=64$  (where the gap’s CI includes zero). This is a statement about the  
859 *level* of  $\Phi$  at the tested cadences, not about the *shape* of  $\Phi(K)$ : consistent with the monotone trend,  
860 we make no claim of an interior extremum (§4.3).

861 The negative within-cadence covariance documented in §4.5 creates the Jensen-inequality gap  
862 between  $\Phi$  and the marginal product-of-means:  $E[X] \cdot E[Y] > E[X \cdot Y]$  when  $\text{Cov}(X, Y) < 0$ . The  
863 mean-of-products is reported as the primary summary because it absorbs the covariance correctly.  
864 Figure 7 shows the K-sweep view of both conditional rates and mean  $t_{\text{succ}}$ ; Figure 8 shows the  
865 per-configuration scatter that the negative covariance summarises.

## 866 F.2 Design decomposition and run-count reconciliation

867 The 35,640 figure decomposes exactly as in Table 9. The base design is the nine admissible  $(N, T)$   
868 scarcity pairs (for  $N=3$ :  $T \in \{2, 3\}$ ;  $N=5$ :  $T \in \{2, 3, 5\}$ ;  $N=8$ :  $T \in \{2, 3, 5, 8\}$ ) crossed with three  
869 layouts, three hazard densities, and 40 seeds, giving  $3,240$  base configurations. The peer architecture  
870 is run at all eight cadences ( $3,240 \times 8 = 25,920$ ); the three cadence-free architectures (independent,  
871 shared bus, central aggregator) contribute  $3,240$  each. The  $25,920$ -run figure used for the peer  
872 cluster-robust analysis (Appendix F.3) is exactly the peer block; the initial  $12,960$ -run exploratory  
873 sweep (§4.3, “Chronology”) and the extended-cadence and  $N=12$  sweeps (Appendix E) are separate  
874 run sets and are not pooled with the main sweep.




Figure 7: Bottleneck shift across eight peer cadences (40-seed sweep). Top:  $P(M^*|R^*)$  falls and  $P(C^*|M^*)$  rises in  $K$  (moderate, cluster-robust Spearman effect sizes  $-0.53$  and  $+0.43$ ; CIs in Appendix F.3). Bottom: mean time-to-success has an interior minimum at  $K = 8$  (7.47 ticks) below the independent-agent baseline (8.02, dashed); the location is robust but its margin over the adjacent fast cadences ( $K=1, 4$ ) is not resolved (Appendix F.3). Error bars and shaded band are 95% bootstrap CIs.

Table 9: Run-count reconciliation for the main multi-agent sweep (35,640 runs). Counts are recomputed directly from the released run log.

| Factor                         | Levels                                                   | Count         |
|--------------------------------|----------------------------------------------------------|---------------|
| $(N, T)$ scarcity pairs        | $(3, 2)(3, 3)(5, 2)(5, 3)(5, 5)(8, 2)(8, 3)(8, 5)(8, 8)$ | 9             |
| Layouts                        | symmetric, asymmetric, random                            | 3             |
| Hazard densities               | 0.0, 0.05, 0.10                                          | 3             |
| Seeds                          |                                                          | 40            |
| Base configurations            | $9 \times 3 \times 3 \times 40$                          | 3,240         |
| Peer (all 8 cadences $K$ )     | $3,240 \times 8$                                         | 25,920        |
| Independent / shared / central | $3 \times 3,240$                                         | 9,720         |
| <b>Total</b>                   |                                                          | <b>35,640</b> |

### 875 F.3 Effect sizes and cluster-robust inference

876 The  $p < 10^{-4}$  figures quoted in §4.5 are reported for completeness but are not the basis of the  
877 claim: at tens of thousands of runs almost any non-zero slope is “significant,” and the runs are not  
878 independent—they cluster within the  $(N, T, \text{layout}, \text{hazard})$  configuration. The full sweep has 25,920  
879 peer runs but only **81 independent design cells**. We therefore report effect sizes with cluster-robust  
880 confidence intervals from a block bootstrap that resamples *configuration cells* (not rows),  $B=4000$   
881 resamples.



Bottleneck-shift signature: negative within-cadence correlation peaks at  $K = 4$  and decays at  $K = 64$




Figure 8: Per-configuration scatter of  $P(M^*|R^*)$  vs.  $P(C^*|M^*)$  at three cadences. The strong negative correlation at  $K = 4$  (Pearson  $-0.61$ ) is the bottleneck-shift signature: the same configurations that succeed at  $M$  tend to fail at  $C$  and vice versa. The correlation decays to noise by  $K = 64$ .

882 The two slopes of Empirical Claim 1 survive as genuine, but moderate, effects: the Spearman  
883 correlation of  $P(M^*|R^*)$  with  $K$  is  $\rho = -0.53$  (95% cluster CI  $[-0.59, -0.47]$ , strong) and of  
884  $P(C^*|M^*)$  with  $K$  is  $\rho = +0.43$  (95% cluster CI  $[+0.37, +0.50]$ , moderate). Both intervals exclude  
885 zero under cell-level resampling, so the bottleneck shift is not an artefact of the correlated run count;  
886 it is a moderate monotone trend, and we describe it as such rather than as a near-deterministic law.

887 The interior optimum on mean time-to-success is robust in *location* but shallow in *margin*. Paired  
888 block bootstrap on the cell-level  $t_{\text{succ}}$  differences gives  $t_{\text{succ}}(K=8)$  below  $t_{\text{succ}}(K=16)$  by 0.40 ticks  
889 (CI  $[0.08, 0.83]$ ) and below  $K=64$  by 1.67 ticks (CI  $[0.75, 2.83]$ ), both resolved; but its advantage  
890 over  $K=4$  (0.49, CI  $[-0.47, 1.51]$ ) and  $K=1$  (0.27, CI  $[-0.43, 0.97]$ ) is within noise.<sup>1</sup> A success-  
891 weighted efficiency metric ( $t_{\text{succ}}$  divided by success rate) has the same interior arg min at  $K=8$ , so  
892 the location of the optimum is robust to the efficiency definition even though its margin over the  
893 adjacent fast cadences is not resolved at this sample size. We read this as support for the *shape*  
894 predicted by the  $M \leftrightarrow C$  trade-off, not for a sharply tuned optimal cadence.

#### 895 F.4 Construct validity: task success $Y$ vs. semantic-path completion $C^*$

896 A success  $Y$  can in principle bypass the semantic path, so a stage bottleneck in  $Q/R/M/C$  need not  
897 explain variation in  $Y$ . Two facts bound this concern. First, in the multi-agent sweep that Empirical  
898 Claim 1 is about, success is *defined* as reaching a `water_source` target, which is exactly the  $C^*$   
899 event; empirically  $Y$  and  $C^*$  coincide run-for-run (correlation 1.00, mean gap 0.00). The central  
900 multi-agent claim is thus not exposed to a parallel-channel objection. Second, the gap  $Y > C^*$  is  
901 a single-agent phenomenon on the easy task families (water, rest), where a fallback controller can  
902 reach the goal without a semantic target lock; it is largest exactly where success is already high and  
903 the diagnosis is least needed, and it vanishes on the hard families (hazard recovery, goal-region)  
904 where the framework does its work. A per-episode, stage-resolved mediation of  $Y$  (rather than of  
905  $C^*$ ) requires logs that tag each success with whether it traversed the semantic path, and is a planned  
906 instrumentation refinement.

#### 907 F.5 Threshold ( $\theta$ ) robustness of the empty-candidate-set diagnosis

908 The single-agent attribution “candidate generation fails, not ranking” (§3.5) rests on a  
909 high empty-candidate-set rate. A reviewer asks whether that is an artefact of the  
910 required-tag threshold  $\theta$  (deployed 0.5; the candidacy gate is `profile[tag]  $\geq \theta$`  in

<sup>1</sup>This paired bootstrap and the effect-size Spearman correlations above are computed on the same sweep execution as Table 8 (25,920 peer runs); the  $t_{\text{succ}}$  statistic is defined only over successful episodes, so the paired cell-level difference (1.67 for  $K=8$  vs.  $K=64$ ) is on that subpopulation and matches the marginal-mean difference in Table 8 ( $9.19 - 7.47 = 1.72$ ) to within rounding. The effect sizes are identical on the independent `paper1_clean` execution ( $-0.531/+0.434$ ), confirming they are not run-specific.

symbolic\_memory.\_required\_matches\_profile). We sweep  $\theta$  via an override on the query builder (experiments/big\_experiment/theta\_ablation.py); the endpoints confirm the gate is live, and the operating range shows the diagnosis is not a threshold effect.

Table 10: Empty-candidate-set rate vs. required-tag threshold  $\theta$  (semantic method, assist off, hazard-recovery / goal-region). The rate is *invariant* across the entire realistic range  $\theta \in [0.05, 0.95]$ ; only the degenerate endpoints move it, and only to the no-candidate floor. The empties are therefore absence of stored evidence, not sub-threshold filtering.

| $\theta$           | empty rate (hazard recovery) | empty rate (goal region) |
|--------------------|------------------------------|--------------------------|
| 0.00 (accept all)  | 0.50                         | —                        |
| 0.05               | 0.533                        | 0.267                    |
| 0.15               | 0.533                        | 0.267                    |
| 0.25               | 0.533                        | 0.267                    |
| 0.50 (deployed)    | 0.533                        | 0.267                    |
| 0.75               | 0.533                        | 0.267                    |
| 0.95               | 0.533                        | 0.267                    |
| 1.00 (accept none) | 0.75                         | —                        |

The reading is decisive (Table 10). At  $\theta=1.0$  nothing passes and the empty rate rises; at  $\theta=0.0$  the required-tag gate accepts every node ( $\text{conf} \geq 0$  always holds), yet the empty rate falls only to 0.50, its floor—meaning that at those decision points there is *no candidate node of any confidence* to retrieve. Across the entire sane range  $\theta \in [0.05, 0.95]$  the rate is flat at 0.533: loosening or tightening the threshold does not recover a candidate, because the required semantic place was never stored. This is exactly the “absence of evidence, not contradiction” regime named in §2.3, now shown to be threshold-invariant, so the memory-accumulation attribution is not an artefact of  $\theta$ . (Absolute rates here are on a small hazard-recovery / goal-region sweep and are not the headline 91% of §3.5, which is on the full benchmark; the point is the  $\theta$ -invariance, not the level.)

## F.6 Robustness to Assumption 1 violations

Definition 1 rests on the conditional stage-Markov property (Assumption 1). The two practically relevant violations—adaptive query rewrites not surfaced to the logger, and hidden between-stage memory updates—are exactly cases where an *unlogged latent* couples non-adjacent stages. We quantify the resulting estimator bias with a controlled simulation ( $4 \times 10^5$  episodes per setting; experiments/big\_experiment/assumption1\_stress\_test.py). Episodes are generated with known mechanism rates  $P(M^*|R^*)=0.70$ ,  $P(C^*|M^*)=0.65$ , and a hidden binary latent  $U$  that shifts *both* the M-rate and the C-rate by  $\pm\delta$ ; the naive estimator conditions only on the previous starred stage, not on  $U$ .

Table 11: Bias of the naive conditional estimators under an unlogged latent of strength  $\delta$  that couples the M and C stages (violating Assumption 1).  $P(M^*|R^*)$  stays unbiased;  $P(C^*|M^*)$  is biased upward, materially only under strong coupling.

| $\delta$ (coupling strength) | bias in $\hat{P}(M^* R^*)$ | bias in $\hat{P}(C^* M^*)$ |
|------------------------------|----------------------------|----------------------------|
| 0.00                         | −0.001                     | −0.002                     |
| 0.05                         | +0.001                     | +0.004                     |
| 0.10                         | +0.001                     | +0.015                     |
| 0.15                         | −0.000                     | +0.031                     |
| 0.20                         | +0.000                     | +0.056                     |
| 0.25                         | −0.001                     | +0.089                     |

Two conclusions (Table 11). The  $R \rightarrow M$  conditional is unbiased at every  $\delta$ : the latent re-weights episodes but does not confound that edge. The  $C \rightarrow M$  conditional acquires a one-signed *upward* bias—the latent makes high-M episodes also high-C, so conditioning on M over-counts completion—but the bias is only +1.5 percentage points at a mild violation ( $\delta \leq 0.10$ ) and reaches +5.6 points only under strong unlogged coupling ( $\delta=0.20$ ). The framework’s exposure to Assumption 1 violations is therefore bounded and its failure direction is known, which is the relevant robustness statement for applying the logger to learned or LLM agents where such couplings are more common.

## 939 F.7 Preliminary continuous-substrate check (VMAS)

940 The most-requested robustness check is whether the  $M \leftrightarrow C$  bottleneck shift is specific to grid sub-  
 941 strates. We report a *preliminary* positive on a continuous-state substrate (VMAS (Bettini et al., 2022)).  
 942 The symbolic peer memory and the operational Q/R/M/C definitions are reused unchanged; a thin  
 943 discretisation bridge maps continuous  $(x, y)$  to coarse cells so the memory applies, and the controller  
 944 acts in continuous space toward the materialised waypoint (experiments/vmas\_portability/).  
 945 Materialisation uses the same occupancy-sensitive lock as the grid planner (a candidate whose target  
 946 is already claimed does not lock).

Table 12: Preliminary VMAS portability: both slope signs of Empirical Claim 1 reproduce on a continuous substrate in the coupling regime ( $N=8$ ,  $T=3$ , 40 seeds, horizon 25).  $\text{Spearman}(P(M^*|R^*), K) = -1.00$  and  $\text{Spearman}(P(C^*|M^*), K) = +1.00$ . Beyond  $K \approx 32$  the system enters the independent regime (rare broadcast, no contention) and  $P(M^*|R^*)$  rebounds, the same regime boundary seen on the grid.

| $K$ | $P(M^* R^*)$ | $P(C^* M^*)$ |
|-----|--------------|--------------|
| 1   | 1.000        | 0.182        |
| 2   | 0.583        | 0.305        |
| 4   | 0.448        | 0.405        |
| 8   | 0.305        | 0.608        |
| 16  | 0.278        | 0.639        |

947 We label this preliminary and do not fold it into the headline claims: it is a single scenario con-  
 948 figuration, the horizon is short, there are no cluster-robust CIs yet, the memory is fed through the  
 949 discretisation bridge (continuous *dynamics* with discretised symbolic *memory*, not fully continuous  
 950 memory), and the reported  $\pm 1.00$  Spearman is a *perfect* rank correlation over only five monotone  
 951 points ( $K \leq 16$ ) — weak evidence on its own, indicating direction rather than a resolved effect size.  
 952 What it establishes is that the two slope signs are not an artefact of grid cells: the same freshness-  
 953 vs-contention mechanism operates when positions and motion are continuous. Promoting this to a  
 954 headline result requires scaling to several scarcity cells and layouts with the cluster-robust bootstrap  
 955 of Appendix F.3.

## 956 F.8 Open problems suggested by the diagnostics

957 Four concrete questions follow directly from the measurements of §3.5 and §4.5. **(1) Adaptive**  
 958 **cadence.** Empirical Claim 1 assumes a fixed  $K$ ; the bottleneck-shift signature suggests an online  
 959 controller  $K_t$  that broadcasts when the (estimated) M-side gain exceeds the C-side occupancy cost.  
 960 **(2) Selective consolidation.** Snapshot exchange merges everything; merge rules that prioritize  
 961 evidence about under-represented semantic tags attack the candidate-generation starvation of §3.5  
 962 directly. **(3) Trust-weighted merging under occupancy.** The central aggregator already performs  
 963 a trust-weighted merge but pays the C-collapse cost; decentralized trust rules that account for peer  
 964 *intent* (not only peer evidence) could decouple freshness from target racing. *Since submission of*  
 965 *the first draft, this problem has been addressed constructively:* a decentralized protocol in which  
 966 an agent’s target lock piggybacks a lightweight *reservation* on its next scheduled broadcast (no  
 967 coordinator, no extra channel) and locks dominated by an earlier reservation are retracted by an  
 968 anti- $M^*$  rollback with exponential backoff. On the scarcity cells this recovers the fast-cadence  
 969 completion loss on *unconditional* metrics — success rate rises at every  $K \in \{1, 2, 4\}$  with disjoint  
 970 bootstrap CIs, and fast sharing with the protocol exceeds the  $K=8$  optimum baseline (0.614–0.620 vs.  
 971 0.583) — with rollback latency bounded by  $\approx K$  ticks. Details in a companion manuscript; the open  
 972 problem stands revised toward its harder variants (adaptive reservation scope, heterogeneous trust).  
 973 **(4) Substrate generality.** The decomposition transfers across MiniGrid, BabyAI, the custom grid,  
 974 and a standard-MiniGrid multi-agent wrapper with identical operational definitions and verified slope  
 975 signs of Empirical Claim 1 (§ 6), and a preliminary check reproduces both slopes on the continuous-  
 976 state VMAS substrate (Appendix F.7). A full continuous-state validation with cluster-robust CIs, and  
 977 richer 3D substrates (Habitat (Savva et al., 2019)), remain the next portability steps.

## F.9 Extended related work

Classical cognitive-map perspectives motivate explicit place abstractions (Tolman, 1948; O’Keefe and Nadel, 1978; Kuipers, 2000); modern embodied AI combines mapping and planning (Gupta et al., 2017) or exposes topological memory (SPTM, Savinov et al., 2018); goal-conditioned RL conditions on explicit goals or embeddings (Andrychowicz et al., 2017; Ghosh et al., 2021), with successor-representation work providing predictive structure (Dayan, 1993; Stachenfeld et al., 2017; Momennejad et al., 2017). Our setting differs in that the target is a semantic pattern retrieved from memory rather than a coordinate. In multi-agent RL, CTDE dominates (Lowe et al., 2017; Foerster et al., 2018) and learned communication adds a per-step exchange variable (Sukhbaatar et al., 2016; Foerster et al., 2016); gossip and consensus protocols (Boyd et al., 2006; Xiao et al., 2007) are the closest network-level analogue of our cadence axis, lifted here to the cognitive-architecture level. LLM-agent systems combine reasoning/acting, reflection, explicit memory, and multi-agent conversation (Yao et al., 2023; Shinn et al., 2023; Wang et al., 2023; Park et al., 2023; Wu et al., 2023; Packer et al., 2023; Chhikara et al., 2025). Our framework is orthogonal: rather than learning a communication policy end-to-end, it *diagnoses* where in the Q/R/M/C chain a given architecture succeeds or fails, and yields a slope-level prediction about cadence that is testable on logs.

## F.10 SOTA baseline and coverage matrix

Table 14 separates direct experimental baselines from broader SOTA references. The local experiments are intentionally modest and fully instrumented: random / graph-only / SPTM-like / learned BC in MiniGrid, CommNet PPO in the multi-agent grid, and a TextNav LLM step-budget study over two local Ollama models (in which Q/R/M/C localizes a completion-limited failure, §4.5). ReAct, Reflexion, Voyager, Generative Agents, AutoGen, MemGPT, and Mem0 are not reimplemented as full systems here; they define the current memory-agent design space against which Q/R/M/C is positioned as a diagnostic protocol.

**Positioning against SOTA memory designs.** The memory-agent design space can be organised on two axes: single- vs. multi-agent, and how memory is represented and shared (Table 13). Vector-embedding memory (MemGPT, Voyager single-agent; Generative Agents multi-agent) retrieves by opaque nearest-neighbour lookup, so a failure is not localisable to a stage. Shared-context designs (MetaGPT, AutoGen) pool everything into one window, dissolving per-agent state. Learned-communication MARL (CommNet, and stronger CTDE methods QMIX (Rashid et al., 2018), MAPPO (Yu et al., 2022), MADDPG (Lowe et al., 2017)) optimises a message channel end-to-end and, as we verify for CommNet PPO (§4.5), emits no separable query/lock/completion events. The cell our system occupies—per-agent *symbolic* memory with an explicit query/retrieve/materialise interface, exchanged at a tunable cadence—is empty in this matrix, and it is exactly the interface that makes the four Q/R/M/C stages observable. Our claim against SOTA is therefore not higher reward but higher *diagnosability at matched task success*: the same interface that a strong learned method trades away for end-to-end optimisation is what lets the framework localise where memory helps.

Table 13: Where the framework sits in the memory-agent design space. Q/R/M/C stage-observability requires an explicit query/lock interface, which the symbolic-per-agent cell provides.

| Memory representation                   | Single-agent                                             | Multi-agent                                                                                                     |
|-----------------------------------------|----------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------|
| Vector-embedding memory                 | MemGPT, Voyager (Packer et al., 2023; Wang et al., 2023) | Generative Agents (Park et al., 2023)                                                                           |
| Shared context window                   | —                                                        | AutoGen (Wu et al., 2023)                                                                                       |
| Learned-comm. MARL                      | —                                                        | CommNet, QMIX, MAPPO, MADDPG (Sukhbaatar et al., 2016; Rashid et al., 2018; Yu et al., 2022; Lowe et al., 2017) |
| Symbolic per-agent + explicit interface | Q/R/M/C single-agent (§3)                                | <b>this work</b> (peer memory, cadence $K$ )                                                                    |

## 1015 **F.11 Multi-agent oracle-intervention numerics**

## 1016 **F.12 Oracle C: definition and reading**

1017 Oracle C is the completion-stage analogue of oracle R and oracle M, added to close the R/M/C triple  
1018 symmetrically. The real Q/R/M pipeline runs unchanged (real memory query, real materialisation),  
1019 and once an agent has locked a target that is a genuine `water_source`, the oracle guarantees  
1020 arrival by placing the agent on that cell; scarcity is preserved because a cell already claimed or  
1021 occupied is not available, so oracle C cannot make more than  $T$  agents succeed. The runner is  
1022 `experiments/big_experiment/exp_oracle_interventions.py` (mode `oracle_C`).

1023 Two readings follow from Table 15. (i) Oracle C produces the *lowest* mean  $t_{\text{succ}}$  of the three  
1024 interventions ( $\sim 6.1$ – $6.4$  vs.  $\sim 7.2$  for R and M), because it removes the entire post-lock travel  
1025 time; this confirms that a real part of the excess time is controller travel to a correctly materialised  
1026 target. (ii) Oracle C yields the *smallest* success-rate gain ( $+0.01$ – $0.02$  over baseline, below both  
1027 oracle R at 0.614 and oracle M at 0.602), because in scarcity the constraint on *whether* an agent  
1028 succeeds is upstream—retrieving and locking a still-available target—not reaching it. The two facts  
1029 together are exactly the framework’s thesis in intervention form: completion is where time is spent,  
1030 retrieval/materialisation is where success is won.

## 1031 **F.13 Stage-agnostic cadence hypotheses and what the framework changes**

1032 A first version of the multi-agent cadence sweep (12,960 runs at 20 seeds; subsumed by the 35,640-run  
1033 sweep in this paper) was originally planned with four operational hypotheses formulated *before*  
1034 Empirical Claim 1. They are retained here because they motivated the eventual framework. The  
1035 framework was formulated after the sweep was complete; it is not claimed to have predicted the  
1036 sweep numbers a priori (see “Chronology” in §4.3).

### 1037 **Original hypotheses (pre-framework).**

1038 **H1.** An interior  $K^* \in \{2, 4, 8\}$  minimising mean time-to-success exists for most  $(N, T, \text{layout})$   
1039 cells.

1040 **H2.**  $K^*$  grows monotonically with scarcity  $\rho = N/T$ .

1041 **H3.** When  $\rho \leq 1$ ,  $K^* = 1$ .

1042 **H4.** Peer-broadcast architectures Pareto-dominate centralized aggregation on (mean time-to-success,  
1043 success rate).

1044 **Stand-alone result on the original 20-seed sweep.** The practical claim was directionally supported:  
1045 peer at the empirically best  $K$  was faster than centralized in the scarcity regime (a modest effect;  
1046 Mann–Whitney  $p=0.044$  on this small standalone sweep, reported as exploratory rather than as  
1047 primary evidence, consistent with the effect-size-first line of Appendix F.3). Among the four  
1048 structural hypotheses, only H4 was clearly supported in the random-layout subset. H1 held in 7 of  
1049 18 cells; H2 had a non-significant negative correlation; H3 was supported in 4 of 9 admissible cells.  
1050 The 35,640-run extended sweep used elsewhere in the paper produces matching stand-alone numbers  
1051 within bootstrap uncertainty.

1052 **What the framework changes.** The hypotheses H1–H4 are stage-agnostic: each ties  $K^*$  to scarcity  
1053  $\rho$  without specifying mechanism. Empirical Claim 1, formulated after the sweep, ties the slopes of  
1054  $P(M^*|R^*)$  and  $P(C^*|M^*)$  to cadence; the slope conditions are testable on logs. The data verify  
1055 both signs as moderate, cluster-robust effect sizes (Appendix F.3). We make no formal claim about  
1056 the shape of the conditional product  $\Phi$ , on which the data are monotone. The interior optimum holds  
1057 on the efficiency metric mean time-to-success.

## 1058 **G Assist on/off analysis**

1059 Across all four single-agent task families, top-line success is essentially unchanged under assist on/off  
1060 ( $\Delta \leq 0.10$ ). The only non-trivial effect is on rest, where assists improve post-retrieval completion  
1061 by  $+0.05$  without changing query/retrieval rates, consistent with assists acting on materialization and  
1062 post-retrieval execution.

## H Retrained retriever and CommNet baseline

### H.1 Retrained candidate-generation MLP (§3.5)

**Dataset extraction.** From the 10-seed hazard-recovery benchmark (MiniGrid-LavaGapS7-v0, method `milestone_state_conditioned_hazard_recovery_v7`), an 18-dimensional feature vector was extracted for every step of every episode of every seed — 3 numeric agent-state scalars (energy, thirst, risk\_budget), 9 semantic-tag confidences, 6 intent one-hot — and a binary label: *was this symbolic node ever selected as a candidate during the run?* Dataset totals: 814 examples, 113 positives (13.9%).

**Split.** By seed: train {2, 3, 4, 5, 6, 7} (505 examples, 9.7% positive), val {8, 9} (170, 22.4% positive), test {10, 11} (139, 18.7% positive).

**Architecture.**  $18 \rightarrow 32 \rightarrow 32 \rightarrow 1$ , ReLU activations, BCE with  $\text{pos\_weight} \approx 9.3$  for class imbalance, Adam  $\text{lr} = 2 \cdot 10^{-3}$  with  $\text{weight\_decay} = 10^{-4}$ , batch 64, 200 epochs, best by val loss.

### Results.




Figure 9: Refined attribution diagnostic. *Left*: 349 retrieval calls aggregated across 10 seeds; 317 (90.8%) returned empty candidate sets, 31 (8.9%) returned at least one satisfied candidate, only 1 (0.3%) returned candidates that were ranked but unsatisfied. The bottleneck is candidate generation, not candidate ranking. *Right*: the trained MLP underperforms a majority-class baseline on all three seed-disjoint splits, confirming that instantaneous state features alone are insufficient to predict candidate eligibility.

The MLP underperforms majority-class on accuracy across all splits and achieves only moderate precision at fixed recall. Instantaneous state features are insufficient to predict candidate eligibility downstream. This is the negative-but-informative finding referenced in §3.5: the framework’s “retrieval-limited” attribution refines to “memory-accumulation-limited” rather than “ranking-limited”.

### H.2 CommNet-style learned communication baseline (§4.5)

**Policy architecture.** Per-agent observation is 4-dim direction one-hot +  $5 \times 5 \times 5$  local cell-type window (125) + 2-dim normalised position = 131. Shared-weights encoder  $131 \rightarrow 64 \rightarrow 64$  (ReLU), 16-dim comm projection, peer mean-pool (excluding self), combine  $80 \rightarrow 64$  (ReLU), actor head  $64 \rightarrow 4$  (TURN\_LEFT, TURN\_RIGHT, FORWARD, NOOP), critic head  $64 \rightarrow 1$ . This is a CommNet-style learned-communication architecture (Sukhbaatar et al., 2016).

**Training.** REINFORCE (Williams, 1992) with value-baseline;  $\text{loss} = -\text{mean}(\log \pi(a) \cdot (R - V)) + 0.5 \cdot \text{MSE}(V, R) - 0.01 \cdot H[\pi]$ ; Adam  $\text{lr} = 3 \cdot 10^{-4}$ , gradient clipping 1.0. Reward shaping: +1.0 on first success per agent,  $-0.01$  step cost,  $-0.1$  per hazard contact. 2000 episodes cycling 200 environment seeds, wall-clock 100 s on CPU. Held-out eval: 50 environment seeds 250–299.

1090 **Training curve.** Episode-level success (last-100 window) stays in  $[0.13, 0.23]$  throughout training;  
 1091 no monotonic improvement after  $\sim 200$  episodes. The policy converges to a local optimum where  
 1092 one of three agents finds water by chance.




Figure 10: Learned-communication baseline diagnostics. *Left:* CommNet training curve (smoothed success rate over a 100-episode window) over 2000 training episodes; success stays below 0.25 throughout. The dashed line is the symbolic peer architecture at  $K=8$  on the same scenario for reference. *Right:* Q/R/M/C profile of the trained CommNet on 50 held-out seeds vs symbolic peer at  $K=8$ .  $Q^*$  is saturated at 1.0 on CommNet (degenerate by construction);  $R/M$  collapse together; only on the symbolic architecture are the conditional rates separated.

1093 **Q/R/M/C evaluation (50 held-out seeds, stochastic action sampling).**

1094 Compare to the symbolic peer architecture at  $K=8$ :  $P(M^*|R^*)=0.69$ ,  $P(C^*|M^*)=0.87$ , success  
 1095  $\sim 0.65$ .

1096 **PPO variant: structural degeneracy verified at competitive success.** A PPO-trained (Schulman et al., 2017) variant of the same CommNet architecture (clip  $\epsilon=0.2$ , generalized advantage estimation (Schulman et al., 2016) with  $\lambda=0.95$ ,  $\gamma=0.99$ , 4 PPO epochs per rollout, 64 rollouts per update, 200 updates,  $\approx 96$  minutes on a single CPU core; full setup in experiments/commnet\_ppo\_baseline/train\_ppo\_commnet.py) reaches 0.667 success on 100 held-out seeds — competitive with the symbolic peer architecture at  $K=8$  (0.65).

1102 The PPO result converts the structural argument from an assertion (“training budget would not change this”) to a measurement:  $Q^*$  remains saturated at 1.00,  $P(M^*|R^*)$  remains 1.00 (R/M collapse), even  
 1103 at competitive success. The Q stage is degenerate by construction in any policy whose communication  
 1104 channel is a real-valued vector with a learned linear projection — the channel is never informatively  
 1105 empty, regardless of training budget — and without an explicit symbolic target-lock interface the M  
 1106 event cannot separate from R. Only architectures with an explicit *query/lock* interface (the symbolic  
 1107 stack of §4.1) expose stage-separable Q/R/M/C events.  
 1108

## 1109 I Reproducibility commands

1110 An anonymized public repository accompanies this submission. The main paper artifacts were  
 1111 generated from the following scripts; PYTHONPATH=. and a Python 3 virtual environment with numpy,  
 1112 scipy, pandas, matplotlib, imageio, gymnasium, and minigrid are assumed throughout.

1113 **One-command reproduction (multi-agent block,  $\approx 22$  minutes on 8 cores)**

1114 bashscripts/reproduce\_paper.sh

## 1115 Clean Paper 1 suite

- 1116 • smoke suite (all local non-LLM blocks): `pythonscripts/run_paper1_clean_experiments.py --modesmoke --skip_llm --out_dir/tmp/paper1_clean_experiments_smoke`
- 1117
- 1118 • full local suite (single-agent, oracle, multi-agent, CommNet, portability/noise):  
1119 `pythonscripts/run_paper1_clean_experiments.py --modefull --skip_`  
1120 `llm --workers8 --out_dir/tmp/paper1_clean_experiments_full`
- 1121 • LLM TextNav step-budget sweep (requires local Ollama), run once per budget  $K' \in \{6, 12, 25\}$ :  
1122 `python-mexperiments.llm_collective.run_model_sweep --modelsqwen3:4bllama3.`  
1123 `1:latest --episodes8 --step_limit<K'> --out_dir/tmp/llm_steplimit_<K'>`

## 1124 Single-agent (§3.2–§3.4)

- 1125 • article benchmark and learned-baseline suite: `scripts/benchmark_symbolic_memory_`  
1126 `article.py`
- 1127 • article analysis and figure generation: `scripts/analyze_symbolic_memory_article.py`
- 1128 • synthetic factorization sanity check: `scripts/validate_qrmc_factorization.py`
- 1129 • oracle-stage intervention benchmark: `scripts/benchmark_oracle_stage_interventions.`  
1130 `py`
- 1131 • BabyAI transfer benchmark: `scripts/compare_semnav_babyai.py`

## 1132 Multi-agent (§4.5)

- 1133 • Base batch — peer at  $K \in \{1, 2, 4, 8, 16\}$  and the three cadence-free architectures (in-  
1134 dependent, shared bus, central aggregator), 40 seeds (25,920 runs = 16,200 peer + 9,720  
1135 cadence-free),  $\approx 18$  min: `pythonexperiments/big_experiment/exp_cadence_phase.`  
1136 `py --modefull --workers8 --out_dir/tmp/big_experiment_qrmc_40`
- 1137 • Extended-cadence batch — peer at  $K \in \{32, 48, 64\}$  only (9,720 runs),  $\approx 4$  min:  
1138 `pythonexperiments/big_experiment/exp_extra_K.py --workers8 --out_dir/tmp/`  
1139 `big_experiment_extraK`
- 1140 *Note.* These two batches partition the 35,640 total by *cadence range*; the architecture-wise  
1141 partition of Appendix F.2 (25,920 peer over all eight cadences vs. 9,720 cadence-free) reuses  
1142 the same two totals for *different* subsets (a coincidence of 3 extended cadences  $\times$  3,240 = 3  
1143 cadence-free architectures  $\times$  3,240). The combined log is `tmp/paper1_clean_experiments_`  
1144 `full/multiagent_cadence_full/runs.csv`.
- 1145 • Q/R/M/C stage analysis and bootstrap CIs: `pythonexperiments/big_experiment/`  
1146 `analyze_qrmc.py --runs_csvtmp/big_experiment_qrmc_40/runs.csv --out_`  
1147 `dirtmp/big_experiment_qrmc_40`
- 1148 • Multi-agent oracle interventions: `pythonexperiments/big_experiment/exp_oracle_`  
1149 `interventions.py --workers8 --out_dir/tmp/big_experiment_oracle`
- 1150 • Phase 1–4 smoke tests: `pythonscripts/run_all_smokes.py --out_dir/tmp/all_smokes`
- 1151 • PPO-trained CommNet baseline (Table 16,  $\approx 96$  min on a single CPU core):  
1152 `python-mexperiments.commnet_ppo_baseline.train_ppo_commnet --total_`  
1153 `updates200 --rollouts_per_update64`
- 1154 • Q/R/M/C eval of the PPO policy on 100 held-out seeds: `pythonexperiments/commnet_`  
1155 `baseline/eval_with_qrmc.py --policy_pathexperiments/commnet_ppo_baseline/`  
1156 `ppo_policy.pt --out_dir/experiments/commnet_ppo_baseline --n_episodes100`
- 1157 • MiniGrid-substrate portability sweep (§6, 100 runs,  $\approx 2$  min): `python-mexperiments.`  
1158 `minigrid_multiagent_wrapper.run_portability_sweep --out_dir/tmp/minigrid_`  
1159 `portability`

## 1160 J Asset and license note

1161 The experiments use MiniGrid (Chevalier-Boisvert et al., 2023) and BabyAI (Chevalier-Boisvert  
1162 et al., 2018), both released under the MIT License. The locally implemented baselines and benchmark



1163 runner scripts introduced in this paper will be released under the MIT License as part of the camera-  
1164 ready artifact package.

Table 14: SOTA coverage for Paper 1. “Instrumented” means the run emits the operational  $Q^*$ ,  $R^*$ ,  $M^*$ ,  $C^*$  events used by Definition 1.

| Family                                        | Representative method                                                         | Status in this paper             | Q/R/M/C comparison point                                                                                                                                                                                                                                                                        |
|-----------------------------------------------|-------------------------------------------------------------------------------|----------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Topological navigation memory                 | SPTM-like graph baseline (Savinov et al., 2018)                               | Direct MiniGrid baseline         | Tests whether topological memory without semantic query/lock exposes non-trivial stage events. Strong success on easy tasks but no explicit query/lock interface, so diagnostics are structurally limited.                                                                                      |
| Behavior cloning                              | Learned grid-observation BC                                                   | Direct MiniGrid baseline         |                                                                                                                                                                                                                                                                                                 |
| Learned communication                         | CommNet PPO (Sukhbaatar et al., 2016; Schulman et al., 2017)                  | Direct multi-agent baseline      | Competitive success; $Q^*$ saturates and R/M collapse, verifying observability loss.                                                                                                                                                                                                            |
| Cooperative MARL (CTDE / value decomposition) | QMIX, MAPPO, MADDPG (Rashid et al., 2018; Yu et al., 2022; Lowe et al., 2017) | Primary-source reference         | Stronger end-to-end baselines than CommNet; they optimise a joint policy/message channel without an explicit query/lock interface, so the same observability-collapse argument applies — the framework quantifies the diagnostic cost of implicit coordination rather than competing on reward. |
| Reasoning/acting LLM agents                   | ReAct (Yao et al., 2023)                                                      | Primary-source reference         | Q/R/M/C can instrument the resulting action traces if queries, retrieved evidence, target locks, and completions are logged.                                                                                                                                                                    |
| Reflective LLM agents                         | Reflexion (Shinn et al., 2023)                                                | Primary-source reference         | Reflection can be treated as upstream Q/R revision; the framework asks whether it improves R, M, or C.                                                                                                                                                                                          |
| Embodied LLM agents                           | Voyager (Wang et al., 2023)                                                   | Primary-source reference         | Skill/memory retrieval maps naturally onto R and M, but full reproduction is outside this paper.                                                                                                                                                                                                |
| Social/memory simulations                     | Generative Agents (Park et al., 2023)                                         | Primary-source reference         | Their observation/memory/reflection loop motivates logging stage-specific memory use rather than only believability.                                                                                                                                                                            |
| Multi-agent LLM frameworks                    | AutoGen (Wu et al., 2023)                                                     | Primary-source reference         | Conversation traces provide candidate Q/R/M/C events; this paper supplies the measurement contract.                                                                                                                                                                                             |
| Long-term memory systems                      | MemGPT / Mem0 (Packer et al., 2023; Chhikara et al., 2025)                    | Primary-source reference         | Persistent memory systems can be evaluated by where they fail: retrieval, materialization, or completion.                                                                                                                                                                                       |
| Local LLM sub-strate                          | qwen3:4b, llama3.1                                                            | Direct TextNav step-budget study | Q/R/M/C localizes a completion-limited failure ( $C^* \rightarrow 0$ while $Q/R/M$ stay saturated) in real LLM agents; see Table 3.                                                                                                                                                             |

Table 15: Multi-agent oracle interventions on the scarcity cases at three cadences. Mean time-to-success (and success rate) averaged across  $(N, T) \in \{(3, 2), (5, 3), (8, 5)\}$ , two layouts (asymmetric, random), 15 seeds. Oracle R returns ground-truth water positions; oracle M locks the nearest real water; **oracle C** guarantees completion of a real materialised lock (Appendix F.12). Oracle C gives the largest *time* reduction (it removes travel) but the smallest *success* gain — completion is a time cost, not the success bottleneck, which sits upstream at R/M.

| $K$                                      | Baseline | Oracle R     | Oracle M | Oracle C    |
|------------------------------------------|----------|--------------|----------|-------------|
| <i>mean <math>t_{\text{succ}}</math></i> |          |              |          |             |
| 1                                        | 10.33    | 7.22         | 7.21     | <b>6.19</b> |
| 4                                        | 10.98    | 7.22         | 7.21     | <b>6.07</b> |
| 16                                       | 8.45     | 7.22         | 7.21     | <b>6.41</b> |
| <i>success rate</i>                      |          |              |          |             |
| 1                                        | 0.591    | <b>0.614</b> | 0.602    | 0.591       |
| 4                                        | 0.575    | <b>0.614</b> | 0.602    | 0.600       |
| 16                                       | 0.578    | <b>0.614</b> | 0.602    | 0.598       |

| Metric                           | Train | Val   | Test  |
|----------------------------------|-------|-------|-------|
| Accuracy                         | 0.584 | 0.553 | 0.590 |
| Majority-class baseline accuracy | 0.903 | 0.776 | 0.813 |
| Precision @ recall = 0.70        | 0.171 | 0.278 | 0.275 |

| Quantity     | Value                      |
|--------------|----------------------------|
| $Q^*$ rate   | <b>1.00</b> (saturated)    |
| $R^*$ rate   | 0.50                       |
| $M^*$ rate   | 0.50 (collapses onto $R$ ) |
| $C^*$ rate   | 0.19                       |
| $P(R   Q)$   | 0.50                       |
| $P(M   R)$   | 1.00                       |
| $P(C   M)$   | 0.42                       |
| Success rate | 0.19                       |
| Mean steps   | 80 (budget exhausted)      |

Table 16: REINFORCE vs PPO CommNet vs symbolic peer ( $K=8$ ) on the scarcity scenario ( $N=3$ ,  $T=2$ , asymmetric, hazard 0.05). PPO is competitive on success rate, yet still exhibits Q-saturation and R/M collapse — the diagnostic claim verified, not asserted.

| Metric         | CommNet REINFORCE | <b>CommNet PPO</b> | Symbolic peer ( $K=8$ ) |
|----------------|-------------------|--------------------|-------------------------|
| Success rate   | 0.19              | <b>0.667</b>       | 0.65                    |
| $Q^*$ rate     | 1.00              | <b>1.00</b>        | 0.99                    |
| $R^*$ rate     | 0.50              | <b>0.997</b>       | 0.69                    |
| $M^*$ rate     | 0.50              | <b>0.997</b>       | 0.69                    |
| $P(M^*   R^*)$ | 1.00              | <b>1.00</b>        | 0.69                    |
| $P(C^*   M^*)$ | 0.42              | <b>0.67</b>        | 0.87                    |
| R/M separated? | no                | <b>no</b>          | <b>yes</b>              |

## 1165 NeurIPS Paper Checklist

### 1166 1. Claims

1167 Question: Do the main claims made in the abstract and introduction accurately reflect the  
1168 paper’s contributions and scope?

1169 Answer: [Yes]

1170 Justification: The abstract and introduction present the paper primarily as a four-stage  
1171 diagnostic framework with a reference symbolic-semantic instantiation, and they explicitly  
1172 scope the evidence to benchmarked navigation environments with a limited portability check  
1173 (Abstract; Sections 1 and 6).

1174 Guidelines:

- 1175 • The answer [N/A] means that the abstract and introduction do not include the claims  
1176 made in the paper.
- 1177 • The abstract and/or introduction should clearly state the claims made, including the  
1178 contributions made in the paper and important assumptions and limitations. A [No] or  
1179 [N/A] answer to this question will not be perceived well by the reviewers.
- 1180 • The claims made should match theoretical and experimental results, and reflect how  
1181 much the results can be expected to generalize to other settings.
- 1182 • It is fine to include aspirational goals as motivation as long as it is clear that these goals  
1183 are not attained by the paper.

### 1184 2. Limitations

1185 Question: Does the paper discuss the limitations of the work performed by the authors?

1186 Answer: [Yes]

1187 Justification: The paper includes an explicit “Limitations and future work” section that dis-  
1188 cusses environment difficulty, the observable-query and observable-occupancy assumptions,  
1189 the slope-level scope of Empirical Claim 1, the learned-baseline scope, and agent-count  
1190 scaling (Section 6).

1191 Guidelines:

- 1192 • The answer [N/A] means that the paper has no limitation while the answer [No] means  
1193 that the paper has limitations, but those are not discussed in the paper.
- 1194 • The authors are encouraged to create a separate “Limitations” section in their paper.
- 1195 • The paper should point out any strong assumptions and how robust the results are to  
1196 violations of these assumptions (e.g., independence assumptions, noiseless settings,  
1197 model well-specification, asymptotic approximations only holding locally). The authors  
1198 should reflect on how these assumptions might be violated in practice and what the  
1199 implications would be.
- 1200 • The authors should reflect on the scope of the claims made, e.g., if the approach was  
1201 only tested on a few datasets or with a few runs. In general, empirical results often  
1202 depend on implicit assumptions, which should be articulated.
- 1203 • The authors should reflect on the factors that influence the performance of the approach.  
1204 For example, a facial recognition algorithm may perform poorly when image resolution  
1205 is low or images are taken in low lighting. Or a speech-to-text system might not be  
1206 used reliably to provide closed captions for online lectures because it fails to handle  
1207 technical jargon.
- 1208 • The authors should discuss the computational efficiency of the proposed algorithms  
1209 and how they scale with dataset size.
- 1210 • If applicable, the authors should discuss possible limitations of their approach to  
1211 address problems of privacy and fairness.
- 1212 • While the authors might fear that complete honesty about limitations might be used by  
1213 reviewers as grounds for rejection, a worse outcome might be that reviewers discover  
1214 limitations that aren’t acknowledged in the paper. The authors should use their best  
1215 judgment and recognize that individual actions in favor of transparency play an impor-  
1216 tant role in developing norms that preserve the integrity of the community. Reviewers  
1217 will be specifically instructed to not penalize honesty concerning limitations.

### 3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: The paper states the conditional stage-Markov property and positivity assumptions in Section 2.2, gives a proof sketch in the main text, and includes a full formal argument for Definition 1 (single-agent stage estimability) in Appendix A together with the multi-agent extension (Definition 2) and the slope-level argument for Empirical Claim 1.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

### 4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The paper specifies the task families, environments, seeds, episode counts, assist modes, baselines, statistical protocol, and benchmark semantics; the single-agent setup is in Section 3.1 and the multi-agent sweep protocol in Section 4.4; the operational stage definitions used to extract  $Q^*$ ,  $R^*$ ,  $M^*$ ,  $C^*$  from runtime logs are in Appendix B.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

1273 (d) We recognize that reproducibility may be tricky in some cases, in which case  
1274 authors are welcome to describe the particular way they provide for reproducibility.  
1275 In the case of closed-source models, it may be that access to the model is limited in  
1276 some way (e.g., to registered users), but it should be possible for other researchers  
1277 to have some path to reproducing or verifying the results.

## 1278 5. Open access to data and code

1279 Question: Does the paper provide open access to the data and code, with sufficient instruc-  
1280 tions to faithfully reproduce the main experimental results, as described in supplemental  
1281 material?

1282 Answer: [Yes]

1283 Justification: The paper describes a unified benchmark runner and source-of-truth artifacts.  
1284 An anonymized public release of the codebase — single-agent benchmark, multi-agent grid  
1285 environment, 35,640-run sweep driver, Q/R/M/C analyser, oracle-intervention scripts, and  
1286 exact commands to reproduce every table and figure — accompanies the submission as  
1287 supplemental material (anonymized GitHub repository; single-command reproduction of  
1288 the main multi-agent sweep on an 8-core machine in  $\approx 22$  minutes, matching Appendix H).

1289 Guidelines:

- 1290 • The answer [N/A] means that paper does not include experiments requiring code.
- 1291 • Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- 1292 • While we encourage the release of code and data, we understand that this might not  
1293 be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not  
1294 including code, unless this is central to the contribution (e.g., for a new open-source  
1295 benchmark).
- 1296 • The instructions should contain the exact command and environment needed to run to  
1297 reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- 1298 • The authors should provide instructions on data access and preparation, including how  
1299 to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- 1300 • The authors should provide scripts to reproduce all experimental results for the new  
1301 proposed method and baselines. If only a subset of experiments are reproducible, they  
1302 should state which ones are omitted from the script and why.
- 1303 • At submission time, to preserve anonymity, the authors should release anonymized  
1304 versions (if applicable).
- 1305 • Providing as much information as possible in supplemental material (appended to the  
1306 paper) is recommended, but including URLs to data and code is permitted.

## 1307 6. Experimental setting/details

1310 Question: Does the paper specify all the training and test details (e.g., data splits, hyperpa-  
1311 rameters, how they were chosen, type of optimizer) necessary to understand the results?

1312 Answer: [Yes]

1313 Justification: The benchmark protocol, task definitions, assist settings, seeds, episode counts,  
1314 baseline families, and the learned baselines are described in the main text together with the  
1315 measurement and statistical protocols: single-agent setup in Section 3.1, multi-agent sweep  
1316 in Section 4.4, and per-baseline hyperparameters in Appendix G.

1317 Guidelines:

- 1318 • The answer [N/A] means that the paper does not include experiments.
- 1319 • The experimental setting should be presented in the core of the paper to a level of detail  
1320 that is necessary to appreciate the results and make sense of them.
- 1321 • The full details can be provided either with the code, in appendix, or as supplemental  
1322 material.

## 1323 7. Experiment statistical significance

1324 Question: Does the paper report error bars suitably and correctly defined or other appropriate  
1325 information about the statistical significance of the experiments?

Answer: [Yes]

Justification: The paper reports 95% bootstrap confidence intervals over seed-level aggregates ( $B=5,000$  resamples; Section 3.1, “Protocol” paragraph) and uses paired Wilcoxon signed-rank tests with Holm-Bonferroni correction for assist comparisons. CIs are shown in Tables 3 and 5 (Best semantic per-task rates; the 35,640-run MOP column) and in Figures 4–6 (bootstrap shaded bands).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

## 8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: The paper discloses the hardware platform (10-core ARM CPU, 32 GB RAM; CPU-only), the scale of each benchmark, and wall-clock costs — single-agent oracle sweep  $< 1$  minute (Section 3.1); main multi-agent sweep (35,640 runs)  $\approx 22$  minutes on 8 cores (Section 4.4);  $N=12$  scale-up (8,640 runs)  $\approx 12$  minutes (Appendix D); PPO CommNet baseline (200 updates  $\times$  64 rollouts)  $\approx 96$  minutes on a single CPU core (Appendix G); MiniGrid portability sweep (100 runs)  $\approx 2$  minutes on a single core (§6).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn’t make it into the paper).

## 9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The work is a simulation-only study in standard benchmark environments, uses anonymized authorship for review, and does not involve human subjects, private data, or deployment claims that conflict with the NeurIPS Code of Ethics.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

## 10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: The paper includes a dedicated “Broader impacts” section that discusses the primary methodological benefits of the work and the indirect risks of improving navigation systems without adequate safeguards (Section 7).

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

## 11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper does not release high-risk models or scraped datasets; it studies navigation policies in simulated environments.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.



- 1432 • We recognize that providing effective safeguards is challenging, and many papers do  
1433 not require this, but we encourage authors to take this into account and make a best  
1434 faith effort.

## 1435 12. Licenses for existing assets

1436 Question: Are the creators or original owners of assets (e.g., code, data, models), used in  
1437 the paper, properly credited and are the license and terms of use explicitly mentioned and  
1438 properly respected?

1439 Answer: [Yes]

1440 Justification: The paper cites the upstream benchmark platforms (MiniGrid, BabyAI) in  
1441 the main references. Appendix I states that both are MIT-licensed and that the benchmark  
1442 scripts introduced here will be released under the MIT License at camera-ready.

1443 Guidelines:

- 1444 • The answer [N/A] means that the paper does not use existing assets.
- 1445 • The authors should cite the original paper that produced the code package or dataset.
- 1446 • The authors should state which version of the asset is used and, if possible, include a  
1447 URL.
- 1448 • The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- 1449 • For scraped data from a particular source (e.g., website), the copyright and terms of  
1450 service of that source should be provided.
- 1451 • If assets are released, the license, copyright information, and terms of use in the  
1452 package should be provided. For popular datasets, `paperswithcode.com/datasets`  
1453 has curated licenses for some datasets. Their licensing guide can help determine the  
1454 license of a dataset.
- 1455 • For existing datasets that are re-packaged, both the original license and the license of  
1456 the derived asset (if it has changed) should be provided.
- 1457 • If this information is not available online, the authors are encouraged to reach out to  
1458 the asset's creators.

## 1459 13. New assets

1460 Question: Are new assets introduced in the paper well documented and is the documentation  
1461 provided alongside the assets?

1462 Answer: [No]

1463 Justification: The submission introduces benchmark scripts and analysis artifacts. A formal  
1464 anonymized release package with documentation is planned for camera-ready; Appendix H  
1465 lists all scripts needed to reproduce each reported result, and Appendix I states the intended  
1466 license for the new assets.

1467 Guidelines:

- 1468 • The answer [N/A] means that the paper does not release new assets.
- 1469 • Researchers should communicate the details of the dataset/code/model as part of their  
1470 submissions via structured templates. This includes details about training, license,  
1471 limitations, etc.
- 1472 • The paper should discuss whether and how consent was obtained from people whose  
1473 asset is used.
- 1474 • At submission time, remember to anonymize your assets (if applicable). You can either  
1475 create an anonymized URL or include an anonymized zip file.

## 1476 14. Crowdsourcing and research with human subjects

1477 Question: For crowdsourcing experiments and research with human subjects, does the paper  
1478 include the full text of instructions given to participants and screenshots, if applicable, as  
1479 well as details about compensation (if any)?

1480 Answer: [N/A]

1481 Justification: The paper does not involve crowdsourcing or human-subject experiments.

1482 Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

#### 15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: The paper does not involve human subjects or human-participant data collection.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

#### 16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: LLMs are not an important, original, or non-standard component of the core methods or experiments reported in the paper.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.