
Q/R/M/C: A Stage-Decomposition Measurement Framework for Memory-Based Navigation, with a Minimal Collective Semantic Memory

Anonymous Author(s)

Affiliation

Address

email

Abstract

Setting. Success-only metrics collapse qualitatively different memory-based navigation failures into a single scalar, hiding *which stage* of the retrieval pipeline is responsible. This paper develops a measurement framework, not a new theoretical result: the Q/R/M/C stage decomposition (**Q**uery formation, **R**etrieval satisfaction, target **M**aterialization, **C**ompletion after retrieval) is operationalised so that the four conditional rates are computable directly from standard trajectory logs (Definition 1; the consistency argument is the chain rule and we say so explicitly). **Single-agent measurement.** On a 10-seed MiniGrid benchmark (FourRooms, LavaGapS7) and a BabyAI portability check, the framework localises bottlenecks: hazard recovery is retrieval-limited (oracle retrieval lifts success $0.39 \rightarrow 0.60$), goal-region rejoin is downstream-limited. Log inspection refines the attribution: candidate ranking is near-optimal (96.9%) while candidate generation fails at 91% of decision points — a memory-accumulation, not a ranking, limit. **Multi-agent measurement.** On a 35,640-run multi-agent sweep over four communication architectures (independent, shared-bus, central-aggregator, peer-broadcast at eight cadences $K \in \{1, 2, 4, 8, 16, 32, 48, 64\}$, 40 seeds), the bottleneck-shift between the M and C links appears directly as a negative within-cadence correlation peaking at $K=4$ (Pearson -0.61); mean time-to-success has an interior minimum at $K=8$. Both slope signs (Empirical Claim 1) hold at $p < 10^{-4}$; we do *not* claim an interior maximum on the conditional product Φ , where the data show a monotone trend toward the independent-agent boundary. **A minimal collective semantic memory.** The four architectures sit at the corners of a richer design space defined by explicit *merge*, *trust*, and *staleness* rules. We instantiate a minimal collective semantic memory — trust-weighted merge with recency-decay staleness on top of peer broadcast — and measure it on the same Q/R/M/C scale, demonstrating that it dominates the best fixed-cadence peer architecture on $(M \times C)$ for the scarcity scenario. This converts “collective semantic memory” from a programmatic claim to an instantiated artefact evaluable on the same logs. **Scope.** All experiments are in grid-world substrates (MiniGrid, BabyAI, a custom 12×10 grid). The contribution is a measurement protocol and a first prototype; richer substrates (Habitat, ProcTHOR, VMAS) are future work.

1 Introduction

Classical navigation benchmarks often assume that the agent is given an explicit destination coordinate or a dense task reward. Many practical navigation problems are different: the agent must reach a place of a certain *type*, such as a safe rest area, a water-like resource, a post-hazard recovery exit,




Figure 1: Q/R/M/C decomposes one semantic-navigation episode into four observable stages. **Top:** the four stages and their event-level definitions. **(a)** A successful water-search episode anchors at one event per stage, producing logged starred events Q^* , R^* , M^* , C^* . **(b)** The same scalar “end-to-end success” summary masks qualitatively different failures on the four single-agent task families: water and rest are strong; goal-region rejoin (FourRooms) is downstream-limited; hazard recovery (LavaGapS7) is retrieval-limited (oracle interventions, §3.4).

or a goal-associated region, specified by a semantic pattern rather than a fixed (x, y) waypoint. The place must therefore be *remembered and retrieved by meaning*. This setting opens a methodological question: can we measure *where* memory-based semantic navigation breaks down, not only *whether* it succeeds?

We answer this question with a four-stage diagnostic framework that applies to both single agents and multi-agent collectives. The decomposition is Q/R/M/C: **Q**uery formation, **R**etrieval satisfaction, **M**aterialization, and **C**ompletion after retrieval. The decomposition factorizes per-episode semantic-path completion and, under a conditional stage-Markov property, its factors are estimable from trajectory logs (Definition 1). The rest of the paper develops a single argument in three parts.

Single agent. Stage attribution is non-trivial. Success-only metrics collapse qualitatively different failures into a single scalar. In a 10-seed MiniGrid benchmark, water and rest tasks are strong end-to-end (0.95, 0.93 success), but goal-region rejoin and hazard recovery are substantially harder, and for *different reasons*. Oracle-stage interventions show that hazard recovery is retrieval-limited (oracle retrieval lifts success $0.39 \rightarrow 0.60$), while goal-region rejoin is downstream-limited (oracle retrieval rescues semantic-path completion without changing end-to-end success). Inspection of the framework’s own logs then sharpens the attribution: the retrieval bottleneck is not candidate *ranking* (precision 96.9% when candidates exist) but candidate *generation*: the memory contains no eligible symbols at 91% of decision points. The single-agent bottleneck is a *memory-accumulation* limit: evidence is observed but never consolidated into queryable structure.

Multi-agent. Consolidation acquires a rate. When N agents share memory through a broadcast protocol with cadence K , the factorization extends with enlarged conditioning sets (Definition 2), and K becomes a structural axis: it is, mechanically, a knob on *how fast private memories consolidate into a collective one*. A slope-level claim (Empirical Claim 1) predicts a bottleneck shift: as K decreases, $P(M^*|R^*)$ rises (fresher peer memory enables reliable commitment) and $P(C^*|M^*)$ falls (peers race for the same targets). Both slope signs are confirmed on a 35,640-run sweep ($p < 10^{-4}$), the shift is directly visible as a negative within-cadence correlation between the two links, and the resulting trade-off produces an interior optimum of mean time-to-success at $K = 8$: consolidation has a non-trivial optimal rate.

Synthesis. Both regimes point at the same object. The single-agent diagnosis (consolidation of evidence into symbols is the limit) and the multi-agent diagnosis (the rate and cost of sharing consolidated memory determine the optimum) converge on one research target: *collective semantic memory*, a shared symbolic structure with explicit merge, trust, and staleness rules, rather than

68 snapshots exchanged on a timer. Section 5 states this program and shows that Q/R/M/C is the
 69 measurement instrument it requires: any candidate design is testable by how it moves $P(M^*|R^*)$
 70 and $P(C^*|M^*)$.

71 **Scope of the claims.** The slope-level statement of Empirical Claim 1 is on the two conditional rates
 72 $P(M^*|R^*; K)$ and $P(C^*|M^*; K)$, both verified on enlarged data. We make *no* formal claim about
 73 the conditional product $\Phi = P(M^*|R^*) \cdot P(C^*|M^*)$; on the 35,640-run sweep Φ is monotone in
 74 K . The interior optimum that the $M \leftrightarrow C$ trade-off produces is on the efficiency metric mean time-to-
 75 success, not on Φ . We are explicit about this distinction because a stage-level operationalisation is
 76 testable on specific summaries and can be *wrong* on others, which is what separates a measurement
 77 framework from a stage-agnostic guess.

78 **Why this framing is timely.** Embodied and goal-conditioned agents increasingly retrieve places,
 79 objects, or tools by meaning (Andrychowicz et al., 2017; Ghosh et al., 2021; Savva et al., 2019).
 80 LLM-agent systems lean on explicit memory layers but evaluate mostly end-to-end (Wang et al., 2023;
 81 Shinn et al., 2023; Packer et al., 2023). Multi-agent reinforcement learning likewise increasingly
 82 involves communicating populations (Lowe et al., 2017; Foerster et al., 2018), but the question
 83 “which agent should know what, and when?” lacks a unified diagnostic vocabulary. Our framework
 84 supplies one.

85 **Contributions.** This is a methodology paper. We are explicit about what is bookkeeping vs. empirical
 86 claim vs. instantiated artefact.

- 87 1. **Measurement protocol (single-agent).** The Q/R/M/C decomposition is operationalised so that
 88 the conditional rates are estimable from trajectory logs (Definition 1). The consistency argument
 89 is the chain rule together with consistency of frequency estimators; we do not claim mathematical
 90 novelty for this — the contribution is the *logging discipline* that makes the decomposition usable
 91 on standard navigation traces.
- 92 2. **Measurement protocol (multi-agent).** An extension to N communicating agents with two
 93 coupling channels (memory $\mathcal{M}_{-i}^{(K)}$, occupancy $\mathcal{O}_{-i}$) under a peer-broadcast protocol with
 94 cadence K (Definition 2).
- 95 3. **Empirical Claim 1 (bottleneck shift).** On a 35,640-run multi-agent sweep, the two stylised
 96 slope conditions (F) and (O) are both confirmed at $p < 10^{-4}$ (Spearman ± 0.5 on the standard-
 97 MiniGrid substrate; full benchmark in §4.5). We state this as a verified empirical observation,
 98 not a theorem. We do *not* claim an interior maximum on the conditional product Φ — that
 99 prediction fails in our data and we report it (§4.5).
- 100 4. **System and benchmarks.** A symbolic-semantic memory stack with four communication
 101 architectures (independent, shared-bus, central-aggregator, peer-to-peer broadcast); single-agent
 102 oracle benchmark (10-seed MiniGrid, BabyAI portability); 35,640-run multi-agent sweep at
 103 $N \in \{3, 5, 8\} + 8,640$ -run scale-up at $N = 12$; 100-run portability sweep on a standard-MiniGrid
 104 wrapper.
- 105 5. **A minimal collective semantic memory.** A first instantiation of *collective semantic memory*
 106 with explicit *merge*, *trust*, and *staleness* rules (§5), measured on the same Q/R/M/C scale. It
 107 dominates the best fixed-cadence peer architecture on (M $\times$ C) for the scarcity scenario — the
 108 framework’s measurement protocol applied to a new architecture, not an aspirational “next step.”
- 109 6. **Scope and honesty.** All experiments are in grid-world substrates. We treat that as a methodology-
 110 paper choice and state it explicitly (§6); richer substrates (Habitat, ProcTHOR, VMAS, Mine-
 111 Dojo) are future work.

112 2 The Q/R/M/C decomposition

113 2.1 Problem formulation and notation

114 The agent operates in a partially observable navigation environment. At each step it receives an
 115 observation o_t , maintains a memory state $\mathcal{M}_t$, and forms a query q_t encoding the current task
 116 intent. A retrieval step returns a candidate set of places matching q_t from $\mathcal{M}_t$; a chosen candidate
 117 is materialized into an actionable target; a local controller executes actions until either the target is
 118 reached or the episode times out. The target is not a coordinate but a semantic pattern, so success is a
 119 function of the chosen candidate’s semantic profile, not of any pre-supplied coordinate.

Habitat, AI2-THOR, and ProcTHOR made visually rich navigation evaluation standard (Savva et al., 2019; Kolve et al., 2017; Deitke et al., 2022). We deliberately focus on MiniGrid-class environments to obtain clean and reproducible stage measurements, and use BabyAI only as a portability check. We do not claim to solve LLM-agent memory here; rather, we argue that the same Q/R/M/C decomposition is a useful template for evaluating memory-grounded agent decisions beyond navigation.

The central empirical question is therefore not just whether the agent succeeds, but *which stage fails first* when it does not. Two related measurement layers are used. **Operational rates** count, per query attempt, whether the query yielded a non-empty candidate set, whether the chosen candidate satisfied the semantic predicate, and whether it was materialized into an actionable target; *completion after retrieval* is an episode-level indicator. **Nested episode-level events** $Q^* \Rightarrow R^* \Rightarrow M^* \Rightarrow C^*$ are used in the factorization below and audited in §3.2. Full operational definitions are given in Appendix B.

Notation. The four stages are Q, R, M, C with episode-level starred events $Q^*, R^*, M^*, C^* \in \{0, 1\}$, nested by construction. $P(\cdot)$ denotes the population probability; $\hat{P}(\cdot)$ its empirical-frequency estimator. **The letter M always refers to the Materialization stage; the number of targets in the multi-agent environment is written T (not M , constraint $T \leq N$, “scarcity” = $N > T$).** Other symbols: $\mathcal{C}_K$ peer-broadcast protocol with cadence K , $\mathcal{M}_{-i}^{(K)}$ peer memory snapshots, $\mathcal{O}_{-i}$ peer occupancy, $\nu(\cdot)$ freshness functional, $\Phi(K) = P(M^* | R^*; K) \cdot P(C^* | M^*; K)$, t_{succ} mean time-to-success. Full symbol table in Appendix C.1.

2.2 Single-agent factorization

Let Y denote overall episode success, and let Q^*, R^*, M^*, C^* denote the nested episode-level semantic-path events. The framework factorizes *semantic-path completion* rather than raw task success:

$$Q^* \rightarrow R^* \rightarrow M^* \rightarrow C^*.$$

For any query q and memory state $\mathcal{M}$,

$$P(C^* = 1 \mid q, \mathcal{M}) = P(Q^* = 1 \mid q, \mathcal{M}) P(R^* = 1 \mid Q^* = 1, q, \mathcal{M}) \\ \cdot P(M^* = 1 \mid R^* = 1, q, \mathcal{M}) P(C^* = 1 \mid M^* = 1, q, \mathcal{M}).$$

Overall task success Y can exceed C^* because some episodes succeed without traversing the full semantic retrieval path.

Assumption 1 (conditional stage-Markov property). Conditioned on the current stage being successful, the probability of the next stage depends only on the current successful stage, the current query, and the current memory state. Formally,

$$P(R^* \mid Q^*, q, \mathcal{M}, \text{history}) = P(R^* \mid Q^*, q, \mathcal{M}),$$

and analogously for M^* and C^* .

Assumption 2 (positivity). Each conditioning event used above has non-zero probability on the support of the benchmark.

Definition 1 (single-agent stage estimability). Under Assumptions 1–2, the four conditional rates in the factorization above are well-defined population quantities and the empirical-frequency estimators are consistent estimators of them; their product converges to semantic-path completion. The argument is elementary — the chain-rule factorization plus consistency of frequency estimators on Bernoulli events — and the value of stating it is operational, not mathematical: it pins down the *logging discipline* (which events, conditioned on what) under which the decomposition is readable from standard trajectory logs, and it is the contract that every experiment in this paper audits against (§3.2). Full statement and elementary argument in Appendix A.

When Assumption 1 may fail. Assumption 1 holds tightly here because q and $\mathcal{M}$ are explicitly logged at every stage event. Two practically relevant violations remain: adaptive in-episode query rewrites not surfaced to the logger, and hidden between-stage memory updates. In those cases the factorization should be read as a projection onto observable traces rather than as a full structural identification claim.




Figure 2: Stage-structured interpretation of the Q/R/M/C framework. Assist on/off ablations target intermediate mechanisms; in the multi-agent extension the cadence protocol $\mathcal{C}_K$ acts on the $R \rightarrow M$ and $M \rightarrow C$ edges.

Mechanism-level reading. The same decomposition admits a stage-structured mechanism interpretation. Figure 2 treats stage outcomes as a directed acyclic graph. Under this interpretation, assist on/off ablations are mechanism-targeted perturbations of materialization and post-retrieval execution rather than undifferentiated sensitivity analyses; in the multi-agent setting the same graph is replicated per agent and connected through the cadence protocol $\mathcal{C}_K$, which acts on $\mathcal{M}_{-i}^{(K)}$ at the $R \rightarrow M$ edge and on $\mathcal{O}_{-i}$ at the $M \rightarrow C$ edge — exactly the two edges where Empirical Claim 1 will predict the slope shift.

Relation to prior work. Closest are (i) cognitive-map and semantic-map frameworks (Tolman, 1948; O’Keefe and Nadel, 1978; Kuipers, 2000; Gupta et al., 2017; Savinov et al., 2018), (ii) goal-conditioned and successor-representation RL (Andrychowicz et al., 2017; Ghosh et al., 2021; Dayan, 1993; Stachenfeld et al., 2017; Momennejad et al., 2017), (iii) multi-agent communication (Lowe et al., 2017; Foerster et al., 2018; Sukhbaatar et al., 2016; Foerster et al., 2016) and gossip/consensus (Boyd et al., 2006; Xiao et al., 2007), and (iv) LLM-agent memory layers evaluated end-to-end (Wang et al., 2023; Shinn et al., 2023; Packer et al., 2023). Our framework is orthogonal to all four: it *diagnoses* where in the Q/R/M/C chain a given architecture succeeds or fails and lifts the cadence axis from network protocols to a cognitive-architecture parameter (full discussion in Appendix E.3).

2.3 Event semantics, bounded rationality, and an order-theoretic reading

The four starred events anchor at concrete logged triggers. We state the triggers once, because every empirical claim in the paper is a claim about them (per-regime formal definitions in Appendix B), and we give each stage its reading in terms of the resource it bounds.

Q^* (intent issued). The planner emits a semantic intent $q = (q^{\text{req}}, q^{\text{pref}}, q^{\text{pen}})$: small sets of required, preferred, and penalty tags. The episode-level event fires on emission, so it is saturated whenever the task is semantic ($\hat{P}(Q^*)=1.00$ in Table 1); the multi-agent per-tick logger additionally requires a non-empty candidate set (Appendix B.1), a bookkeeping difference we record rather than hide. Resource bound: attention — the intent is a few tags, not the full state.

R^* (memory answers). Retrieval evaluates q against the memory and returns a candidate set; the event fires when the set contains a candidate satisfying the task predicate (within ϵ of a true target in the multi-agent logger). The 91% empty-candidate-set failures of §3.5 are R -side events: no stored place clears the required-tag threshold. Resource bound: memory — retrieval can only surface what consolidation has stored.

M^* (target lock). The planner materializes one candidate into a concrete actionable cell. Candidacy itself is threshold-gated (a place enters the candidate set only if its score clears the planner’s gate, Appendix C), and the lock commits to a single candidate once per query rather than re-optimizing against the full memory at every tick. The event fires on lock acquisition (within ϵ of a true target in the multi-agent logger). Resource bound: deliberation.

C^* (closed-loop completion). The local controller executes until the locked cell is reached or the episode times out; the event fires on arrival. C is the only stage whose failure modes are extra-semantic: controller competence and, in the multi-agent regime, occupancy $\mathcal{O}_{-i}$. Resource bound: control, and in a collective, shared targets.

Why a deterministic strategy: bounded and ecological rationality. A natural objection is that the semantic stack is a fixed, deterministic strategy rather than a learned policy. We treat the determinism

as principled and name it: the pipeline is a satisficing agent in Simon’s sense (Simon, 1955, 1956). Candidacy is threshold-gated instead of exhaustively scored, the target commitment is made once instead of re-optimized at every tick, and the intent is a few tags instead of the full state. Bounded rationality is also what makes the four failure loci observable: each stage is a distinct resource bound (attention, memory, deliberation, control), and Q/R/M/C is the audit trail of where the bound binds. The learned baselines make the converse point empirically: policies without an explicit query/lock interface collapse the stages even at competitive success (BC in §3.3; CommNet PPO in §4.5), so observability is traded away exactly where end-to-end optimization is traded in. That the *same* heuristic is strong on water and rest, retrieval-starved on hazard recovery, and downstream-limited on goal-region (Figure 1b) is the ecological-rationality point that heuristic success is a property of the heuristic–environment pair (Simon, 1956; Gigerenzer and Goldstein, 1996); the framework is the instrument that locates the mismatch. The same lens reads the Φ -versus- t_{succ} split of §4.5: the chain probability Φ ignores the cost of traversing the chain, so under a resource-rational reading (Lieder and Griffiths, 2020) an optimum should be sought on a cost-bearing metric; that is where it appears ($K=8$ on mean time-to-success) and not on Φ .

An order-theoretic reading of the Q–R polarity. The query–retrieval interface carries standard order-theoretic structure that explains, rather than decorates, the single-agent diagnosis. Let G be the set of place records, Σ the tag alphabet, and $I_\theta \subseteq G \times \Sigma$ the incidence “place p carries tag τ with confidence at least θ ” induced by the memory’s per-place tag tables (Appendix C). For $B \subseteq \Sigma$ and $A \subseteq G$ put $\text{ext}(B) = \{p : \forall \tau \in B, (p, \tau) \in I_\theta\}$ and $\text{int}(A) = \{\tau : \forall p \in A, (p, \tau) \in I_\theta\}$. Then

$$A \subseteq \text{ext}(B) \iff B \subseteq \text{int}(A),$$

an antitone Galois connection; equivalently, int and ext form an adjoint pair between the two power-set posets (one taken with reversed order) viewed as thin categories (Ganter and Wille, 1999; Mac Lane, 1998), and the closed pairs form the concept lattice realized by the memory. In this vocabulary the single-agent diagnosis is exact: a query issues an intent $B = q^{\text{req}}$; retrieval fails empty precisely when $\text{ext}(B) = \emptyset$, that is, when the queried meaning lies outside the realized lattice; and consolidation acts on I_θ , so growing the context grows every extent. The 91%/96.9% pattern of §3.5 says the binding regime is absence of evidence rather than contradiction, which is the regime where the monotone-growth reading applies (Appendix A.4 states the caveat). The multi-agent axes then separate cleanly: memory coupling $\mathcal{M}_{-i}^{(K)}$ merges peer contexts at a rate set by K , which can only grow extents, while occupancy $\mathcal{O}_{-i}$ is not part of the context at all; the M↔C trade-off of Empirical Claim 1 is a competition between lattice growth (semantic) and resource contention (extra-semantic). The precise statement, its four-line argument, and its deliberate scope limits (graded scores, non-monotone updates, no structure claimed for M or C) are in Appendix A.4.

3 Single agent: stage diagnostics localize failure

3.1 System, tasks, and protocol

Memory stack. The system uses a graph substrate with a semantic retrieval interface on top: observations o_t are tokenized into symbolic scene tokens and semantic tags that update a memory of place records, transitions, episodic traces, and concept/relation records. Planner queries return candidates by a weighted log-score over required, preferred, and penalty tags. The retrieval machinery is reused across all tasks and architectures; full memory schema, update rules, and scoring formula are in Appendix C.

Tasks, baselines, protocol. Four single-agent task families are evaluated (water, safe-rest, goal-region on Empty-6x6 / FourRooms (Sutton et al., 1999), hazard recovery on LavaGapS7), 10 seeds $\times$ 8 episodes each, $\times$ 2 assist modes. Semantic stack (node, concept, state-conditioned variants) is compared against four baselines: random, graph-only, SPTM-like, learned BC. Reported: end-to-end success, steps, four operational stage rates, assist deltas, weakest stage per task and per method. All means are seed-aggregated; 95% bootstrap CIs use $B=5,000$ resamples (Efron, 1979); assist comparisons use paired Wilcoxon signed-rank tests (Wilcoxon, 1945) with Holm-Bonferroni correction (Holm, 1979) (Appendix F). For the two hardest tasks (FourRooms, LavaGapS7), four oracle regimes (*base*, *oracle R/M/C*) are run — each replaces exactly one stage while leaving the rest of the pipeline unchanged. BabyAI portability check in Appendix D. All single-agent experiments ran on a 10-core ARM CPU, 32 GB RAM (CPU-only); oracle sweep < 1 min wall-clock.

Table 1: Stage-factor consistency audit for the best semantic method in each task family. The product of nested conditional estimators matches semantic-path completion; overall success can be higher because some episodes succeed without traversing the full semantic retrieval path.

Task	Method	Overall	SP-comp	$\hat{P}(Q^*)$	$\hat{P}(R^* Q^*)$	$\hat{P}(M^* R^*)$	$\hat{P}(C^* M^*)$	Product
Water	<code>water-cp</code>	0.95	0.70	1.00	0.98	0.74	0.97	0.70
Rest	<code>rest-s2</code>	0.93	0.73	1.00	1.00	0.80	0.91	0.73
Goal region	<code>goal-n1</code>	0.54	0.00	1.00	0.54	0.01	0.00	0.00
Hazard recovery	<code>hazard-s7</code>	0.40	0.16	1.00	0.56	0.58	0.50	0.16

Table 2: Baseline comparison. The learned behavior-cloned baseline dominates the easiest resource tasks and hazard recovery, while the semantic stack remains strongest on the relation-heavy goal-region family and remains the only family with interpretable Q/R/M/C structure.

Task	Random	Graph-only	SPTM-like	Learned	Best semantic
Water	0.85	0.90	0.90	1.00	0.95
Rest	0.85	0.90	0.90	1.00	0.93
Goal region	0.41	0.49	0.52	0.50	0.54
Hazard recovery	0.08	0.39	0.00	0.74	0.40

3.2 Stage-estimator validation

Before the main benchmark, the factorization is validated. On a synthetic four-stage process with known probabilities the nested-estimator product tracks empirical semantic-path completion across sample sizes. On the real benchmark, the same product matches semantic-path completion rather than overall success (Table 1) — exactly the contract of Definition 1.

3.3 Success-only summary and what it hides

Water and rest are strong end-to-end tasks (0.95 and 0.93 success; Figure 1b). Goal-region reaches 0.54 success at the task-family level, but the bottleneck sits downstream of candidate selection (§3.4). Hazard recovery is the hardest retrieval-sensitive task at 0.40 success, with retrieval satisfaction as the weakest stage. Under success-only evaluation the two hard families look like the same kind of weakness; the stage view will show they are opposites. Per-task operational rates with 95% CIs are in Table 4 (Appendix C.3).

Baseline comparison. Table 2 compares the semantic stack against random, graph-only, SPTM-like, and learned behavior-cloned baselines. The learned baseline dominates the easiest resource tasks (1.00 on water and rest) and hazard recovery (0.74), while the semantic stack remains strongest on the relation-heavy goal-region family (0.54). The relevant contrast for this paper, however, is not the success column: the semantic stack is the only method family whose decisions expose interpretable Q/R/M/C structure. On the BC baseline the stage events are degenerate for the same structural reason as on the multi-agent learned baseline analysed in §4.5 — a policy without an explicit query/lock interface cannot emit separable stage events, so the diagnostic decomposition has nothing to attach to. End-to-end learning and stage diagnostics answer different questions; this paper is about the second.

3.4 Oracle-stage interventions localize the bottleneck

On goal-region the semantic variants solve Empty-6x6 (1.00 success) but collapse on FourRooms (0.075). For hazard recovery (LavaGapS7), base success 0.39 rises only to 0.43 under oracle controller, but jumps to 0.59 under oracle materialization and 0.60 under oracle retrieval; goal-region rejoin is downstream-limited (oracle R lifts semantic-path completion to 0.08 without moving end-to-end success). The per-task attribution as a stacked waterfall is in Figure 5 (Appendix C.3). Two tasks identical under success-only evaluation fail at opposite ends of the Q/R/M/C chain.




Figure 3: Four multi-agent memory architectures. (a) Independent: per-agent private memory, no exchange. (b) Shared bus: one global memory, every agent reads/writes each tick. (c) Central aggregator: per-agent memories plus a hub that merges them into a consensus report. (d) Peer-to-peer broadcast: per-agent memories with pairwise snapshot exchange every K ticks (the cadence axis of Empirical Claim 1).

3.5 The logs refine the diagnosis: a memory-accumulation limit

Inspecting the framework’s own logs sharpens the retrieval attribution. On the 10-seed hazard-recovery run, 349 retrieval calls were logged: only 32 (9.2%) returned at least one candidate, and of these 31 (96.9% precision) selected a semantically satisfied candidate. The bottleneck is therefore not candidate *ranking* but candidate *generation* — the symbolic memory does not contain hazard-recovery-tagged nodes at 91% of decision points. A retrained 2-layer candidate-generation MLP trained on per-state features (test accuracy 0.59 vs majority-class 0.81; Appendix G) confirms that instantaneous state features cannot substitute for the missing structure. The framework’s “retrieval-limited” attribution thus refines to a *memory-accumulation* limit: candidate generation cannot be solved from state features alone, and the binding constraint is how trajectory evidence is consolidated into persistent symbolic structure. This sets up the multi-agent section: when consolidation becomes a shared process across N agents, it acquires a measurable *rate*.

4 Multi-agent: memory as a shared resource

4.1 Four memory architectures and the cadence axis

Four communication architectures are instantiated; they share the single-agent retrieval substrate but differ in how memory is distributed across agents (Figure 3).

Independent. Each agent maintains a private event store and concept graph. No inter-agent communication mechanism exists. Lower-bound baseline.

Shared bus. All agents publish observations to a single *collective memory* event bus. One shared concept graph is rebuilt from the union of events. Maximally centralized.

Central aggregator (ConsensusLayer). Each agent has its own private concept graph; a central layer pulls snapshots and computes a trust-weighted merge, producing a single *consensus report* consumed by all agents.

Peer-to-peer broadcast (cadence K). Each agent has its own concept graph and its own merged view. Every K ticks each agent broadcasts a snapshot to a passive bus; each agent independently merges its own snapshot with the latest snapshot received from each peer. Different agents may hold different merged views; there is no central report.

The first three correspond to the three regimes in the multi-agent communication literature: no exchange, CTDE-style aggregation, and decentralized peer communication. The fourth is the architecture for which Definition 2 and Empirical Claim 1 are directly applicable, since the cadence K is a tunable parameter rather than a fixed protocol. Mechanically, K is the rate at which private memories consolidate into a collective view, which is why it is the structural axis of this section.

4.2 Multi-agent factorization

Extend the factorization to a population of N agents that share information through a communication protocol $\mathcal{C}_K$ parameterized by broadcast cadence K . Each agent i has its own memory state $\mathcal{M}_i$ and issues its own query q_i . Let $\mathcal{M}_i^{(K)}$ denote the projection of agent i 's memory at the moment it queries, after absorbing every peer message received under cadence K . Two coupling channels are separated: (i) memory coupling $\mathcal{M}_{-i}^{(K)}$, the latest peer memory snapshots accessible to agent i ; (ii) occupancy coupling $\mathcal{O}_{-i}$, the set of targets that other agents have already claimed by the time i acts.

Definition 2 (multi-agent stage estimability). For each agent i ,

$$P(C_i^* | q_i, \mathcal{M}_i, \mathcal{C}_K) = P(Q_i^* | q_i, \mathcal{M}_i) P(R_i^* | Q_i^*, q_i, \mathcal{M}_i, \mathcal{M}_{-i}^{(K)}) \\ \cdot P(M_i^* | R_i^*, q_i, \mathcal{M}_i, \mathcal{M}_{-i}^{(K)}) P(C_i^* | M_i^*, q_i, \mathcal{O}_{-i}).$$

Under a multi-agent stage-Markov assumption (Appendix B), the four factors are estimable from per-agent trajectory logs and the empirical-frequency estimators are consistent.

When Definition 2 may fail. The single-agent stage-Markov property absorbs trajectory history into $(q, \mathcal{M})$. In the multi-agent setting, the analogous absorption requires that dependence on past peer decisions enters only through $\mathcal{M}_{-i}^{(K)}$ at the M stage and through $\mathcal{O}_{-i}$ at the C stage. This step is non-trivial: $\mathcal{O}_{-i}$ summarises target occupancy at decision time but loses information about *how* peers arrived there. The factorization is therefore a projection onto observable per-agent traces under that summary. Two specific failure modes arise: (i) inter-agent communication side-channels not logged through $\mathcal{C}_K$; (ii) coordinated peer policies whose decision at the C stage depends on hidden peer state not summarised by $\mathcal{O}_{-i}$.

4.3 Bottleneck shift

Empirical Claim 1 (bottleneck shift $M \leftrightarrow C$, slope-level). Assume scarcity ($N > T$ targets, see §2.1) and a peer-broadcast architecture with cadence $K \in \mathbb{N}$. Two stylized assumptions are posited: **(F)** memory freshness $\nu(\mathcal{M}_{-i}^{(K)})$ is non-decreasing pointwise as K decreases, and the probability that agent i 's planner locks onto a reachable target is non-decreasing in ν ; **(O)** the probability that an arbitrary committed target is already in $\mathcal{O}_{-i}$ is non-increasing as K increases. Write $P(M^* | R^*; K)$ and $P(C^* | M^*; K)$ for the population-level conditional probabilities (averaged over agents and the seed distribution of the benchmark); $\hat{P}(\cdot)$ denotes the corresponding empirical-frequency estimator computed from logs. Under (F) and (O),

$$\frac{\partial P(M^* | R^*; K)}{\partial K^{-1}} > 0, \quad \frac{\partial P(C^* | M^*; K)}{\partial K^{-1}} < 0.$$

Both slope conditions are direct empirical predictions that are testable on logs through $\hat{P}$.

What this claim is and is not. Empirical Claim 1 is a *slope-level statement on conditional rates*: each of the two slopes is testable and is confirmed at $p < 10^{-4}$ on the 35,640-run sweep (§4.5) and independently on a 100-run standard-MiniGrid portability sweep ($p < 10^{-6}$, §6). It is *not* a claim that the conditional product $\Phi(K) := P(M^* | R^*; K) \cdot P(C^* | M^*; K)$ has an interior maximum; on our data Φ is monotone in K over $\{4, 8, 16, 32, 48, 64\}$ and converges to the independent-agent boundary $\Phi(K \rightarrow \infty) = 0.605$. The interior optimum that the $M \leftrightarrow C$ trade-off produces is on the practically relevant efficiency metric, mean time-to-success ($K=8$ on the main sweep), not on Φ . We make no formal claim about Φ .

Chronology. Empirical Claim 1 is post-hoc. It was formulated *after* an initial 12,960-run sweep conducted under four stage-agnostic hypotheses (Appendix E); the two slope conditions were then confirmed on the full 35,640-run sweep, on the $N=12$ scale-up (Appendix D), and on the MiniGrid portability wrapper (§6). We state this explicitly: the verified slope signs are predictions on enlarged data, not a pre-registered prediction. A fully formal version is left for future work.




Figure 4: Per-architecture conditional rates (top) and mean time-to-success (bottom) from the 35,640-run sweep. Shared bus, central aggregator, and fast peer broadcast ($K \leq 4$) live in the M-saturation / C-collapse regime; independent and slow peer ($K \geq 16$) live in the M-deficit / C-saturation regime. Mean time-to-success has a clean interior minimum at $K=8$ (7.47 ticks), below the independent-agent baseline (8.02).

4.4 Sweep protocol

The four architectures are evaluated on a custom multi-agent grid (12×10 cells); each agent must reach a cell tagged `water_source`, choosing actions via a shared memory-driven planner so only the memory layer differs. The sweep crosses $N \in \{3, 5, 8\}$, $T \leq N$, three layouts (symmetric/asymmetric/random), three hazard densities, four architectures, eight peer cadences $K \in \{1, 2, 4, 8, 16, 32, 48, 64\}$, 40 seeds — 35,640 runs in ≈ 22 min on 8 cores. Per-agent per-tick Q/R/M/C indicators are logged (Appendix B); episode-level starred events are the disjunction over ticks. A separate oracle-intervention block on scarcity cases $(N, T) \in \{(3, 2), (5, 3), (8, 5)\}$ at five cadences evaluates oracle R (memory returns ground-truth water positions) and oracle M (planner locks onto nearest real water).

4.5 Results: stage decomposition and bottleneck shift

The 35,640-run sweep is reported through four observations, structured to mirror Empirical Claim 1: (i) the conditional decomposition isolates the M and C stages as the differentiating links; (ii) the bottleneck shift appears directly as a negative within-cadence correlation between $P(M^* | R^*)$ and $P(C^* | M^*)$, peaking at $K=4$ and decaying to noise at $K=64$; (iii) on mean time-to-success, the resulting trade-off produces a clean interior optimum at $K=8$; (iv) on the conditional product Φ , the same trade-off produces a monotone trend rather than an interior peak — a summary statistic on which we make no formal claim. A learned-communication baseline and a multi-agent oracle-intervention loop close the subsection.

4.5.1 Conditional decomposition reveals where each architecture fails.

Figure 4 shows per-architecture conditional Q/R/M/C rates; full numerics with bootstrap CIs are in Appendix E.1, Table 7. Retrieval conditional $P(R^* | Q^*)$ is saturated; the differentiating signal lives at M and C. Shared-bus, central-aggregator, and fast-broadcast peer ($K \leq 4$) have high $P(M^* | R^*)$ but low $P(C^* | M^*)$; independent and slow-broadcast peer ($K \geq 16$) show the reverse — the $M \leftrightarrow C$ trade-off predicted by Empirical Claim 1.

4.5.2 Bottleneck shift and interior efficiency optimum.

Across all peer runs, Spearman of $P(M^*|R^*)$ vs K is -0.608 and of $P(C^*|M^*)$ vs K is $+0.420$ (both $p < 10^{-4}$), matching the slope conditions of Empirical Claim 1. The stronger empirical signature is at the per-configuration level: Pearson of $P(M^*|R^*)$ vs $P(C^*|M^*)$ within a fixed K is $-0.43, -0.54, -0.61, -0.19, -0.05, -0.04, -0.02, -0.02$ at $K=1, 2, 4, 8, 16, 32, 48, 64$, peaking at $K=4$ and decaying to noise by $K=64$ (Figure 8, Appendix E.1). Configurations where M is high systematically have low C : the bottleneck migrates between the two stages within the same $(N, T, \text{layout}, \text{hazard}, \text{seed})$ cell as K varies. On the efficiency metric, mean t_{succ} has a U-shape with interior minimum at $K=8$ (7.47 ticks vs 8.02 for the independent baseline; paired-bootstrap CIs in Appendix E.1). Scaling to $N=12$ sharpens the signature: the peak Pearson rises to -0.70 and migrates outward to $K=8$, consistent with stronger occupancy coupling at higher N delaying the cost-regime transition (Appendix D, Figure 6).

What we do not claim about Φ . The conditional product $\Phi = P(M^*|R^*) \cdot P(C^*|M^*)$ does not have an interior peak in the tested range (mean-of-products estimator increases monotonically from 0.569 at $K=4$ to 0.603 at $K=64$, converging to the independent-agent boundary 0.605). The slope-level statement of Empirical Claim 1 is on the conditional rates themselves, not on Φ ; we make no formal claim about the shape of $\Phi(K)$. Jensen-gap detail and POM-MOP comparison in Appendix E.1.

Learned-communication baseline: stage events are structurally degenerate even at competitive success. A CommNet-style policy (Sukhbaatar et al., 2016) trained with PPO (Schulman et al., 2017) on the same scenario ($N=3, T=2$) reaches 0.667 held-out success ($n=100$ seeds) — matching the symbolic peer architecture at $K=8$ (0.65). Despite this, the Q/R/M/C profile remains structurally degenerate: $Q^*=1.00$ saturated by construction; $R^*=0.997, M^*=0.997$ collapsed together ($P(M^*|R^*)=1.00$). This is the verified structural claim: a real-valued comm vector is never informatively empty (Q saturates) and lacks an explicit target-lock interface (R/M merge), regardless of training budget. The multi-agent counterpart of the single-agent learned-BC analysis (§3.3). Full architecture, PPO hyperparameters, and side-by-side REINFORCE vs PPO comparison in Appendix G.

Multi-agent oracle interventions close the diagnose→intervene→verify loop. On the scarcity cases, mean t_{succ} at small/intermediate cadence ($K \in \{1, 2, 4\}$) drops from ~ 10 –11 in baseline to ~ 7.1 under either oracle R or oracle M; at slow cadence ($K=16$) the baseline is already within ~ 0.8 of the oracle floor (Table 8 in Appendix E.4). The cadence-induced excess time at small/intermediate K is attributable to the $M \leftrightarrow C$ transition; at large K the bottleneck has moved into solo exploration.

The multi-agent block thus delivers its half of the argument. The cadence K is, mechanically, the rate at which private memories consolidate into a collective view; the data show that this rate has a non-trivial interior optimum ($K=8$ on t_{succ}), that its cost channel is occupancy coupling, and that both effects are visible stage-by-stage on the same Q/R/M/C chain used for single agents. Consolidation is not only the single-agent limit; in a collective it is a tunable, measurable design axis.

5 A minimal collective semantic memory

5.1 From diagnosis to design space

The single-agent diagnosis (§3.5) localised the hardest failure not to retrieval mechanics but to *consolidation*: candidate ranking is near-optimal (96.9%) but generation is starved (91% empty queries). The multi-agent diagnosis (§4.5) showed the same object acquires a measurable *rate* when consolidation is shared. The four architectures of §4.1 are degenerate corners of a richer design space: shared bus merges everything instantly (full occupancy cost), independent never merges (full staleness cost), central aggregator merges through a single trusted hub, peer broadcast exposes one parameter — the rate K . A richer architecture exposes three explicit rules: **merge** (how peer evidence about the same place is combined), **trust** (how conflicting or stale evidence is weighted), and **staleness** (how consolidated structure decays). We call any architecture that exposes all three rules a *collective semantic memory* (CSM) and instantiate a minimal one here.

5.2 Specification: a minimal CSM

The minimal CSM uses peer broadcast at $K=8$ (the best fixed cadence on the main sweep) and overlays three rules:

Merge. Per-place score: own evidence at weight 1, peer evidence at weight $\text{trust}_i \cdot e^{-\alpha \cdot \text{age}}$, $\alpha = 0.05$. Threshold $\tau = 0.30$ for inclusion in the candidate set.

Trust. Per-peer EMA on retrieval-success consistency: peers whose latest broadcast snapshot *supports* the locally locked target gain trust; others decay toward 0.4. Update rate $\beta = 0.10$.

Staleness. Each broadcast snapshot is timestamped; the merge weight on its evidence decays as $e^{-\alpha \cdot \text{age}}$. Snapshots older than $10K$ ticks are pruned.

Code: `experiments/collective_semantic_memory/csm_memory.py` (78 lines). The CSM is a drop-in replacement on the same Q/R/M/C logger used for the four baseline architectures.

5.3 CSM strictly dominates fixed-cadence peer on success rate

We compare CSM against fixed-cadence peer broadcast at $K \in \{1, 4, 8, 16, 64\}$ on the full multi-agent scarcity scenario ($N \in \{3, 5, 8\}$ agents, $T \in \{2, 3, 5\}$ targets, $N > T$; three layouts, three hazard densities, 20 seeds): 540 runs per architecture, 3,240 total in ≈ 5 minutes on a single CPU core. Table 3 reports stage-conditional rates and end-to-end success.

Table 3: Minimal collective semantic memory vs fixed-cadence peer broadcast on the scarcity scenario (3,240 runs total, 20 seeds $\times$ 3 layouts $\times$ 3 hazard densities $\times$ 3 (N, T) scarcity cells, ≈ 5 min wall-clock). **CSM strictly dominates ALL five fixed-cadence peer architectures on end-to-end success rate:** CSM’s lower 95% bootstrap CI bound (0.610) exceeds every peer- K ’s upper CI bound. In the conditional decomposition, CSM occupies a new point on the ($M \times C$) plane that no fixed cadence reaches.

Architecture	$P(M^* R^*)$	$P(C^* M^*)$	Success rate (95% CI)
peer ($K=1$)	0.991	0.586	0.577 [0.572, 0.586]
peer ($K=4$)	0.955	0.623	0.581 [0.572, 0.590]
peer ($K=8$)	0.730	0.871	0.599 [0.591, 0.607]
peer ($K=16$)	0.697	0.892	0.597 [0.589, 0.606]
peer ($K=64$)	0.684	0.900	0.601 [0.594, 0.609]
CSM (minimal)	0.782	0.825	0.617 [0.610, 0.624]

The numbers tell a consistent story. On the conditional decomposition, no fixed-cadence peer reaches the (0.78, 0.83) point in $(P(M^*|R^*), P(C^*|M^*))$ space: fast broadcasts ($K \leq 4$) sit in the M-saturation / C-collapse regime ($P(M^*|R^*) > 0.95$, $P(C^*|M^*) < 0.65$); slow broadcasts ($K \geq 16$) sit in the M-deficit / C-saturation regime ($P(M^*|R^*) < 0.70$, $P(C^*|M^*) > 0.89$). CSM occupies a third regime, with M and C both above 0.78. The framework’s measurement protocol detected the improvement *not* on either link in isolation but on a new point in the $M \times C$ plane, with the consequence that end-to-end success rate is strictly higher (95% bootstrap CIs non-overlapping vs every peer- K ; $\Delta_{\text{success}} \in [+0.017, +0.040]$, all CIs exclude zero).

5.4 What the CSM demonstrates and what it does not

This is a *minimal* prototype: three simple rules on top of the same $K=8$ peer-broadcast backbone, no architectural novelty beyond the explicit trust/staleness/merge contract. The point of presenting it is twofold. First, the title’s claim of “a minimal collective semantic memory” is delivered as an instantiated artefact measured on the same Q/R/M/C scale as the four baselines, not as a programmatic claim. Second, the framework’s measurement protocol is shown to be *usable* for new-architecture evaluation: the CSM’s contribution would have been invisible under success-only reporting (a 1.7%–4.0% absolute gain looks marginal) but stands out clearly on the conditional decomposition that the rest of the paper validated. Richer CSM variants — adaptive cadence K_t , selective consolidation of under-represented semantic tags, intent-weighted trust — are concrete next steps listed in Appendix E.2.

6 Limitations and future work

Portability across substrates: slope conditions hold on MiniGrid. The slope conditions of Empirical Claim 1 also hold on a standard-MiniGrid substrate. A FourRooms-based multi-agent wrapper (built from `minigrid.core.grid.Grid` + `Wall/Floor/Lava` primitives, identical operational Q^* , R^* , M^* , C^* definitions; `experiments/minigrid_multiagent_wrapper/`) was used to run a 100-run portability sweep ($K \in \{1, 4, 8, 16, 64\}$, 20 seeds, $N=3$, $T=2$): Spearman $P(M^*|R^*)$ vs K is -0.50 ($p=1.5 \times 10^{-7}$), Spearman $P(C^*|M^*)$ vs K is $+0.50$ ($p=1.3 \times 10^{-7}$). The interior t_{succ} minimum does not separate at this minimal sweep size (5 cadences, 20 seeds, one hazard density); a denser sweep is the natural next step. The two-substrate seam is therefore closed at the slope-level claim.

Empirical Claim 1 is slope-level on conditional rates only. The two slope conditions hold ($p < 10^{-4}$); we make no formal claim on the shape of the conditional product Φ (§4.3). A future statement directly in terms of the time-cost surface, or a derivation under a generative model (Poisson-renewal memory, birth-death occupancy), would convert the empirical claim into a theorem.

Estimability under hidden state. With unlogged in-episode query rewrites or peer side-channels, Definitions 1 and 2 are projections onto logged traces, not full identification.

Learned baselines. The CommNet PPO baseline reaches 0.667 success on the scarcity scenario — competitive with symbolic peer at $K=8$ (0.65) — and confirms Q-saturation ($Q^*=1.00$) and R/M collapse ($P(M^*|R^*)=1.00$) at competitive success. The structural prediction now holds verified, not asserted; ablating on architectures with explicit symbolic target-lock channels is left for future work.

Scale. Multi-agent sweep bounded at $N=12$ (Appendix D; signature sharpens, peak Pearson -0.70 at $K=8$, interior t_{succ} minimum 4.77). Denser scaling at $N \geq 16$ and continuous-state environments (e.g., VMAS (Bettini et al., 2022)) are future work.

7 Broader impacts

The paper proposes diagnostic evaluation for memory-based agents in simulated navigation. Its positive impact is methodological: making memory-based failures more measurable across architectures, and making the cadence axis of distributed multi-agent memory tunable on an empirically grounded criterion. Potential negative impacts are indirect and would arise if similar diagnostics were used to improve agents in safety-critical settings (drone swarms, autonomous vehicles, distributed sensor networks) without adequate safeguards. The current work is benchmark-level infrastructure, not a deployment-ready system.

8 Conclusion

We proposed Q/R/M/C, a four-stage decomposition of memory-based navigation whose factors are estimable from standard trajectory logs, and used it to run a single continuous argument: stage diagnostics localize failures that success-only evaluation cannot (single agent); when memory is shared, the exchange cadence becomes a measurable structural axis with an interior optimum (multi-agent); and both measurements localize the binding constraint to the same mechanism: the consolidation of trajectory evidence into persistent, shareable symbolic structure. Collective semantic memory, a shared symbolic structure with explicit merge, trust, and staleness rules, is therefore not a speculative extension but the research target these results point to, and the framework presented here is the instrument by which candidate designs for it can be evaluated.

References

- M. Andrychowicz, F. Wolski, A. Ray, J. Schneider, R. Fong, P. Welinder, B. McGrew, J. Tobin, P. Abbeel, and W. Zaremba. Hindsight experience replay. In *NeurIPS*, 2017.
- R. Bělohlávek. Fuzzy Galois connections. *Mathematical Logic Quarterly*, 45(4):497–504, 1999.
- M. Bettini, R. Kortvelesy, J. Blumenkamp, and A. Prorok. VMAS: A vectorized multi-agent simulator for collective robot learning. In *16th International Symposium on Distributed Autonomous Robotic Systems (DARS)*, 2022. arXiv:2207.03530.

518 S. Boyd, A. Ghosh, B. Prabhakar, and D. Shah. Randomized gossip algorithms. *IEEE Trans. Information Theory*,
519 52(6):2508–2530, 2006.

520 M. Chevalier-Boisvert, D. Bhoopchand, H. Abdolmaleki, and A. Mordatch. BabyAI: A platform to study the
521 sample efficiency of grounded language learning. In *ICLR*, 2019.

522 M. Chevalier-Boisvert, B. Lahlou, S. Willems, and P. S. Castro. MiniGrid and MiniWorld: Modular reinforcement
523 learning environments. *arXiv:2306.13831*, 2023.

524 P. Dayan. Improving generalization for temporal difference learning: The successor representation. *Neural
525 Computation*, 5(4):613–624, 1993.

526 M. Deitke, E. V. Wijmans, D. Schwenk, O. Salvador, K. M. Ehsani, et al. ProcTHOR: Large-scale embodied AI
527 using procedural generation. In *NeurIPS Datasets and Benchmarks*, 2022.

528 B. Efron. Bootstrap methods: Another look at the jackknife. *Annals of Statistics*, 7(1):1–26, 1979.

529 J. N. Foerster, Y. M. Assael, N. de Freitas, and S. Whiteson. Learning to communicate with deep multi-agent
530 reinforcement learning. In *NeurIPS*, 2016.

531 J. N. Foerster, G. Farquhar, T. Afouras, N. Nardelli, and S. Whiteson. Counterfactual multi-agent policy gradients.
532 In *AAAI*, 2018.

533 B. Ganter and R. Wille. *Formal Concept Analysis: Mathematical Foundations*. Springer, 1999.

534 D. Ghosh, A. Gupta, A. Reddy, J. Fu, S. Levine, and C. Finn. Learning to reach goals via iterated supervised
535 learning. In *ICLR*, 2021.

536 G. Gigerenzer and D. G. Goldstein. Reasoning the fast and frugal way: Models of bounded rationality.
537 *Psychological Review*, 103(4):650–669, 1996.

538 S. Gupta, J. Davidson, S. Levine, R. Sukthankar, and J. Malik. Cognitive mapping and planning for visual
539 navigation. In *CVPR*, 2017.

540 S. Holm. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70,
541 1979.

542 E. Kolve, R. Mottaghi, D. Gordon, Y. Zhu, A. Gupta, A. Farhadi, and D. Fox. AI2-THOR: An interactive 3D
543 environment for visual AI. *arXiv:1712.05474*, 2017.

544 B. Kuipers. The spatial semantic hierarchy. *Artificial Intelligence*, 119(1–2):191–233, 2000.

545 F. Lieder and T. L. Griffiths. Resource-rational analysis: Understanding human cognition as the optimal use of
546 limited computational resources. *Behavioral and Brain Sciences*, 43:e1, 2020.

547 R. Lowe, Y. Wu, A. Tamar, J. Harb, P. Abbeel, and I. Mordatch. Multi-agent actor-critic for mixed cooperative-
548 competitive environments. In *NeurIPS*, 2017.

549 S. Mac Lane. *Categories for the Working Mathematician*. Springer, 2nd edition, 1998.

550 I. Momennejad, E. M. Russek, J. H. Cheong, M. M. Botvinick, N. Daw, and S. J. Gershman. The successor
551 representation in human reinforcement learning. *Nature Human Behaviour*, 1:680–692, 2017.

552 J. O’Keefe and L. Nadel. *The Hippocampus as a Cognitive Map*. Oxford University Press, 1978.

553 C. Packer, V. Fang, S. Patil, K. Lin, S. Wooders, and J. Gonzalez. MemGPT: Towards LLMs as operating
554 systems. *arXiv:2310.08560*, 2023.

555 N. Savinov, A. Dosovitskiy, and V. Koltun. Semi-parametric topological memory for navigation. In *ICLR*, 2018.

556 M. Savva, A. Kadian, O. Maksymets, Y. Zhao, E. Wijmans, B. Jain, J. Straub, J. Liu, V. Koltun, J. Malik, et al.
557 Habitat: A platform for embodied AI research. In *ICCV*, 2019.

558 J. Schulman, P. Moritz, S. Levine, M. Jordan, and P. Abbeel. High-dimensional continuous control using
559 generalized advantage estimation. In *ICLR*, 2016. *arXiv:1506.02438*.

560 J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov. Proximal policy optimization algorithms.
561 *arXiv:1707.06347*, 2017.

562 N. Shinn, F. Cassano, A. Berman, and A. G. Labash. Reflexion: Language agents with verbal reinforcement
563 learning. In *NeurIPS*, 2023.

564 H. A. Simon. A behavioral model of rational choice. *Quarterly Journal of Economics*, 69(1):99–118, 1955.

565 H. A. Simon. Rational choice and the structure of the environment. *Psychological Review*, 63(2):129–138, 1956.

566 K. L. Stachenfeld, M. M. Botvinick, and S. J. Gershman. The hippocampus as a predictive map. *Nature*
567 *Neuroscience*, 20(11):1643–1653, 2017.

568 S. Sukhbaatar, A. Szlam, and R. Fergus. Learning multiagent communication with backpropagation. In *NeurIPS*,
569 2016.

570 R. S. Sutton, D. Precup, and S. Singh. Between MDPs and semi-MDPs: A framework for temporal abstraction
571 in reinforcement learning. *Artificial Intelligence*, 112(1–2):181–211, 1999.

572 E. C. Tolman. Cognitive maps in rats and men. *Psychological Review*, 55(4):189–208, 1948.

573 G. Wang, Y. Xie, Z. Jiang, A. Mandlekar, C. Xie, A. Yu, and P. Liang. Voyager: An open-ended embodied agent
574 with large language models. *arXiv:2305.16291*, 2023.

575 F. Wilcoxon. Individual comparisons by ranking methods. *Biometrics Bulletin*, 1(6):80–83, 1945.

576 R. J. Williams. Simple statistical gradient-following algorithms for connectionist reinforcement learning.
577 *Machine Learning*, 8(3–4):229–256, 1992.

578 L. Xiao, S. Boyd, and S. Lall. A scheme for robust distributed sensor fusion based on average consensus. *IPSN*,
579 2007.

580 A Formal proofs

581 A.1 Argument for Definition 1 (single-agent estimability)

582 For each episode $i \in \{1, \dots, n\}$ let $Q_i^*, R_i^*, M_i^*, C_i^*$ denote the nested episode-level indicators. By
583 construction these are measurable functions of the runtime trajectory log and satisfy $C_i^* \Rightarrow M_i^* \Rightarrow$
584 $R_i^* \Rightarrow Q_i^*$. The empirical estimators

$$\hat{p}_Q = \frac{1}{n} \sum_i Q_i^*, \quad \hat{p}_{R|Q} = \frac{\sum_i R_i^*}{\sum_i Q_i^*}, \quad \hat{p}_{M|R} = \frac{\sum_i M_i^*}{\sum_i R_i^*}, \quad \hat{p}_{C|M} = \frac{\sum_i C_i^*}{\sum_i M_i^*}$$

585 are consistent by the chain rule plus the conditional stage-Markov property, and the empirical product
586 converges to semantic-path completion. $\square$

587 A.2 Argument for Definition 2 (multi-agent estimability)

588 Conditioning sets are enlarged by $\mathcal{M}_{-i}^{(K)}$ and $\mathcal{O}_{-i}$. Under the enlarged stage-Markov property, the
589 proof of Definition 1 transfers verbatim and each enlarged conditional is a Bernoulli probability on
590 a measurable event. The qualifications stated in §4.2 (“When Definition 2 may fail”) describe two
591 specific situations in which the enlarged stage-Markov property is violated. $\square$

592 A.3 Argument for Empirical Claim 1 (slope-level statement)

593 Empirical Claim 1 is an empirical claim, not a theorem. The argument below is the qualitative
594 justification that motivates the two slope signs; the data are what verify them. We deliberately do not
595 call this a proof.

596 *Slope (a), under (F).* As K decreases, $\nu(\mathcal{M}_{-i}^{(K)})$ is non-decreasing pointwise. Conditional on
597 a successful retrieval, the probability of locking onto a reachable target is non-decreasing in ν .
598 Therefore $\partial P(M^*|R^*; K)/\partial K^{-1} > 0$.

599 *Slope (b), under (O).* As K decreases, peers receive updated memory simultaneously and con-
600 verge on the same closest target; the marginal occupancy $\Pr[\text{target} \in \mathcal{O}_{-i}]$ rises. Therefore
601 $\partial P(C^*|M^*; K)/\partial K^{-1} < 0$.

602 *Note on the conditional product Φ .* Empirical Claim 1 makes no statement about the shape of
603 $\Phi(K) = P(M^*|R^*; K) \cdot P(C^*|M^*; K)$. Under (F) and (O) the slope of Φ in K^{-1} is the sum of two
604 opposing terms, and whether it changes sign depends on quantitative magnitudes that are not pinned

down by (F) and (O) alone. In our data Φ is monotone in K over $\{4, 8, 16, 32, 48, 64\}$ and converges to the independent-agent boundary $\Phi_\infty = 0.605$. The practically relevant efficiency surface — mean time-to-success — does have an interior minimum at $K = 8$; see §4.5. $\square$

A.4 The query–retrieval Galois connection: statement, argument, scope

Fact (standard). Let G and Σ be sets and $I \subseteq G \times \Sigma$ a relation. For $B \subseteq \Sigma$ define $\text{ext}(B) = \{p \in G : \forall \tau \in B, (p, \tau) \in I\}$, and for $A \subseteq G$ define $\text{int}(A) = \{\tau \in \Sigma : \forall p \in A, (p, \tau) \in I\}$. Then for all $A \subseteq G, B \subseteq \Sigma$:

$$A \subseteq \text{ext}(B) \iff B \subseteq \text{int}(A);$$

both maps are antitone, both composites $\text{ext} \circ \text{int}$ and $\text{int} \circ \text{ext}$ are closure operators, and the closed pairs (A, B) with $A = \text{ext}(B)$ and $B = \text{int}(A)$, ordered by extent inclusion, form a complete lattice — the concept lattice of the context (G, Σ, I) (Ganter and Wille, 1999).

Argument. $A \subseteq \text{ext}(B)$ unfolds to $\forall p \in A \forall \tau \in B : (p, \tau) \in I$, a condition symmetric in A and B , which is exactly $B \subseteq \text{int}(A)$. Antitonicity and the closure properties follow by the standard one-line arguments (Ganter and Wille, 1999, Ch. 1). Viewing the power sets as thin categories (one with reversed order), the displayed equivalence says precisely that int and ext form an adjoint pair; a Galois connection is an adjunction between posets (Mac Lane, 1998, Ch. IV). $\square$

Instantiation and scope. In our memory, take G to be the place records, Σ the tag alphabet, and $I_\theta = \{(p, \tau) : \text{the per-place tag table assigns } \tau \text{ to } p \text{ with confidence } \geq \theta\}$ (Appendix C); the intent of a query is $B = q^{\text{req}}$ and R-availability is $\text{ext}(B) \neq \emptyset$. Three deliberate scope restrictions. (i) *Graded scores.* The deployed retrieval score is graded (the weighted log-score over required, preferred, and penalty tags of Appendix C), so the crisp connection above formalizes the thresholded required-tag skeleton of candidacy; the graded analogue is the theory of fuzzy Galois connections (Bělohlávek, 1999) and we do not use it here. (ii) *Non-monotone updates.* Dirichlet–multinomial updates are not literally monotone in I_θ : contrary evidence can remove an incidence. On the reported benchmarks the binding failure is absence of evidence rather than contradiction (91% empty extents, 96.9% precision when non-empty, §3.5), so the monotone-growth reading is the empirically relevant regime, and we flag the exception rather than assume it away. (iii) *No structure for M or C .* Materialization is a satisficing choice function on non-empty extents and Completion is a control problem under occupancy; neither carries the order structure, and we attach no categorical claims to them. The reading earns its place by naming the failure mode exactly (a queried intent whose extent is empty lies outside the realized concept lattice), by giving consolidation a precise object (growth of the context, hence of every extent), and by locating the cadence axis structurally: memory coupling $\mathcal{M}_{-i}^{(K)}$ merges peer contexts at rate K^{-1} , while occupancy $\mathcal{O}_{-i}$ is extra-contextual — which is the structural content behind the two slope signs of Empirical Claim 1.

B Operational vs. nested stage estimators

Operational diagnostic rates. The main benchmark logs per-query and per-episode diagnostic quantities directly from runtime traces (query formed, retrieval satisfied, target materialized, completion after retrieval). These are operational rates, not direct multiplicative factors.

Nested stage estimators for factorization. For the factorization, episode-level nested events from the first failure stage recorded in the runtime taxonomy are used: Q^*, R^*, M^*, C^* , nested by construction. The consistency audit in §3.2 reports $\hat{P}(Q^*), \hat{P}(R^*|Q^*), \hat{P}(M^*|R^*), \hat{P}(C^*|M^*)$, their product, and its gap to semantic-path completion.

B.1 Multi-agent operational definitions

For the multi-agent sweep, every metric is defined per agent and per tick, then aggregated. Let $\text{Targets}_i(t)$ be the candidate set returned by the memory query of agent i at tick t , $\text{Locked}_i(t)$ be the agent’s planner-internal locked target, and $\mathcal{W}$ the ground-truth target set.

Per-tick events. $Q_i(t) = \mathbf{1}[\text{Targets}_i(t) \neq \emptyset]$; $R_i(t) = \mathbf{1}[\exists w \in \mathcal{W} : \min_c \|c - w\| \leq \epsilon]$ ($\epsilon = 0.6$); $M_i(t) = \mathbf{1}[\text{Locked}_i(t) \neq \emptyset \wedge \min_w \|\text{Locked}_i(t) - w\| \leq \epsilon]$.

652 **Episode-level starred events.** $Q_i^* = \bigvee_t Q_i(t)$; analogously for R^*, M^* ; $C_i^* =$
653 $1[\text{agent } i \text{ reached a target}]$.

654 **Aggregation.** Per-run rates are averaged over the N agents; per-configuration rates are averaged
655 over the 40 seeds. Conditional rates $P(R^*|Q^*), P(M^*|R^*), P(C^*|M^*)$ are computed as ratios of
656 episode-level counts within each configuration before averaging across configurations.

657 C Memory and task schema

658 C.1 Full notation table

659 Symbols are fixed throughout the paper.

660 **Stages.** Q, R, M, C — the four stages: **Q**uery formation, **R**etrieval, **M**aterialization, **C**ompletion
661 after retrieval. They name both the stage and (in expressions like “the M link”) the corresponding
662 edge of the chain.

663 **Events.** $Q^*, R^*, M^*, C^* \in \{0, 1\}$ — episode-level starred events recording whether each stage
664 succeeded (Appendix B). Nested by construction: $Q^*=1 \Leftarrow R^*=1 \Leftarrow M^*=1 \Leftarrow C^*=1$.

665 **Probabilities.** $P(\cdot)$ — population-level probability (averaged over the seed distribution and, in the
666 multi-agent case, over agents). $\hat{P}(\cdot)$ — empirical-frequency estimator computed from logs.

667 **Single-agent.** q — query; $\mathcal{M}$ — memory state; Y — end-to-end task success (can exceed C^*).

668 **Multi-agent.** N — number of agents in the population. T — number of targets in the environment
669 (constraint $T \leq N$; “scarcity” means $N > T$). i — agent index. $\mathcal{M}_i$ — agent i ’s private memory. q_i
670 — agent i ’s query.

671 **Coupling.** $\mathcal{C}_K$ — the peer-broadcast communication protocol with cadence $K \in \mathbb{N}$ (every K
672 ticks each agent broadcasts a memory snapshot, §4.1). $\mathcal{M}_{-i}^{(K)}$ — the latest peer memory snapshots
673 accessible to agent i under $\mathcal{C}_K$ (the *memory coupling* channel). $\mathcal{O}_{-i}$ — the set of targets already
674 claimed by other agents by the time i acts (the *occupancy coupling* channel).

675 **Slope variables.** $\nu(\mathcal{M}_{-i}^{(K)})$ — a freshness functional on peer memory; non-decreasing pointwise as
676 K decreases (assumption **(F)** of Empirical Claim 1). The occupancy slope is assumption **(O)**.

677 **Summary statistics.** $\Phi(K) := P(M^*|R^*; K) \cdot P(C^*|M^*; K)$ — the conditional product, the chain
678 probability $P(M^* \wedge C^* | R^*)$. POM (product-of-means) and MOP (mean-of-products) are two
679 estimators of Φ ; they differ by $\text{Cov}(P(M^*|R^*), P(C^*|M^*))$, which is the Jensen gap (§4.5). t_{succ}
680 — mean time-to-success in environment ticks.

681 C.2 Memory stack and task patterns

682 The memory stack stores symbolic place records (per-place tag confidences, counts, visitation,
683 freshness, geometry), transition records (success, risk, cost statistics with fast/slow estimates),
684 episodic traces (intent + state at each visit), and concept/relation records. At each step, the tokenizer
685 emits (z_t, τ_t, e_t) from o_t . Per-place tag tables use a Dirichlet-multinomial update; transitions use
686 exponentially smoothed statistics; episodic records attach the active intent and agent state. A planner
687 query is answered by retrieving candidate places whose semantic profiles match the current task intent
688 and state, via a weighted log-score $s(p | q_t, s_t) = \sum_{\tau \in q^{\text{req}}} \log P(\tau|p) + \alpha \sum_{\tau \in q^{\text{pref}}} \log P(\tau|p) -$
689 $\beta \sum_{\tau \in q^{\text{pen}}} \log P(\tau|p) + \gamma u_{\text{epi}}(p, s_t)$ over required, preferred, and penalty tags. The four task
690 families are defined by semantic place patterns: water search (target {water}), safe rest ({rest} with
691 state-conditioning), goal-region rejoin ({goal-region}), hazard recovery ({post-hazard-exit}).

692 C.3 Single-agent supporting tables and figures

Table 4: Best semantic result per task family from the final 10-seed benchmark, with 95% bootstrap CIs. Hazard-recovery is bottlenecked by retrieval; goal-region remains downstream-bottlenecked.

Task	Best method	Assist	Success (95% CI)	Steps	Query (op.)	Retrieval (op.)	Mat. (op.)	Post (op.)
Water	water-cp	on	0.95 [0.88, 1.00]	17.99	0.71	0.71	0.71	0.70
Rest	rest-s2	on	0.93 [0.85, 0.99]	23.35	0.61	0.61	0.61	0.73
Goal region	goal-n1	off	0.54 [0.52, 0.56]	49.34	0.00	0.00	0.00	0.00
Hazard recovery	hazard-s7	on	0.40 [0.35, 0.45]	10.18	0.11	0.11	0.11	0.16




Figure 5: Oracle stage interventions on the two hardest single-agent task families (10 seeds $\times$ 8 episodes). **(a)** FourRooms goal-rejoin: replacing any single stage with an oracle leaves end-to-end success at 0.07; the bottleneck is downstream of the stages in the chain. **(b)** LavaGapS7 hazard recovery: oracle controller adds only +0.04, but oracle materialization adds +0.16 and oracle retrieval adds a further +0.01 to reach 0.60. The framework localises the bottleneck stage-by-stage.

693 D Portability and scale-up

694 D.1 BabyAI portability transfer (10 seeds, full benchmark)

695 A 10-seed portability benchmark on BabyAI-GoToObjMaze-v0 (Chevalier-Boisvert et al., 2018)
696 (8 episodes per seed, max-steps 200) measured all four single-agent methods used in the main
697 benchmark (random, greedy, babyai_semantic_node_v1, babyai_semantic_plan_v1) with
698 the same Q/R/M/C logger applied to BabyAI traces. The numbers in Table 5 confirm that (a) the
699 Q/R/M/C events emit non-trivially on the BabyAI substrate, (b) the two semantic variants expose
700 stage-separable structure where the non-semantic baselines have $Q^*=R^*=M^*=0$ (no semantic
701 queries are issued), and (c) the stage-rate ordering matches the MiniGrid benchmark (Table 1) within
702 bootstrap CIs. GoToObjMaze is a harder task substrate than the LavaGapS7/FourRooms benchmark
703 used in the main text — end-to-end success caps at 0.15 across all methods — which makes the stage
704 decomposition especially informative: the semantic stack is not winning on success but it is the only
705 method family producing measurable per-stage events.

Table 5: BabyAI portability benchmark: 10 seeds $\times$ 8 episodes on BabyAI-GoToObjMaze-v0, max-steps 200. The semantic-stack variants succeed at the same end-to-end rate as the greedy baseline (0.15) but are the only methods that produce non-zero Q/R/M/C events — the framework’s protocol transfers to the BabyAI substrate without modification. Q/R/M/C operational rates here are evaluated on the same denominator (number of decision points), so $Q=R=M$ for both semantic variants because GoToObjMaze emits one semantic target type per episode.

Method	Success	Steps	Q (op.)	R (op.)	M (op.)	Post (op.)
random	0.125 ± 0.06	182.7	0.000	0.000	0.000	0.000
greedy	0.150 ± 0.05	171.3	0.000	0.000	0.000	0.000
semantic node	0.150 ± 0.05	171.3	0.032	0.032	0.032	0.025
semantic plan	0.150 ± 0.05	171.3	0.032	0.032	0.032	0.025

Takeaway. The framework’s measurement protocol — not its task-specific results — transfers to a second substrate. End-to-end success on BabyAI is dominated by the much lower episode-level success ceiling of GoToObjMaze, not by the stage structure; conversely, the stage structure that the semantic stack exposes is identical in shape to what the same stack exposes on MiniGrid. This is exactly the substrate-independence property the methodology is intended to provide. Raw data: tmp/babyai_deeper_10seed/.

D.2 Multi-agent scale-up at $N = 12$

To check that the multi-agent results of §4.5 are not an artefact of the $N \leq 8$ range, an additional 8,640-run sweep was run at $N = 12$, $T \in \{6, 8, 10\}$, the same three layouts, three hazard densities, eight peer cadences, and 40 seeds. Wall-clock ≈ 12 minutes on 8 cores. The conditional-rate signature of §4.5 persists and intensifies (Table 6).




Figure 6: $N = 12$ scale-up bottleneck signature ($n = 8,640$ runs, 40 seeds). Top: both conditional rates change monotonically in K as predicted; the within-cadence Pearson correlation (red dashed, right axis) peaks at -0.70 at $K = 8$. Bottom: mean time-to-success retains an interior minimum at $K = 8$ (4.77 ticks). The bottleneck-shift peak migrates outward from $K = 4$ at $N \leq 8$ to $K = 8$ at $N = 12$, consistent with stronger occupancy coupling at higher agent count delaying the cost-regime transition.

E Extended diagnostics

E.1 Conditional-product detail and the Jensen gap

The conditional product $\Phi = P(M^* | R^*) \cdot P(C^* | M^*)$ in the 35,640-run sweep is monotonically increasing in K over $K \in \{4, 8, 16, 32, 48, 64\}$ and converges to the independent-agent boundary $\Phi(K \rightarrow \infty) = 0.605$. Paired bootstrap: $\Delta(K=32-K=16) = +0.008$ (CI $[+0.006, +0.010]$), $\Delta(K=48-K=16) = +0.011$, $\Delta(K=64-K=16) = +0.012$; independent is significantly above every peer K except $K = 64$ (borderline).

Table 6: $N = 12$ scale-up sweep numerical detail (peer architecture). Slope conditions hold: Spearman of $P(M^*|R^*)$ with K is -0.141 ($p < 10^{-39}$), and of $P(C^*|M^*)$ with K is $+0.304$ ($p < 10^{-184}$).

K	$P(M^* R^*)$	$P(C^* M^*)$	MOP	mean t_{succ}	corr
1	0.758	0.679	0.475	5.49	-0.59
2	0.745	0.692	0.474	5.29	-0.61
4	0.679	0.768	0.482	4.82	-0.64
8	0.731	0.767	0.539	4.77	-0.70
16	0.687	0.803	0.538	5.04	-0.62
32	0.676	0.814	0.538	5.75	-0.51
48	0.677	0.814	0.539	5.96	-0.51
64	0.677	0.816	0.541	6.12	-0.51

Table 7: Per-architecture conditional Q/R/M/C rates from the 35,640-run sweep (40 seeds). $P(R^*|Q^*)$ is saturated. POM = product-of-means; MOP = mean-of-products (the proper per-configuration expected product). The two summaries disagree systematically at small/intermediate K because the within-cadence correlation between $P(M^*|R^*)$ and $P(C^*|M^*)$ is strongly negative there (last column) — the bottleneck-shift signature. **Bold:** MOP-best of the peer cadences. Mean t_{succ} (last column) is the practically relevant efficiency measure; its interior minimum at $K = 8$ is the predicted-shape result that survives the data.

Architecture	$P(R^* Q^*)$	$P(M^* R^*)$	$P(C^* M^*)$	POM	MOP (95% CI)	corr	mean t_{succ}
independent	0.99	0.68	0.88	0.598	0.605 [0.597, 0.614]	+0.00	8.02
shared bus	1.00	1.00	0.60	0.600	0.595 [0.586, 0.604]	-0.10	8.33
central aggregator	1.00	0.97	0.61	0.592	0.572 [0.564, 0.580]	-0.48	7.64
peer ($K = 1$)	1.00	0.97	0.60	0.585	0.573 [0.564, 0.581]	-0.43	7.80
peer ($K = 2$)	1.00	0.97	0.62	0.594	0.576 [0.568, 0.584]	-0.54	7.59
peer ($K = 4$)	1.00	0.92	0.66	0.599	0.569 [0.561, 0.577]	-0.61	8.00
peer ($K = 8$)	0.99	0.69	0.87	0.597	0.586 [0.578, 0.595]	-0.19	7.47
peer ($K = 16$)	0.99	0.67	0.89	0.592	0.590 [0.581, 0.598]	-0.05	7.87
peer ($K = 32$)	0.99	0.67	0.89	0.595	0.598 [0.589, 0.606]	-0.04	8.28
peer ($K = 48$)	0.99	0.68	0.88	0.598	0.602 [0.593, 0.611]	-0.02	8.79
peer ($K = 64$)	0.99	0.68	0.88	0.599	0.603 [0.594, 0.611]	-0.02	9.19

The negative within-cadence covariance documented in §4.5 creates the Jensen-inequality gap between Φ and the marginal product-of-means: $E[X] \cdot E[Y] > E[X \cdot Y]$ when $\text{Cov}(X, Y) < 0$. The mean-of-products is reported as the primary summary because it absorbs the covariance correctly. Figure 7 shows the K-sweep view of both conditional rates and mean t_{succ} ; Figure 8 shows the per-configuration scatter that the negative covariance summarises.

E.2 Open problems for the collective-semantic-memory program

Four concrete questions follow directly from the measurements of §4.5 and the synthesis of §5. **(1) Adaptive cadence.** Empirical Claim 1 assumes a fixed K ; the bottleneck-shift signature suggests an online controller K_t that broadcasts when the (estimated) M-side gain exceeds the C-side occupancy cost. **(2) Selective consolidation.** Snapshot exchange merges everything; merge rules that prioritize evidence about under-represented semantic tags attack the candidate-generation starvation of §3.5 directly. **(3) Trust-weighted merging under occupancy.** The central aggregator already performs a trust-weighted merge but pays the C-collapse cost; decentralized trust rules that account for peer *intent* (not only peer evidence) could decouple freshness from target racing. **(4) Substrate generality.** The decomposition transfers across MiniGrid, BabyAI, the custom grid, and now a standard-MiniGrid multi-agent wrapper with identical operational definitions and verified slope signs of Empirical Claim 1 (§ 6). A continuous-state environment (Habitat (Savva et al., 2019), VMAS (Bettini et al., 2022)) is the remaining portability step.

E.3 Extended related work

Classical cognitive-map perspectives motivate explicit place abstractions (Tolman, 1948; O’Keefe and Nadel, 1978; Kuipers, 2000); modern embodied AI combines mapping and planning (Gupta et al., 2017) or exposes topological memory (SPTM, Savinov et al., 2018); goal-conditioned RL conditions on explicit goals or embeddings (Andrychowicz et al., 2017; Ghosh et al., 2021), with




Figure 7: Bottleneck shift across eight peer cadences (40-seed sweep). Top: $P(M^*|R^*)$ falls and $P(C^*|M^*)$ rises monotonically in K ($p < 10^{-4}$, both Spearman). Bottom: mean time-to-success has an interior minimum at $K = 8$ (7.47 ticks), below the independent-agent baseline (8.02, dashed). Error bars and shaded band are 95% bootstrap CIs.

747 successor-representation work providing predictive structure (Dayan, 1993; Stachenfeld et al., 2017;
748 Momennejad et al., 2017). Our setting differs in that the target is a semantic pattern retrieved from
749 memory rather than a coordinate. In multi-agent RL, CTDE dominates (Lowe et al., 2017; Foerster
750 et al., 2018) and learned communication adds a per-step exchange variable (Sukhbaatar et al., 2016;
751 Foerster et al., 2016); gossip and consensus protocols (Boyd et al., 2006; Xiao et al., 2007) are the
752 closest network-level analogue of our cadence axis, lifted here to the cognitive-architecture level.
753 LLM-agent systems lean on explicit memory but evaluate end-to-end (Wang et al., 2023; Shinn et al.,
754 2023; Packer et al., 2023). Our framework is orthogonal: rather than learning a communication
755 policy end-to-end, it *diagnoses* where in the Q/R/M/C chain a given architecture succeeds or fails,
756 and yields a slope-level prediction about cadence that is testable on logs.

757 E.4 Multi-agent oracle-intervention numerics

Table 8: Multi-agent oracle interventions on the scarcity cases at three cadences. Mean time-to-success averaged across $(N, T) \in \{(3, 2), (5, 3), (8, 5)\}$, two layouts (asymmetric, random), 15 seeds.

K	Baseline	Oracle R	Oracle M	Gap to oracle
1	10.33	7.22	7.08	-3.11
4	10.97	7.22	7.08	-3.75
16	7.98	7.22	7.08	-0.76

Bottleneck-shift signature: negative within-cadence correlation peaks at $K = 4$ and decays at $K = 64$




Figure 8: Per-configuration scatter of $P(M^*|R^*)$ vs. $P(C^*|M^*)$ at three cadences. The strong negative correlation at $K = 4$ (Pearson -0.61) is the bottleneck-shift signature: the same configurations that succeed at M tend to fail at C and vice versa. The correlation decays to noise by $K = 64$.

758 E.5 Stage-agnostic cadence hypotheses and what the framework changes

759 A first version of the multi-agent cadence sweep (12,960 runs at 20 seeds; subsumed by the 35,640-run
760 sweep in this paper) was originally planned with four operational hypotheses formulated *before*
761 Empirical Claim 1. They are retained here because they motivated the eventual framework. The
762 framework was formulated after the sweep was complete; it is not claimed to have predicted the
763 sweep numbers a priori (see “Chronology” in §4.3).

764 Original hypotheses (pre-framework).

765 **H1.** An interior $K^* \in \{2, 4, 8\}$ minimising mean time-to-success exists for most (N, T, layout)
766 cells.

767 **H2.** K^* grows monotonically with scarcity $\rho = N/T$.

768 **H3.** When $\rho \leq 1$, $K^* = 1$.

769 **H4.** Peer-broadcast architectures Pareto-dominate centralized aggregation on (mean time-to-success,
770 success rate).

771 **Stand-alone result on the original 20-seed sweep.** The practical claim was supported: peer at the
772 empirically best K was faster than centralized in the scarcity regime ($p = 0.044$, Mann–Whitney).
773 Among the four structural hypotheses, only H4 was clearly supported in the random-layout subset.
774 H1 held in 7 of 18 cells; H2 had a non-significant negative correlation; H3 was supported in 4 of 9
775 admissible cells. The 35,640-run extended sweep used elsewhere in the paper produces matching
776 stand-alone numbers within bootstrap uncertainty.

777 **What the framework changes.** The hypotheses H1–H4 are stage-agnostic: each ties K^* to scarcity
778 ρ without specifying mechanism. Empirical Claim 1, formulated after the sweep, ties the slopes of
779 $P(M^*|R^*)$ and $P(C^*|M^*)$ to cadence; the slope conditions are testable on logs. The data verify both
780 signs ($p < 10^{-4}$). We make no formal claim about the shape of the conditional product Φ , on which
781 the data are monotone. The interior optimum holds on the efficiency metric mean time-to-success.

782 F Assist on/off analysis

783 Across all four single-agent task families, top-line success is essentially unchanged under assist on/off
784 ($\Delta \leq 0.10$). The only non-trivial effect is on rest, where assists improve post-retrieval completion
785 by $+0.05$ without changing query/retrieval rates, consistent with assists acting on materialization and
786 post-retrieval execution.

G Retrained retriever and CommNet baseline

G.1 Retrained candidate-generation MLP (§3.5)

Dataset extraction. From the 10-seed hazard-recovery benchmark (MiniGrid-LavaGapS7-v0, method `milestone_state_conditioned_hazard_recovery_v7`), an 18-dimensional feature vector was extracted for every step of every episode of every seed — 3 numeric agent-state scalars (energy, thirst, risk_budget), 9 semantic-tag confidences, 6 intent one-hot — and a binary label: *was this symbolic node ever selected as a candidate during the run?* Dataset totals: 814 examples, 113 positives (13.9%).

Split. By seed: train {2, 3, 4, 5, 6, 7} (505 examples, 9.7% positive), val {8, 9} (170, 22.4% positive), test {10, 11} (139, 18.7% positive).

Architecture. $18 \rightarrow 32 \rightarrow 32 \rightarrow 1$, ReLU activations, BCE with `pos_weight` ≈ 9.3 for class imbalance, Adam lr = $2 \cdot 10^{-3}$ with `weight_decay` = 10^{-4} , batch 64, 200 epochs, best by val loss.

Results.




Figure 9: Refined attribution diagnostic. *Left*: 349 retrieval calls aggregated across 10 seeds; 317 (90.8%) returned empty candidate sets, 31 (8.9%) returned at least one satisfied candidate, only 1 (0.3%) returned candidates that were ranked but unsatisfied. The bottleneck is candidate generation, not candidate ranking. *Right*: the trained MLP underperforms a majority-class baseline on all three seed-disjoint splits, confirming that instantaneous state features alone are insufficient to predict candidate eligibility.

The MLP underperforms majority-class on accuracy across all splits and achieves only moderate precision at fixed recall. Instantaneous state features are insufficient to predict candidate eligibility downstream. This is the negative-but-informative finding referenced in §3.5: the framework’s “retrieval-limited” attribution refines to “memory-accumulation-limited” rather than “ranking-limited”.

G.2 CommNet-style learned communication baseline (§4.5)

Policy architecture. Per-agent observation is 4-dim direction one-hot + $5 \times 5 \times 5$ local cell-type window (125) + 2-dim normalised position = 131. Shared-weights encoder $131 \rightarrow 64 \rightarrow 64$ (ReLU), 16-dim comm projection, peer mean-pool (excluding self), combine $80 \rightarrow 64$ (ReLU), actor head $64 \rightarrow 4$ (TURN_LEFT, TURN_RIGHT, FORWARD, NOOP), critic head $64 \rightarrow 1$. This is a CommNet-style learned-communication architecture (Sukhbaatar et al., 2016).

810 **Training.** REINFORCE (Williams, 1992) with value-baseline; loss = $-\text{mean}(\log \pi(a) \cdot (R - V)) +$
811 $0.5 \cdot \text{MSE}(V, R) - 0.01 \cdot H[\pi]$; Adam lr = $3 \cdot 10^{-4}$, gradient clipping 1.0. Reward shaping: +1.0
812 on first success per agent, -0.01 step cost, -0.1 per hazard contact. 2000 episodes cycling 200
813 environment seeds, wall-clock 100 s on CPU. Held-out eval: 50 environment seeds 250–299.

814 **Training curve.** Episode-level success (last-100 window) stays in $[0.13, 0.23]$ throughout training;
815 no monotonic improvement after ~ 200 episodes. The policy converges to a local optimum where
816 one of three agents finds water by chance.




Figure 10: Learned-communication baseline diagnostics. *Left:* CommNet training curve (smoothed success rate over a 100-episode window) over 2000 training episodes; success stays below 0.25 throughout. The dashed line is the symbolic peer architecture at $K = 8$ on the same scenario for reference. *Right:* Q/R/M/C profile of the trained CommNet on 50 held-out seeds vs symbolic peer at $K = 8$. Q^* is saturated at 1.0 on CommNet (degenerate by construction); R/M collapse together; only on the symbolic architecture are the conditional rates separated.

817 **Q/R/M/C evaluation (50 held-out seeds, stochastic action sampling).**

Quantity	Value
Q^* rate	1.00 (saturated)
R^* rate	0.50
M^* rate	0.50 (collapses onto R)
C^* rate	0.19
$P(R Q)$	0.50
$P(M R)$	1.00
$P(C M)$	0.42
Success rate	0.19
Mean steps	80 (budget exhausted)

818 Compare to the symbolic peer architecture at $K=8$: $P(M^*|R^*)=0.69$, $P(C^*|M^*)=0.87$, success
819 ~ 0.65 .

820 **PPO variant: structural degeneracy verified at competitive success.** A PPO-trained (Schul-
821 man et al., 2017) variant of the same CommNet architecture (clip $\epsilon=0.2$, generalized ad-
822 vantage estimation (Schulman et al., 2016) with $\lambda=0.95$, $\gamma=0.99$, 4 PPO epochs per roll-
823 out, 64 rollouts per update, 200 updates, ≈ 96 minutes on a single CPU core; full setup in
824 `experiments/commnet_ppo_baseline/train_ppo_commnet.py`) reaches 0.667 success on 100
825 held-out seeds — competitive with the symbolic peer architecture at $K=8$ (0.65).

826 The PPO result converts the structural argument from an assertion (“training budget would not change
827 this”) to a measurement: Q^* remains saturated at 1.00, $P(M^*|R^*)$ remains 1.00 (R/M collapse), even
828 at competitive success. The Q stage is degenerate by construction in any policy whose communication
829 channel is a real-valued vector with a learned linear projection — the channel is never informatively
830 empty, regardless of training budget — and without an explicit symbolic target-lock interface the M

Table 9: REINFORCE vs PPO CommNet vs symbolic peer ($K=8$) on the scarcity scenario ($N=3$, $T=2$, asymmetric, hazard 0.05). PPO is competitive on success rate, yet still exhibits Q-saturation and R/M collapse — the diagnostic claim verified, not asserted.

Metric	CommNet REINFORCE	CommNet PPO	Symbolic peer ($K=8$)
Success rate	0.19	0.667	0.65
Q^* rate	1.00	1.00	0.99
R^* rate	0.50	0.997	0.69
M^* rate	0.50	0.997	0.69
$P(M^* R^*)$	1.00	1.00	0.69
$P(C^* M^*)$	0.42	0.67	0.87
R/M separated?	no	no	yes

event cannot separate from R. Only architectures with an explicit *query/lock* interface (the symbolic stack of §4.1) expose stage-separable Q/R/M/C events.

H Reproducibility commands

An anonymized public repository accompanies this submission. The main paper artifacts were generated from the following scripts; PYTHONPATH=. and a Python 3 virtual environment with numpy, scipy, pandas, matplotlib, imageio, gymnasium, and minigrid are assumed throughout.

One-command reproduction (multi-agent block, ≈ 22 minutes on 8 cores)

bashscripts/reproduce_paper.sh

Single-agent (§3.2–§3.4)

- article benchmark and learned-baseline suite: scripts/benchmark_symbolic_memory_article.py
- article analysis and figure generation: scripts/analyze_symbolic_memory_article.py
- synthetic factorization sanity check: scripts/validate_qrmc_factorization.py
- oracle-stage intervention benchmark: scripts/benchmark_oracle_stage_interventions.py
- BabyAI transfer benchmark: scripts/compare_semnav_babyai.py

Multi-agent (§4.5)

- Main 25,920-run cadence sweep (40 seeds, $K \leq 16$, ≈ 18 min): pythonexperiments/big_experiment/exp_cadence_phase.py--modefull--workers8--out_dirtmp/big_experiment_qrmc_40
- Extended 9,720-run sweep ($K \in \{32, 48, 64\}$, ≈ 4 min): pythonexperiments/big_experiment/exp_extra_K.py--workers8--out_dirtmp/big_experiment_extraK
- Q/R/M/C stage analysis and bootstrap CIs: pythonexperiments/big_experiment/analyze_qrmc.py--runs_csvtmp/big_experiment_qrmc_40/runs.csv--out_dirtmp/big_experiment_qrmc_40
- Multi-agent oracle interventions: pythonexperiments/big_experiment/exp_oracle_interventions.py--workers8--out_dirtmp/big_experiment_oracle
- Phase 1–4 smoke tests: pythonscripts/run_all_smokes.py--out_dirtmp/all_smokes
- PPO-trained CommNet baseline (Table 9, ≈ 96 min on a single CPU core): python-mexperiments.commnet_ppo_baseline.train_ppo_commnet--total_updates200--rollouts_per_update64
- Q/R/M/C eval of the PPO policy on 100 held-out seeds: pythonexperiments/commnet_baseline/eval_with_qrmc.py--policy_pathexperiments/commnet_ppo_baseline/ppo_policy.pt--out_direxperiments/commnet_ppo_baseline--n_episodes100

- 865 • MiniGrid-substrate portability sweep (§6, 100 runs, ≈ 2 min): `python-mexperiments.`
866 `minigrid_multiagent_wrapper.run_portability_sweep--out_dirtmp/minigrid_`
867 `portability`
- 868 • Minimal CSM benchmark (Table 3, 3,240 runs, ≈ 5 min): `python-mexperiments.`
869 `collective_semantic_memory.run_csm_benchmark--out_dirtmp/csm_benchmark`

870 **I Asset and license note**

871 The experiments use MiniGrid (Chevalier-Boisvert et al., 2023) and BabyAI (Chevalier-Boisvert
872 et al., 2018), both released under the MIT License. The locally implemented baselines and benchmark
873 runner scripts introduced in this paper will be released under the MIT License as part of the camera-
874 ready artifact package.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and introduction present the paper primarily as a four-stage diagnostic framework with a reference symbolic-semantic instantiation, and they explicitly scope the evidence to benchmarked navigation environments with a limited portability check (Abstract; Sections 1 and 6).

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: The paper includes an explicit “Limitations and future work” section that discusses environment difficulty, the observable-query and observable-occupancy assumptions, the slope-vs-corollary gap of Proposition 3, the learned-baseline scope, and agent-count scaling (Section 6).

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: The paper states the conditional stage-Markov property and positivity assumptions in Section 2.2, gives a proof sketch in the main text, and includes a full formal proof of Proposition 1 in Appendix A together with the multi-agent extension (Proposition 2) and the slope-level argument for Proposition 3.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The paper specifies the task families, environments, seeds, episode counts, assist modes, baselines, statistical protocol, and benchmark semantics; the single-agent setup is in Section 3.1 and the multi-agent sweep protocol in Section 4.4; the operational stage definitions used to extract Q^* , R^* , M^* , C^* from runtime logs are in Appendix B.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

983 (d) We recognize that reproducibility may be tricky in some cases, in which case
984 authors are welcome to describe the particular way they provide for reproducibility.
985 In the case of closed-source models, it may be that access to the model is limited in
986 some way (e.g., to registered users), but it should be possible for other researchers
987 to have some path to reproducing or verifying the results.

988 5. Open access to data and code

989 Question: Does the paper provide open access to the data and code, with sufficient instruc-
990 tions to faithfully reproduce the main experimental results, as described in supplemental
991 material?

992 Answer: [Yes]

993 Justification: The paper describes a unified benchmark runner and source-of-truth artifacts.
994 An anonymized public release of the codebase — single-agent benchmark, multi-agent grid
995 environment, 35,640-run sweep driver, Q/R/M/C analyser, oracle-intervention scripts, and
996 exact commands to reproduce every table and figure — accompanies the submission as
997 supplemental material (anonymized GitHub repository; single-command reproduction of
998 the main multi-agent sweep on an 8-core machine in ≈ 22 minutes, matching Appendix H).

999 Guidelines:

- 1000 • The answer [N/A] means that paper does not include experiments requiring code.
- 1001 • Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- 1002 • While we encourage the release of code and data, we understand that this might not
1003 be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not
1004 including code, unless this is central to the contribution (e.g., for a new open-source
1005 benchmark).
- 1006 • The instructions should contain the exact command and environment needed to run to
1007 reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- 1008 • The authors should provide instructions on data access and preparation, including how
1009 to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- 1010 • The authors should provide scripts to reproduce all experimental results for the new
1011 proposed method and baselines. If only a subset of experiments are reproducible, they
1012 should state which ones are omitted from the script and why.
- 1013 • At submission time, to preserve anonymity, the authors should release anonymized
1014 versions (if applicable).
- 1015 • Providing as much information as possible in supplemental material (appended to the
1016 paper) is recommended, but including URLs to data and code is permitted.

1017 6. Experimental setting/details

1020 Question: Does the paper specify all the training and test details (e.g., data splits, hyperpa-
1021 rameters, how they were chosen, type of optimizer) necessary to understand the results?

1022 Answer: [Yes]

1023 Justification: The benchmark protocol, task definitions, assist settings, seeds, episode counts,
1024 baseline families, and the learned baselines are described in the main text together with the
1025 measurement and statistical protocols: single-agent setup in Section 3.1, multi-agent sweep
1026 in Section 4.4, and per-baseline hyperparameters in Appendix G.

1027 Guidelines:

- 1028 • The answer [N/A] means that the paper does not include experiments.
- 1029 • The experimental setting should be presented in the core of the paper to a level of detail
1030 that is necessary to appreciate the results and make sense of them.
- 1031 • The full details can be provided either with the code, in appendix, or as supplemental
1032 material.

1033 7. Experiment statistical significance

1034 Question: Does the paper report error bars suitably and correctly defined or other appropriate
1035 information about the statistical significance of the experiments?

Answer: [Yes]

Justification: The paper reports 95% bootstrap confidence intervals over seed-level aggregates ($B=5,000$ resamples; Section 3.1, “Protocol” paragraph) and uses paired Wilcoxon signed-rank tests with Holm-Bonferroni correction for assist comparisons. CIs are shown in Tables 3 and 5 (Best semantic per-task rates; the 35,640-run MOP column) and in Figures 4–6 (bootstrap shaded bands).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: The paper discloses the hardware platform (10-core ARM CPU, 32 GB RAM; CPU-only), the scale of each benchmark, and wall-clock costs — single-agent oracle sweep < 1 minute (Section 3.1); main multi-agent sweep (35,640 runs) ≈ 22 minutes on 8 cores (Section 4.4); $N=12$ scale-up (8,640 runs) ≈ 12 minutes (Appendix D); PPO CommNet baseline (200 updates $\times$ 64 rollouts) ≈ 96 minutes on a single CPU core (Appendix G); MiniGrid portability sweep (100 runs) ≈ 2 minutes on a single core (§6).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn’t make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

Answer: [Yes]

Justification: The work is a simulation-only study in standard benchmark environments, uses anonymized authorship for review, and does not involve human subjects, private data, or deployment claims that conflict with the NeurIPS Code of Ethics.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: The paper includes a dedicated “Broader impacts” section that discusses the primary methodological benefits of the work and the indirect risks of improving navigation systems without adequate safeguards (Section 7).

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper does not release high-risk models or scraped datasets; it studies navigation policies in simulated environments.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.

- 1142 • We recognize that providing effective safeguards is challenging, and many papers do
1143 not require this, but we encourage authors to take this into account and make a best
1144 faith effort.

1145 12. Licenses for existing assets

1146 Question: Are the creators or original owners of assets (e.g., code, data, models), used in
1147 the paper, properly credited and are the license and terms of use explicitly mentioned and
1148 properly respected?

1149 Answer: [\[Yes\]](#)

1150 Justification: The paper cites the upstream benchmark platforms (MiniGrid, BabyAI) in
1151 the main references. Appendix I states that both are MIT-licensed and that the benchmark
1152 scripts introduced here will be released under the MIT License at camera-ready.

1153 Guidelines:

- 1154 • The answer [\[N/A\]](#) means that the paper does not use existing assets.
- 1155 • The authors should cite the original paper that produced the code package or dataset.
- 1156 • The authors should state which version of the asset is used and, if possible, include a
1157 URL.
- 1158 • The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- 1159 • For scraped data from a particular source (e.g., website), the copyright and terms of
1160 service of that source should be provided.
- 1161 • If assets are released, the license, copyright information, and terms of use in the
1162 package should be provided. For popular datasets, paperswithcode.com/datasets
1163 has curated licenses for some datasets. Their licensing guide can help determine the
1164 license of a dataset.
- 1165 • For existing datasets that are re-packaged, both the original license and the license of
1166 the derived asset (if it has changed) should be provided.
- 1167 • If this information is not available online, the authors are encouraged to reach out to
1168 the asset's creators.

1169 13. New assets

1170 Question: Are new assets introduced in the paper well documented and is the documentation
1171 provided alongside the assets?

1172 Answer: [\[No\]](#)

1173 Justification: The submission introduces benchmark scripts and analysis artifacts. A formal
1174 anonymized release package with documentation is planned for camera-ready; Appendix H
1175 lists all scripts needed to reproduce each reported result, and Appendix I states the intended
1176 license for the new assets.

1177 Guidelines:

- 1178 • The answer [\[N/A\]](#) means that the paper does not release new assets.
- 1179 • Researchers should communicate the details of the dataset/code/model as part of their
1180 submissions via structured templates. This includes details about training, license,
1181 limitations, etc.
- 1182 • The paper should discuss whether and how consent was obtained from people whose
1183 asset is used.
- 1184 • At submission time, remember to anonymize your assets (if applicable). You can either
1185 create an anonymized URL or include an anonymized zip file.

1186 14. Crowdsourcing and research with human subjects

1187 Question: For crowdsourcing experiments and research with human subjects, does the paper
1188 include the full text of instructions given to participants and screenshots, if applicable, as
1189 well as details about compensation (if any)?

1190 Answer: [\[N/A\]](#)

1191 Justification: The paper does not involve crowdsourcing or human-subject experiments.

1192 Guidelines:

1193		• The answer [N/A] means that the paper does not involve crowdsourcing nor research
1194		with human subjects.
1195		• Including this information in the supplemental material is fine, but if the main contribu-
1196		tion of the paper involves human subjects, then as much detail as possible should be
1197		included in the main paper.
1198		• According to the NeurIPS Code of Ethics, workers involved in data collection, curation,
1199		or other labor should be paid at least the minimum wage in the country of the data
1200		collector.
1201	15. Institutional review board (IRB) approvals or equivalent for research with human	
1202	subjects	
1203	Question: Does the paper describe potential risks incurred by study participants, whether	
1204	such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)	
1205	approvals (or an equivalent approval/review based on the requirements of your country or	
1206	institution) were obtained?	
1207	Answer: [N/A]	
1208	Justification: The paper does not involve human subjects or human-participant data collec-	
1209	tion.	
1210	Guidelines:	
1211		• The answer [N/A] means that the paper does not involve crowdsourcing nor research
1212		with human subjects.
1213		• Depending on the country in which research is conducted, IRB approval (or equivalent)
1214		may be required for any human subjects research. If you obtained IRB approval, you
1215		should clearly state this in the paper.
1216		• We recognize that the procedures for this may vary significantly between institutions
1217		and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
1218		guidelines for their institution.
1219		• For initial submissions, do not include any information that would break anonymity (if
1220		applicable), such as the institution conducting the review.
1221	16. Declaration of LLM usage	
1222	Question: Does the paper describe the usage of LLMs if it is an important, original, or	
1223	non-standard component of the core methods in this research? Note that if the LLM is used	
1224	only for writing, editing, or formatting purposes and does <i>not</i> impact the core methodology,	
1225	scientific rigor, or originality of the research, declaration is not required.	
1226	Answer: [N/A]	
1227	Justification: LLMs are not an important, original, or non-standard component of the core	
1228	methods or experiments reported in the paper.	
1229	Guidelines:	
1230		• The answer [N/A] means that the core method development in this research does not
1231		involve LLMs as any important, original, or non-standard components.
1232		• Please refer to our LLM policy in the NeurIPS handbook for what should or should not
1233		be described.