
Q/R/M/C: Stage Decomposition for Memory-Based Navigation in Single Agents and Cadence-Coupled Collectives

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Success-only metrics collapse qualitatively different memory-based navigation fail-
2 ures into a single scalar, hiding *which stage* of the retrieval pipeline is responsible.
3 We propose the Q/R/M/C decomposition — **Q**uery formation, **R**etrieval satisfac-
4 tion, target **M**aterialization, **C**ompletion after retrieval — and show that, under a
5 conditional stage-Markov assumption, the four stage factors are estimable directly
6 from standard trajectory logs (Proposition 1). **Single-agent.** On a 10-seed Mini-
7 Grid benchmark the framework localizes bottlenecks: hazard recovery is retrieval-
8 limited (oracle retrieval lifts success $0.39 \rightarrow 0.60$) while goal-region rejoin is
9 downstream-limited (oracle retrieval rescues semantic-path completion without
10 moving end-to-end success). Inspection of the framework’s own logs refines the at-
11 tribution further: candidate *ranking* is near-optimal while candidate *generation* fails
12 at 91% of decision points, exposing a memory-accumulation limit. **Multi-agent.**
13 On a 35,640-run sweep with broadcast cadence $K \in \{1, 2, 4, 8, 16, 32, 48, 64\}$
14 and 40 seeds, the bottleneck-shift signature appears as a strongly negative within-
15 configuration correlation between $P(M^*|R^*)$ and $P(C^*|M^*)$ (Pearson -0.61 at
16 $K = 4$, decaying to -0.02 at $K = 64$), and mean time-to-success has an interior
17 minimum at $K = 8$. **Honest limit.** The strict interior-maximum corollary on the
18 conditional product $P(M^*|R^*) \cdot P(C^*|M^*)$ does *not* hold — the product mono-
19 tonically approaches the independent-agent boundary. The slope-level claim of
20 Proposition 3 is supported while one of its corollaries is not, and we report the
21 gap explicitly. Across both regimes the diagnostics converge on a single binding
22 constraint — how trajectory evidence is consolidated into persistent, shareable
23 symbolic structure — motivating *collective semantic memory* as the engineering
24 target that the framework is built to measure.

1 Introduction

26 Classical navigation benchmarks often assume that the agent is given an explicit destination coordinate
27 or a dense task reward. Many practical navigation problems are different: the agent must reach a
28 place of a certain *type* — a safe rest area, a water-like resource, a post-hazard recovery exit, or a
29 goal-associated region — specified by a semantic pattern rather than a fixed (x, y) waypoint. The
30 place must therefore be *remembered and retrieved by meaning*. This setting opens a methodological
31 question: can we measure *where* memory-based semantic navigation breaks down, not only *whether*
32 it succeeds?

33 We answer this question with a four-stage diagnostic framework that applies to both single agents
34 and multi-agent collectives. The decomposition is Q/R/M/C: **Q**uery formation, **R**etrieval satisfaction,
35 target **M**aterialization, and **C**ompletion after retrieval. The decomposition factorizes per-episode




Figure 1: Q/R/M/C decomposes one semantic-navigation episode into four observable stages. **Top:** the four stages and their event-level definitions. **(a)** A successful water-search episode anchors at one event per stage, producing logged starred events Q^* , R^* , M^* , C^* . **(b)** The same scalar “end-to-end success” summary masks qualitatively different failures on the four single-agent task families: water and rest are strong; goal-region rejoin (FourRooms) is downstream-limited; hazard recovery (LavaGapS7) is retrieval-limited (oracle interventions, §3.4).

semantic-path completion and, under a conditional stage-Markov property, its factors are estimable from trajectory logs (Proposition 1). The rest of the paper is one continuous argument in three moves.

Move 1 (single agent): stage attribution is non-trivial. Success-only metrics collapse qualitatively different failures into a single scalar. In a 10-seed MiniGrid benchmark, water and rest tasks are strong end-to-end (0.95, 0.93 success), but goal-region rejoin and hazard recovery are substantially harder — and for *different reasons*. Oracle-stage interventions show that hazard recovery is retrieval-limited (oracle retrieval lifts success $0.39 \rightarrow 0.60$), while goal-region rejoin is downstream-limited (oracle retrieval rescues semantic-path completion without changing end-to-end success). Inspection of the framework’s own logs then sharpens the attribution: the retrieval bottleneck is not candidate *ranking* (precision 96.9% when candidates exist) but candidate *generation* — the memory contains no eligible symbols at 91% of decision points. The single-agent bottleneck is a *memory-accumulation* limit: evidence is observed but never consolidated into queryable structure.

Move 2 (multi-agent): consolidation acquires a rate. When N agents share memory through a broadcast protocol with cadence K , the factorization extends with enlarged conditioning sets (Proposition 2), and K becomes a structural axis: it is, mechanically, a knob on *how fast private memories consolidate into a collective one*. A slope-level claim (Proposition 3) predicts a bottleneck shift: as K decreases, $P(M^*|R^*)$ rises (fresher peer memory enables reliable commitment) and $P(C^*|M^*)$ falls (peers race for the same targets). Both slope signs are confirmed on a 35,640-run sweep ($p < 10^{-4}$), the shift is directly visible as a negative within-cadence correlation between the two links, and the resulting trade-off produces an interior optimum of mean time-to-success at $K = 8$: consolidation has a non-trivial optimal rate.

Move 3 (synthesis): both regimes point at the same object. The single-agent diagnosis (consolidation of evidence into symbols is the limit) and the multi-agent diagnosis (the rate and cost of sharing consolidated memory determine the optimum) converge on one research target: *collective semantic memory* — a shared symbolic structure with explicit merge, trust, and staleness rules, rather than snapshots exchanged on a timer. Section 5 states this program and shows that Q/R/M/C is the measurement instrument it requires: any candidate design is testable by how it moves $P(M^*|R^*)$ and $P(C^*|M^*)$.

What the data do and do not support. A naive reading of the two slopes of Proposition 3 predicts an interior maximum on the conditional product $P(M^*|R^*) \cdot P(C^*|M^*)$. In this environment that prediction fails: the product is monotone in K and approaches the independent-agent boundary. The interior optimum appears instead on mean time-to-success. The gap is reported explicitly — it is the

68 difference between a stage-level operationalisation that can be *wrong* on specific summaries and a
69 stage-agnostic guess that is unfalsifiable.

70 **Why this framing is timely.** Embodied and goal-conditioned agents increasingly retrieve places,
71 objects, or tools by meaning (Andrychowicz et al., 2017; Ghosh et al., 2021; Savva et al., 2019).
72 LLM-agent systems lean on explicit memory layers but evaluate mostly end-to-end (Wang et al., 2023;
73 Shinn et al., 2023; Packer et al., 2023). Multi-agent reinforcement learning likewise increasingly
74 involves communicating populations (Lowe et al., 2017; Foerster et al., 2018), but the question
75 “which agent should know what, and when?” lacks a unified diagnostic vocabulary. Our framework
76 supplies one.

77 Contributions.

- 78 1. **Framework.** A stage-decomposition of memory-based navigation with a probabilistic factoriza-
79 tion whose factors are estimable from standard trajectory logs, for single agents (Proposition 1)
80 and multi-agent collectives (Proposition 2).
- 81 2. **Slope-level claim.** A bottleneck-shift claim (Proposition 3) about how cadence K trades off
82 the M and C stages, verified empirically ($p < 10^{-4}$ on both slopes), with its chronology stated
83 explicitly (§4.3).
- 84 3. **System.** A symbolic-semantic memory stack instantiating the framework, with four communi-
85 cation architectures (independent, shared-bus, central-aggregator, peer-to-peer broadcast).
- 86 4. **Benchmarks.** (a) 10-seed single-agent MiniGrid benchmark with oracle stage interventions; (b)
87 35,640-run multi-agent sweep at $N \in \{3, 5, 8\}$ plus an 8,640-run scale-up at $N = 12$.
- 88 5. **Honest accounting.** Explicit report of what fails: the strict interior-maximum corollary on the
89 conditional product does not hold in this environment; the interior optimum reappears on the
90 efficiency metric.
- 91 6. **Research program.** A data-grounded case that the binding constraint in both regimes is
92 consolidation of evidence into shareable symbolic structure, motivating collective semantic
93 memory as the next target and Q/R/M/C as its evaluation instrument (§5).

94 2 The Q/R/M/C decomposition

95 2.1 Problem formulation and notation

96 The agent operates in a partially observable navigation environment. At each step it receives an
97 observation o_t , maintains a memory state $\mathcal{M}_t$, and forms a query q_t encoding the current task
98 intent. A retrieval step returns a candidate set of places matching q_t from $\mathcal{M}_t$; a chosen candidate
99 is materialized into an actionable target; a local controller executes actions until either the target is
100 reached or the episode times out. The target is not a coordinate but a semantic pattern, so success is a
101 function of the chosen candidate’s semantic profile, not of any pre-supplied coordinate.

102 Habitat, AI2-THOR, and ProcTHOR made visually rich navigation evaluation standard (Savva
103 et al., 2019; Kolve et al., 2017; Deitke et al., 2022). We deliberately focus on MiniGrid-class
104 environments to obtain clean and reproducible stage measurements, and use BabyAI only as a
105 portability check. We do not claim to solve LLM-agent memory here; rather, we argue that the
106 same Q/R/M/C decomposition is a useful template for evaluating memory-grounded agent decisions
107 beyond navigation.

108 The central empirical question is therefore not just whether the agent succeeds, but *which stage fails*
109 *first* when it does not. Two related measurement layers are used. **Operational rates** count, per query
110 attempt, whether the query yielded a non-empty candidate set, whether the chosen candidate satisfied
111 the semantic predicate, and whether it was materialized into an actionable target; *completion after*
112 *retrieval* is an episode-level indicator. **Nested episode-level events** $Q^* \Rightarrow R^* \Rightarrow M^* \Rightarrow C^*$ are used
113 in the factorization below and audited in §3.2. Full operational definitions are given in Appendix B.

114 **Notation.** The four stages are Q, R, M, C with episode-level starred events $Q^*, R^*, M^*, C^* \in$
115 $\{0, 1\}$, nested by construction. $P(\cdot)$ denotes the population probability; $\hat{P}(\cdot)$ its empirical-frequency
116 estimator. **The letter M always refers to the Materialization stage; the number of targets in**
117 **the multi-agent environment is written T (not M , constraint $T \leq N$, “scarcity” = $N > T$).** Other

118 symbols: $\mathcal{C}_K$ peer-broadcast protocol with cadence K , $\mathcal{M}_{-i}^{(K)}$ peer memory snapshots, $\mathcal{O}_{-i}$ peer
 119 occupancy, $\nu(\cdot)$ freshness functional, $\Phi(K)=P(M^*|R^*;K)\cdot P(C^*|M^*;K)$, t_{succ} mean time-to-
 120 success. Full symbol table in Appendix C.1.

121 2.2 Single-agent factorization

122 Let Y denote overall episode success, and let Q^*, R^*, M^*, C^* denote the nested episode-level
 123 semantic-path events. The framework factorizes *semantic-path completion* rather than raw task
 124 success:

$$Q^* \rightarrow R^* \rightarrow M^* \rightarrow C^*.$$

125 For any query q and memory state $\mathcal{M}$,

$$P(C^* = 1 \mid q, \mathcal{M}) = P(Q^* = 1 \mid q, \mathcal{M}) P(R^* = 1 \mid Q^* = 1, q, \mathcal{M}) \\ \cdot P(M^* = 1 \mid R^* = 1, q, \mathcal{M}) P(C^* = 1 \mid M^* = 1, q, \mathcal{M}).$$

126 Overall task success Y can exceed C^* because some episodes succeed without traversing the full
 127 semantic retrieval path.

128 **Assumption 1 (conditional stage-Markov property).** Conditioned on the current stage being
 129 successful, the probability of the next stage depends only on the current successful stage, the current
 130 query, and the current memory state. Formally,

$$P(R^* \mid Q^*, q, \mathcal{M}, \text{history}) = P(R^* \mid Q^*, q, \mathcal{M}),$$

131 and analogously for M^* and C^* .

132 **Assumption 2 (positivity).** Each conditioning event used above has non-zero probability on the
 133 support of the benchmark.

134 **Proposition 1 (single-agent estimability).** Under Assumptions 1–2, the four stage factors are
 135 estimable from trajectory observations and the empirical-frequency estimators are consistent; their
 136 product converges to semantic-path completion. Proof in Appendix A.

137 The mathematical content is elementary — a chain-rule factorization plus consistency of frequency
 138 estimators. The value of the statement is operational: it pins down the exact logging discipline (which
 139 events, conditioned on what) under which the decomposition can be read off standard trajectory logs,
 140 and it is the contract that every experiment in this paper audits against (§3.2).

141 **When Assumption 1 may fail.** Assumption 1 holds tightly here because q and $\mathcal{M}$ are explicitly
 142 logged at every stage event. Two practically relevant violations remain: adaptive in-episode query
 143 rewrites not surfaced to the logger, and hidden between-stage memory updates. In those cases the
 144 factorization should be read as a projection onto observable traces rather than as a full structural
 145 identification claim.

146 **Mechanism-level reading.** The same decomposition admits a stage-structured mechanism interpre-
 147 tation. Figure 2 treats stage outcomes as a directed acyclic graph. Under this interpretation, assist
 148 on/off ablations are mechanism-targeted perturbations of materialization and post-retrieval execution
 149 rather than undifferentiated sensitivity analyses; in the multi-agent setting the same graph is replicated
 150 per agent and connected through the cadence protocol $\mathcal{C}_K$, which acts on $\mathcal{M}_{-i}^{(K)}$ at the R→M edge
 151 and on $\mathcal{O}_{-i}$ at the M→C edge — exactly the two edges where Proposition 3 will predict the slope
 152 shift.

153 **Relation to prior work.** Closest are (i) cognitive-map and semantic-map frameworks (Tolman,
 154 1948; O’Keefe and Nadel, 1978; Kuipers, 2000; Gupta et al., 2017; Savinov et al., 2018), (ii) goal-
 155 conditioned and successor-representation RL (Andrychowicz et al., 2017; Ghosh et al., 2021; Dayan,
 156 1993; Stachenfeld et al., 2017; Momennejad et al., 2017), (iii) multi-agent communication (Lowe
 157 et al., 2017; Foerster et al., 2018; Sukhbaatar et al., 2016; Foerster et al., 2016) and gossip/consensus
 158 (Boyd et al., 2006; Xiao et al., 2007), and (iv) LLM-agent memory layers evaluated end-to-end (Wang
 159 et al., 2023; Shinn et al., 2023; Packer et al., 2023). Our framework is orthogonal to all four: it
 160 *diagnoses* where in the Q/R/M/C chain a given architecture succeeds or fails and lifts the cadence
 161 axis from network protocols to a cognitive-architecture parameter (full discussion in Appendix E.3).




Figure 2: Stage-structured interpretation of the Q/R/M/C framework. Assist on/off ablations target intermediate mechanisms; in the multi-agent extension the cadence protocol $\mathcal{C}_K$ acts on the $R \rightarrow M$ and $M \rightarrow C$ edges.

Table 1: Stage-factor consistency audit for the best semantic method in each task family. The product of nested conditional estimators matches semantic-path completion; overall success can be higher because some episodes succeed without traversing the full semantic retrieval path.

Task	Method	Overall	SP-comp	$\hat{P}(Q^*)$	$\hat{P}(R^* Q^*)$	$\hat{P}(M^* R^*)$	$\hat{P}(C^* M^*)$	Product
Water	water-cp	0.95	0.70	1.00	0.98	0.74	0.97	0.70
Rest	rest-s2	0.93	0.73	1.00	1.00	0.80	0.91	0.73
Goal region	goal-n1	0.54	0.00	1.00	0.54	0.01	0.00	0.00
Hazard recovery	hazard-s7	0.40	0.16	1.00	0.56	0.58	0.50	0.16

3 Single agent: stage diagnostics localize failure

3.1 System, tasks, and protocol

Memory stack. The system uses a graph substrate with a semantic retrieval interface on top: observations o_t are tokenized into symbolic scene tokens and semantic tags that update a memory of place records, transitions, episodic traces, and concept/relation records. Planner queries return candidates by a weighted log-score over required, preferred, and penalty tags. The retrieval machinery is reused across all tasks and architectures; full memory schema, update rules, and scoring formula are in Appendix C.

Tasks, baselines, protocol. Four single-agent task families are evaluated (water, safe-rest, goal-region on Empty-6x6 / FourRooms, hazard recovery on LavaGapS7), 10 seeds $\times$ 8 episodes each, $\times$ 2 assist modes. Semantic stack (node, concept, state-conditioned variants) is compared against four baselines: random, graph-only, SPTM-like, learned BC. Reported: end-to-end success, steps, four operational stage rates, assist deltas, weakest stage per task and per method. All means are seed-aggregated; 95% bootstrap CIs use $B=5,000$ resamples; assist comparisons use paired Wilcoxon + Holm-Bonferroni (Appendix F). For the two hardest tasks (FourRooms, LavaGapS7), four oracle regimes (*base*, *oracle R/M/C*) are run — each replaces exactly one stage while leaving the rest of the pipeline unchanged. BabyAI portability check in Appendix D. All single-agent experiments ran on a 10-core ARM CPU, 32 GB RAM (CPU-only); oracle sweep < 1 min wall-clock.

3.2 Stage-estimator validation

Before the main benchmark, the factorization is validated. On a synthetic four-stage process with known probabilities the nested-estimator product tracks empirical semantic-path completion across sample sizes. On the real benchmark, the same product matches semantic-path completion rather than overall success (Table 1) — exactly the contract of Proposition 1.

3.3 Success-only summary and what it hides

Water and rest are strong end-to-end tasks (0.95 and 0.93 success; Figure 1b). Goal-region reaches 0.54 success at the task-family level, but the bottleneck sits downstream of candidate selection (§3.4). Hazard recovery is the hardest retrieval-sensitive task at 0.40 success, with retrieval satisfaction as the weakest stage. Under success-only evaluation the two hard families look like the same kind of

Table 2: Baseline comparison. The learned behavior-cloned baseline dominates the easiest resource tasks and hazard recovery, while the semantic stack remains strongest on the relation-heavy goal-region family and remains the only family with interpretable Q/R/M/C structure.

Task	Random	Graph-only	SPTM-like	Learned	Best semantic
Water	0.85	0.90	0.90	1.00	0.95
Rest	0.85	0.90	0.90	1.00	0.93
Goal region	0.41	0.49	0.52	0.50	0.54
Hazard recovery	0.08	0.39	0.00	0.74	0.40

weakness; the stage view will show they are opposites. Per-task operational rates with 95% CIs are in Table 3 (Appendix C.3).

Baseline comparison. Table 2 compares the semantic stack against random, graph-only, SPTM-like, and learned behavior-cloned baselines. The learned baseline dominates the easiest resource tasks (1.00 on water and rest) and hazard recovery (0.74), while the semantic stack remains strongest on the relation-heavy goal-region family (0.54). The relevant contrast for this paper, however, is not the success column: the semantic stack is the only method family whose decisions expose interpretable Q/R/M/C structure. On the BC baseline the stage events are degenerate for the same structural reason as on the multi-agent learned baseline analysed in §4.5 — a policy without an explicit query/lock interface cannot emit separable stage events, so the diagnostic decomposition has nothing to attach to. End-to-end learning and stage diagnostics answer different questions; this paper is about the second.

3.4 Oracle-stage interventions localize the bottleneck

On goal-region the semantic variants solve Empty-6x6 (1.00 success) but collapse on FourRooms (0.075). For hazard recovery (LavaGapS7), base success 0.39 rises only to 0.43 under oracle controller, but jumps to 0.59 under oracle materialization and 0.60 under oracle retrieval; goal-region rejoin is downstream-limited (oracle R lifts semantic-path completion to 0.08 without moving end-to-end success). The per-task attribution as a stacked waterfall is in Figure 5 (Appendix C.3). Two tasks identical under success-only evaluation fail at opposite ends of the Q/R/M/C chain.

3.5 The logs refine the diagnosis: a memory-accumulation limit

Inspecting the framework’s own logs sharpens the retrieval attribution. On the 10-seed hazard-recovery run, 349 retrieval calls were logged: only 32 (9.2%) returned at least one candidate, and of these 31 (96.9% precision) selected a semantically satisfied candidate. The bottleneck is therefore not candidate *ranking* but candidate *generation* — the symbolic memory does not contain hazard-recovery-tagged nodes at 91% of decision points. A retrained 2-layer candidate-generation MLP trained on per-state features (test accuracy 0.59 vs majority-class 0.81; Appendix G) confirms that instantaneous state features cannot substitute for the missing structure. The framework’s “retrieval-limited” attribution thus refines to a *memory-accumulation* limit: candidate generation cannot be solved from state features alone, and the binding constraint is how trajectory evidence is consolidated into persistent symbolic structure. This sets up the multi-agent section: when consolidation becomes a shared process across N agents, it acquires a measurable *rate*.

4 Multi-agent: memory as a shared resource

4.1 Four memory architectures and the cadence axis

Four communication architectures are instantiated; they share the single-agent retrieval substrate but differ in how memory is distributed across agents (Figure 3).

Independent. Each agent maintains a private event store and concept graph. No inter-agent communication mechanism exists. Lower-bound baseline.

Shared bus. All agents publish observations to a single *collective memory* event bus. One shared concept graph is rebuilt from the union of events. Maximally centralized.




Figure 3: Four multi-agent memory architectures. **(a) Independent:** per-agent private memory, no exchange. **(b) Shared bus:** one global memory, every agent reads/writes each tick. **(c) Central aggregator:** per-agent memories plus a hub that merges them into a consensus report. **(d) Peer-to-peer broadcast:** per-agent memories with pairwise snapshot exchange every K ticks (the cadence axis of Proposition 3).

Central aggregator (ConsensusLayer). Each agent has its own private concept graph; a central layer pulls snapshots and computes a trust-weighted merge, producing a single *consensus report* consumed by all agents.

Peer-to-peer broadcast (cadence K). Each agent has its own concept graph and its own merged view. Every K ticks each agent broadcasts a snapshot to a passive bus; each agent independently merges its own snapshot with the latest snapshot received from each peer. Different agents may hold different merged views; there is no central report.

The first three correspond to the three regimes in the multi-agent communication literature: no exchange, CTDE-style aggregation, and decentralized peer communication. The fourth is the architecture for which Propositions 2 and 3 are directly applicable, since the cadence K is a tunable parameter rather than a fixed protocol. Mechanically, K is the rate at which private memories consolidate into a collective view — which is why it is the structural axis of this section.

4.2 Multi-agent factorization

Extend the factorization to a population of N agents that share information through a communication protocol $\mathcal{C}_K$ parameterized by broadcast cadence K . Each agent i has its own memory state $\mathcal{M}_i$ and issues its own query q_i . Let $\mathcal{M}_i^{(K)}$ denote the projection of agent i 's memory at the moment it queries, after absorbing every peer message received under cadence K . Two coupling channels are separated: (i) memory coupling $\mathcal{M}_{-i}^{(K)}$, the latest peer memory snapshots accessible to agent i ; (ii) occupancy coupling $\mathcal{O}_{-i}$, the set of targets that other agents have already claimed by the time i acts.

Proposition 2 (multi-agent factorization, estimability). For each agent i ,

$$P(C_i^* | q_i, \mathcal{M}_i, \mathcal{C}_K) = P(Q_i^* | q_i, \mathcal{M}_i) P(R_i^* | Q_i^*, q_i, \mathcal{M}_i, \mathcal{M}_{-i}^{(K)}) \\ \cdot P(M_i^* | R_i^*, q_i, \mathcal{M}_i, \mathcal{M}_{-i}^{(K)}) P(C_i^* | M_i^*, q_i, \mathcal{O}_{-i}).$$

Under a multi-agent stage-Markov assumption (Appendix B), the four factors are estimable from per-agent trajectory logs and the empirical-frequency estimators are consistent.

When Proposition 2 may fail. The single-agent stage-Markov property absorbs trajectory history into $(q, \mathcal{M})$. In the multi-agent setting, the analogous absorption requires that dependence on past peer decisions enters only through $\mathcal{M}_{-i}^{(K)}$ at the M stage and through $\mathcal{O}_{-i}$ at the C stage. This step is non-trivial: $\mathcal{O}_{-i}$ summarises target occupancy at decision time but loses information about *how* peers arrived there. The factorization is therefore a projection onto observable per-agent traces under that summary. Two specific failure modes arise: (i) inter-agent communication side-channels not logged through $\mathcal{C}_K$; (ii) coordinated peer policies whose decision at the C stage depends on hidden peer state not summarised by $\mathcal{O}_{-i}$.

4.3 Bottleneck shift

Proposition 3 (bottleneck shift $M \leftrightarrow C$, slope-level claim). Assume scarcity ($N > T$ targets, see §2.1) and a peer-broadcast architecture with cadence $K \in \mathbb{N}$. Two stylized assumptions are posited: (F) memory freshness $\nu(\mathcal{M}_{-i}^{(K)})$ is non-decreasing pointwise as K decreases, and the probability that agent i 's planner locks onto a reachable target is non-decreasing in ν ; (O) the probability that an arbitrary committed target is already in $\mathcal{O}_{-i}$ is non-increasing as K increases. Write $P(M^*|R^*; K)$ and $P(C^*|M^*; K)$ for the population-level conditional probabilities (averaged over agents and the seed distribution of the benchmark); $\hat{P}(\cdot)$ denotes the corresponding empirical-frequency estimator computed from logs. Under (F) and (O),

$$\frac{\partial P(M^*|R^*; K)}{\partial K^{-1}} > 0, \quad \frac{\partial P(C^*|M^*; K)}{\partial K^{-1}} < 0.$$

Both slope conditions are direct empirical predictions that are testable on logs through $\hat{P}$.

Strict interior-maximum corollary (NOT empirically supported). A naive reading of (F) and (O) suggests that the product $\Phi(K) := P(M^*|R^*; K) \cdot P(C^*|M^*; K)$ admits an interior maximum at some $K^* \in \{2, \dots, K_{\max} - 1\}$. In the 35,640-run sweep (§4.5) this corollary does *not* hold: Φ is monotonically increasing in K over $K \in \{4, 8, 16, 32, 48, 64\}$ and continues to increase toward the independent-agent boundary value $\Phi(K \rightarrow \infty) = 0.605$. The slopes (F) and (O) are real, but the M-side gain in $P(M^*|R^*)$ at small K is not large enough to overcome the C-side cost; the two effects do not produce an interior peak in Φ in this environment. The interior optimum that the framework intuitively predicts does appear, however, on the practically relevant efficiency metric mean time-to-success: this metric has a clear minimum at $K = 8$ (§4.5), with significantly worse performance at both faster and slower cadences. The conditional product Φ captures the chain probability $P(M^* \wedge C^* | R^*)$ but ignores the cost of getting through the chain.

Status and chronology. Assumptions (F) and (O) are intuitive but not derived from a generative model; a formal proof would require, e.g., modelling memory updates as a Poisson-renewal process with rate λ and occupancy as a birth-death process whose rate scales with K^{-1} . Proposition 3 is therefore treated as a slope-level empirical claim verified on data, separating the verified slope conditions from the corollary that does not survive the data. We also state the chronology explicitly: Proposition 3 was formulated *after* an initial 12,960-run sweep conducted under four stage-agnostic hypotheses (Appendix E); its testable slope consequences were then confirmed on the full 35,640-run sweep and, independently, on the $N = 12$ scale-up (Appendix D). It is a post-hoc structural explanation whose predictions were verified on enlarged data, not a pre-registered prediction. A fully formal version is left for future work.

4.4 Sweep protocol

The four architectures are evaluated on a custom multi-agent grid (12×10 cells); each agent must reach a cell tagged `water_source`, choosing actions via a shared memory-driven planner so only the memory layer differs. The sweep crosses $N \in \{3, 5, 8\}$, $T \leq N$, three layouts (symmetric/asymmetric/random), three hazard densities, four architectures, eight peer cadences $K \in \{1, 2, 4, 8, 16, 32, 48, 64\}$, 40 seeds — 35,640 runs in ≈ 22 min on 8 cores. Per-agent per-tick Q/R/M/C indicators are logged (Appendix B); episode-level starred events are the disjunction over ticks. A separate oracle-intervention block on scarcity cases $(N, T) \in \{(3, 2), (5, 3), (8, 5)\}$ at five cadences evaluates oracle R (memory returns ground-truth water positions) and oracle M (planner locks onto nearest real water).

4.5 Results: stage decomposition and bottleneck shift

The 35,640-run sweep is reported through four observations, structured to mirror Proposition 3: (i) the conditional decomposition isolates the M and C stages as the differentiating links; (ii) the bottleneck shift appears directly as a negative within-cadence correlation between $P(M^*|R^*)$ and $P(C^*|M^*)$, peaking at $K = 4$ and decaying to noise at $K = 64$; (iii) on mean time-to-success, the resulting trade-off produces a clean interior optimum at $K = 8$; (iv) on the conditional product Φ , the same trade-off produces a monotone trend rather than an interior peak — the corollary that does not




Figure 4: Per-architecture conditional rates (top) and mean time-to-success (bottom) from the 35,640-run sweep. Shared bus, central aggregator, and fast peer broadcast ($K \leq 4$) live in the M-saturation / C-collapse regime; independent and slow peer ($K \geq 16$) live in the M-deficit / C-saturation regime. Mean time-to-success has a clean interior minimum at $K=8$ (7.47 ticks), below the independent-agent baseline (8.02).

307 survive the data. A learned-communication baseline and a multi-agent oracle-intervention loop close
 308 the subsection.

309 4.5.1 Conditional decomposition reveals where each architecture fails.

310 Figure 4 shows per-architecture conditional $Q/R/M/C$ rates; full numerics with bootstrap CIs are in
 311 Appendix E.1, Table 5. Retrieval conditional $P(R^*|Q^*)$ is saturated; the differentiating signal lives
 312 at M and C. Shared-bus, central-aggregator, and fast-broadcast peer ($K \leq 4$) have high $P(M^*|R^*)$
 313 but low $P(C^*|M^*)$; independent and slow-broadcast peer ($K \geq 16$) show the reverse — the $M \leftrightarrow C$
 314 trade-off predicted by Proposition 3.

315 4.5.2 Bottleneck shift and interior efficiency optimum.

316 Across all peer runs, Spearman of $P(M^*|R^*)$ vs K is -0.608 and of $P(C^*|M^*)$ vs K is $+0.420$
 317 (both $p < 10^{-4}$), matching the slope conditions of Proposition 3. The stronger empirical signa-
 318 ture is at the per-configuration level: Pearson of $P(M^*|R^*)$ vs $P(C^*|M^*)$ within a fixed K is
 319 $-0.43, -0.54, -0.61, -0.19, -0.05, -0.04, -0.02, -0.02$ at $K=1, 2, 4, 8, 16, 32, 48, 64$, peaking
 320 at $K=4$ and decaying to noise by $K=64$ (Figure 8, Appendix E.1). Configurations where M is
 321 high systematically have low C: the bottleneck migrates between the two stages within the same
 322 ($N, T, \text{layout}, \text{hazard}, \text{seed}$) cell as K varies. On the efficiency metric, mean t_{succ} has a U-shape with
 323 interior minimum at $K=8$ (7.47 ticks vs 8.02 for the independent baseline; paired-bootstrap CIs
 324 in Appendix E.1). Scaling to $N=12$ sharpens the signature: the peak Pearson rises to -0.70 and
 325 migrates outward to $K=8$, consistent with stronger occupancy coupling at higher N delaying the
 326 cost-regime transition (Appendix D, Figure 6).

327 **Honest limit on the conditional product.** The conditional product $\Phi = P(M^*|R^*) \cdot P(C^*|M^*)$
 328 does *not* have an interior peak in the tested range: the mean-of-products estimator increases mono-
 329 tonically from 0.569 at $K=4$ to 0.603 at $K=64$ and converges to the independent-agent boundary
 330 0.605. The slopes of Proposition 3 are supported; the additional interior-maximum corollary is not,
 331 in this environment. Jensen-gap detail and POM-MOP comparison in Appendix E.1.

332 **Learned-communication baseline: stage events are structurally degenerate even at competitive**
 333 **success.** A CommNet-style policy (Sukhbaatar et al., 2016) trained with PPO on the same scenario

($N=3, T=2$) reaches 0.667 held-out success ($n=100$ seeds) — matching the symbolic peer architecture at $K=8$ (0.65). Despite this, the Q/R/M/C profile remains structurally degenerate: $Q^*=1.00$ saturated by construction; $R^*=0.997$, $M^*=0.997$ collapsed together ($P(M^*|R^*)=1.00$). This is the verified structural claim: a real-valued comm vector is never informatively empty (Q saturates) and lacks an explicit target-lock interface (R/M merge), regardless of training budget. The multi-agent counterpart of the single-agent learned-BC analysis (§3.3). Full architecture, PPO hyperparameters, and side-by-side REINFORCE vs PPO comparison in Appendix G.

Multi-agent oracle interventions close the diagnose→intervene→verify loop. On the scarcity cases, mean t_{succ} at small/intermediate cadence ($K \in \{1, 2, 4\}$) drops from ~ 10 –11 in baseline to ~ 7.1 under either oracle R or oracle M; at slow cadence ($K=16$) the baseline is already within ~ 0.8 of the oracle floor (Table 6 in Appendix E.4). The cadence-induced excess time at small/intermediate K is attributable to the $M \leftrightarrow C$ transition; at large K the bottleneck has moved into solo exploration.

Bridge. The multi-agent block thus delivers its half of the argument. The cadence K is, mechanically, the rate at which private memories consolidate into a collective view; the data show that this rate has a non-trivial interior optimum ($K=8$ on t_{succ}), that its cost channel is occupancy coupling, and that both effects are visible stage-by-stage on the same Q/R/M/C chain used for single agents. Consolidation is not only the single-agent limit — in a collective it is a tunable, measurable design axis.

5 Synthesis: toward collective semantic memory

The single-agent and multi-agent halves of this paper run on different substrates, with different agents and different questions, yet they return the same answer. The single-agent diagnosis localized the hardest failure not to retrieval mechanics but to *consolidation*: candidate ranking is near-optimal (96.9% precision when candidates exist), but candidate generation is starved (91% empty candidate sets), and instantaneous state features cannot substitute for the missing structure (§3.5). The multi-agent diagnosis showed that when consolidation becomes a shared process it acquires a *rate* (the cadence K) with a benefit channel at the $R \rightarrow M$ edge and a cost channel at the $M \rightarrow C$ edge, whose trade-off produces an interior optimum neither fully centralized nor fully independent architectures can reach. The binding constraint in both regimes is the consolidation of evidence into shared symbolic structure.

We therefore propose *collective semantic memory* — a shared symbolic structure with explicit **merge**, **trust**, and **staleness** rules — as the research target the data select. The architectures of §4.1 are degenerate corners of this design space (shared bus merges everything instantly and pays full occupancy cost; independent never merges and pays full staleness; peer-broadcast exposes exactly one design parameter, the rate K). The framework built in this paper is the measurement instrument any candidate design requires: a good consolidation rule must raise $P(M^*|R^*)$ without collapsing $P(C^*|M^*)$ — that is, reach the upper-right region of Figure 8 that no fixed cadence reaches — and reduce the empty-candidate-set fraction of §3.5. Concrete open problems (adaptive cadence, selective consolidation, intent-weighted trust, substrate generality) are listed in Appendix E.2.

6 Limitations and future work

Portability across substrates: slope conditions hold on MiniGrid. The slope conditions of Proposition 3 also hold on a standard-MiniGrid substrate. A FourRooms-based multi-agent wrapper (built from `minigrid.core.grid.Grid` + `Wall/Floor/Lava` primitives, identical operational Q^*, R^*, M^*, C^* definitions; `experiments/minigrid_multiagent_wrapper/`) was used to run a 100-run portability sweep ($K \in \{1, 4, 8, 16, 64\}$, 20 seeds, $N=3, T=2$): Spearman $P(M^*|R^*)$ vs K is -0.50 ($p=1.5 \times 10^{-7}$), Spearman $P(C^*|M^*)$ vs K is $+0.50$ ($p=1.3 \times 10^{-7}$). The interior t_{succ} minimum does not separate at this minimal sweep size (5 cadences, 20 seeds, one hazard density); a denser sweep is the natural next step. The two-substrate seam is therefore closed at the slope-level claim.

Proposition 3 at slope level only. The two slope conditions hold ($p < 10^{-4}$); the strict interior-maximum corollary on Φ does not (§4.3). A version of Proposition 3 stated directly in time-cost would

close this gap. Proof under a generative model (Poisson-renewal memory, birth-death occupancy) is future work.

Estimability under hidden state. With unlogged in-episode query rewrites or peer side-channels, Propositions 1 and 2 are projections onto logged traces, not full identification.

Learned baselines. The CommNet PPO baseline reaches 0.667 success on the scarcity scenario — competitive with symbolic peer at $K=8$ (0.65) — and confirms Q-saturation ($Q^*=1.00$) and R/M collapse ($P(M^*|R^*)=1.00$) at competitive success. The structural prediction now holds verified, not asserted; ablating on architectures with explicit symbolic target-lock channels is left for future work.

Scale. Multi-agent sweep bounded at $N=12$ (Appendix D; signature sharpens, peak Pearson -0.70 at $K=8$, interior t_{succ} minimum 4.77). Denser scaling at $N \geq 16$ and continuous-state environments are future work.

7 Broader impacts

The paper proposes diagnostic evaluation for memory-based agents in simulated navigation. Its positive impact is methodological: making memory-based failures more measurable across architectures, and making the cadence axis of distributed multi-agent memory tunable on an empirically grounded criterion. Potential negative impacts are indirect and would arise if similar diagnostics were used to improve agents in safety-critical settings (drone swarms, autonomous vehicles, distributed sensor networks) without adequate safeguards. The current work is benchmark-level infrastructure, not a deployment-ready system.

8 Conclusion

We proposed Q/R/M/C, a four-stage decomposition of memory-based navigation whose factors are estimable from standard trajectory logs, and used it to run a single continuous argument: stage diagnostics localize failures that success-only evaluation cannot (single agent); when memory is shared, the exchange cadence becomes a measurable structural axis with an interior optimum (multi-agent); and both measurements localize the binding constraint to the same mechanism — the consolidation of trajectory evidence into persistent, shareable symbolic structure. Collective semantic memory — a shared symbolic structure with explicit merge, trust, and staleness rules — is therefore not a speculative extension but the research target that the data select, and the framework presented here is the instrument by which candidate designs for it can be evaluated.

References

- E. C. Tolman. Cognitive maps in rats and men. *Psychological Review*, 55(4):189–208, 1948.
- J. O’Keefe and L. Nadel. *The Hippocampus as a Cognitive Map*. Oxford University Press, 1978.
- B. Kuipers. The spatial semantic hierarchy. *Artificial Intelligence*, 119(1–2):191–233, 2000.
- S. Gupta, J. Davidson, S. Levine, R. Sukthankar, and J. Malik. Cognitive mapping and planning for visual navigation. In *CVPR*, 2017.
- P. Dayan. Improving generalization for temporal difference learning: The successor representation. *Neural Computation*, 5(4):613–624, 1993.
- M. Andrychowicz, F. Wolski, A. Ray, J. Schneider, R. Fong, P. Welinder, B. McGrew, J. Tobin, P. Abbeel, and W. Zaremba. Hindsight experience replay. In *NeurIPS*, 2017.
- N. Savinov, A. Dosovitskiy, and V. Koltun. Semi-parametric topological memory for navigation. In *ICLR*, 2018.
- K. L. Stachenfeld, M. M. Botvinick, and S. J. Gershman. The hippocampus as a predictive map. *Nature Neuroscience*, 20(11):1643–1653, 2017.
- I. Momennejad, E. M. Russek, J. H. Cheong, M. M. Botvinick, N. Daw, and S. J. Gershman. The successor representation in human reinforcement learning. *Nature Human Behaviour*, 1:680–692, 2017.
- D. Ghosh, A. Gupta, A. Reddy, J. Fu, S. Levine, and C. Finn. Learning to reach goals via iterated supervised learning. In *ICLR*, 2021.

- 430 E. Kolve, R. Mottaghi, D. Gordon, Y. Zhu, A. Gupta, A. Farhadi, and D. Fox. AI2-THOR: An interactive 3D
431 environment for visual AI. *arXiv:1712.05474*, 2017.
- 432 M. Savva, A. Kadian, O. Maksymets, Y. Zhao, E. Wijmans, B. Jain, J. Straub, J. Liu, V. Koltun, J. Malik, et al.
433 Habitat: A platform for embodied AI research. In *ICCV*, 2019.
- 434 M. Deitke, E. V. Wijmans, D. Schwenk, O. Salvador, K. M. Ehsani, et al. ProcTHOR: Large-scale embodied AI
435 using procedural generation. In *NeurIPS Datasets and Benchmarks*, 2022.
- 436 G. Wang, Y. Xie, Z. Jiang, A. Mandlekar, C. Xie, A. Yu, and P. Liang. Voyager: An open-ended embodied agent
437 with large language models. *arXiv:2305.16291*, 2023.
- 438 N. Shinn, F. Cassano, A. Berman, and A. G. Labash. Reflexion: Language agents with verbal reinforcement
439 learning. In *NeurIPS*, 2023.
- 440 C. Packer, V. Fang, S. Patil, K. Lin, S. Wooders, and J. Gonzalez. MemGPT: Towards LLMs as operating
441 systems. *arXiv:2310.08560*, 2023.
- 442 M. Chevalier-Boisvert, D. Bhoopchand, H. Abdolmaleki, and A. Mordatch. BabyAI: A platform to study the
443 sample efficiency of grounded language learning. In *ICLR*, 2019.
- 444 M. Chevalier-Boisvert, B. Lahlou, S. Willems, and P. S. Castro. MiniGrid and MiniWorld: Modular reinforcement
445 learning environments. *arXiv:2306.13831*, 2023.
- 446 R. Lowe, Y. Wu, A. Tamar, J. Harb, P. Abbeel, and I. Mordatch. Multi-agent actor-critic for mixed cooperative-
447 competitive environments. In *NeurIPS*, 2017.
- 448 J. N. Foerster, G. Farquhar, T. Afouras, N. Nardelli, and S. Whiteson. Counterfactual multi-agent policy gradients.
449 In *AAAI*, 2018.
- 450 S. Sukhbaatar, A. Szlam, and R. Fergus. Learning multiagent communication with backpropagation. In *NeurIPS*,
451 2016.
- 452 J. N. Foerster, Y. M. Assael, N. de Freitas, and S. Whiteson. Learning to communicate with deep multi-agent
453 reinforcement learning. In *NeurIPS*, 2016.
- 454 S. Boyd, A. Ghosh, B. Prabhakar, and D. Shah. Randomized gossip algorithms. *IEEE Trans. Information Theory*,
455 52(6):2508–2530, 2006.
- 456 L. Xiao, S. Boyd, and S. Lall. A scheme for robust distributed sensor fusion based on average consensus. *IPSN*,
457 2007.

458 A Formal proofs

459 A.1 Proof of Proposition 1

460 For each episode $i \in \{1, \dots, n\}$ let $Q_i^*, R_i^*, M_i^*, C_i^*$ denote the nested episode-level indicators. By
461 construction these are measurable functions of the runtime trajectory log and satisfy $C_i^* \Rightarrow M_i^* \Rightarrow$
462 $R_i^* \Rightarrow Q_i^*$. The empirical estimators

$$\hat{p}_Q = \frac{1}{n} \sum_i Q_i^*, \quad \hat{p}_{R|Q} = \frac{\sum_i R_i^*}{\sum_i Q_i^*}, \quad \hat{p}_{M|R} = \frac{\sum_i M_i^*}{\sum_i R_i^*}, \quad \hat{p}_{C|M} = \frac{\sum_i C_i^*}{\sum_i M_i^*}$$

463 are consistent by the chain rule plus the conditional stage-Markov property, and the empirical product
464 converges to semantic-path completion. $\square$

465 A.2 Proof of Proposition 2

466 Conditioning sets are enlarged by $\mathcal{M}_{-i}^{(K)}$ and $\mathcal{O}_{-i}$. Under the enlarged stage-Markov property, the
467 proof of Proposition 1 transfers verbatim and each enlarged conditional is a Bernoulli probability on
468 a measurable event. The qualifications stated in §4.2 (“When Proposition 2 may fail”) describe two
469 specific situations in which the enlarged stage-Markov property is violated. $\square$

470 A.3 Justification of Proposition 3

471 No full proof of Proposition 3 is given here; the slope-level argument under (F) and (O) is recorded
472 below.

473 *Slope (a), under (F).* As K decreases, $\nu(\mathcal{M}_{-i}^{(K)})$ is non-decreasing pointwise. Conditional on
474 a successful retrieval, the probability of locking onto a reachable target is non-decreasing in ν .
475 Therefore $\partial P(M^*|R^*; K)/\partial K^{-1} > 0$.

476 *Slope (b), under (O).* As K decreases, peers receive updated memory simultaneously and con-
477 verge on the same closest target; the marginal occupancy $\Pr[\text{target} \in \mathcal{O}_{-i}]$ rises. Therefore
478 $\partial P(C^*|M^*; K)/\partial K^{-1} < 0$.

479 *Interior maximum corollary fails in our data.* A naive reading would predict that $\Phi(K) =$
480 $P(M^*|R^*; K) \cdot P(C^*|M^*; K)$ admits an interior maximum. In the 35,640-run sweep Φ is monotoni-
481 cally increasing over $K \in \{4, 8, 16, 32, 48, 64\}$ and converges to the independent-agent boundary
482 value $\Phi_\infty = 0.605$. Under (F) and (O) the slope of Φ in K^{-1} is the sum of two opposing terms;
483 in our environment the M-side gain at small K is not large enough to flip the sign of the sum. A
484 formal version of the corollary would require quantitative bounds on the relative magnitudes, e.g. via
485 a generative model in which freshness $\nu(K)$ saturates at a known rate and occupancy scales with
486 K^{-1} . That is left to future work; the empirical truth is reported as found: slope conditions verified;
487 interior peak on Φ not found. $\square$

488 B Operational vs. nested stage estimators

489 **Operational diagnostic rates.** The main benchmark logs per-query and per-episode diagnos-
490 tic quantities directly from runtime traces (query formed, retrieval satisfied, target materialized,
491 completion after retrieval). These are operational rates, not direct multiplicative factors.

492 **Nested stage estimators for factorization.** For the factorization, episode-level nested events
493 from the first failure stage recorded in the runtime taxonomy are used: Q^*, R^*, M^*, C^* , nested by
494 construction. The consistency audit in §3.2 reports $\hat{P}(Q^*), \hat{P}(R^*|Q^*), \hat{P}(M^*|R^*), \hat{P}(C^*|M^*)$,
495 their product, and its gap to semantic-path completion.

496 B.1 Multi-agent operational definitions

497 For the multi-agent sweep, every metric is defined per agent and per tick, then aggregated. Let
498 $\text{Targets}_i(t)$ be the candidate set returned by the memory query of agent i at tick t , $\text{Locked}_i(t)$ be the
499 agent’s planner-internal locked target, and $\mathcal{W}$ the ground-truth target set.

500 **Per-tick events.** $Q_i(t) = \mathbf{1}[\text{Targets}_i(t) \neq \emptyset]$; $R_i(t) = \mathbf{1}[\exists w \in \mathcal{W} : \min_c \|c - w\| \leq \epsilon]$ ($\epsilon = 0.6$);
501 $M_i(t) = \mathbf{1}[\text{Locked}_i(t) \neq \emptyset \wedge \min_w \|\text{Locked}_i(t) - w\| \leq \epsilon]$.

502 **Episode-level starred events.** $Q_i^* = \bigvee_t Q_i(t)$; analogously for R^*, M^* ; $C_i^* =$
503 $\mathbf{1}[\text{agent } i \text{ reached a target}]$.

504 **Aggregation.** Per-run rates are averaged over the N agents; per-configuration rates are averaged
505 over the 40 seeds. Conditional rates $P(R^*|Q^*), P(M^*|R^*), P(C^*|M^*)$ are computed as ratios of
506 episode-level counts within each configuration before averaging across configurations.

507 C Memory and task schema

508 C.1 Full notation table

509 Symbols are fixed throughout the paper.

510 **Stages.** R, M, C — the four stages: **Q**uery formation, **R**etrieval, **M**aterialization, **C**ompletion after
511 retrieval. They name both the stage and (in expressions like “the M link”) the corresponding edge
512 of the chain.

513 **Events.** $R^*, M^*, C^* \in \{0, 1\}$ — episode-level starred events recording whether each stage succeeded
514 (Appendix B). Nested by construction: $Q^*=1 \Leftarrow R^*=1 \Leftarrow M^*=1 \Leftarrow C^*=1$.

515 $\text{Prob}(\cdot)$ — population-level probability (averaged over the seed distribution and, in the multi-agent
 516 a - case, over agents). $\hat{P}(\cdot)$ — empirical-frequency estimator computed from logs.
 517 q — query; $\mathcal{M}$ — memory state; Y — end-to-end task success (can exceed C^*).
 518 N — number of agents in the population. T — number of targets in the environment (constraint
 519 $T \leq N$; “scarcity” means $N > T$). i — agent index. $\mathcal{M}_i$ — agent i ’s private memory. q_i —
 520 a - agent i ’s query.
 521 $\mathcal{C}_K$ — the peer-broadcast communication protocol with cadence $K \in \mathbb{N}$ (every K ticks each agent
 522 pl broadcasts a memory snapshot, §4.1). $\mathcal{M}_{-i}^{(K)}$ — the latest peer memory snapshots accessible to
 523 agent i under $\mathcal{C}_K$ (the *memory coupling* channel). $\mathcal{O}_{-i}$ — the set of targets already claimed by
 524 other agents by the time i acts (the *occupancy coupling* channel).
 525 $\text{Slope}(\mathcal{M}_{-i}^{(K)})$ — a freshness functional on peer memory; non-decreasing pointwise as K decreases
 526 var assumption (F) of Proposition 3). The occupancy slope is assumption (O).
 527 $\text{Chain}(K) := P(M^*|R^*; K) \cdot P(C^*|M^*; K)$ — the conditional product, the chain probability
 528 mar $(M^* \wedge C^* | R^*)$. POM (product-of-means) and MOP (mean-of-products) are two estimators of
 529 $stat$ $\text{Chain}(K)$; they differ by $\text{Cov}(P(M^*|R^*), P(C^*|M^*))$, which is the Jensen gap (§4.5). t_{succ} — mean
 530 $tics$ time-to-success in environment ticks.

531 C.2 Memory stack and task patterns

532 The memory stack stores symbolic place records (per-place tag confidences, counts, visitation,
 533 freshness, geometry), transition records (success, risk, cost statistics with fast/slow estimates),
 534 episodic traces (intent + state at each visit), and concept/relation records. At each step, the tokenizer
 535 emits (z_t, τ_t, e_t) from o_t . Per-place tag tables use a Dirichlet-multinomial update; transitions use
 536 exponentially smoothed statistics; episodic records attach the active intent and agent state. A planner
 537 query is answered by retrieving candidate places whose semantic profiles match the current task intent
 538 and state, via a weighted log-score $s(p | q_t, s_t) = \sum_{\tau \in q^{\text{req}}} \log P(\tau|p) + \alpha \sum_{\tau \in q^{\text{pref}}} \log P(\tau|p) -$
 539 $\beta \sum_{\tau \in q^{\text{pen}}} \log P(\tau|p) + \gamma u_{\text{epi}}(p, s_t)$ over required, preferred, and penalty tags. The four task
 540 families are defined by semantic place patterns: water search (target {water}), safe rest ({rest} with
 541 state-conditioning), goal-region rejoin ({goal-region}), hazard recovery ({post-hazard-exit}).

542 C.3 Single-agent supporting tables and figures

Table 3: Best semantic result per task family from the final 10-seed benchmark, with 95% bootstrap CIs. Hazard-recovery is bottlenecked by retrieval; goal-region remains downstream-bottlenecked.

Task	Best method	Assist	Success (95% CI)	Steps	Query (op.)	Retrieval (op.)	Mat. (op.)	Post (op.)
Water	water-cp	on	0.95 [0.88, 1.00]	17.99	0.71	0.71	0.71	0.70
Rest	rest-s2	on	0.93 [0.85, 0.99]	23.35	0.61	0.61	0.61	0.73
Goal region	goal-n1	off	0.54 [0.52, 0.56]	49.34	0.00	0.00	0.00	0.00
Hazard recovery	hazard-s7	on	0.40 [0.35, 0.45]	10.18	0.11	0.11	0.11	0.16

Oracle stage interventions localise the bottleneck (single-agent, 10 seeds $\times$ 8 episodes)




Figure 5: Oracle stage interventions on the two hardest single-agent task families (10 seeds $\times$ 8 episodes). **(a)** FourRooms goal-rejoin: replacing any single stage with an oracle leaves end-to-end success at 0.07; the bottleneck is downstream of the stages in the chain. **(b)** LavaGapS7 hazard recovery: oracle controller adds only +0.04, but oracle materialization adds +0.16 and oracle retrieval adds a further +0.01 to reach 0.60. The framework localises the bottleneck stage-by-stage.

D Portability and scale-up

D.1 BabyAI portability transfer

A 5-seed portability check on BabyAI-GoToObjMaze-v0 (8 episodes/seed) confirms that the same Q/R/M/C events are recoverable from BabyAI traces and that the stage-rate ordering matches the MiniGrid benchmark within bootstrap CIs.

D.2 Multi-agent scale-up at $N = 12$

To check that the multi-agent results of §4.5 are not an artefact of the $N \leq 8$ range, an additional 8,640-run sweep was run at $N = 12$, $T \in \{6, 8, 10\}$, the same three layouts, three hazard densities, eight peer cadences, and 40 seeds. Wall-clock $\approx$ 12 minutes on 8 cores. The conditional-rate signature of §4.5 persists and intensifies (Table 4).

Table 4: $N = 12$ scale-up sweep numerical detail (peer architecture). Slope conditions hold: Spearman of $P(M^*|R^*)$ with K is -0.141 ($p < 10^{-39}$), and of $P(C^*|M^*)$ with K is $+0.304$ ($p < 10^{-184}$).

K	$P(M^* R^*)$	$P(C^* M^*)$	MOP	mean t_{succ}	corr
1	0.758	0.679	0.475	5.49	-0.59
2	0.745	0.692	0.474	5.29	-0.61
4	0.679	0.768	0.482	4.82	-0.64
8	0.731	0.767	0.539	4.77	-0.70
16	0.687	0.803	0.538	5.04	-0.62
32	0.676	0.814	0.538	5.75	-0.51
48	0.677	0.814	0.539	5.96	-0.51
64	0.677	0.816	0.541	6.12	-0.51

E Extended diagnostics

E.1 Conditional-product detail and the Jensen gap

The conditional product $\Phi = P(M^*|R^*) \cdot P(C^*|M^*)$ in the 35,640-run sweep is monotonically increasing in K over $K \in \{4, 8, 16, 32, 48, 64\}$ and converges to the independent-agent boundary $\Phi(K \rightarrow \infty) = 0.605$. Paired bootstrap: $\Delta(K=32-K=16) = +0.008$ (CI $[+0.006, +0.010]$), $\Delta(K=48-K=16) = +0.011$, $\Delta(K=64-K=16) = +0.012$; independent is significantly above every peer K except $K=64$ (borderline).




Figure 6: $N = 12$ scale-up bottleneck signature ($n = 8,640$ runs, 40 seeds). Top: both conditional rates change monotonically in K as predicted; the within-cadence Pearson correlation (red dashed, right axis) peaks at -0.70 at $K = 8$. Bottom: mean time-to-success retains an interior minimum at $K = 8$ (4.77 ticks). The bottleneck-shift peak migrates outward from $K = 4$ at $N \leq 8$ to $K = 8$ at $N = 12$, consistent with stronger occupancy coupling at higher agent count delaying the cost-regime transition.

560 The negative within-cadence covariance documented in §4.5 creates the Jensen-inequality gap
561 between Φ and the marginal product-of-means: $E[X] \cdot E[Y] > E[X \cdot Y]$ when $\text{Cov}(X, Y) < 0$. The
562 mean-of-products is reported as the primary summary because it absorbs the covariance correctly.
563 Figure 7 shows the K -sweep view of both conditional rates and mean t_{succ} ; Figure 8 shows the
564 per-configuration scatter that the negative covariance summarises.

565 E.2 Open problems for the collective-semantic-memory program

566 Four concrete questions follow directly from the measurements of §4.5 and the synthesis of §5. **(1)**
567 **Adaptive cadence.** Proposition 3 assumes a fixed K ; the bottleneck-shift signature suggests an
568 online controller K_t that broadcasts when the (estimated) M-side gain exceeds the C-side occupancy
569 cost. **(2) Selective consolidation.** Snapshot exchange merges everything; merge rules that prioritize
570 evidence about under-represented semantic tags attack the candidate-generation starvation of §3.5
571 directly. **(3) Trust-weighted merging under occupancy.** The central aggregator already performs
572 a trust-weighted merge but pays the C-collapse cost; decentralized trust rules that account for peer
573 *intent* (not only peer evidence) could decouple freshness from target racing. **(4) Substrate generality.**
574 The decomposition transfers across MiniGrid, BabyAI, the custom grid, and now a standard-MiniGrid
575 multi-agent wrapper with identical operational definitions and verified slope signs of Proposition 3
576 (§ 6). A continuous-state environment (Habitat, VMAS) is the remaining portability step.

Table 5: Per-architecture conditional Q/R/M/C rates from the 35,640-run sweep (40 seeds). $P(R^*|Q^*)$ is saturated. POM = product-of-means; MOP = mean-of-products (the proper per-configuration expected product). The two summaries disagree systematically at small/intermediate K because the within-cadence correlation between $P(M^*|R^*)$ and $P(C^*|M^*)$ is strongly negative there (last column) — the bottleneck-shift signature. **Bold:** MOP-best of the peer cadences. Mean t_{succ} (last column) is the practically relevant efficiency measure; its interior minimum at $K = 8$ is the predicted-shape result that survives the data.

Architecture	$P(R^* Q^*)$	$P(M^* R^*)$	$P(C^* M^*)$	POM	MOP (95% CI)	corr	mean t_{succ}
independent	0.99	0.68	0.88	0.598	0.605 [0.597, 0.614]	+0.00	8.02
shared bus	1.00	1.00	0.60	0.600	0.595 [0.586, 0.604]	−0.10	8.33
central aggregator	1.00	0.97	0.61	0.592	0.572 [0.564, 0.580]	−0.48	7.64
peer ($K = 1$)	1.00	0.97	0.60	0.585	0.573 [0.564, 0.581]	−0.43	7.80
peer ($K = 2$)	1.00	0.97	0.62	0.594	0.576 [0.568, 0.584]	−0.54	7.59
peer ($K = 4$)	1.00	0.92	0.66	0.599	0.569 [0.561, 0.577]	−0.61	8.00
peer ($K = 8$)	0.99	0.69	0.87	0.597	0.586 [0.578, 0.595]	−0.19	7.47
peer ($K = 16$)	0.99	0.67	0.89	0.592	0.590 [0.581, 0.598]	−0.05	7.87
peer ($K = 32$)	0.99	0.67	0.89	0.595	0.598 [0.589, 0.606]	−0.04	8.28
peer ($K = 48$)	0.99	0.68	0.88	0.598	0.602 [0.593, 0.611]	−0.02	8.79
peer ($K = 64$)	0.99	0.68	0.88	0.599	0.603 [0.594, 0.611]	−0.02	9.19

577 E.3 Extended related work

578 Classical cognitive-map perspectives motivate explicit place abstractions (Tolman, 1948; O’Keefe
579 and Nadel, 1978; Kuipers, 2000); modern embodied AI combines mapping and planning (Gupta
580 et al., 2017) or exposes topological memory (SPTM, Savinov et al., 2018); goal-conditioned RL
581 conditions on explicit goals or embeddings (Andrychowicz et al., 2017; Ghosh et al., 2021), with
582 successor-representation work providing predictive structure (Dayan, 1993; Stachenfeld et al., 2017;
583 Momennejad et al., 2017). Our setting differs in that the target is a semantic pattern retrieved from
584 memory rather than a coordinate. In multi-agent RL, CTDE dominates (Lowe et al., 2017; Foerster
585 et al., 2018) and learned communication adds a per-step exchange variable (Sukhbaatar et al., 2016;
586 Foerster et al., 2016); gossip and consensus protocols (Boyd et al., 2006; Xiao et al., 2007) are the
587 closest network-level analogue of our cadence axis, lifted here to the cognitive-architecture level.
588 LLM-agent systems lean on explicit memory but evaluate end-to-end (Wang et al., 2023; Shinn et al.,
589 2023; Packer et al., 2023). Our framework is orthogonal: rather than learning a communication
590 policy end-to-end, it *diagnoses* where in the Q/R/M/C chain a given architecture succeeds or fails,
591 and yields a slope-level prediction about cadence that is testable on logs.

592 E.4 Multi-agent oracle-intervention numerics

Table 6: Multi-agent oracle interventions on the scarcity cases at three cadences. Mean time-to-success averaged across $(N, T) \in \{(3, 2), (5, 3), (8, 5)\}$, two layouts (asymmetric, random), 15 seeds.

K	Baseline	Oracle R	Oracle M	Gap to oracle
1	10.33	7.22	7.08	−3.11
4	10.97	7.22	7.08	−3.75
16	7.98	7.22	7.08	−0.76

593 E.5 Stage-agnostic cadence hypotheses and what the framework changes

594 A first version of the multi-agent cadence sweep (12,960 runs at 20 seeds; subsumed by the 35,640-run
595 sweep in this paper) was originally planned with four operational hypotheses formulated *before*
596 Proposition 3. They are retained here because they motivated the eventual framework. The framework
597 was formulated after the sweep was complete; it is not claimed to have predicted the sweep numbers
598 a priori (see “Status and chronology” in §4.3).

599 Original hypotheses (pre-framework).

600 **H1.** An interior $K^* \in \{2, 4, 8\}$ minimising mean time-to-success exists for most (N, T, layout)
601 cells.




Figure 7: Bottleneck shift across eight peer cadences (40-seed sweep). Top: $P(M^*|R^*)$ falls and $P(C^*|M^*)$ rises monotonically in K ($p < 10^{-4}$, both Spearman). Bottom: mean time-to-success has an interior minimum at $K = 8$ (7.47 ticks), below the independent-agent baseline (8.02, dashed). Error bars and shaded band are 95% bootstrap CIs.

602 **H2.** K^* grows monotonically with scarcity $\rho = N/T$.

603 **H3.** When $\rho \leq 1$, $K^* = 1$.

604 **H4.** Peer-broadcast architectures Pareto-dominate centralized aggregation on (mean time-to-success,
605 success rate).

606 **Stand-alone result on the original 20-seed sweep.** The practical claim was supported: peer at the
607 empirically best K was faster than centralized in the scarcity regime ($p = 0.044$, Mann–Whitney).
608 Among the four structural hypotheses, only H4 was clearly supported in the random-layout subset.
609 H1 held in 7 of 18 cells; H2 had a non-significant negative correlation; H3 was supported in 4 of 9
610 admissible cells. The 35,640-run extended sweep used elsewhere in the paper produces matching
611 stand-alone numbers within bootstrap uncertainty.

612 **What the framework changes.** The hypotheses H1–H4 are stage-agnostic: each ties K^* to scarcity
613 ρ without specifying mechanism. Proposition 3, formulated after the sweep, ties the slopes of
614 $P(M^*|R^*)$ and $P(C^*|M^*)$ to cadence; the slope conditions are testable. The data verify both signs
615 ($p < 10^{-4}$), but the corollary about an interior maximum on the conditional product fails. The interior
616 optimum holds on the efficiency metric mean time-to-success.

617 F Assist on/off analysis

618 Across all four single-agent task families, top-line success is essentially unchanged under assist on/off
619 ($\Delta \leq 0.10$). The only non-trivial effect is on rest, where assists improve post-retrieval completion
620 by +0.05 without changing query/retrieval rates, consistent with assists acting on materialization and
621 post-retrieval execution.

Bottleneck-shift signature: negative within-cadence correlation peaks at $K = 4$ and decays at $K = 64$




Figure 8: Per-configuration scatter of $P(M^*|R^*)$ vs. $P(C^*|M^*)$ at three cadences. The strong negative correlation at $K = 4$ (Pearson -0.61) is the bottleneck-shift signature: the same configurations that succeed at M tend to fail at C and vice versa. The correlation decays to noise by $K = 64$.

G Retrained retriever and CommNet baseline

G.1 Retrained candidate-generation MLP (§3.5)

Dataset extraction. From the 10-seed hazard-recovery benchmark (MiniGrid-LavaGapS7-v0, method `milestone_state_conditioned_hazard_recovery_v7`), an 18-dimensional feature vector was extracted for every step of every episode of every seed — 3 numeric agent-state scalars (energy, thirst, risk_budget), 9 semantic-tag confidences, 6 intent one-hot — and a binary label: *was this symbolic node ever selected as a candidate during the run?* Dataset totals: 814 examples, 113 positives (13.9%).

Split. By seed: train $\{2, 3, 4, 5, 6, 7\}$ (505 examples, 9.7% positive), val $\{8, 9\}$ (170, 22.4% positive), test $\{10, 11\}$ (139, 18.7% positive).

Architecture. $18 \rightarrow 32 \rightarrow 32 \rightarrow 1$, ReLU activations, BCE with $\text{pos_weight} \approx 9.3$ for class imbalance, Adam lr = $2 \cdot 10^{-3}$ with $\text{weight_decay} = 10^{-4}$, batch 64, 200 epochs, best by val loss.

Results.

Metric	Train	Val	Test
Accuracy	0.584	0.553	0.590
Majority-class baseline accuracy	0.903	0.776	0.813
Precision @ recall = 0.70	0.171	0.278	0.275

The MLP underperforms majority-class on accuracy across all splits and achieves only moderate precision at fixed recall. Instantaneous state features are insufficient to predict candidate eligibility downstream. This is the negative-but-informative finding referenced in §3.5: the framework’s “retrieval-limited” attribution refines to “memory-accumulation-limited” rather than “ranking-limited”.

G.2 CommNet-style learned communication baseline (§4.5)

Policy architecture. Per-agent observation is 4-dim direction one-hot + $5 \times 5 \times 5$ local cell-type window (125) + 2-dim normalised position = 131. Shared-weights encoder $131 \rightarrow 64 \rightarrow 64$ (ReLU), 16-dim comm projection, peer mean-pool (excluding self), combine $80 \rightarrow 64$ (ReLU), actor head $64 \rightarrow 4$ (TURN_LEFT, TURN_RIGHT, FORWARD, NOOP), critic head $64 \rightarrow 1$.

Training. REINFORCE with value-baseline; loss = $-\text{mean}(\log \pi(a) \cdot (R - V)) + 0.5 \cdot \text{MSE}(V, R) - 0.01 \cdot H[\pi]$; Adam lr = $3 \cdot 10^{-4}$, gradient clipping 1.0. Reward shaping: +1.0 on first success per agent, -0.01 step cost, -0.1 per hazard contact. 2000 episodes cycling 200 environment seeds, wall-clock 100 s on CPU. Held-out eval: 50 environment seeds 250–299.




Figure 9: Refined attribution diagnostic. *Left*: 349 retrieval calls aggregated across 10 seeds; 317 (90.8%) returned empty candidate sets, 31 (8.9%) returned at least one satisfied candidate, only 1 (0.3%) returned candidates that were ranked but unsatisfied. The bottleneck is candidate generation, not candidate ranking. *Right*: the trained MLP underperforms a majority-class baseline on all three seed-disjoint splits, confirming that instantaneous state features alone are insufficient to predict candidate eligibility.

648 **Training curve.** Episode-level success (last-100 window) stays in $[0.13, 0.23]$ throughout training;
649 no monotonic improvement after ~ 200 episodes. The policy converges to a local optimum where
650 one of three agents finds water by chance.




Figure 10: Learned-communication baseline diagnostics. *Left*: CommNet training curve (smoothed success rate over a 100-episode window) over 2000 training episodes; success stays below 0.25 throughout. The dashed line is the symbolic peer architecture at $K=8$ on the same scenario for reference. *Right*: Q/R/M/C profile of the trained CommNet on 50 held-out seeds vs symbolic peer at $K=8$. Q^* is saturated at 1.0 on CommNet (degenerate by construction); R/M collapse together; only on the symbolic architecture are the conditional rates separated.

651 **Q/R/M/C evaluation (50 held-out seeds, stochastic action sampling).**

652 Compare to the symbolic peer architecture at $K=8$: $P(M^*|R^*)=0.69$, $P(C^*|M^*)=0.87$, success
653 ~ 0.65 .

654 **PPO variant: structural degeneracy verified at competitive success.** A PPO-trained variant
655 of the same CommNet architecture (clip $\epsilon=0.2$, GAE $\lambda=0.95$, $\gamma=0.99$, 4 PPO epochs per roll-
656 out, 64 rollouts per update, 200 updates, ≈ 96 minutes on a single CPU core; full setup in

Quantity	Value
Q^* rate	1.00 (saturated)
R^* rate	0.50
M^* rate	0.50 (collapses onto R)
C^* rate	0.19
$P(R Q)$	0.50
$P(M R)$	1.00
$P(C M)$	0.42
Success rate	0.19
Mean steps	80 (budget exhausted)

657 experiments/commnet_ppo_baseline/train_ppo_commnet.py) reaches 0.667 success on 100
658 held-out seeds — competitive with the symbolic peer architecture at $K=8$ (0.65).

Table 7: REINFORCE vs PPO CommNet vs symbolic peer ($K=8$) on the scarcity scenario ($N=3$, $T=2$, asymmetric, hazard 0.05). PPO is competitive on success rate, yet still exhibits Q-saturation and R/M collapse — the diagnostic claim verified, not asserted.

Metric	CommNet REINFORCE	CommNet PPO	Symbolic peer ($K=8$)
Success rate	0.19	0.667	0.65
Q^* rate	1.00	1.00	0.99
R^* rate	0.50	0.997	0.69
M^* rate	0.50	0.997	0.69
$P(M^* R^*)$	1.00	1.00	0.69
$P(C^* M^*)$	0.42	0.67	0.87
R/M separated?	no	no	yes

659 The PPO result converts the structural argument from an assertion (“training budget would not change
660 this”) to a measurement: Q^* remains saturated at 1.00, $P(M^*|R^*)$ remains 1.00 (R/M collapse), even
661 at competitive success. The Q stage is degenerate by construction in any policy whose communication
662 channel is a real-valued vector with a learned linear projection — the channel is never informatively
663 empty, regardless of training budget — and without an explicit symbolic target-lock interface the M
664 event cannot separate from R. Only architectures with an explicit *query/lock* interface (the symbolic
665 stack of §4.1) expose stage-separable Q/R/M/C events.

666 H Reproducibility commands

667 An anonymized public repository accompanies this submission. The main paper artifacts were
668 generated from the following scripts; PYTHONPATH=. and a Python 3 virtual environment with numpy,
669 scipy, pandas, matplotlib, imageio, gymnasium, and minigrid are assumed throughout.

670 One-command reproduction (multi-agent block, ≈ 22 minutes on 8 cores)

671 bashscripts/reproduce_paper.sh

672 Single-agent (§3.2–§3.4)

- 673 • article benchmark and learned-baseline suite: scripts/benchmark_symbolic_memory_
- 674 article.py
- 675 • article analysis and figure generation: scripts/analyze_symbolic_memory_article.py
- 676 • synthetic factorization sanity check: scripts/validate_qrmc_factorization.py
- 677 • oracle-stage intervention benchmark: scripts/benchmark_oracle_stage_interventions.
- 678 py
- 679 • BabyAI transfer benchmark: scripts/compare_semnai_babyai.py

680 Multi-agent (§4.5)

- 681 • Main 25,920-run cadence sweep (40 seeds, $K \leq 16$, ≈ 18 min): `pythonexperiments/`
682 `big_experiment/exp_cadence_phase.py--modefull--workers8--out_dirtmp/big_`
683 `experiment_qrmc_40`
- 684 • Extended 9,720-run sweep ($K \in \{32, 48, 64\}$, ≈ 4 min): `pythonexperiments/big_`
685 `experiment/exp_extra_K.py--workers8--out_dirtmp/big_experiment_extraK`
- 686 • Q/R/M/C stage analysis and bootstrap CIs: `pythonexperiments/big_experiment/`
687 `analyze_qrmc.py--runs_csvtmp/big_experiment_qrmc_40/runs.csv--out_`
688 `dirtmp/big_experiment_qrmc_40`
- 689 • Multi-agent oracle interventions: `pythonexperiments/big_experiment/exp_oracle_`
690 `interventions.py--workers8--out_dirtmp/big_experiment_oracle`
- 691 • Phase 1–4 smoke tests: `pythonscripts/run_all_smokes.py--out_dirtmp/all_smokes`
- 692 • PPO-trained CommNet baseline (Table 7, ≈ 96 min on a single CPU core):
693 `python-mexperiments.commnet_ppo_baseline.train_ppo_commnet--total_`
694 `updates200--rollouts_per_update64`
- 695 • Q/R/M/C eval of the PPO policy on 100 held-out seeds: `pythonexperiments/commnet_`
696 `baseline/eval_with_qrmc.py--policy_pathexperiments/commnet_ppo_baseline/`
697 `ppo_policy.pt--out_direxperiments/commnet_ppo_baseline--n_episodes100`
- 698 • MiniGrid-substrate portability sweep (§6, 100 runs, ≈ 2 min): `python-mexperiments.`
699 `minigrid_multiagent_wrapper.run_portability_sweep--out_dirtmp/minigrid_`
700 `portability`

701 I Asset and license note

702 The experiments use MiniGrid and BabyAI, both released under the MIT License. The locally
703 implemented baselines and benchmark runner scripts introduced in this paper will be released under
704 the MIT License as part of the camera-ready artifact package.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: The abstract and introduction present the paper primarily as a four-stage diagnostic framework with a reference symbolic-semantic instantiation, and they explicitly scope the evidence to benchmarked navigation environments with a limited portability check (Abstract; Sections 1 and 6).

Guidelines:

- The answer [\[N/A\]](#) means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [\[No\]](#) or [\[N/A\]](#) answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: The paper includes an explicit “Limitations and future work” section that discusses environment difficulty, the observable-query and observable-occupancy assumptions, the slope-vs-corollary gap of Proposition 3, the learned-baseline scope, and agent-count scaling (Section 6).

Guidelines:

- The answer [\[N/A\]](#) means that the paper has no limitation while the answer [\[No\]](#) means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: The paper states the conditional stage-Markov property and positivity assumptions in Section 2.2, gives a proof sketch in the main text, and includes a full formal proof of Proposition 1 in Appendix A together with the multi-agent extension (Proposition 2) and the slope-level argument for Proposition 3.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The paper specifies the task families, environments, seeds, episode counts, assist modes, baselines, statistical protocol, and benchmark semantics; the single-agent setup is in Section 3.1 and the multi-agent sweep protocol in Section 4.4; the operational stage definitions used to extract Q^* , R^* , M^* , C^* from runtime logs are in Appendix B.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The paper describes a unified benchmark runner and source-of-truth artifacts. An anonymized public release of the codebase — single-agent benchmark, multi-agent grid environment, 35,640-run sweep driver, Q/R/M/C analyser, oracle-intervention scripts, and exact commands to reproduce every table and figure — accompanies the submission as supplemental material (anonymized GitHub repository; single-command reproduction of the main multi-agent sweep on an 8-core machine in ≈ 22 minutes, matching Appendix H).

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: The benchmark protocol, task definitions, assist settings, seeds, episode counts, baseline families, and the learned baselines are described in the main text together with the measurement and statistical protocols: single-agent setup in Section 3.1, multi-agent sweep in Section 4.4, and per-baseline hyperparameters in Appendix G.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: The paper reports 95% bootstrap confidence intervals over seed-level aggregates ($B=5,000$ resamples; Section 3.1, “Protocol” paragraph) and uses paired Wilcoxon signed-rank tests with Holm-Bonferroni correction for assist comparisons. CIs are shown in Tables 3 and 5 (Best semantic per-task rates; the 35,640-run MOP column) and in Figures 4–6 (bootstrap shaded bands).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: The paper discloses the hardware platform (10-core ARM CPU, 32 GB RAM; CPU-only), the scale of each benchmark, and wall-clock costs — single-agent oracle sweep < 1 minute (Section 3.1); main multi-agent sweep (35,640 runs) ≈ 22 minutes on 8 cores (Section 4.4); $N=12$ scale-up (8,640 runs) ≈ 12 minutes (Appendix D); PPO CommNet baseline (200 updates $\times$ 64 rollouts) ≈ 96 minutes on a single CPU core (Appendix G); MiniGrid portability sweep (100 runs) ≈ 2 minutes on a single core (§6).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn’t make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The work is a simulation-only study in standard benchmark environments, uses anonymized authorship for review, and does not involve human subjects, private data, or deployment claims that conflict with the NeurIPS Code of Ethics.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: The paper includes a dedicated “Broader impacts” section that discusses the primary methodological benefits of the work and the indirect risks of improving navigation systems without adequate safeguards (Section 7).

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper does not release high-risk models or scraped datasets; it studies navigation policies in simulated environments.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.

- 972 • We recognize that providing effective safeguards is challenging, and many papers do
973 not require this, but we encourage authors to take this into account and make a best
974 faith effort.

975 **12. Licenses for existing assets**

976 Question: Are the creators or original owners of assets (e.g., code, data, models), used in
977 the paper, properly credited and are the license and terms of use explicitly mentioned and
978 properly respected?

979 Answer: [\[Yes\]](#)

980 Justification: The paper cites the upstream benchmark platforms (MiniGrid, BabyAI) in
981 the main references. Appendix I states that both are MIT-licensed and that the benchmark
982 scripts introduced here will be released under the MIT License at camera-ready.

983 Guidelines:

- 984 • The answer [\[N/A\]](#) means that the paper does not use existing assets.
- 985 • The authors should cite the original paper that produced the code package or dataset.
- 986 • The authors should state which version of the asset is used and, if possible, include a
987 URL.
- 988 • The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- 989 • For scraped data from a particular source (e.g., website), the copyright and terms of
990 service of that source should be provided.
- 991 • If assets are released, the license, copyright information, and terms of use in the
992 package should be provided. For popular datasets, paperswithcode.com/datasets
993 has curated licenses for some datasets. Their licensing guide can help determine the
994 license of a dataset.
- 995 • For existing datasets that are re-packaged, both the original license and the license of
996 the derived asset (if it has changed) should be provided.
- 997 • If this information is not available online, the authors are encouraged to reach out to
998 the asset's creators.

999 **13. New assets**

1000 Question: Are new assets introduced in the paper well documented and is the documentation
1001 provided alongside the assets?

1002 Answer: [\[No\]](#)

1003 Justification: The submission introduces benchmark scripts and analysis artifacts. A formal
1004 anonymized release package with documentation is planned for camera-ready; Appendix H
1005 lists all scripts needed to reproduce each reported result, and Appendix I states the intended
1006 license for the new assets.

1007 Guidelines:

- 1008 • The answer [\[N/A\]](#) means that the paper does not release new assets.
- 1009 • Researchers should communicate the details of the dataset/code/model as part of their
1010 submissions via structured templates. This includes details about training, license,
1011 limitations, etc.
- 1012 • The paper should discuss whether and how consent was obtained from people whose
1013 asset is used.
- 1014 • At submission time, remember to anonymize your assets (if applicable). You can either
1015 create an anonymized URL or include an anonymized zip file.

1016 **14. Crowdsourcing and research with human subjects**

1017 Question: For crowdsourcing experiments and research with human subjects, does the paper
1018 include the full text of instructions given to participants and screenshots, if applicable, as
1019 well as details about compensation (if any)?

1020 Answer: [\[N/A\]](#)

1021 Justification: The paper does not involve crowdsourcing or human-subject experiments.

1022 Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: The paper does not involve human subjects or human-participant data collection.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: LLMs are not an important, original, or non-standard component of the core methods or experiments reported in the paper.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.