
Q/R/M/C: A Stage-Decomposition Measurement Framework for Memory-Based Navigation

Anonymous Author(s)

Affiliation

Address

email

Abstract

Setting. Success-only metrics collapse qualitatively different memory-based navigation failures into a single scalar, hiding *which stage* of the retrieval pipeline is responsible. This paper contributes an *evaluation protocol*: the Q/R/M/C stage decomposition (Query formation, Retrieval satisfaction, target Materialization, Completion after retrieval), operationalised so that the four conditional rates are computable directly from standard trajectory logs (Definition 1). The arithmetic is the chain rule, deliberately so — just as decomposing accuracy into precision and recall added no mathematics but changed what the field could *see*, the value here is identifiability from logs and what the decomposition reveals, not a theorem. **What it decomposes** is important to state up front: the product of the four rates converges to *semantic-path completion* C^* , not to raw task success Y ; because an episode can succeed without traversing the retrieval path, Y can exceed C^* , so Q/R/M/C localises *where the semantic route fails*, not every route to success. **Single-agent measurement.** On a 10-seed MiniGrid benchmark (FourRooms, LavaGapS7) and a BabyAI portability check, the framework localises bottlenecks: hazard recovery is retrieval-limited (oracle retrieval lifts success $0.39 \rightarrow 0.60$), goal-region rejoin is downstream-limited. Log inspection refines the attribution: candidate ranking is near-optimal (96.9%) while candidate generation fails at 91% of decision points — a memory-accumulation, not a ranking, limit. **Multi-agent measurement.** On a 35,640-run multi-agent sweep over four communication architectures (independent, shared-bus, central-aggregator, peer-broadcast at eight cadences $K \in \{1, 2, 4, 8, 16, 32, 48, 64\}$, 40 seeds), the bottleneck-shift between the M and C links appears directly as a negative within-cadence correlation peaking at $K=4$ (Pearson -0.61); mean time-to-success has an interior minimum at $K=8$ (robust in location but shallow in margin — not resolved from the adjacent fast cadences). Both slope signs (Empirical Claim 1) hold as moderate, cluster-robust effect sizes (Spearman -0.53 and $+0.43$ vs. K , 95% CIs $[-0.59, -0.47]$ and $[+0.37, +0.50]$ over 81 configuration cells; Appendix F.3). Non-claims (notably on the conditional product Φ) are consolidated in one place (§1, Scope and non-claims). **Diagnostic use.** The same four rates also expose why learned policies with implicit communication can be competitive on success while collapsing Q/R/M/C observability: their communication vectors are never informatively empty and they lack an explicit target-lock interface. We scope this claim explicitly: Q/R/M/C is separately observable only for systems with an explicit query/lock (materialization) interface; for opaque learned policies the stages are not individually recoverable, so the framework is *complementary to*, not a replacement for, probing, causal-mediation, and message-ablation diagnostics. On two local LLM agents (qwen3:4b, llama3.1 on TextNav), the same logger localizes a completion-limited failure: under a tight step budget both agents form the query, retrieve, and lock the target ($Q^*=R^*=M^*=1$) yet none complete ($C^*=0$), so an otherwise opaque




Figure 1: Q/R/M/C decomposes one semantic-navigation episode into four observable stages. **Top:** the four stages and their event-level definitions. **(a)** A successful water-search episode anchors at one event per stage, producing logged starred events Q^* , R^* , M^* , C^* . **(b)** The same scalar “end-to-end success” summary masks qualitatively different failures on the four single-agent task families: water and rest are strong; goal-region rejoin (FourRooms) is downstream-limited; hazard recovery (LavaGapS7) is retrieval-limited (oracle interventions, §3.4).

success collapse is resolved into a specific failing stage rather than a leaderboard number. Thus the contribution is a measurement protocol for making memory failures comparable across single- and multi-agent symbolic and LLM systems wherever a query/lock interface is exposed, and for measuring stage observability itself on opaque learned policies. **Scope.** The primary experiments are in grid-world substrates (MiniGrid, BabyAI, a custom 12×10 grid). The contribution is the diagnostic framework and its validation; a preliminary continuous-state check on VMAS reproduces both slope signs directionally (Appendix F.7), and a full continuous/3D validation with cluster-robust CIs (Habitat, ProcTHOR) is future work.

1 Introduction

Classical navigation benchmarks often assume that the agent is given an explicit destination coordinate or a dense task reward. Many practical navigation problems are different: the agent must reach a place of a certain *type*, such as a safe rest area, a water-like resource, a post-hazard recovery exit, or a goal-associated region, specified by a semantic pattern rather than a fixed (x, y) waypoint. The place must therefore be *remembered and retrieved by meaning*. This setting opens a methodological question: can we measure *where* memory-based semantic navigation breaks down, not only *whether* it succeeds?

We answer this question with a four-stage diagnostic framework that applies to both single agents and multi-agent collectives. The decomposition is Q/R/M/C: **Q**uery formation, **R**etrieval satisfaction, **M**aterialization, and **C**ompletion after retrieval. The decomposition factorizes per-episode semantic-path completion and, under a conditional stage-Markov property, its factors are estimable from trajectory logs (Definition 1). The rest of the paper develops a single argument in three parts.

Single agent. Stage attribution is non-trivial. Success-only metrics collapse qualitatively different failures into a single scalar. In a 10-seed MiniGrid benchmark, water and rest tasks are strong end-to-end (0.95, 0.93 success), but goal-region rejoin and hazard recovery are substantially harder, and for *different reasons*. Oracle-stage interventions show that hazard recovery is retrieval-limited (oracle retrieval lifts success $0.39 \rightarrow 0.60$), while goal-region rejoin is downstream-limited (oracle retrieval rescues semantic-path completion without changing end-to-end success). Inspection of the framework’s own logs then sharpens the attribution: the retrieval bottleneck is not candidate *ranking*

(precision 96.9% when candidates exist) but candidate *generation*: the memory contains no eligible symbols at 91% of decision points. The single-agent bottleneck is a *memory-accumulation* limit: evidence is observed but never consolidated into queryable structure.

Multi-agent. Consolidation acquires a rate. When N agents share memory through a broadcast protocol with cadence K , the factorization extends with enlarged conditioning sets (Definition 2), and K becomes a structural axis: it is, mechanically, a knob on *how fast private memories consolidate into a collective one*. A slope-level claim (Empirical Claim 1) predicts a bottleneck shift: as K decreases, $P(M^*|R^*)$ rises (fresher peer memory enables reliable commitment) and $P(C^*|M^*)$ falls (peers race for the same targets). Both slope signs are confirmed on a 35,640-run sweep as moderate, cluster-robust effect sizes (Appendix F.3), the shift is directly visible as a negative within-cadence correlation between the two links, and the resulting trade-off produces an interior optimum of mean time-to-success at $K=8$: consolidation has a non-trivial optimal rate.

Diagnostic synthesis. Both regimes point at the same measurement object. The single-agent diagnosis says that evidence consolidation into symbols can be the binding constraint; the multi-agent diagnosis says that the rate and cost of sharing consolidated memory determine the optimum. Q/R/M/C is the common instrument: candidate systems can be compared by how they move $P(R^*|Q^*)$, $P(M^*|R^*)$, and $P(C^*|M^*)$, even when their end-to-end success rates look similar.

Scope and non-claims (consolidated). All caveats of this paper in one place; the running text no longer repeats them. Each limitation is stated together with the reason the results survive it. (i) The contribution is the operationalisation: log-level definitions under which the four conditional rates are identifiable on any system exposing a query/lock interface (Definitions 1–2). The consistency argument itself is the chain rule plus frequency estimation, and no mathematical novelty is claimed for it. (ii) Empirical Claim 1 is a slope-level empirical observation rather than a theorem, stated deliberately so that each slope is separately falsifiable on logs; both slopes were confirmed with cluster-robust confidence intervals (Appendix F.3). (iii) We make no claim about the shape of the conditional product Φ ; in our data it is monotone in K , and the interior optimum lives on mean time-to-success (§4.5). (iv) Primary substrates are grid worlds because they admit oracle interventions and exhaustive sweeps that richer substrates do not; VMAS is a directional continuous-state check, and the LLM studies are small and controlled (two local models on TextNav; an $N=3$ collective; a three-framework external validation at $n=20$ per cell with two models, §5). (v) Multi-agent place identity assumes a shared coordinate frame for the main sweep (relaxed to semantic, frame-free matching in Appendix E.3); the assumption affects the agents’ merge, not the event definitions. (vi) The four-rate product converges to *semantic-path completion* C^* , not to raw task success Y : Y can exceed C^* when an episode succeeds without traversing the retrieval path. This is by design — Q/R/M/C localises where the *semantic route* fails, which a success-only metric conflates with every other route to success. (vii) Stage-separability requires an explicit query/lock interface; on opaque learned policies the stages are not individually observable and the framework is *complementary to*, not a replacement for, probing, causal-mediation, and message-ablation diagnostics. A stage-level operationalisation is testable on specific summaries and can be *wrong* on others — (iii) is an instance we report — which is what separates a measurement protocol from a stage-agnostic guess.

Why this framing is timely. Embodied and goal-conditioned agents increasingly retrieve places, objects, or tools by meaning (Andrychowicz et al., 2017; Ghosh et al., 2021; Savva et al., 2019). LLM-agent systems lean on explicit memory layers but evaluate mostly end-to-end (Wang et al., 2023; Shinn et al., 2023; Packer et al., 2023). Multi-agent reinforcement learning likewise increasingly involves communicating populations (Lowe et al., 2017; Foerster et al., 2018), but the question “which agent should know what, and when?” lacks a unified diagnostic vocabulary. Our framework supplies one.

Contributions. This is a methodology paper. We are explicit about what is bookkeeping vs. empirical claim vs. instantiated artefact.

1. Measurement protocol (single-agent). The Q/R/M/C decomposition is operationalised so that the conditional rates are estimable from trajectory logs (Definition 1); the contribution is the *logging discipline* that makes the decomposition usable on standard navigation traces (non-claims: see box above).

Table 1: One contract, five system classes. The same four events (Q^* , R^* , M^* , C^*), defined once, instrument each class without modification — and localize a different failure in each.

System class	Substrate	What the same four events revealed
Symbolic single agent	MiniGrid, BabyAI	ranking near-optimal (96.9%) while candidate generation is empty at 91% of decisions: a memory- <i>accumulation</i> limit (§3)
Symbolic collective (4 architectures $\times$ cadence K)	grid, 35,640 runs	$M \leftrightarrow C$ bottleneck shift; interior t_{succ} optimum at $K=8$; provenance mechanism (§4)
Learned policies (BC; CommNet-PPO)	grid	competitive success with structurally collapsed stage observability (§3, §4)
LLM agents (2 models; $N=3$ collective)	TextNav	an opaque $1.00 \rightarrow 0.00$ collapse resolved to pure C-budget failure; sharing lowers success $2/3 \rightarrow 1/3$ (§4)
Third-party stacks (OpenAI SDK, LangGraph, AutoGen) $\times$ 2 models	tool traces	structural failures localize identically (20/20); first-commitment accuracy separates models; frameworks reshape it (§5)

2. **Measurement protocol (multi-agent).** An extension to N communicating agents with two coupling channels (memory $\mathcal{M}_{-i}^{(K)}$, occupancy $\mathcal{O}_{-i}$) under a peer-broadcast protocol with cadence K (Definition 2).
 3. **Empirical Claim 1 (bottleneck shift).** On a 35,640-run multi-agent sweep, the two stylised slope conditions (F) and (O) are both confirmed as moderate, cluster-robust effect sizes (Spearman -0.53 and $+0.43$ vs. K , 95% CIs excluding zero over 81 configuration cells; Appendix F.3; full benchmark in §4.5), with a provenance-level mechanism (§4.5).
 4. **System and benchmarks.** A symbolic-semantic memory stack with four communication architectures (independent, shared-bus, central-aggregator, peer-to-peer broadcast); single-agent oracle benchmark (10-seed MiniGrid, BabyAI portability); 35,640-run multi-agent sweep at $N \in \{3, 5, 8\} + 8,640$ -run scale-up at $N=12$; 100-run portability sweep on a standard-MiniGrid wrapper.
 5. **Diagnostic comparison across architectures.** The same logging contract separates symbolic memory systems from learned policies with implicit communication: the latter can be competitive on success while collapsing Q/R/M/C stage observability. This is an evaluation result, not a claim that one architecture class should always dominate.
 6. **Scope and honesty.** The primary experiments are in grid-world substrates. We treat that as a methodology-paper choice and state it explicitly (§6); a preliminary VMAS check on a continuous substrate is included (Appendix F.7), and full continuous/3D validation (Habitat, ProcTHOR, MineDojo) is future work.
- Portability overview.** The central claim is not any single number below; it is that one logging contract transfers across system classes and returns a located diagnosis in each. Table 1 summarises this correspondence; each subsequent section instantiates one row.

2 The Q/R/M/C decomposition

2.1 Problem formulation and notation

The agent operates in a partially observable navigation environment. At each step it receives an observation o_t , maintains a memory state $\mathcal{M}_t$, and forms a query q_t encoding the current task intent. A retrieval step returns a candidate set of places matching q_t from $\mathcal{M}_t$; a chosen candidate is materialized into an actionable target; a local controller executes actions until either the target is reached or the episode times out. The target is not a coordinate but a semantic pattern, so success is a function of the chosen candidate’s semantic profile, not of any pre-supplied coordinate.

Habitat, AI2-THOR, and ProcTHOR made visually rich navigation evaluation standard (Savva et al., 2019; Kolve et al., 2017; Deitke et al., 2022). We deliberately focus on MiniGrid-class environments to obtain clean and reproducible stage measurements, and use BabyAI only as a

portability check. We do not claim to solve LLM-agent memory here; rather, we argue that the same Q/R/M/C decomposition is a useful template for evaluating memory-grounded agent decisions beyond navigation.

The central empirical question is therefore not just whether the agent succeeds, but *which stage fails first* when it does not. Two related measurement layers are used. **Operational rates** count, per query attempt, whether the query yielded a non-empty candidate set, whether the chosen candidate satisfied the semantic predicate, and whether it was materialized into an actionable target; *completion after retrieval* is an episode-level indicator. **Nested episode-level events** $Q^* \Rightarrow R^* \Rightarrow M^* \Rightarrow C^*$ are used in the factorization below and audited in §3.2. Full operational definitions are given in Appendix C.

Notation. The four stages are Q, R, M, C with episode-level starred events $Q^*, R^*, M^*, C^* \in \{0, 1\}$, nested by construction. $P(\cdot)$ denotes the population probability; $\hat{P}(\cdot)$ its empirical-frequency estimator. **The letter M always refers to the Materialization stage; the number of targets in the multi-agent environment is written T (not M , constraint $T \leq N$, “scarcity” = $N > T$).** Other symbols: $\mathcal{C}_K$ peer-broadcast protocol with cadence K , $\mathcal{M}_{-i}^{(K)}$ peer memory snapshots, $\mathcal{O}_{-i}$ peer occupancy, $\nu(\cdot)$ freshness functional, $\Phi(K) = P(M^*|R^*; K) \cdot P(C^*|M^*; K)$, t_{succ} mean time-to-success. Full symbol table in Appendix D.1.

2.2 Single-agent factorization

Let Y denote overall episode success, and let Q^*, R^*, M^*, C^* denote the nested episode-level semantic-path events. The framework factorizes *semantic-path completion* rather than raw task success:

$$Q^* \rightarrow R^* \rightarrow M^* \rightarrow C^*.$$

For any query q and memory state $\mathcal{M}$,

$$P(C^* = 1 \mid q, \mathcal{M}) = P(Q^* = 1 \mid q, \mathcal{M}) P(R^* = 1 \mid Q^* = 1, q, \mathcal{M}) \\ \cdot P(M^* = 1 \mid R^* = 1, q, \mathcal{M}) P(C^* = 1 \mid M^* = 1, q, \mathcal{M}).$$

Overall task success Y can exceed C^* because some episodes succeed without traversing the full semantic retrieval path.

Assumption 1 (conditional stage-Markov property). Conditioned on the current stage being successful, the probability of the next stage depends only on the current successful stage, the current query, and the current memory state. Formally,

$$P(R^* \mid Q^*, q, \mathcal{M}, \text{history}) = P(R^* \mid Q^*, q, \mathcal{M}),$$

and analogously for M^* and C^* .

Assumption 2 (positivity). Each conditioning event used above has non-zero probability on the support of the benchmark.

Definition 1 (single-agent stage estimability). Under Assumptions 1–2, the four conditional rates in the factorization above are well-defined population quantities and the empirical-frequency estimators are consistent estimators of them; their product converges to semantic-path completion. The argument is elementary — the chain-rule factorization plus consistency of frequency estimators on Bernoulli events — and the value of stating it is operational, not mathematical: it pins down the *logging discipline* (which events, conditioned on what) under which the decomposition is readable from standard trajectory logs, and it is the contract that every experiment in this paper audits against (§3.2). Full statement and elementary argument in Appendix A.

When Assumption 1 may fail. Assumption 1 holds tightly here because q and $\mathcal{M}$ are explicitly logged at every stage event. Two practically relevant violations remain: adaptive in-episode query rewrites not surfaced to the logger, and hidden between-stage memory updates. In those cases the factorization should be read as a projection onto observable traces rather than as a full structural identification claim. Appendix F.6 quantifies the estimator bias under a controlled violation of this kind: it is negligible for mild unlogged couplings and one-signed (an upward bias on $P(C^*|M^*)$), which bounds the framework’s exposure when it is applied to learned or LLM agents.




Figure 2: Stage-structured interpretation of the Q/R/M/C framework. Assist on/off ablations target intermediate mechanisms; in the multi-agent extension the cadence protocol $\mathcal{C}_K$ acts on the $R \rightarrow M$ and $M \rightarrow C$ edges.

Mechanism-level reading. The same decomposition admits a stage-structured mechanism interpretation. Figure 2 treats stage outcomes as a directed acyclic graph. Under this interpretation, assist on/off ablations are mechanism-targeted perturbations of materialization and post-retrieval execution rather than undifferentiated sensitivity analyses; in the multi-agent setting the same graph is replicated per agent and connected through the cadence protocol $\mathcal{C}_K$, which acts on $\mathcal{M}_{-i}^{(K)}$ at the $R \rightarrow M$ edge and on $\mathcal{O}_{-i}$ at the $M \rightarrow C$ edge — exactly the two edges where Empirical Claim 1 will predict the slope shift.

Relation to prior work. Closest are (i) cognitive-map and semantic-map frameworks (Tolman, 1948; O’Keefe and Nadel, 1978; Kuipers, 2000; Gupta et al., 2017; Savinov et al., 2018), (ii) goal-conditioned and successor-representation RL (Andrychowicz et al., 2017; Ghosh et al., 2021; Dayan, 1993; Stachenfeld et al., 2017; Momennejad et al., 2017), (iii) multi-agent communication (Lowe et al., 2017; Foerster et al., 2018; Sukhbaatar et al., 2016; Foerster et al., 2016) and gossip/consensus (Boyd et al., 2006; Xiao et al., 2007), and (iv) LLM-agent memory layers evaluated end-to-end (Wang et al., 2023; Shinn et al., 2023; Packer et al., 2023). Our framework is orthogonal to all four: it *diagnoses* where in the Q/R/M/C chain a given architecture succeeds or fails and lifts the cadence axis from network protocols to a cognitive-architecture parameter (full discussion in Appendix F.10).

2.3 Event semantics, bounded rationality, and an order-theoretic reading

The four starred events anchor at concrete logged triggers. We state the triggers once, because every empirical claim in the paper is a claim about them (per-regime formal definitions in Appendix C), and we give each stage its reading in terms of the resource it bounds.

Q^* (intent issued). The planner emits a semantic intent $q = (q^{\text{req}}, q^{\text{pref}}, q^{\text{pen}})$: small sets of required, preferred, and penalty tags. The episode-level event fires on emission, so it is saturated whenever the task is semantic ($\hat{P}(Q^*) = 1.00$ in Table 2); the multi-agent per-tick logger additionally requires a non-empty candidate set (Appendix C.1), a bookkeeping difference we record rather than hide. Resource bound: attention — the intent is a few tags, not the full state.

R^* (memory answers). Retrieval evaluates q against the memory and returns a candidate set; the event fires when the set contains a candidate satisfying the task predicate (within ϵ of a true target in the multi-agent logger). The 91% empty-candidate-set failures of §3.5 are R -side events: no stored place clears the required-tag threshold. Resource bound: memory — retrieval can only surface what consolidation has stored.

M^* (target lock). The planner materializes one candidate into a concrete actionable cell. Candidacy itself is threshold-gated (a place enters the candidate set only if its score clears the planner’s gate, Appendix D), and the lock commits to a single candidate once per query rather than re-optimizing against the full memory at every tick. The event fires on lock acquisition (within ϵ of a true target in the multi-agent logger). Resource bound: deliberation.

C^* (closed-loop completion). The local controller executes until the locked cell is reached or the episode times out; the event fires on arrival. C is the only stage whose failure modes are extra-semantic: controller competence and, in the multi-agent regime, occupancy $\mathcal{O}_{-i}$. Resource bound: control, and in a collective, shared targets.

Interpretive readings (all deferred). The stages admit three readings, none of which is a contribution and all of which now live in the appendices: the pipeline as a *satisficing*, *boundedly rational*

agent whose four resource bounds (attention, memory, deliberation, control) are exactly the four failure loci (Appendix B); an *order-theoretic* account of empty retrieval — the 91% regime — via the query–extent Galois connection (Appendix A.4); and an *options/semi-MDP* account of why optima appear on time-bearing metrics rather than on the duration-blind chain probability Φ (Appendix B).

3 Single agent: stage diagnostics localize failure

This section instantiates the first row of Table 1: a single symbolic agent, the simplest system class, where the question is whether stage attribution adds information that end-to-end success does not.

3.1 System, tasks, and protocol

Memory stack. The system uses a graph substrate with a semantic retrieval interface on top: observations o_t are tokenized into symbolic scene tokens and semantic tags that update a memory of place records, transitions, episodic traces, and concept/relation records. Planner queries return candidates by a weighted log-score over required, preferred, and penalty tags. The retrieval machinery is reused across all tasks and architectures; full memory schema, update rules, and scoring formula are in Appendix D.

Tasks, baselines, protocol. Four single-agent task families are evaluated (water, safe-rest, goal-region on Empty-6x6 / FourRooms (Sutton et al., 1999), hazard recovery on LavaGapS7), 10 seeds $\times$ 8 episodes each, $\times$ 2 assist modes. Semantic stack (node, concept, state-conditioned variants) is compared against four baselines: random, graph-only, SPTM-like, learned BC. Reported: end-to-end success, steps, four operational stage rates, assist deltas, weakest stage per task and per method. All means are seed-aggregated; 95% bootstrap CIs use $B=5,000$ resamples (Efron, 1979); assist comparisons use paired Wilcoxon signed-rank tests (Wilcoxon, 1945) with Holm-Bonferroni correction (Holm, 1979) (Appendix G). For the two hardest tasks (FourRooms, LavaGapS7), four oracle regimes (*base*, *oracle R/M/C*) are run — each replaces exactly one stage while leaving the rest of the pipeline unchanged. BabyAI portability check in Appendix E. All single-agent experiments ran on a 10-core ARM CPU, 32 GB RAM (CPU-only); oracle sweep < 1 min wall-clock.

LLM diagnostic block. To check that Q/R/M/C is not tied to the symbolic planner implementation, we also run a fixed TextNav substrate with an LLM-driven observe/query/retrieve/decide loop. The TextNav task, prompts, parser contract, and event definitions are held fixed while only the local Ollama model changes; the backend uses model-compatible completion settings when an Ollama chat template otherwise hides short structured answers in a separate thinking field. This block is a controlled step-budget study (Table 4): sweeping the completion budget drives both LLM agents into a failure that Q/R/M/C localizes to the *C* stage while *Q/R/M* stay saturated, exhibiting the framework’s diagnostic property on real language agents. Reproducing full ReAct/Reflexion/Voyager/MemGPT-style agents is a separate systems contribution; the comparison to those systems is a coverage comparison in Appendix F.11.

3.2 Stage-estimator validation

Before the main benchmark, the factorization is validated. On a synthetic four-stage process with known probabilities the nested-estimator product tracks empirical semantic-path completion across sample sizes. On the real benchmark, the same product matches semantic-path completion rather than overall success (Table 2) — exactly the contract of Definition 1.

3.3 Success-only summary and what it hides

Water and rest are strong end-to-end tasks (0.95 and 0.93 success; Figure 1b). Goal-region reaches 0.54 success at the task-family level — an average that pools a solved environment (Empty-6x6, 1.00) with a hard one (FourRooms, 0.075; §3.4), so the family number is bimodal rather than a uniform mid-difficulty — but the bottleneck sits downstream of candidate selection (§3.4). Hazard recovery is the hardest retrieval-sensitive task at 0.40 success, with retrieval satisfaction as the weakest stage. Under success-only evaluation the two hard families look like the same kind of weakness;

Table 2: Stage-factor consistency audit for the best semantic method in each task family. The product of nested conditional estimators matches semantic-path completion; overall success can be higher because some episodes succeed without traversing the full semantic retrieval path.

Task	Method	Overall	SP-comp	$\hat{P}(Q^*)$	$\hat{P}(R^* Q^*)$	$\hat{P}(M^* R^*)$	$\hat{P}(C^* M^*)$	Product
Water	water-cp	0.95	0.70	1.00	0.98	0.74	0.97	0.70
Rest	rest-s2	0.93	0.73	1.00	1.00	0.80	0.91	0.73
Goal region	goal-n1	0.54	0.00	1.00	0.54	0.01	0.00	0.00
Hazard recovery	hazard-s7	0.40	0.16	1.00	0.56	0.58	0.50	0.16

Table 3: Baseline comparison. The learned behavior-cloned baseline dominates the easiest resource tasks and hazard recovery, while the semantic stack remains strongest on the relation-heavy goal-region family and remains the only family with interpretable Q/R/M/C structure.

Task	Random	Graph-only	SPTM-like	Learned	Best semantic
Water	0.85	0.90	0.90	1.00	0.95
Rest	0.85	0.90	0.90	1.00	0.93
Goal region	0.41	0.49	0.52	0.50	0.54
Hazard recovery	0.08	0.39	0.00	0.74	0.40

the stage view will show they are opposites. Per-task operational rates with 95% CIs are in Table 8 (Appendix D.3).

Baseline comparison. Table 3 compares the semantic stack against random, graph-only, SPTM-like, and learned behavior-cloned baselines. The learned baseline dominates the easiest resource tasks (1.00 on water and rest) and hazard recovery (0.74), while the semantic stack remains strongest on the relation-heavy goal-region family (0.54). The relevant contrast for this paper, however, is not the success column: the semantic stack is the only method family whose decisions expose interpretable Q/R/M/C structure. On the BC baseline the stage events are degenerate for the same structural reason as on the multi-agent learned baseline analysed in §4.5 — a policy without an explicit query/lock interface cannot emit separable stage events, so the diagnostic decomposition has nothing to attach to. End-to-end learning and stage diagnostics answer different questions; this paper is about the second.

3.4 Oracle-stage interventions localize the bottleneck

On goal-region the semantic variants solve Empty-6x6 (1.00 success) but collapse on FourRooms (0.075). For hazard recovery (LavaGapS7), base success 0.39 (bootstrap 95% CI [0.35, 0.43]) rises only to 0.43 under oracle controller, but jumps to 0.59 under oracle materialization and 0.60 under oracle retrieval (CI [0.51, 0.69]; per-seed paired deltas positive in 9/10 seeds, exact sign test $p=0.004$); goal-region rejoin is downstream-limited (oracle R lifts semantic-path completion to 0.08 without moving end-to-end success). The per-task attribution as a stacked waterfall is in Figure 6 (Appendix D.3). Two tasks identical under success-only evaluation fail at opposite ends of the Q/R/M/C chain.

3.5 The logs refine the diagnosis: a memory-accumulation limit

Inspecting the framework’s own logs sharpens the retrieval attribution. On the 10-seed hazard-recovery run, 349 retrieval calls were logged: only 32 (9.2%) returned at least one candidate, and of these 31 (96.9% precision) selected a semantically satisfied candidate. The bottleneck is therefore not candidate *ranking* but candidate *generation* — the symbolic memory does not contain hazard-recovery-tagged nodes at 91% of decision points. A retrained 2-layer candidate-generation MLP trained on per-state features (test accuracy 0.59 vs majority-class 0.81; Appendix H) confirms that instantaneous state features cannot substitute for the missing structure. The framework’s “retrieval-limited” attribution thus refines to a *memory-accumulation* limit: candidate generation cannot be solved from state features alone, and the binding constraint is how trajectory evidence is consolidated into persistent symbolic structure. This attribution is not an artefact of the required-tag threshold




Figure 3: Four multi-agent memory architectures. **(a) Independent:** per-agent private memory, no exchange. **(b) Shared bus:** one global memory, every agent reads/writes each tick. **(c) Central aggregator:** per-agent memories plus a hub that merges them into a consensus report. **(d) Peer-to-peer broadcast:** per-agent memories with pairwise snapshot exchange every K ticks (the cadence axis of Empirical Claim 1).

299 θ : sweeping θ across the full range leaves the empty-candidate-set rate invariant (Appendix F.5),
 300 because the empties are absence of any stored candidate, not sub-threshold filtering. Nor is it an
 301 artefact of the weighting: the empty rate is 0.908 pooled over calls, 0.909 seed-weighted, and 0.891
 302 episode-weighted (mean over the 80 episodes of each episode’s own empty share), so long failing
 303 episodes contributing more decision points do not inflate the figure. Accumulation itself is not one of
 304 the four stages: the protocol localised the failure to the R link, and this within-link forensics named
 305 the cause — we read Q/R/M/C as coarse localisation with a designated drill-down, not as a complete
 306 causal ontology of memory failure. This sets up the multi-agent section: when consolidation becomes
 307 a shared process across N agents, it acquires a measurable *rate*.

308 4 Multi-agent: memory as a shared resource

309 This section instantiates rows two to four of Table 1: the same logged events, unchanged, instrument
 310 a communicating collective — and, within the same section, learned policies and LLM agents —
 311 where the additional failure axis is contention for the targets that memory designates.

312 4.1 Four memory architectures and the cadence axis

313 Four communication architectures are instantiated; they share the single-agent retrieval substrate but
 314 differ in how memory is distributed across agents (Figure 3).

315 **Independent.** Each agent maintains a private event store and concept graph. No inter-agent commu-
 316 nication mechanism exists. Lower-bound baseline.

317 **Shared bus.** All agents publish observations to a single *collective memory* event bus. One shared
 318 concept graph is rebuilt from the union of events. Maximally centralized.

319 **Central aggregator (ConsensusLayer).** Each agent has its own private concept graph; a central
 320 layer pulls snapshots and computes a trust-weighted merge, producing a single *consensus report*
 321 consumed by all agents.

322 **Peer-to-peer broadcast (cadence K).** Each agent has its own concept graph and its own merged
 323 view. Every K ticks each agent broadcasts a snapshot to a passive bus; each agent independently
 324 merges its own snapshot with the latest snapshot received from each peer. Different agents may
 325 hold different merged views; there is no central report.

326 The first three correspond to the three regimes in the multi-agent communication literature: no
 327 exchange, CTDE-style aggregation, and decentralized peer communication. The fourth is the
 328 architecture for which Definition 2 and Empirical Claim 1 are directly applicable, since the cadence
 329 K is a tunable parameter rather than a fixed protocol. Mechanically, K is the rate at which private
 330 memories consolidate into a collective view, which is why it is the structural axis of this section.

4.2 Multi-agent factorization

Extend the factorization to a population of N agents that share information through a communication protocol $\mathcal{C}_K$ parameterized by broadcast cadence K . Each agent i has its own memory state $\mathcal{M}_i$ and issues its own query q_i . Let $\mathcal{M}_i^{(K)}$ denote the projection of agent i 's memory at the moment it queries, after absorbing every peer message received under cadence K . Two coupling channels are separated: (i) memory coupling $\mathcal{M}_{-i}^{(K)}$, the latest peer memory snapshots accessible to agent i ; (ii) occupancy coupling $\mathcal{O}_{-i}$, the set of targets that other agents have already claimed by the time i acts.

Definition 2 (multi-agent stage estimability). For each agent i ,

$$P(C_i^* | q_i, \mathcal{M}_i, \mathcal{C}_K) = P(Q_i^* | q_i, \mathcal{M}_i) P(R_i^* | Q_i^*, q_i, \mathcal{M}_i, \mathcal{M}_{-i}^{(K)}) \\ \cdot P(M_i^* | R_i^*, q_i, \mathcal{M}_i, \mathcal{M}_{-i}^{(K)}) P(C_i^* | M_i^*, q_i, \mathcal{O}_{-i}).$$

Under a multi-agent stage-Markov assumption (Appendix C), the four factors are estimable from per-agent trajectory logs and the empirical-frequency estimators are consistent.

When Definition 2 may fail. The single-agent stage-Markov property absorbs trajectory history into $(q, \mathcal{M})$. In the multi-agent setting, the analogous absorption requires that dependence on past peer decisions enters only through $\mathcal{M}_{-i}^{(K)}$ at the M stage and through $\mathcal{O}_{-i}$ at the C stage. This step is non-trivial: $\mathcal{O}_{-i}$ summarises target occupancy at decision time but loses information about *how* peers arrived there. The factorization is therefore a projection onto observable per-agent traces under that summary. Two specific failure modes arise: (i) inter-agent communication side-channels not logged through $\mathcal{C}_K$; (ii) coordinated peer policies whose decision at the C stage depends on hidden peer state not summarised by $\mathcal{O}_{-i}$.

4.3 Bottleneck shift

Empirical Claim 1 (bottleneck shift $\mathbf{M} \leftrightarrow \mathbf{C}$, slope-level). Assume scarcity ($N > T$ targets, see §2.1) and a peer-broadcast architecture with cadence $K \in \mathbb{N}$. Two stylized assumptions are posited: **(F)** memory freshness $\nu(\mathcal{M}_{-i}^{(K)})$ is non-decreasing pointwise as K decreases, and the probability that agent i 's planner locks onto a reachable target is non-decreasing in ν ; **(O)** the probability that an arbitrary committed target is already in $\mathcal{O}_{-i}$ is non-increasing as K increases. Write $P(M^* | R^*; K)$ and $P(C^* | M^*; K)$ for the population-level conditional probabilities (averaged over agents and the seed distribution of the benchmark); $\hat{P}(\cdot)$ denotes the corresponding empirical-frequency estimator computed from logs. Under (F) and (O),

$$\frac{\partial P(M^* | R^*; K)}{\partial K^{-1}} > 0, \quad \frac{\partial P(C^* | M^*; K)}{\partial K^{-1}} < 0.$$

Both slope conditions are direct empirical predictions that are testable on logs through $\hat{P}$.

What this claim is and is not. Empirical Claim 1 is a *slope-level statement on conditional rates*: each of the two slopes is testable and is confirmed on the 35,640-run sweep (§4.5) and independently on a 100-run standard-MiniGrid portability sweep (§6). We report the slopes as *effect sizes with cluster-robust confidence intervals* (resampled over the 81 independent configuration cells rather than the correlated run count) in Appendix F.3. The claim's boundary with respect to the conditional product Φ is stated once, in the Scope-and-non-claims box (§1) and in §4.5. The shift is a signature of scarce-target contention, not of the grid design: under abundance ($T > N$) both slopes flatten and the $\mathbf{M} \leftrightarrow \mathbf{C}$ anticorrelation attenuates (Appendix F.8), i.e. the prediction breaks exactly where the contention mechanism is removed.

Chronology. Empirical Claim 1 is post-hoc. It was formulated *after* an initial 12,960-run sweep conducted under four stage-agnostic hypotheses (Appendix F); the two slope conditions were then confirmed on the full 35,640-run sweep, on the $N = 12$ scale-up (Appendix E), and on the MiniGrid portability wrapper (§6). We state this explicitly: the verified slope signs are predictions on enlarged data, not a pre-registered prediction. A fully formal version is left for future work.

374 4.4 Sweep protocol

375 The four architectures are evaluated on a custom multi-agent grid (12×10 cells); each agent
 376 must reach a cell tagged `water_source`, choosing actions via a shared memory-driven plan-
 377 ner so only the memory layer differs. The sweep crosses $N \in \{3, 5, 8\}$, $T \leq N$, three layouts
 378 (symmetric/asymmetric/random), three hazard densities, four architectures, eight peer cadences
 379 $K \in \{1, 2, 4, 8, 16, 32, 48, 64\}$, 40 seeds — 35,640 runs in ≈ 22 min on 8 cores. Per-agent per-tick
 380 Q/R/M/C indicators are logged (Appendix C); episode-level starred events are the disjunction over
 381 ticks. A separate oracle-intervention block on scarcity cases $(N, T) \in \{(3, 2), (5, 3), (8, 5)\}$ at five
 382 cadences evaluates oracle R (memory returns ground-truth water positions), oracle M (planner
 383 locks onto nearest real water), and oracle C (completion of a real materialised lock is guaranteed;
 384 Appendix F.13).

385 4.5 Results: stage decomposition and bottleneck shift

386 The 35,640-run sweep is reported through four observations, structured to mirror Empirical Claim 1:
 387 (i) the conditional decomposition isolates the M and C stages as the differentiating links; (ii) the
 388 bottleneck shift appears directly as a negative within-cadence correlation between $P(M^*|R^*)$ and
 389 $P(C^*|M^*)$, peaking at $K = 4$ and decaying to noise at $K = 64$; (iii) on mean time-to-success, the
 390 resulting trade-off produces an interior optimum at $K = 8$ — robust in location but shallow in margin,
 391 with the advantage over the adjacent fast cadences not resolved at this sample size (Appendix F.3);
 392 (iv) on the conditional product Φ , the same trade-off produces a monotone trend rather than an interior
 393 peak — a summary statistic on which we make no formal claim. A learned-communication baseline
 394 and a multi-agent oracle-intervention loop close the subsection.

395 4.5.1 Conditional decomposition reveals where each architecture fails.

396 Figure 4 shows per-architecture conditional Q/R/M/C rates; full numerics with bootstrap CIs are in
 397 Appendix F.1, Table 11. Retrieval conditional $P(R^*|Q^*)$ is saturated; the differentiating signal lives
 398 at M and C. Shared-bus, central-aggregator, and fast-broadcast peer ($K \leq 4$) have high $P(M^*|R^*)$
 399 but low $P(C^*|M^*)$; independent and slow-broadcast peer ($K \geq 16$) show the reverse — the $M \leftrightarrow C$
 400 trade-off predicted by Empirical Claim 1.

401 4.5.2 Bottleneck shift and interior efficiency optimum.

402 Across all peer runs, Spearman of $P(M^*|R^*)$ vs K is -0.53 (95% cluster-robust CI $[-0.59, -0.47]$)
 403 and of $P(C^*|M^*)$ vs K is $+0.43$ (CI $[+0.37, +0.50]$); both are moderate effect sizes whose CIs
 404 exclude zero under resampling of the 81 configuration cells (Appendix F.3), matching the slope
 405 conditions of Empirical Claim 1. We report these rather than p -values, which are uninforma-
 406 tive at this correlated sample size. Both slope signs are also robust to the two free thresholds
 407 of the logger — candidacy θ and lock tolerance ϵ (Appendix F.5). The stronger empirical signa-
 408 ture is at the per-configuration level: Pearson of $P(M^*|R^*)$ vs $P(C^*|M^*)$ within a fixed K is
 409 $-0.43, -0.54, -0.61, -0.19, -0.05, -0.04, -0.02, -0.02$ at $K=1, 2, 4, 8, 16, 32, 48, 64$, peaking
 410 at $K=4$ and decaying to noise by $K=64$ (Figure 9, Appendix F.1). Configurations where M is
 411 high systematically have low C: the bottleneck migrates between the two stages within the same
 412 $(N, T, \text{layout}, \text{hazard}, \text{seed})$ cell as K varies. On the efficiency metric, mean t_{succ} has a U-shape
 413 with interior minimum at $K=8$ (7.47 ticks vs 8.02 for the independent baseline); the minimum’s
 414 location is robust but shallow — its margin over the adjacent fast cadences ($K=1, 4$) is not re-
 415 solved under cell-level resampling (paired-bootstrap CIs in Appendices F.1 and F.3). Both state-
 416 ments condition on success. Under a censoring-aware expected time that charges failed runs the
 417 full step limit, the fast flank reverses and the independent baseline is marginally best (cell-paired
 418 $\Delta(K=8 - \text{independent}) = +0.32$, CI $[+0.05, +0.63]$; Appendix F.3): at fast cadence the completion
 419 decline reduces the success rate by more than the materialization gain increases it, so the interior
 420 optimum is a statement about conditional speed, not unconditional efficiency — and closing the
 421 unconditional gap is what the reservation protocol of open problem (3) achieves (Appendix F.9).
 422 Scaling to $N=12$ *sharpens* the signature: the peak Pearson rises to -0.70 and migrates outward to
 423 $K=8$, consistent with stronger occupancy coupling at higher N delaying the cost-regime transition
 424 (Appendix E, Figure 7).




Figure 4: Per-architecture conditional rates (top) and mean time-to-success among successes (bottom) from the 35,640-run sweep. Error bars and band: 95% cluster-bootstrap CIs over the 81 design cells; points are cell-weighted means (7.63 at $K=8$ vs. 8.04 independent; the run-pooled values of Table 11 are 7.47/8.02). Shared bus, central aggregator, and fast peer broadcast ($K \leq 4$) live in the M-saturation / C-collapse regime; independent and slow peer ($K \geq 16$) live in the M-deficit / C-saturation regime. The interior minimum at $K=8$ is robust in location but shallow in margin (Appendix F.3); the censoring-aware contrast is in the text.

425 **Provenance stratification: the mechanism behind the shift.** Annotating each M^* lock with the
426 *provenance* of the evidence behind it makes the bottleneck shift mechanistic rather than correlational.
427 Let φ be the foreign share of the evidence mass supporting the locked candidate, computed from the
428 same trust/support weights the merge itself used and frozen at lock time, and call a lock a *warp* (W^*)
429 when $\varphi \geq 0.8$: the agent is committing to a target it knows essentially only from peers. On a dedicated
430 peer sweep over the scarcity cells (random layouts, 20 seeds, $K \in \{1, 2, 4, 8, 16\}$) the warp share
431 $P(W^* | M^*)$ falls from 0.55 at $K=1$ to 0.11 at $K=8$ while the own-evidence stratum completes at a
432 flat 0.26–0.34 and the warp stratum completes below 0.05 throughout: *the C-collapse of fast sharing*
433 *is concentrated in the warp stratum*. What cadence controls, mechanically, is how much of the collec-
434 tive’s materialisation runs on foreign evidence — and that is precisely the stratum whose completions
435 are lost to occupancy contention. The stratification is not an artefact of the 0.8 cutoff: recomputed
436 from the raw per-lock φ values (23,237 locks) at thresholds $\varphi \in \{0.5, 0.6, 0.7, 0.8, 0.9, 0.95\}$, the
437 foreign stratum completes at ≤ 0.03 while the own stratum stays at 0.26–0.50, at every threshold and
438 cadence. The two coupling channels also separate interventionally with arms already in the design:
439 the independent arm removes peer evidence while leaving target contention in place, and the reserva-
440 tion protocol of open problem (3) (Appendix F.9) keeps peer evidence while managing contention —
441 under it, unconditional success at fast cadence exceeds the $K=8$ baseline (0.614–0.620 vs. 0.583),
442 identifying occupancy racing rather than evidence quality as the cost channel of fast sharing. The
443 same stratification reproduces on the continuous VMAS bridge (0.72 $\rightarrow$ 0.00 over $K=1/4/16$) and
444 on an LLM text substrate (0.42 $\rightarrow$ 0.08 over $K=2/8$), with the curves nearly coinciding across
445 substrates. A full provenance-conditioned analysis (counterfactual attribution by masked replay,
446 a closed-form staleness horizon with a registered 6/6-exact hold-out, and a decentralized repair
447 protocol) appears in a companion manuscript; here the stratification serves as the mechanistic reading
448 of Empirical Claim 1.

449 **What we do not claim about Φ .** The conditional product $\Phi = P(M^* | R^*) \cdot P(C^* | M^*)$ does
450 *not* have an interior peak in the tested range (mean-of-products estimator increases monotonically
451 from 0.569 at $K=4$ to 0.603 at $K=64$, converging to the independent-agent boundary 0.605). The
452 slope-level statement of Empirical Claim 1 is on the conditional rates themselves, not on Φ ; we

make no formal claim about the shape of $\Phi(K)$. Jensen-gap detail and POM-MOP comparison in Appendix F.1.

Learned-communication baseline: stage events are structurally degenerate even at competitive success. A CommNet-style policy (Sukhbaatar et al., 2016) trained with PPO (Schulman et al., 2017) on the same scenario ($N=3$, $T=2$) reaches 0.667 held-out success ($n=100$ seeds) — matching the symbolic peer architecture at $K=8$ (0.65). Despite this, the Q/R/M/C profile remains structurally degenerate: $Q^*=1.00$ saturated by construction; $R^*=0.997$, $M^*=0.997$ collapsed together ($P(M^*|R^*)=1.00$). This is the verified structural claim: a real-valued comm vector is never informatively empty (Q saturates) and lacks an explicit target-lock interface (R/M merge), regardless of training budget. The multi-agent counterpart of the single-agent learned-BC analysis (§3.3). Full architecture, PPO hyperparameters, and side-by-side REINFORCE vs PPO comparison in Appendix H.

LLM model sweep: the logger localizes a real failure in language agents. Table 4 runs the same Q/R/M/C event contract on two local LLM agents (qwen3:4b and llama3.1 via Ollama) on TextNav, sweeping the step budget over $\{6, 12, 25\}$ (8 episodes each; the task needs ~ 10 low-level steps). At a generous budget (12, 25) both agents succeed and all four stages fire. At a budget below the requirement (6), success drops to 0/8 for *both* models — and the decomposition shows *where*: $Q^*=R^*=M^*=1.00$ (the agents form the query, retrieve, and lock the correct target every episode) while $C^*=0.00$ (none complete in time). The framework localizes the failure precisely to a budget-limited completion (C) rather than to query formation, retrieval, or materialization; that both models fail identically shows the failure is structural, not a model artifact. This is the diagnostic property the framework claims, now exhibited on *real language agents*: an otherwise opaque success collapse ($1.00 \rightarrow 0.00$) is resolved into a specific failing stage. (Instrumentation note: a naive qwen3:4b/Ollama chat-template call initially returned empty response strings because the model spent the short structured-output budget in a separate thinking field; we use raw structured completion for qwen3:4b and a first-line parser, recorded for reproducibility.)

The multi-agent phenomenology reproduces on a live LLM collective. Beyond the single-agent budget study, the same event contract instruments a collective of $N=3$ LLM agents (llama3.1) sharing symbolic place memory by peer broadcast under scarcity ($M=2$ takeable targets, four layouts, $K \in \{2, 8\}$), with all predictions registered before the runs. The cadence phenomenology of §4.5 carries over to the language substrate: a paired comparison against the no-communication baseline on identical layouts shows raw sharing *lowering* success from 2/3 to 1/3 per episode — the occupancy-racing C-collapse, now produced by agents whose queries, tag extraction, and commitments are LLM-made — and target collisions appear in 7/8 shared-memory episodes. Attaching the reservation-and-rollback repair protocol (the constructive answer to open problem 3 of Appendix F.9) recovers part of the loss ($0.333 \rightarrow 0.417$ at both cadences, rollback latency again $\approx K$). Two interface findings transfer to any LLM-over-symbolic-memory design: LLM query formers free-generate retrieval tags outside the memory’s vocabulary (requiring tolerant retrieval gates), and small local models are reliable at semantic commitment but not at coordinate arithmetic or systematic exploration (requiring layered execution). Full protocol and provenance analysis in the companion manuscript.

Scope: these are controlled studies on local models; scaling the same contract to full Voyager/MemGPT-style agents on richer tasks with multiple failure modes is developed next (external-framework validation) and in Appendix F.11.

Multi-agent oracle interventions close the diagnose→intervene→verify loop. On the scarcity cases, mean t_{succ} at small/intermediate cadence ($K \in \{1, 2, 4\}$) drops from ~ 10 –11 in baseline to ~ 7.2 under either oracle R or oracle M, and further to ~ 6.1 under oracle C, which additionally removes post-lock travel; at slow cadence ($K=16$) the baseline is already within ~ 1.2 of the R/M oracle floor (Table 19 in Appendix F.12). Crucially, oracle C gives the largest time reduction but the smallest success-rate gain (Appendix F.13): in scarcity, completion costs time while retrieval and materialisation decide success. The cadence-induced excess time at small/intermediate K is attributable to the $M \leftrightarrow C$ transition; at large K the bottleneck has moved into solo exploration.

The multi-agent block thus delivers its half of the argument. The cadence K is, mechanically, the rate at which private memories consolidate into a collective view; the data show that this rate has a non-trivial interior optimum ($K=8$ on t_{succ}), that its cost channel is occupancy coupling, and that both

Table 4: LLM TextNav step-budget sweep: Q/R/M/C localizes a completion-limited failure in two local Ollama agents. Eight episodes per (model, step limit), temperature 0.0; the task needs ~ 10 low-level steps. At a budget below that requirement (step limit 6) success drops to 0.00 for both models, yet $Q^*=R^*=M^*=1.00$ (the agents form the query, retrieve, and lock the correct target every episode) while $C^*=0.00$: the framework localizes the failure to the C stage, not to Q/R/M. Both models behave identically, so the failure is structural (budget), not model-specific. Artifacts: `tmp/llm_steplimit_{6,12,25}`.

Model	Step limit	Success	Q^*	R^*	M^*	C^*
qwen3:4b	6	0.00	1.00	1.00	1.00	0.00
llama3.1	6	0.00	1.00	1.00	1.00	0.00
qwen3:4b	12	1.00	1.00	1.00	1.00	1.00
llama3.1	12	1.00	1.00	1.00	1.00	1.00
qwen3:4b	25	1.00	1.00	1.00	1.00	1.00
llama3.1	25	1.00	1.00	1.00	1.00	1.00




Figure 5: Stage profiles of three third-party agent stacks on the four engineered variants (llama3.1, $n=20$ per cell). Query and retrieval saturate identically everywhere; the two structural failures produce exact zero columns on every stack ($R^*=0$ under the consolidation gap, $C^*=0$ under the tight budget); the behavioral spread between frameworks lives entirely in the M and C stages.

507 effects are visible stage-by-stage on the same Q/R/M/C chain used for single agents. Consolidation is
508 not only the single-agent limit; in a collective it is a tunable, measurable design axis.

509 5 The contract on third-party agent stacks

510 This section instantiates the last row of Table 1. We further evaluate the protocol on systems built
511 by others. The contract has two layers: an *invariant core* — four event roles (a memory query is
512 issued; returned evidence satisfies the task; the agent commits to a target; the task completes) — and
513 a *system-specific adapter* that maps logged actions to those roles. We instrument three third-party
514 agent stacks — a function-calling loop on the OpenAI SDK, LangGraph’s prebuilt ReAct agent, and
515 AutoGen’s AssistantAgent/UserProxy pair, all driving the same local model (llama3.1, temperature 0)
516 — with one tool-trace adapter shared by all three: an episodic household memory is exposed through
517 `recall/goto/take` tools, and Q^* fires on querying memory, R^* on *attribution* (a returned record
518 names the item at its true place), M^* on a commitment (a `goto` to the true place; first-commitment
519 accuracy M_1 is logged as an auxiliary lock-quality diagnostic), and C^* on task completion within a
520 tool budget. A tool trace exposes no lock chain, so these are *marginal event frequencies* $f(\cdot)$ in the
521 sense of Appendix C: an agent can commit to the true place without a supporting retrieval, the events
522 are not nested, and no factorisation is claimed — the four frequencies serve as separate diagnostics
523 and are never substituted into the chain-rule product. What varies across system classes is the adapter,
524 not the event roles. Three variants engineer one failure each — a consolidation gap (the location fact
525 was never written to memory), a stale-distractor ambiguity, a tool budget below the task’s requirement
526 — plus a clean control; all predictions were registered before the runs, and the two revisions of the
527 adapter (an attribution leak; an over-tightened M^*) are recorded in the registration file.

528 Three findings (Table 5, Figure 5). **(1) Structural localizations are exact and framework-**
529 **independent.** When the memory lacks the fact, $R^*=0.00$ on every stack; when the budget cannot
530 fit completion, $C^*=0.00$ at $Q^*=R^*=1$ on every stack (20/20 episodes each). The contract needs

Table 5: Marginal Q/R/M/C event frequencies on three third-party agent frameworks (llama3.1, 20 episodes per cell, identical episodes/tools/budgets). Each entry is the fraction of episodes in which the event occurred: Q^* — the agent queried memory at least once; R^* — a returned record named the item at its true place; M^* — the agent committed (goto) to the true place, with or without retrieval support; C^* — take succeeded within the tool budget; M_1 : first goto targets the true place; stale_1 : first goto targets the stale distractor. The events are not nested and no product is claimed (Appendix C): under the consolidation gap, budgeted blind search over the four places sometimes commits to and reaches the true place, hence $M^*, C^* > 0$ at $R^*=0$. Structural failures localize identically everywhere: the consolidation gap is $R^*=0.00$ and the budget failure is $C^*=0.00$ at $Q^*=R^*=1$, for all three stacks. On the *clean* control the frameworks spread from 0.50 to 0.75 end-to-end — and the decomposition attributes the spread (see text).

Framework	Variant	$f(Q^*)$	$f(R^*)$	$f(M^*)$	$f(C^*)$	$f(M_1)$	$f(\text{stale}_1)$
OpenAI SDK	control	1.00	1.00	0.85	0.75	0.25	—
	consolidation gap	1.00	0.00	0.50	0.40	0.25	—
	stale ambiguity	1.00	1.00	0.90	0.30	0.25	0.20
	tight budget	1.00	1.00	0.25	0.00	0.25	—
LangGraph	control	1.00	1.00	0.70	0.50	0.50	—
	consolidation gap	1.00	0.00	0.20	0.15	0.10	—
	stale ambiguity	1.00	1.00	0.65	0.60	0.50	0.35
	tight budget	1.00	1.00	0.10	0.00	0.10	—
AutoGen	control	1.00	1.00	0.85	0.55	0.45	—
	consolidation gap	1.00	0.00	0.30	0.20	0.15	—
	stale ambiguity	1.00	1.00	0.50	0.45	0.35	0.40
	tight budget	1.00	1.00	0.15	0.00	0.15	—

no per-framework adaptation. **(2) The clean control separates the frameworks, and the decomposition says how.** End-to-end success on the SAME clean episodes ranges 0.50 (LangGraph) to 0.75 (OpenAI SDK). The stage profiles attribute the spread to a failure mode common to all three stacks combined with framework-specific recovery: first-commitment accuracy is low everywhere ($M_1=0.25\text{--}0.50$ — the agent’s first goto frequently ignores its own just-retrieved evidence), episode-level materialization recovers to 0.70–0.85 through re-query loops, and the frameworks then differ in how much of that recovery survives to completion (C^*/M^* from 0.88 for the OpenAI SDK loop down to 0.65 for AutoGen). **(3) The stale distractor pulls, measurably and unevenly.** The outdated record captures the first commitment in 0.20–0.40 of ambiguous episodes (most strongly for AutoGen), and its completion cost is framework-specific: the OpenAI SDK loop pays hardest ($C^* : 0.75 \rightarrow 0.30$) while LangGraph and AutoGen absorb the detour ($0.50 \rightarrow 0.60$; $0.55 \rightarrow 0.45$) — the LLM-memory analogue of the staleness gating that the collective experiments showed to be decisive. Honestly failed registrations: the behavioral predictions that assumed a clean control (all stages ≥ 0.75 ; a uniform ambiguity-induced completion drop) broke *because the frameworks themselves fail a clean memory task 25–50% of the time and pay for ambiguity in different places* — which is finding (2)–(3), not noise; all revisions are recorded as registered failures with mechanisms.

Model robustness: the structure transfers, the behaviour is legible. A prediction registered before the runs — the structural localizations transfer to a second model on all frameworks, while behavioural rates are model-dependent and only reported — was confirmed exactly by repeating all 240 episodes with qwen3:4b (Table 7): $R^*=0.00$ under the consolidation gap and $C^*=0.00$ under the tight budget, 20/20 on every stack, for both models. The behavioural contrast is itself a diagnostic result: qwen3:4b completes every clean episode ($C^*=1.00$ across stacks, vs. 0.50–0.75 for llama3.1) because its *first-commitment accuracy* is high ($M_1=0.90\text{--}0.95$ vs. 0.25–0.50) and its susceptibility to the stale distractor low (0.05–0.30) — except under AutoGen, where qwen’s M_1 collapses to 0.40 while remaining ≥ 0.90 on the other stacks. The stage decomposition thus separates three influences a leaderboard conflates: the task’s structural limits (model- and framework-independent), the model’s commitment quality, and the framework’s prompt/loop shaping that can degrade it.

Blinded fault-localization meta-evaluation (registered). The engineered variants verify localization of failures whose stage is known by construction; the reviewer-grade test treats the protocol as a *classifier*. We injected ten fault types — two to three per stage, struc-

561 tural and behavioural (Table 6) — and fixed a first-failing-link decision rule, with all thresh-
562 olds, *before* any run (tmp/qrmc_external/registered_meta.json; rule and thresholds in
563 experiments/qrmc_external/run_meta_evaluation.py). The rule receives only anonymised
564 marginal profiles plus a labelled control cell for calibration; a held-out control (seeds 20–39) enters the
565 blind set with ground truth “none”, measuring the false-positive rate; a manipulation check excludes
566 faults that do not bind (the prompt-only fault moved qwen but not llama). Stage-level accuracy:
567 26/30 scored cells for llama3.1 (8/10, 8/10, 10/10 per stack) and 29/33 for qwen3:4b (9/11, 10/11,
568 10/11); the five structural faults localize 30/30 across both models and all stacks; held-out controls
569 read “none” in 5/6 cells (one false positive at $n=20$); cross-framework label agreement is 7/11
570 (llama) and 10/11 (qwen). The registered model-contrast prediction was confirmed in its named
571 direction: M-family faults hide behind llama’s weak baseline commitment ($M_1=0.25\text{--}0.50$) and
572 become visible on qwen’s clean baseline — the unmarked duplicate drops M_1 by 0.90/0.80/0.20
573 across stacks — while its accuracy clause held on two of three stacks (AutoGen: 10/11 vs. llama’s
574 10/10, a one-cell difference). One registered risk materialised instructively: the *marked* stale distrac-
575 tor changes llama3.1’s behaviour but reads “none” on qwen3:4b, which takes the staleness marker
576 into account — the protocol measures the behavioural effect of a fault, not its presence, and a fault a
577 model neutralizes is reported as no failure.

Table 6: Blinded fault localization: injected fault, ground-truth stage, and correct-localization count over the three stacks, per model. “Neutralized”: the fault does not change qwen’s behaviour (the staleness marker is heeded), so the blind rule reads “none”.

Injected fault	Stage	llama3.1	qwen3:4b
recall tool absent	Q	3/3	3/3
prompt: memory unreliable	Q	non-binding	3/3
fact never consolidated	R	3/3	3/3
silent corruption	R	3/3	3/3
flaky recall ($p=0.7$)	R	3/3	3/3
stale distractor (marked)	M	1/3	0/3 (neutralized)
duplicate (unmarked)	M	2/3	3/3
budget below requirement	C	3/3	3/3
flaky goto ($p=0.5$)	C	2/3	3/3
take always fails	C	3/3	3/3
held-out control	none	3/3	2/3

Table 7: Second-model replication (qwen3:4b, 20 episodes per cell, identical episodes/tools/budgets; thinking disabled request-level for budget parity, recorded in the registration). Entries are marginal event frequencies $f(E)$ as in Table 5 — not chain-rule links. Structural zeros transfer exactly; behavioural rates differ from llama3.1 (Table 5) in the direction the M_1 diagnostic explains.

Framework	Variant	$f(Q^*)$	$f(R^*)$	$f(M^*)$	$f(C^*)$	$f(M_1)$	$f(\text{stale}_1)$
OpenAI SDK	control	1.00	1.00	1.00	1.00	0.95	—
	consolidation gap	1.00	0.00	0.45	0.45	0.20	—
	stale ambiguity	1.00	1.00	1.00	0.95	0.90	0.05
	tight budget	1.00	1.00	0.30	0.00	0.30	—
LangGraph	control	1.00	1.00	1.00	1.00	0.90	—
	consolidation gap	1.00	0.00	0.60	0.60	0.30	—
	stale ambiguity	1.00	1.00	0.95	0.95	0.75	0.20
	tight budget	1.00	1.00	0.15	0.00	0.15	—
AutoGen	control	1.00	1.00	1.00	1.00	0.40	—
	consolidation gap	1.00	0.00	0.55	0.50	0.30	—
	stale ambiguity	1.00	1.00	1.00	0.90	0.40	0.30
	tight budget	1.00	1.00	0.15	0.00	0.15	—

578 Scope: three stacks at $n=20$ per cell on two local models; Letta/CrewAI (Python ≥ 3.10) and
579 API-scale models are the natural extension.

6 Limitations and future work

Portability across substrates: slope conditions hold on MiniGrid. The slope conditions of Empirical Claim 1 also hold on a standard-MiniGrid substrate. A FourRooms-based multi-agent wrapper (built from `minigrid.core.grid.Grid + Wall/Floor/Lava` primitives, identical operational Q^*, R^*, M^*, C^* definitions; `experiments/minigrid_multiagent_wrapper/`) was used to run a 100-run portability sweep ($K \in \{1, 4, 8, 16, 64\}$, 20 seeds, $N=3, T=2$): Spearman $P(M^*|R^*)$ vs K is -0.50 and Spearman $P(C^*|M^*)$ vs K is $+0.50$ (moderate rank correlations; this is a small 100-run sweep, so we report the effect sizes and do not lean on significance). The interior t_{succ} minimum does not separate at this minimal sweep size (5 cadences, 20 seeds, one hazard density); a denser sweep is the natural next step. The two-substrate seam is therefore closed at the slope-level claim.

Empirical Claim 1 is slope-level on conditional rates only. The two slope conditions hold as moderate, cluster-robust effect sizes (Appendix F.3); we make no formal claim on the shape of the conditional product Φ (§4.3). A future statement directly in terms of the time-cost surface, or a derivation under a generative model (Poisson-renewal memory, birth-death occupancy), would convert the empirical claim into a theorem.

Estimability under hidden state. With unlogged in-episode query rewrites or peer side-channels, Definitions 1 and 2 are projections onto logged traces, not full identification.

Shared coordinate frame. All multi-agent results assume agents interpret places in a common coordinate frame: cross-agent place identity is coordinate proximity, with semantic profiles as a tie-breaker inside a spatial radius. The place-correspondence problem — whether two agents’ records denote the same place at all — is therefore bracketed here, not solved. A companion manuscript removes the assumption constructively (fingerprint-based place identity with consensus frame recovery up to translation and 90° rotation) and shows the Q/R/M/C measurements, including the cadence slopes and the staleness horizon, unchanged on top of the recovered frames; the harder variants (continuous $SE(2)$, appearance noise, symbol alignment) remain open.

The evaluator’s correspondence oracle. A deeper form of the same dependence survives every representational relaxation, and we state it explicitly. The protocol unifies *observation points*, not internal representations — but scoring the two middle stages presupposes an *entity correspondence* between memory content and world referents: R^* asks whether a returned candidate denotes the true target, M^* whether the lock does. In every study here that oracle is supplied by an instrumented environment (ground-truth cells on the grids; the task’s true location in the tool-trace contract). The dependence is graded, which is what makes the protocol degrade gracefully rather than break: Q^* requires no correspondence at all (did the agent query its memory), C^* is behaviourally observable (did the task complete), while R^* and M^* — precisely the diagnostic payload — are undefined without it. The severity of the requirement depends on the population under evaluation: interchangeable agents (a shared perception interface, possibly different weights) situated in a single environment inherit the correspondence from the shared substrate — the setting of every study in this paper and the common case in multi-agent evaluation. It becomes a genuinely open problem for heterogeneous agents in open worlds, where deploying the protocol requires either an instrumented task or a correspondence mechanism with its own error model; the companion’s meaning-based place identity is a candidate for grounding the *evaluator’s* correspondence, not only the agents’, and we leave that construction open.

Learned baselines. The CommNet PPO baseline reaches 0.667 success on the scarcity scenario — competitive with symbolic peer at $K=8$ (0.65) — and confirms Q-saturation ($Q^*=1.00$) and R/M collapse ($P(M^*|R^*)=1.00$) at competitive success. The structural prediction now holds verified, not asserted; ablating on architectures with explicit symbolic target-lock channels is left for future work.

Scale. Multi-agent sweep bounded at $N=12$ (Appendix E; signature sharpens, peak Pearson -0.70 at $K=8$, interior t_{succ} minimum 4.77). Denser scaling at $N \geq 16$ is future work. **Continuous substrate.** A preliminary VMAS (Bettini et al., 2022) check reproduces *both* slope signs of Empirical Claim 1 in the coupling regime (Spearman -1.00 and $+1.00$ over $K \leq 16$, 40 seeds; Appendix F.7), reusing the symbolic memory through a discretisation bridge. This is a directional signal only — a perfect rank correlation over five monotone points, one configuration, discretised memory — so it indicates the shift is not an artefact of grid cells, but a full continuous-state validation with cluster-robust CIs remains future work.

7 Broader impacts

The paper proposes diagnostic evaluation for memory-based agents in simulated navigation. Its positive impact is methodological: making memory-based failures more measurable across architectures, and making the cadence axis of distributed multi-agent memory tunable on an empirically grounded criterion. Potential negative impacts are indirect and would arise if similar diagnostics were used to improve agents in safety-critical settings (drone swarms, autonomous vehicles, distributed sensor networks) without adequate safeguards. The current work is benchmark-level infrastructure, not a deployment-ready system.

8 Conclusion

We proposed Q/R/M/C, a four-stage decomposition of memory-based navigation whose factors are estimable from standard trajectory logs, and used it as a diagnostic instrument across symbolic, learned, LLM-driven, single-agent, and multi-agent traces. The main result is not a new agent architecture: it is a reproducible logging and estimation protocol that separates query formation, retrieval satisfaction, target materialization, and post-retrieval completion. Across the experiments, this protocol localizes failures that success-only evaluation cannot, exposes degenerate observability in implicit learned policies, and makes memory-sharing cadence measurable as a stage-level design variable. The resulting rates are directly comparable across systems that expose a query/lock interface; for opaque learned policies the protocol measures stage observability itself.

References

- M. Andrychowicz, F. Wolski, A. Ray, J. Schneider, R. Fong, P. Welinder, B. McGrew, J. Tobin, P. Abbeel, and W. Zaremba. Hindsight experience replay. In *NeurIPS*, 2017.
- R. Bělohávek. Fuzzy Galois connections. *Mathematical Logic Quarterly*, 45(4):497–504, 1999.
- J. Carmona, B. van Dongen, A. Solti, and M. Weidlich. *Conformance Checking: Relating Processes and Models*. Springer, 2018.
- S. Es, J. James, L. Espinosa-Anke, and S. Schockaert. RAGAs: Automated evaluation of retrieval augmented generation. In *EACL System Demonstrations*, 2024.
- H. Lightman, V. Kosaraju, Y. Burda, H. Edwards, B. Baker, T. Lee, J. Leike, J. Schulman, I. Sutskever, and K. Cobbe. Let’s verify step by step. In *ICLR*, 2024.
- C. Ma, J. Zhang, Z. Zhu, C. Yang, Y. Yang, Y. Jin, Z. Lan, L. Kong, and J. He. AgentBoard: An analytical evaluation board of multi-turn LLM agents. In *NeurIPS*, 2024.
- D. Ru, L. Qiu, X. Hu, T. Zhang, P. Shi, S. Chang, et al. RAGChecker: A fine-grained framework for diagnosing retrieval-augmented generation. In *NeurIPS Datasets and Benchmarks*, 2024.
- W. M. P. van der Aalst. *Process Mining: Data Science in Action*. Springer, 2016.
- M. Bettini, R. Kortvelesy, J. Blumenkamp, and A. Prorok. VMAS: A vectorized multi-agent simulator for collective robot learning. In *16th International Symposium on Distributed Autonomous Robotic Systems (DARS)*, 2022. arXiv:2207.03530.
- S. Boyd, A. Ghosh, B. Prabhakar, and D. Shah. Randomized gossip algorithms. *IEEE Trans. Information Theory*, 52(6):2508–2530, 2006.
- M. Chevalier-Boisvert, D. Bhoopchand, H. Abdolmaleki, and A. Mordatch. BabyAI: A platform to study the sample efficiency of grounded language learning. In *ICLR*, 2019.
- M. Chevalier-Boisvert, B. Lahlou, S. Willems, and P. S. Castro. MiniGrid and MiniWorld: Modular reinforcement learning environments. arXiv:2306.13831, 2023.
- P. Chhikara, D. Khant, S. Aryan, T. Singh, and D. Yadav. Mem0: Building production-ready AI agents with scalable long-term memory. arXiv:2504.19413, 2025.
- P. Dayan. Improving generalization for temporal difference learning: The successor representation. *Neural Computation*, 5(4):613–624, 1993.

681 M. Deitke, E. V. Wijmans, D. Schwenk, O. Salvador, K. M. Ehsani, et al. ProcTHOR: Large-scale embodied AI
682 using procedural generation. In *NeurIPS Datasets and Benchmarks*, 2022.

683 B. Efron. Bootstrap methods: Another look at the jackknife. *Annals of Statistics*, 7(1):1–26, 1979.

684 J. N. Foerster, Y. M. Assael, N. de Freitas, and S. Whiteson. Learning to communicate with deep multi-agent
685 reinforcement learning. In *NeurIPS*, 2016.

686 J. N. Foerster, G. Farquhar, T. Afouras, N. Nardelli, and S. Whiteson. Counterfactual multi-agent policy gradients.
687 In *AAAI*, 2018.

688 B. Ganter and R. Wille. *Formal Concept Analysis: Mathematical Foundations*. Springer, 1999.

689 D. Ghosh, A. Gupta, A. Reddy, J. Fu, S. Levine, and C. Finn. Learning to reach goals via iterated supervised
690 learning. In *ICLR*, 2021.

691 G. Gigerenzer and D. G. Goldstein. Reasoning the fast and frugal way: Models of bounded rationality.
692 *Psychological Review*, 103(4):650–669, 1996.

693 S. Gupta, J. Davidson, S. Levine, R. Sukthankar, and J. Malik. Cognitive mapping and planning for visual
694 navigation. In *CVPR*, 2017.

695 S. Holm. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70,
696 1979.

697 E. Kolve, R. Mottaghi, D. Gordon, Y. Zhu, A. Gupta, A. Farhadi, and D. Fox. AI2-THOR: An interactive 3D
698 environment for visual AI. *arXiv:1712.05474*, 2017.

699 B. Kuipers. The spatial semantic hierarchy. *Artificial Intelligence*, 119(1–2):191–233, 2000.

700 F. Lieder and T. L. Griffiths. Resource-rational analysis: Understanding human cognition as the optimal use of
701 limited computational resources. *Behavioral and Brain Sciences*, 43:e1, 2020.

702 R. Lowe, Y. Wu, A. Tamar, J. Harb, P. Abbeel, and I. Mordatch. Multi-agent actor-critic for mixed cooperative-
703 competitive environments. In *NeurIPS*, 2017.

704 S. Mac Lane. *Categories for the Working Mathematician*. Springer, 2nd edition, 1998.

705 I. Momennejad, E. M. Russek, J. H. Cheong, M. M. Botvinick, N. Daw, and S. J. Gershman. The successor
706 representation in human reinforcement learning. *Nature Human Behaviour*, 1:680–692, 2017.

707 J. O’Keefe and L. Nadel. *The Hippocampus as a Cognitive Map*. Oxford University Press, 1978.

708 C. Packer, V. Fang, S. Patil, K. Lin, S. Wooders, and J. Gonzalez. MemGPT: Towards LLMs as operating
709 systems. *arXiv:2310.08560*, 2023.

710 J. S. Park, J. C. O’Brien, C. J. Cai, M. R. Morris, P. Liang, and M. S. Bernstein. Generative agents: Interactive
711 simulacra of human behavior. *arXiv:2304.03442*, 2023.

712 T. Rashid, M. Samvelyan, C. Schroeder de Witt, G. Farquhar, J. Foerster, and S. Whiteson. QMIX: Monotonic
713 value function factorisation for deep multi-agent reinforcement learning. In *ICML*, 2018.

714 N. Savinov, A. Dosovitskiy, and V. Koltun. Semi-parametric topological memory for navigation. In *ICLR*, 2018.

715 M. Savva, A. Kadian, O. Maksymets, Y. Zhao, E. Wijmans, B. Jain, J. Straub, J. Liu, V. Koltun, J. Malik, et al.
716 Habitat: A platform for embodied AI research. In *ICCV*, 2019.

717 J. Schulman, P. Moritz, S. Levine, M. Jordan, and P. Abbeel. High-dimensional continuous control using
718 generalized advantage estimation. In *ICLR*, 2016. *arXiv:1506.02438*.

719 J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov. Proximal policy optimization algorithms.
720 *arXiv:1707.06347*, 2017.

721 N. Shinn, F. Cassano, A. Berman, and A. G. Labash. Reflexion: Language agents with verbal reinforcement
722 learning. In *NeurIPS*, 2023.

723 H. A. Simon. A behavioral model of rational choice. *Quarterly Journal of Economics*, 69(1):99–118, 1955.

724 H. A. Simon. Rational choice and the structure of the environment. *Psychological Review*, 63(2):129–138, 1956.

- 725 K. L. Stachenfeld, M. M. Botvinick, and S. J. Gershman. The hippocampus as a predictive map. *Nature*
726 *Neuroscience*, 20(11):1643–1653, 2017.
- 727 S. Sukhbaatar, A. Szlam, and R. Fergus. Learning multiagent communication with backpropagation. In *NeurIPS*,
728 2016.
- 729 R. S. Sutton, D. Precup, and S. Singh. Between MDPs and semi-MDPs: A framework for temporal abstraction
730 in reinforcement learning. *Artificial Intelligence*, 112(1–2):181–211, 1999.
- 731 E. C. Tolman. Cognitive maps in rats and men. *Psychological Review*, 55(4):189–208, 1948.
- 732 G. Wang, Y. Xie, Z. Jiang, A. Mandlekar, C. Xie, A. Yu, and P. Liang. Voyager: An open-ended embodied agent
733 with large language models. *arXiv:2305.16291*, 2023.
- 734 Q. Wu, G. Bansal, J. Zhang, Y. Wu, B. Li, E. Zhu, L. Jiang, X. Zhang, S. Zhang, J. Liu, A. H. Awadallah, R. W.
735 White, D. Burger, and C. Wang. AutoGen: Enabling next-gen LLM applications via multi-agent conversation.
736 *arXiv:2308.08155*, 2023.
- 737 F. Wilcoxon. Individual comparisons by ranking methods. *Biometrics Bulletin*, 1(6):80–83, 1945.
- 738 R. J. Williams. Simple statistical gradient-following algorithms for connectionist reinforcement learning.
739 *Machine Learning*, 8(3–4):229–256, 1992.
- 740 S. Yao, J. Zhao, D. Yu, N. Du, I. Shafran, K. Narasimhan, and Y. Cao. ReAct: Synergizing reasoning and acting
741 in language models. In *ICLR*, 2023.
- 742 L. Xiao, S. Boyd, and S. Lall. A scheme for robust distributed sensor fusion based on average consensus. *IPSN*,
743 2007.
- 744 C. Yu, A. Velu, E. Vinitzky, J. Gao, Y. Wang, A. Bayen, and Y. Wu. The surprising effectiveness of PPO in
745 cooperative multi-agent games. In *NeurIPS Datasets and Benchmarks Track*, 2022.

746 A Formal proofs

747 A.1 Argument for Definition 1 (single-agent estimability)

748 For each episode $i \in \{1, \dots, n\}$ let $Q_i^*, R_i^*, M_i^*, C_i^*$ denote the nested episode-level indicators. By
749 construction these are measurable functions of the runtime trajectory log and satisfy $C_i^* \Rightarrow M_i^* \Rightarrow$
750 $R_i^* \Rightarrow Q_i^*$. The empirical estimators

$$\hat{p}_Q = \frac{1}{n} \sum_i Q_i^*, \quad \hat{p}_{R|Q} = \frac{\sum_i R_i^*}{\sum_i Q_i^*}, \quad \hat{p}_{M|R} = \frac{\sum_i M_i^*}{\sum_i R_i^*}, \quad \hat{p}_{C|M} = \frac{\sum_i C_i^*}{\sum_i M_i^*}$$

751 are consistent by the chain rule plus the conditional stage-Markov property, and the empirical product
752 converges to semantic-path completion. $\square$

753 A.2 Argument for Definition 2 (multi-agent estimability)

754 Conditioning sets are enlarged by $\mathcal{M}_{-i}^{(K)}$ and $\mathcal{O}_{-i}$. Under the enlarged stage-Markov property, the
755 proof of Definition 1 transfers verbatim and each enlarged conditional is a Bernoulli probability on
756 a measurable event. The qualifications stated in §4.2 (“When Definition 2 may fail”) describe two
757 specific situations in which the enlarged stage-Markov property is violated. $\square$

758 A.3 Argument for Empirical Claim 1 (slope-level statement)

759 Empirical Claim 1 is an empirical claim, not a theorem. The argument below is the qualitative
760 justification that motivates the two slope signs; the data are what verify them. We deliberately do not
761 call this a proof.

762 *Slope (a), under (F).* As K decreases, $\nu(\mathcal{M}_{-i}^{(K)})$ is non-decreasing pointwise. Conditional on
763 a successful retrieval, the probability of locking onto a reachable target is non-decreasing in ν .
764 Therefore $\partial P(M^*|R^*; K)/\partial K^{-1} > 0$.

765 *Slope (b), under (O).* As K decreases, peers receive updated memory simultaneously and con-
 766 verge on the same closest target; the marginal occupancy $\Pr[\text{target} \in \mathcal{O}_{-i}]$ rises. Therefore
 767 $\partial P(C^*|M^*; K)/\partial K^{-1} < 0$.

768 *Note on the conditional product Φ .* Empirical Claim 1 makes no statement about the shape of
 769 $\Phi(K) = P(M^*|R^*; K) \cdot P(C^*|M^*; K)$. Under (F) and (O) the slope of Φ in K^{-1} is the sum of two
 770 opposing terms, and whether it changes sign depends on quantitative magnitudes that are not pinned
 771 down by (F) and (O) alone. In our data Φ is monotone in K over $\{4, 8, 16, 32, 48, 64\}$ and converges
 772 to the independent-agent boundary $\Phi_\infty = 0.605$. The practically relevant efficiency surface — mean
 773 time-to-success — does have an interior minimum at $K = 8$; see §4.5. $\square$

774 A.4 The query–retrieval Galois connection: statement, argument, scope

775 **Fact (standard).** Let G and Σ be sets and $I \subseteq G \times \Sigma$ a relation. For $B \subseteq \Sigma$ define $\text{ext}(B) =$
 776 $\{p \in G : \forall \tau \in B, (p, \tau) \in I\}$, and for $A \subseteq G$ define $\text{int}(A) = \{\tau \in \Sigma : \forall p \in A, (p, \tau) \in I\}$.
 777 Then for all $A \subseteq G, B \subseteq \Sigma$:

$$A \subseteq \text{ext}(B) \iff B \subseteq \text{int}(A);$$

778 both maps are antitone, both composites $\text{ext} \circ \text{int}$ and $\text{int} \circ \text{ext}$ are closure operators, and the closed
 779 pairs (A, B) with $A = \text{ext}(B)$ and $B = \text{int}(A)$, ordered by extent inclusion, form a complete lattice
 780 — the concept lattice of the context (G, Σ, I) (Ganter and Wille, 1999).

781 **Argument.** $A \subseteq \text{ext}(B)$ unfolds to $\forall p \in A \forall \tau \in B : (p, \tau) \in I$, a condition symmetric in A
 782 and B , which is exactly $B \subseteq \text{int}(A)$. Antitonicity and the closure properties follow by the standard
 783 one-line arguments (Ganter and Wille, 1999, Ch. 1). Viewing the power sets as thin categories (one
 784 with reversed order), the displayed equivalence says precisely that int and ext form an adjoint pair; a
 785 Galois connection is an adjunction between posets (Mac Lane, 1998, Ch. IV). $\square$

786 **Instantiation and scope.** In our memory, take G to be the place records, Σ the tag alphabet,
 787 and $I_\theta = \{(p, \tau) : \text{the per-place tag table assigns } \tau \text{ to } p \text{ with confidence } \geq \theta\}$ (Appendix D); the
 788 intent of a query is $B = q^{\text{req}}$ and R-availability is $\text{ext}(B) \neq \emptyset$. Three deliberate scope restrictions.
 789 (i) *Graded scores.* The deployed retrieval score is graded (the weighted log-score over required,
 790 preferred, and penalty tags of Appendix D), so the crisp connection above formalizes the thresholded
 791 required-tag skeleton of candidacy; the graded analogue is the theory of fuzzy Galois connections
 792 (Bělohlávek, 1999) and we do not use it here. (ii) *Non-monotone updates.* Dirichlet–multinomial
 793 updates are not literally monotone in I_θ : contrary evidence can remove an incidence. On the
 794 reported benchmarks the binding failure is absence of evidence rather than contradiction (91% empty
 795 extents, 96.9% precision when non-empty, §3.5), so the monotone-growth reading is the empirically
 796 relevant regime, and we flag the exception rather than assume it away. (iii) *No structure for M or C .*
 797 Materialization is a satisficing choice function on non-empty extents and Completion is a control
 798 problem under occupancy; neither carries the order structure, and we attach no categorical claims to
 799 them. The reading earns its place by naming the failure mode exactly (a queried intent whose extent
 800 is empty lies outside the realized concept lattice), by giving consolidation a precise object (growth of
 801 the context, hence of every extent), and by locating the cadence axis structurally: memory coupling
 802 $\mathcal{M}_{-i}^{(K)}$ merges peer contexts at rate K^{-1} , while occupancy $\mathcal{O}_{-i}$ is extra-contextual — which is the
 803 structural content behind the two slope signs of Empirical Claim 1.

804 B An options / hierarchical-MDP reading of the stages

805 Why a deterministic strategy: bounded and ecological rationality

806 A natural objection is that the semantic stack is a fixed, deterministic strategy rather than a learned
 807 policy. We treat the determinism as principled and name it: the pipeline is a satisficing agent in
 808 Simon’s sense (Simon, 1955, 1956). Candidacy is threshold-gated instead of exhaustively scored,
 809 the target commitment is made once instead of re-optimized at every tick, and the intent is a few
 810 tags instead of the full state. Bounded rationality is also what makes the four failure loci observable:
 811 each stage is a distinct resource bound (attention, memory, deliberation, control), and Q/R/M/C is
 812 the audit trail of where the bound binds. The learned baselines make the converse point empirically:
 813 policies without an explicit query/lock interface collapse the stages even at competitive success

(BC in §3.3; CommNet PPO in §4.5), so observability is traded away exactly where end-to-end optimization is traded in. That the *same* heuristic is strong on water and rest, retrieval-starved on hazard recovery, and downstream-limited on goal-region (Figure 1b) is the ecological-rationality point that heuristic success is a property of the heuristic–environment pair (Simon, 1956; Gigerenzer and Goldstein, 1996); the framework is the instrument that locates the mismatch. The same lens reads the Φ -versus- t_{succ} split of §4.5: the chain probability Φ ignores the cost of traversing the chain, so under a resource-rational reading (Lieder and Griffiths, 2020) an optimum should be sought on a cost-bearing metric; that is where it appears ($K=8$ on mean time-to-success) and not on Φ .

This appendix also develops the dynamic reading pointed to from §2.3. It introduces no new algorithm: it names the object the pipeline already computes and makes the Φ -versus- t_{succ} split a property of that object. Where the order-theoretic reading (Appendix A.4) fixes the *static* content of the stages, this one addresses time and cost.

Q/R/M/C as an audit of an option. Each query instantiates one temporally extended action—an option $o = \langle \mathcal{I}, \pi, \beta \rangle$ in the sense of Sutton et al. (1999): an initiation set $\mathcal{I}$, an internal policy π , and a termination condition β . Standard trajectory-level analysis treats an option as a black box scored by its return. The Q/R/M/C decomposition is instead an audit of the option’s *internal* structure, and it lines up edge-for-edge with the factorization $Q^* \rightarrow R^* \rightarrow M^* \rightarrow C^*$:

$Q^* \rightarrow R^*$ (**initiation**). The option “reach a place satisfying q ” is well-initiated only when its initiation set is non-empty. In the vocabulary of §2.3 this is exactly $\text{ext}(q^{\text{req}}) \neq \emptyset$: an empty extent is not a bad choice of option but the absence of any admissible one, so R fails before a target is ever locked.

$R^* \rightarrow M^*$ (**policy over options**). Materialization is the high-level policy that commits to a single option out of the admissible set—the target lock $g^* = \chi(\text{ext}(q^{\text{req}}))$ under the satisficing choice function χ of §2.3.

$M^* \rightarrow C^*$ (**execution to termination**). Completion is the roll-out of the internal policy π until the termination condition β fires (arrival at g^* , or timeout). This is the only edge whose failure modes are extra-semantic: controller competence and, in a collective, occupancy $\mathcal{O}_{-i}$.

So Q/R/M/C is, on this reading, a stage-resolved instrument for the internal structure of a hierarchical-MDP option generator, rather than a bespoke logger for a single environment.

Why the optimum lives on t_{succ} and not on Φ : a semi-MDP reading. In the options / semi-MDP formalism, an option is evaluated by a return accumulated over a *random* execution duration τ ,

$$R(s, o) = \mathbb{E} \left[\sum_{t=0}^{\tau-1} \gamma^t r_{t+1} \mid s_0 = s, o \right],$$

where τ is the number of low-level steps until β fires (Sutton et al., 1999). The chain probability $\Phi = P(M^*|R^*) \cdot P(C^*|M^*)$ records only *whether* the option-audit succeeds; it is by construction blind to τ , which is the cost of traversing it. The cost-bearing metric t_{succ} is τ . This is the structural reason—and not merely an empirical coincidence—that the resource-rational reading of §2.3 (Lieder and Griffiths, 2020) locates an interior optimum on t_{succ} ($K=8$) while Φ stays monotone: the two metrics measure different faces of the same option, success versus duration.

The bottleneck shift as multi-agent option collision. The dynamic reading makes the $M \leftrightarrow C$ trade-off of Empirical Claim 1 physically concrete. Faster broadcast (smaller K) merges peer contexts $\mathcal{M}_{-i}^{(K)}$ more often, which can only enlarge each agent’s extent and therefore raises the probability that a target is admissible and locked: $P(M^*|R^*)$ increases. But because agents share a spatial basis and a tag alphabet (§2.1), fresher shared memory also *synchronizes* the high-level policy: several agents commit to the same option over the same scarce target. Under occupancy $\mathcal{O}_{-i}$ this inflates the execution duration τ and depresses $P(C^*|M^*)$. Cadence K thus regulates the dispersion of the agents’ option distribution: rare exchange desynchronizes option selection and shortens τ , giving the minimum of t_{succ} at an interior K . The bottleneck is therefore grounded in a competition between initiation-set growth (semantic, contextual) and option-collision cost (extra-semantic, occupancy), exactly the two edges the cadence protocol $\mathcal{C}_K$ acts on in Figure 2.

Belief–desire–intention gloss and scope. The high-level state audited above can be named in belief–desire–intention terms, which is the form used by the collective-memory companion: *belief* is

the realized concept lattice of §2.3 (what is retrievable), *desire* is the state-conditioned predicate that selects the query q (e.g. a thirst variable switching the required tags to `water_source`), and *intention* is the locked option g^* . The hierarchical state is then $h = (\text{lattice}, \text{desire}, g^*, s)$ with s the low-level physical state, and the three failure loci separate as before: an empty extent is a belief/ R failure, a non-empty extent with no stable lock is an M failure, and a lock unreachable under occupancy is a C failure. Two scope caveats. We claim no options-*learning* result: the pipeline is a fixed satisficing option generator (Simon, 1955) and Q/R/M/C audits it. And we attach the option structure only to the semantic path— $\mathcal{I}$ to $Q \rightarrow R$, (π, β) to $M \rightarrow C$ —while the semi-MDP return is used as an explanatory lens for the Φ -versus- t_{succ} split, not fitted to the data. Neither this options reading nor the order-theoretic reading of Appendix A.4 is claimed as a contribution: both are standard constructions used as interpretive lenses on the measurement protocol, which is the paper’s contribution. Their value is that each makes a testable statement on the logs—an empty extent localises the R failure, and the semi-MDP duration τ predicts that the optimum appears on t_{succ} and not on Φ .

C Operational vs. nested stage estimators

Operational diagnostic rates. The main benchmark logs per-query and per-episode diagnostic quantities directly from runtime traces (query formed, retrieval satisfied, target materialized, completion after retrieval). These are operational rates, not direct multiplicative factors.

Nested stage estimators for factorization. For the factorization, episode-level nested events from the first failure stage recorded in the runtime taxonomy are used: Q^*, R^*, M^*, C^* , nested by construction. The consistency audit in §3.2 reports $\hat{P}(Q^*), \hat{P}(R^*|Q^*), \hat{P}(M^*|R^*), \hat{P}(C^*|M^*)$, their product, and its gap to semantic-path completion.

C.1 Multi-agent operational definitions

For the multi-agent sweep, every metric is defined per agent and per tick, then aggregated. Let $\text{Targets}_i(t)$ be the candidate set returned by the memory query of agent i at tick t , $\text{Locked}_i(t)$ be the agent’s planner-internal locked target, and $\mathcal{W}$ the ground-truth target set.

Per-tick events. $Q_i(t) = \mathbf{1}[\text{Targets}_i(t) \neq \emptyset]$; $R_i(t) = \mathbf{1}[\exists w \in \mathcal{W} : \min_c \|c - w\| \leq \epsilon]$ ($\epsilon = 0.6$); $M_i(t) = \mathbf{1}[\text{Locked}_i(t) \neq \emptyset \wedge \min_w \|\text{Locked}_i(t) - w\| \leq \epsilon]$.

Episode-level starred events. $Q_i^* = \bigvee_t Q_i(t)$; analogously for R^*, M^* ; $C_i^* = \mathbf{1}[\text{agent } i \text{ reached a target}]$.

Aggregation. Per-run rates are averaged over the N agents; per-configuration rates are averaged over the 40 seeds. Conditional rates $P(R^*|Q^*), P(M^*|R^*), P(C^*|M^*)$ are computed as ratios of episode-level counts within each configuration before averaging across configurations.

D Memory and task schema

D.1 Full notation table

Symbols are fixed throughout the paper.

Stages. Q, R, M, C — the four stages: **Q**uery formation, **R**etrieval, **M**aterialization, **C**ompletion after retrieval. They name both the stage and (in expressions like “the M link”) the corresponding edge of the chain.

Events. $Q^*, R^*, M^*, C^* \in \{0, 1\}$ — episode-level starred events recording whether each stage succeeded (Appendix C). Nested by construction: $Q^*=1 \Leftarrow R^*=1 \Leftarrow M^*=1 \Leftarrow C^*=1$.

Probabilities. $P(\cdot)$ — population-level probability (averaged over the seed distribution and, in the multi-agent case, over agents). $\hat{P}(\cdot)$ — empirical-frequency estimator computed from logs.

Single-agent. q — query; $\mathcal{M}$ — memory state; Y — end-to-end task success (can exceed C^*).

907 **Multi-agent.** N — number of agents in the population. T — number of targets in the environment
 908 (constraint $T \leq N$; “scarcity” means $N > T$). i — agent index. $\mathcal{M}_i$ — agent i ’s private memory. q_i
 909 — agent i ’s query.

910 **Coupling.** $\mathcal{C}_K$ — the peer-broadcast communication protocol with cadence $K \in \mathbb{N}$ (every K
 911 ticks each agent broadcasts a memory snapshot, §4.1). $\mathcal{M}_{-i}^{(K)}$ — the latest peer memory snapshots
 912 accessible to agent i under $\mathcal{C}_K$ (the *memory coupling* channel). $\mathcal{O}_{-i}$ — the set of targets already
 913 claimed by other agents by the time i acts (the *occupancy coupling* channel).

914 **Slope variables.** $\nu(\mathcal{M}_{-i}^{(K)})$ — a freshness functional on peer memory; non-decreasing pointwise as
 915 K decreases (assumption **(F)** of Empirical Claim 1). The occupancy slope is assumption **(O)**.

916 **Summary statistics.** $\Phi(K) := P(M^*|R^*; K) \cdot P(C^*|M^*; K)$ — the conditional product, the chain
 917 probability $P(M^* \wedge C^* | R^*)$. POM (product-of-means) and MOP (mean-of-products) are two
 918 estimators of Φ ; they differ by $\text{Cov}(P(M^*|R^*), P(C^*|M^*))$, which is the Jensen gap (§4.5). t_{succ}
 919 — mean time-to-success in environment ticks.

920 D.2 Memory stack and task patterns

921 The memory stack stores symbolic place records (per-place tag confidences, counts, visitation,
 922 freshness, geometry), transition records (success, risk, cost statistics with fast/slow estimates),
 923 episodic traces (intent + state at each visit), and concept/relation records. At each step, the tokenizer
 924 emits (z_t, τ_t, e_t) from o_t . Per-place tag tables use a Dirichlet-multinomial update; transitions use
 925 exponentially smoothed statistics; episodic records attach the active intent and agent state. A planner
 926 query is answered by retrieving candidate places whose semantic profiles match the current task intent
 927 and state, via a weighted log-score $s(p | q_t, s_t) = \sum_{\tau \in q^{\text{req}}} \log P(\tau|p) + \alpha \sum_{\tau \in q^{\text{pref}}} \log P(\tau|p) -$
 928 $\beta \sum_{\tau \in q^{\text{pen}}} \log P(\tau|p) + \gamma u_{\text{epi}}(p, s_t)$ over required, preferred, and penalty tags. The four task
 929 families are defined by semantic place patterns: water search (target {water}), safe rest ({rest} with
 930 state-conditioning), goal-region rejoin ({goal-region}), hazard recovery ({post-hazard-exit}).

931 D.3 Single-agent supporting tables and figures

Table 8: Best semantic result per task family from the final 10-seed benchmark, with 95% bootstrap CIs. Hazard-recovery is bottlenecked by retrieval; goal-region remains downstream-bottlenecked.

Task	Best method	Assist	Success (95% CI)	Steps	Query (op.)	Retrieval (op.)	Mat. (op.)	Post (op.)
Water	water-cp	on	0.95 [0.88, 1.00]	17.99	0.71	0.71	0.71	0.70
Rest	rest-s2	on	0.93 [0.85, 0.99]	23.35	0.61	0.61	0.61	0.73
Goal region	goal-n1	off	0.54 [0.52, 0.56]	49.34	0.00	0.00	0.00	0.00
Hazard recovery	hazard-s7	on	0.40 [0.35, 0.45]	10.18	0.11	0.11	0.11	0.16

Oracle stage interventions localise the bottleneck (single-agent, 10 seeds $\times$ 8 episodes)




Figure 6: Oracle stage interventions on the two hardest single-agent task families (10 seeds $\times$ 8 episodes). **(a)** FourRooms goal-rejoin: replacing any single stage with an oracle leaves end-to-end success at 0.07; the bottleneck is downstream of the stages in the chain. **(b)** LavaGapS7 hazard recovery: oracle controller adds only +0.04, but oracle materialization adds +0.16 and oracle retrieval adds a further +0.01 to reach 0.60. The framework localises the bottleneck stage-by-stage.

932 **Seed robustness (30 seeds).** The single-agent benchmark of §3 uses 10 seeds. To check the
933 diagnosis is not a small-sample artefact, we reran the two diagnostic task families at 30 seeds
934 (`-num_seeds 30`; $n=240$ hazard-recovery and 480 goal-region agent-episodes). The *qualitative*
935 attribution is unchanged: hazard recovery stays retrieval-starved (empty-candidate-set rate ≈ 0.62 –
936 0.65 , low success) while goal-region is far less retrieval-limited (empty-set ≈ 0.37 , success ≈ 0.61)
937 — the same stage *ordering* the framework reports at 10 seeds. Absolute rates depend on the exact
938 benchmark configuration (episode budget, assist and env-change settings) and so differ from the
939 headline 0.39/0.91 of the 10-seed oracle configuration; what 30 seeds confirms is the ordering the
940 diagnosis rests on, not a particular level.

941 E Portability and scale-up

942 E.1 BabyAI portability transfer (10 seeds, full benchmark)

943 A 10-seed portability benchmark on BabyAI-GoToObjMaze-v0 (Chevalier-Boisvert et al., 2018)
944 (8 episodes per seed, max-steps 200) measured all four single-agent methods used in the main
945 benchmark (random, greedy, `babyai_semantic_node_v1`, `babyai_semantic_plan_v1`) with
946 the same Q/R/M/C logger applied to BabyAI traces. The numbers in Table 9 confirm that (a) the
947 Q/R/M/C events emit non-trivially on the BabyAI substrate, (b) the two semantic variants expose
948 stage-separable structure where the non-semantic baselines have $Q^*=R^*=M^*=0$ (no semantic
949 queries are issued), and (c) the stage-rate ordering matches the MiniGrid benchmark (Table 2) within
950 bootstrap CIs. GoToObjMaze is a harder task substrate than the LavaGapS7/FourRooms benchmark
951 used in the main text — end-to-end success caps at 0.15 across all methods — which makes the stage
952 decomposition especially informative: the semantic stack is not winning on success but it is the only
953 method family producing measurable per-stage events.

954 **Takeaway.** The framework’s measurement protocol — not its task-specific results — transfers to
955 a second substrate. End-to-end success on BabyAI is dominated by the much lower episode-level
956 success ceiling of GoToObjMaze, not by the stage structure; conversely, the stage structure that
957 the semantic stack exposes is identical in shape to what the same stack exposes on MiniGrid. This
958 is exactly the substrate-independence property the methodology is intended to provide. Raw data:
959 `tmp/babyai_deeper_10seed/`.

960 E.2 Multi-agent scale-up at $N = 12$

961 To check that the multi-agent results of §4.5 are not an artefact of the $N \leq 8$ range, an additional
962 8,640-run sweep was run at $N = 12$, $T \in \{6, 8, 10\}$, the same three layouts, three hazard densities,

Table 9: BabyAI portability benchmark: 10 seeds $\times$ 8 episodes on BabyAI-GoToObjMaze-v0, max-steps 200. The semantic-stack variants succeed at the same end-to-end rate as the greedy baseline (0.15) but are the only methods that produce non-zero Q/R/M/C events — the framework’s protocol transfers to the BabyAI substrate without modification. Q/R/M/C operational rates here are evaluated on the same denominator (number of decision points), so $Q=R=M$ for both semantic variants because GoToObjMaze emits one semantic target type per episode.

Method	Success	Steps	Q (op.)	R (op.)	M (op.)	Post (op.)
random	0.125 ± 0.06	182.7	0.000	0.000	0.000	0.000
greedy	0.150 ± 0.05	171.3	0.000	0.000	0.000	0.000
semantic node	0.150 ± 0.05	171.3	0.032	0.032	0.032	0.025
semantic plan	0.150 ± 0.05	171.3	0.032	0.032	0.032	0.025

963 eight peer cadences, and 40 seeds. Wall-clock $\approx$ 12 minutes on 8 cores. The conditional-rate
964 signature of §4.5 persists and intensifies (Table 10).




Figure 7: $N = 12$ scale-up bottleneck signature ($n = 8,640$ runs, 40 seeds). Top: the per-cadence mean rates move in the predicted directions in K ; the within-cadence Pearson correlation (red dashed, right axis) peaks at -0.70 at $K = 8$ (this within-cadence peak is what sharpens at $N=12$; the M-side K -slope itself is not cluster-robust at this agent count, Table 10). Bottom: mean time-to-success retains an interior minimum at $K = 8$ (4.77 ticks); as on the main sweep, the location is robust but the margin is shallow — the gap to the adjacent $K = 4$ (4.82 ticks) is ≈ 0.05 tick and is not resolved at this sample size. The bottleneck-shift peak migrates outward from $K = 4$ at $N \leq 8$ to $K = 8$ at $N = 12$, consistent with stronger occupancy coupling at higher agent count delaying the cost-regime transition.

Table 10: $N=12$ scale-up sweep numerical detail (peer architecture; 27 configuration cells). Effect sizes with cluster-robust 95% CIs (block bootstrap over cells): Spearman of $P(C^*|M^*)$ with K is $+0.31$ (CI $[+0.21, +0.41]$, excludes zero), while Spearman of $P(M^*|R^*)$ with K is only -0.14 (CI $[-0.28, +0.02]$, *not* resolved from zero at $N=12$). The C-slope of Empirical Claim 1 holds; the M-slope is directionally consistent but not cluster-robust at this agent count. What sharpens at $N=12$ is the within-cadence correlation and the t_{succ} minimum (Figure 7), not the M-side K -slope.

K	$P(M^* R^*)$	$P(C^* M^*)$	MOP	mean t_{succ}	corr
1	0.758	0.679	0.475	5.49	-0.59
2	0.745	0.692	0.474	5.29	-0.61
4	0.679	0.768	0.482	4.82	-0.64
8	0.731	0.767	0.539	4.77	-0.70
16	0.687	0.803	0.538	5.04	-0.62
32	0.676	0.814	0.538	5.75	-0.51
48	0.677	0.814	0.539	5.96	-0.51
64	0.677	0.816	0.541	6.12	-0.51

965 E.3 Q/R/M/C under semantic, frame-free place matching

966 The main sweep resolves place identity by shared (x, y) keys. To check that the decomposition is
 967 not tied to a shared coordinate substrate, we rerun a multi-agent navigation loop in which every
 968 agent holds a *private* coordinate frame (a secret per-agent translation) and places are matched
 969 *by meaning*: local semantic fingerprint constellations recover the inter-agent frame and transport a
 970 peer’s target evidence into the receiver’s frame (experiments/warp/semantic_peer_memory.py;
 971 runner experiments/warp/exp_semantic_cadence_qrmc.py). Agents navigate with the same
 972 rule-based planner toward the transported target; Q/R/M/C are logged unchanged, with M^* the lock
 973 of a candidate within ϵ of a true target ($N=6$, $T=2$ scarce water sources, 8 seeds).

974 The stage events remain well-defined without a shared frame, and the framework localises the
 975 coordinate assumption’s failure. The *semantic* arm locks a *true* target on its first materialisation
 976 in 100% of agent-episodes (at every cadence); the *coordinate* arm — which reads a peer’s private
 977 coordinates at face value — first-locks a true target only 25–52% of the time (rising with K as agents
 978 accumulate their own observations), i.e. it *fails open*, materialising phantom cells that are not water.
 979 End-to-end success is scarcity-limited and comparable across arms ($\approx 0.29 \approx T/N$), as expected.
 980 This is the diagnostic property of interest: M^* under coordinate identity fires on non-targets, and the
 981 stage logger exposes it.

982 We do *not* claim the $M \leftrightarrow C$ cadence shift here: with an all-to-all broadcast and a 140-tick horizon
 983 the semantic frame recovers at every tested K , so $P(M^*|R^*)$ is cadence-invariant in this setup
 984 (1.00 throughout). The experiment isolates the *place-identity substrate* question — Q/R/M/C is
 985 architecture- and coordinate-independent, and detects fail-open materialisation — rather than the
 986 cadence result, which stands on the main coordinate-keyed sweep. Folding semantic matching into
 987 the scarcity regime that produces the shift is the natural next integration.

988 F Extended diagnostics

989 F.1 Conditional-product detail and the Jensen gap

990 The conditional product $\Phi = P(M^*|R^*) \cdot P(C^*|M^*)$ in the 35,640-run sweep is monotonically
 991 increasing in K over $K \in \{4, 8, 16, 32, 48, 64\}$ and converges to the independent-agent boundary
 992 $\Phi(K \rightarrow \infty) = 0.605$. Paired bootstrap: $\Delta(K=32-K=16) = +0.008$ (CI $[+0.006, +0.010]$),
 993 $\Delta(K=48-K=16) = +0.011$, $\Delta(K=64-K=16) = +0.012$; the independent-agent level sits
 994 above every peer K except $K=64$ (where the gap’s CI includes zero). This is a statement about the
 995 *level* of Φ at the tested cadences, not about the *shape* of $\Phi(K)$: consistent with the monotone trend,
 996 we make no claim of an interior extremum (§4.3).

997 The negative within-cadence covariance documented in §4.5 creates the Jensen-inequality gap
 998 between Φ and the marginal product-of-means: $E[X] \cdot E[Y] > E[X \cdot Y]$ when $\text{Cov}(X, Y) < 0$. The
 999 mean-of-products is reported as the primary summary because it absorbs the covariance correctly.

Table 11: Per-architecture conditional Q/R/M/C rates from the 35,640-run sweep (40 seeds). $P(R^*|Q^*)$ is saturated. POM = product-of-means; MOP = mean-of-products (the proper per-configuration expected product). The two summaries disagree systematically at small/intermediate K because the within-cadence correlation between $P(M^*|R^*)$ and $P(C^*|M^*)$ is strongly negative there (last column) — the bottleneck-shift signature. **Bold:** MOP-best of the peer cadences. Mean t_{succ} (last column) is the practically relevant efficiency measure; its interior minimum at $K = 8$ is the predicted-shape result, robust in location but shallow in margin (adjacent fast cadences not resolved; Appendix F.3).

Architecture	$P(R^* Q^*)$	$P(M^* R^*)$	$P(C^* M^*)$	POM	MOP (95% CI)	corr	mean t_{succ}
independent	0.99	0.68	0.88	0.598	0.605 [0.597, 0.614]	+0.00	8.02
shared bus	1.00	1.00	0.60	0.600	0.595 [0.586, 0.604]	−0.10	8.33
central aggregator	1.00	0.97	0.61	0.592	0.572 [0.564, 0.580]	−0.48	7.64
peer ($K = 1$)	1.00	0.97	0.60	0.585	0.573 [0.564, 0.581]	−0.43	7.80
peer ($K = 2$)	1.00	0.97	0.62	0.594	0.576 [0.568, 0.584]	−0.54	7.59
peer ($K = 4$)	1.00	0.92	0.66	0.599	0.569 [0.561, 0.577]	−0.61	8.00
peer ($K = 8$)	0.99	0.69	0.87	0.597	0.586 [0.578, 0.595]	−0.19	7.47
peer ($K = 16$)	0.99	0.67	0.89	0.592	0.590 [0.581, 0.598]	−0.05	7.87
peer ($K = 32$)	0.99	0.67	0.89	0.595	0.598 [0.589, 0.606]	−0.04	8.28
peer ($K = 48$)	0.99	0.68	0.88	0.598	0.602 [0.593, 0.611]	−0.02	8.79
peer ($K = 64$)	0.99	0.68	0.88	0.599	0.603 [0.594, 0.611]	−0.02	9.19

Figure 8 shows the K-sweep view of both conditional rates and mean t_{succ} ; Figure 9 shows the per-configuration scatter that the negative covariance summarises.

F.2 Design decomposition and run-count reconciliation

The 35,640 figure decomposes exactly as in Table 12. The base design is the nine admissible (N, T) scarcity pairs (for $N=3$: $T \in \{2, 3\}$; $N=5$: $T \in \{2, 3, 5\}$; $N=8$: $T \in \{2, 3, 5, 8\}$) crossed with three layouts, three hazard densities, and 40 seeds, giving 3,240 base configurations. The peer architecture is run at all eight cadences ($3,240 \times 8 = 25,920$); the three cadence-free architectures (independent, shared bus, central aggregator) contribute 3,240 each. The 25,920-run figure used for the peer cluster-robust analysis (Appendix F.3) is exactly the peer block; the initial 12,960-run exploratory sweep (§4.3, “Chronology”) and the extended-cadence and $N=12$ sweeps (Appendix E) are separate run sets and are not pooled with the main sweep.

Table 12: Run-count reconciliation for the main multi-agent sweep (35,640 runs). Counts are recomputed directly from the released run log.

Factor	Levels	Count
(N, T) scarcity pairs	$(3, 2)(3, 3)(5, 2)(5, 3)(5, 5)(8, 2)(8, 3)(8, 5)(8, 8)$	9
Layouts	symmetric, asymmetric, random	3
Hazard densities	0.0, 0.05, 0.10	3
Seeds		40
Base configurations	$9 \times 3 \times 3 \times 40$	3,240
Peer (all 8 cadences K)	$3,240 \times 8$	25,920
Independent / shared / central	$3 \times 3,240$	9,720
Total		35,640

F.3 Effect sizes and cluster-robust inference

The $p < 10^{-4}$ figures quoted in §4.5 are reported for completeness but are not the basis of the claim: at tens of thousands of runs almost any non-zero slope is “significant,” and the runs are not independent—they cluster within the $(N, T, \text{layout}, \text{hazard})$ configuration. The full sweep has 25,920 peer runs but only **81 independent design cells**. We therefore report effect sizes with cluster-robust confidence intervals from a block bootstrap that resamples *configuration cells* (not rows), $B=4000$ resamples.

The two slopes of Empirical Claim 1 survive as genuine, but moderate, effects: the Spearman correlation of $P(M^*|R^*)$ with K is $\rho = -0.53$ (95% cluster CI $[-0.59, -0.47]$, strong) and of




Figure 8: Bottleneck shift across eight peer cadences (40-seed sweep). Top: $P(M^*|R^*)$ falls and $P(C^*|M^*)$ rises in K (moderate, cluster-robust Spearman effect sizes -0.53 and $+0.43$; CIs in Appendix F.3). Bottom: mean time-to-success has an interior minimum at $K=8$ (7.47 ticks) below the independent-agent baseline (8.02, dashed); the location is robust but its margin over the adjacent fast cadences ($K=1, 4$) is not resolved (Appendix F.3). Error bars and shaded band are 95% bootstrap CIs.

1020 $P(C^*|M^*)$ with K is $\rho = +0.43$ (95% cluster CI $[+0.37, +0.50]$, moderate). Both intervals exclude
 1021 zero under cell-level resampling, so the bottleneck shift is not an artefact of the correlated run count;
 1022 it is a moderate monotone trend, and we describe it as such rather than as a near-deterministic law.

1023 The interior optimum on mean time-to-success is robust in *location* but shallow in *margin*. Paired
 1024 block bootstrap on the cell-level t_{succ} differences gives $t_{succ}(K=8)$ below $t_{succ}(K=16)$ by 0.40 ticks
 1025 (CI $[0.08, 0.83]$) and below $K=64$ by 1.67 ticks (CI $[0.75, 2.83]$), both resolved; but its advantage
 1026 over $K=4$ (0.49, CI $[-0.47, 1.51]$) and $K=1$ (0.27, CI $[-0.43, 0.97]$) is within noise.¹ A success-
 1027 weighted efficiency metric (t_{succ} divided by success rate) has the same interior arg min at $K=8$,
 1028 so the location of the optimum is robust to the efficiency definition even though its margin over
 1029 the adjacent fast cadences is not resolved at this sample size. We read this as support for the
 1030 *shape* predicted by the $M \leftrightarrow C$ trade-off, not for a sharply tuned optimal cadence. Two further
 1031 cell-paired contrasts complete the picture. On t_{succ} , $K=8$ is also resolved against the independent
 1032 baseline (-0.40 , CI $[-0.59, -0.24]$). On a censoring-aware expected time $\mathbb{E}[T] = s t_{succ} + (1 - s) T_{max}$
 1033 (failed runs charged the step limit $T_{max}=120$), the fast flank reverses: $K=8$ beats $K=4$
 1034 (-2.10 , CI $[-3.61, -0.72]$) but loses to $K=16$ ($+0.69$ $[+0.09, +1.43]$) and to independent ($+0.32$
 1035 $[+0.05, +0.63]$). The interior optimum is thus a property of conditional speed among successes;
 1036 on unconditional efficiency the completion collapse of fast sharing dominates, which is the gap the
 1037 reservation protocol of open problem (3) closes.

¹This paired bootstrap and the effect-size Spearman correlations above are computed on the same sweep execution as Table 11 (25,920 peer runs); the t_{succ} statistic is defined only over successful episodes, so the paired cell-level difference (1.67 for $K=8$ vs. $K=64$) is on that subpopulation and matches the marginal-mean difference in Table 11 ($9.19 - 7.47 = 1.72$) to within rounding. The effect sizes are identical on the independent `paper1_clean` execution ($-0.531/+0.434$), confirming they are not run-specific.

Bottleneck-shift signature: negative within-cadence correlation peaks at $K = 4$ and decays at $K = 64$




Figure 9: Per-configuration scatter of $P(M^*|R^*)$ vs. $P(C^*|M^*)$ at three cadences. The strong negative correlation at $K = 4$ (Pearson -0.61) is the bottleneck-shift signature: the same configurations that succeed at M tend to fail at C and vice versa. The correlation decays to noise by $K = 64$.

1038 F.4 Construct validity: task success Y vs. semantic-path completion C^*

1039 A success Y can in principle bypass the semantic path, so a stage bottleneck in $Q/R/M/C$ need not
 1040 explain variation in Y . Two facts bound this concern. First, in the multi-agent sweep that Empirical
 1041 Claim 1 is about, success is *defined* as reaching a `water_source` target, which is exactly the C^*
 1042 event; empirically Y and C^* coincide run-for-run (correlation 1.00, mean gap 0.00). The central
 1043 multi-agent claim is thus not exposed to a parallel-channel objection. Second, the gap $Y > C^*$ is
 1044 a single-agent phenomenon on the easy task families (water, rest), where a fallback controller can
 1045 reach the goal without a semantic target lock; it is largest exactly where success is already high and
 1046 the diagnosis is least needed, and it vanishes on the hard families (hazard recovery, goal-region)
 1047 where the framework does its work. A per-episode, stage-resolved mediation of Y (rather than of
 1048 C^*) requires logs that tag each success with whether it traversed the semantic path, and is a planned
 1049 instrumentation refinement.

1050 F.5 Threshold sensitivity of the $Q/R/M/C$ event definitions (θ and ϵ)

1051 The single-agent attribution “candidate generation fails, not ranking” (§3.5) rests on a
 1052 high empty-candidate-set rate. A reviewer asks whether that is an artefact of the
 1053 required-tag threshold θ (deployed 0.5; the candidacy gate is `profile[tag] ≥ θ` in
 1054 `symbolic_memory._required_matches_profile`). We sweep θ via an override on the query
 1055 builder (`experiments/big_experiment/theta_ablation.py`); the endpoints confirm the gate
 1056 is live, and the operating range shows the diagnosis is not a threshold effect.

Table 13: Empty-candidate-set rate vs. required-tag threshold θ (semantic method, assist off, hazard-recovery / goal-region). The rate is *invariant* across the entire realistic range $\theta \in [0.05, 0.95]$; only the degenerate endpoints move it, and only to the no-candidate floor. The empties are therefore absence of stored evidence, not sub-threshold filtering.

θ	empty rate (hazard recovery)	empty rate (goal region)
0.00 (accept all)	0.50	—
0.05	0.533	0.267
0.15	0.533	0.267
0.25	0.533	0.267
0.50 (deployed)	0.533	0.267
0.75	0.533	0.267
0.95	0.533	0.267
1.00 (accept none)	0.75	—

The reading is decisive (Table 13). At $\theta=1.0$ nothing passes and the empty rate rises; at $\theta=0.0$ the required-tag gate accepts every node ($\text{conf} \geq 0$ always holds), yet the empty rate falls only to 0.50, its floor—meaning that at those decision points there is *no candidate node of any confidence* to retrieve. Across the entire sane range $\theta \in [0.05, 0.95]$ the rate is flat at 0.533: loosening or tightening the threshold does not recover a candidate, because the required semantic place was never stored. This is exactly the “absence of evidence, not contradiction” regime named in §2.3, now shown to be threshold-invariant, so the memory-accumulation attribution is not an artefact of θ . (Absolute rates here are on a small hazard-recovery / goal-region sweep and are not the headline 91% of §3.5, which is on the full benchmark; the point is the θ -invariance, not the level.)

The match/lock tolerance ϵ . The other free constant is ϵ , the tolerance with which a materialised lock counts as “within a true target” (the R^*/M^* match and lock definition, §2.3). Sweeping $\epsilon \in \{0.3, 0.6, 1.0, 1.5\}$ on the scarcity cell ($N=8$, $T=3$, 15 seeds, peer, experiments/big_experiment/exp_eps_sensitivity.py) leaves the two slope signs of Empirical Claim 1 intact: $\text{Spearman}(P(M^*|R^*), K)$ is $-1.00, -1.00, -0.87, -0.63$ and $\text{Spearman}(P(C^*|M^*), K)$ is $+0.50, +0.50, +0.87, +0.87$ across the four tolerances. A looser ϵ raises absolute $P(M^*|R^*)$ (more locks count as on-target) but does not flip either slope. So the bottleneck shift is robust to both free thresholds of the logger: θ (candidacy/gating, above) and ϵ (match/lock).

F.6 Robustness to Assumption 1 violations

Definition 1 rests on the conditional stage-Markov property (Assumption 1). The two practically relevant violations—adaptive query rewrites not surfaced to the logger, and hidden between-stage memory updates—are exactly cases where an *unlogged latent* couples non-adjacent stages. We quantify the resulting estimator bias with a controlled simulation (4×10^5 episodes per setting; experiments/big_experiment/assumption1_stress_test.py). Episodes are generated with known mechanism rates $P(M^*|R^*)=0.70$, $P(C^*|M^*)=0.65$, and a hidden binary latent U that shifts *both* the M-rate and the C-rate by $\pm\delta$; the naive estimator conditions only on the previous starred stage, not on U .

Table 14: Bias of the naive conditional estimators under an unlogged latent of strength δ that couples the M and C stages (violating Assumption 1). $P(M^*|R^*)$ stays unbiased; $P(C^*|M^*)$ is biased upward, materially only under strong coupling.

δ (coupling strength)	bias in $\hat{P}(M^* R^*)$	bias in $\hat{P}(C^* M^*)$
0.00	−0.001	−0.002
0.05	+0.001	+0.004
0.10	+0.001	+0.015
0.15	−0.000	+0.031
0.20	+0.000	+0.056
0.25	−0.001	+0.089

Two conclusions (Table 14). The $R \rightarrow M$ conditional is unbiased at every δ : the latent re-weights episodes but does not confound that edge. The $C \rightarrow M$ conditional acquires a one-signed *upward* bias—the latent makes high-M episodes also high-C, so conditioning on M over-counts completion—but the bias is only +1.5 percentage points at a mild violation ($\delta \leq 0.10$) and reaches +5.6 points only under strong unlogged coupling ($\delta=0.20$). The framework’s exposure to Assumption 1 violations is therefore bounded and its failure direction is known, which is the relevant robustness statement for applying the logger to learned or LLM agents where such couplings are more common.

F.7 Preliminary continuous-substrate check (VMAS)

The most-requested robustness check is whether the $M \leftrightarrow C$ bottleneck shift is specific to grid substrates. We report a *preliminary* positive on a continuous-state substrate (VMAS (Bettini et al., 2022)). The symbolic peer memory and the operational Q/R/M/C definitions are reused unchanged; a thin discretisation bridge maps continuous (x, y) to coarse cells so the memory applies, and the controller acts in continuous space toward the materialised waypoint (experiments/vmas_portability/).

1097 Materialisation uses the same occupancy-sensitive lock as the grid planner (a candidate whose target
1098 is already claimed does not lock).

Table 15: Preliminary VMAS portability: both slope signs of Empirical Claim 1 reproduce on a continuous substrate in the coupling regime ($N=8$, $T=3$, 40 seeds, horizon 25). $\text{Spearman}(P(M^*|R^*), K) = -1.00$ and $\text{Spearman}(P(C^*|M^*), K) = +1.00$. Beyond $K \approx 32$ the system enters the independent regime (rare broadcast, no contention) and $P(M^*|R^*)$ rebounds, the same regime boundary seen on the grid.

K	$P(M^* R^*)$	$P(C^* M^*)$
1	1.000	0.182
2	0.583	0.305
4	0.448	0.405
8	0.305	0.608
16	0.278	0.639

1099 We label this preliminary and do not fold it into the headline claims: it is a single scenario con-
1100 figuration, the horizon is short, there are no cluster-robust CIs yet, the memory is fed through the
1101 discretisation bridge (continuous *dynamics* with discretised symbolic *memory*, not fully continuous
1102 memory), and the reported ± 1.00 Spearman is a *perfect* rank correlation over only five monotone
1103 points ($K \leq 16$) — weak evidence on its own, indicating direction rather than a resolved effect size.
1104 What it establishes is that the two slope signs are not an artefact of grid cells: the same freshness-
1105 vs-contention mechanism operates when positions and motion are continuous. Promoting this to a
1106 headline result requires scaling to several scarcity cells and layouts with the cluster-robust bootstrap
1107 of Appendix F.3.

1108 F.8 Cadence-robustness: the shift tracks scarcity, not grid design

1109 A reviewer may object that the bottleneck shift is built into a scarce-target grid with locks. If instead
1110 it is a signature of contention over *scarce shared* targets, its slope conditions must break when targets
1111 are abundant. We rerun the peer cadence sweep in three regimes on the same env ($N=8$, random
1112 layout, hazard 0.05, 20 seeds; `experiments/big_experiment/exp_cadence_robustness.py`):
1113 scarcity ($T=3$), borderline ($T=N=8$), and abundance ($T=14 > N$).

Table 16: The two slope conditions of Empirical Claim 1 (§4.3) hold under scarcity and flatten under abundance. $P(M^*|R^*)$ should fall and $P(C^*|M^*)$ should rise across K ; “corr” is the within-cadence Pearson correlation of the two links at $K=64$ (the shift signature). $P(C^*|M^*)$ is an operational ratio (capped at 1; abundance completes without a formal lock).

Regime	$P(M^* R^*): K=1 \rightarrow 64$	$P(C^* M^*): K=1 \rightarrow 64$	corr(M, C) @ $K=64$	t_{succ} argmin
Scarcity ($T=3$)	0.86 $\rightarrow$ 0.49 (falls)	0.51 $\rightarrow$ 0.79 (rises)	-0.86	$K=4$ (interior)
Borderline ($T=N=8$)	0.79 $\rightarrow$ 0.72	0.77 $\rightarrow$ 0.95 (rises)	-0.79	$K=8$ (interior)
Abundance ($T=14$)	0.74 $\rightarrow$ 0.77 (flat)	1.00 $\rightarrow$ 0.99 (flat)	-0.29	$K=4$

1114 Both slopes are present under scarcity, weaken at the border, and *flatten* under abundance: with
1115 more targets than agents there is no contention, so fast sharing neither over-saturates materialization
1116 ($P(M^*|R^*)$ stays flat at ~ 0.75 instead of spiking then falling) nor collapses completion ($P(C^*|M^*)$
1117 stays at ~ 1 even at $K=1$ instead of dropping to 0.51). The within-cadence $M \leftrightarrow C$ anticorrelation
1118 attenuates monotonically with abundance ($-0.86 \rightarrow -0.79 \rightarrow -0.29$). The shift therefore tracks
1119 the *contention mechanism* and vanishes when it is removed — it is not a fixed artefact of the grid-
1120 plus-locks design. We note honestly that the interior t_{succ} optimum is more general (present in all
1121 three regimes): it reflects the exploration/travel cost of very fast *or* very slow sharing, of which target
1122 contention is one contributor but not the only one.

1123 F.9 Open problems suggested by the diagnostics

1124 Four concrete questions follow directly from the measurements of §3.5 and §4.5. **(1) Adaptive**
1125 **cadence.** Empirical Claim 1 assumes a fixed K ; the bottleneck-shift signature suggests an online
1126 controller K_t that broadcasts when the (estimated) M-side gain exceeds the C-side occupancy cost.

1127 **(2) Selective consolidation.** Snapshot exchange merges everything; merge rules that prioritize
1128 evidence about under-represented semantic tags attack the candidate-generation starvation of §3.5
1129 directly. **(3) Trust-weighted merging under occupancy.** The central aggregator already performs
1130 a trust-weighted merge but pays the C-collapse cost; decentralized trust rules that account for peer
1131 *intent* (not only peer evidence) could decouple freshness from target racing. *Since submission of*
1132 *the first draft, this problem has been addressed constructively:* a decentralized protocol in which
1133 an agent’s target lock piggybacks a lightweight *reservation* on its next scheduled broadcast (no
1134 coordinator, no extra channel) and locks dominated by an earlier reservation are retracted by an
1135 anti- M^* rollback with exponential backoff. On the scarcity cells this recovers the fast-cadence
1136 completion loss on *unconditional* metrics — success rate rises at every $K \in \{1, 2, 4\}$ with disjoint
1137 bootstrap CIs, and fast sharing with the protocol exceeds the $K=8$ optimum baseline (0.614–0.620
1138 vs. 0.583) — with rollback latency bounded by $\approx K$ ticks. Details in a companion manuscript; the
1139 open problem stands revised toward its harder variants (adaptive reservation scope, heterogeneous
1140 trust). **(4) Substrate generality.** The decomposition transfers across MiniGrid, BabyAI, the custom
1141 grid, and a standard-MiniGrid multi-agent wrapper with identical operational definitions and verified
1142 slope signs of Empirical Claim 1 (§ 6), and a preliminary check reproduces both slopes on the
1143 continuous-state VMAS substrate (Appendix F.7). The decomposition is also coordinate-independent:
1144 under *semantic*, frame-free place matching (agents in private frames, places matched by meaning) the
1145 stage events remain well-defined and expose the coordinate assumption’s fail-open materialisation
1146 (Appendix E.3). A full continuous-state validation with cluster-robust CIs, and richer 3D substrates
1147 (Habitat (Savva et al., 2019)), remain the next portability steps.

1148 **F.10 Extended related work**

1149 Classical cognitive-map perspectives motivate explicit place abstractions (Tolman, 1948; O’Keefe
1150 and Nadel, 1978; Kuipers, 2000); modern embodied AI combines mapping and planning (Gupta
1151 et al., 2017) or exposes topological memory (SPTM, Savinov et al., 2018); goal-conditioned RL
1152 conditions on explicit goals or embeddings (Andrychowicz et al., 2017; Ghosh et al., 2021), with
1153 successor-representation work providing predictive structure (Dayan, 1993; Stachenfeld et al., 2017;
1154 Momennejad et al., 2017). Our setting differs in that the target is a semantic pattern retrieved from
1155 memory rather than a coordinate. In multi-agent RL, CTDE dominates (Lowe et al., 2017; Foerster
1156 et al., 2018) and learned communication adds a per-step exchange variable (Sukhbaatar et al., 2016;
1157 Foerster et al., 2016); gossip and consensus protocols (Boyd et al., 2006; Xiao et al., 2007) are
1158 the closest network-level analogue of our cadence axis, lifted here to the cognitive-architecture
1159 level. LLM-agent systems combine reasoning/acting, reflection, explicit memory, and multi-agent
1160 conversation (Yao et al., 2023; Shinn et al., 2023; Wang et al., 2023; Park et al., 2023; Wu et al.,
1161 2023; Packer et al., 2023; Chhikara et al., 2025). Our framework is orthogonal: rather than learning
1162 a communication policy end-to-end, it *diagnoses* where in the Q/R/M/C chain a given architecture
1163 succeeds or fails, and yields a slope-level prediction about cadence that is testable on logs.

1164 **Stage-wise evaluation.** The closest line decomposes retrieval-augmented generation: RAGAS
1165 scores retrieval and generation quality separately (Es et al., 2024), RAGChecker refines this to claim-
1166 level diagnostics meta-evaluated against human judgments (Ru et al., 2024), and AgentBoard replaces
1167 final success with analytic progress rates for multi-turn LLM agents (Ma et al., 2024). Q/R/M/C
1168 differs on three axes: its events are logged actions of a task-executing agent, not LLM-judged
1169 text quality; it models commitment (M) and the multi-agent coupling channels (shared memory,
1170 occupancy) that RAG metrics do not represent; and its diagnoses are validated interventionally (oracle
1171 stages; engineered structural faults) rather than by correlation with human ratings. Process mining
1172 and conformance checking compare event logs against stage models (van der Aalst, 2016; Carmona
1173 et al., 2018); the protocol shares that log-level discipline and adds agent-specific event semantics with
1174 intervention-based validation. Process reward models score intermediate reasoning steps to supervise
1175 training (Lightman et al., 2024); Q/R/M/C is an evaluation-time instrument for memory pipelines.

1176 **F.11 SOTA baseline and coverage matrix**

1177 Table 18 separates direct experimental baselines from broader SOTA references. The local experi-
1178 ments are intentionally modest and fully instrumented: random / graph-only / SPTM-like / learned
1179 BC in MiniGrid, CommNet PPO in the multi-agent grid, and a TextNav LLM step-budget study over
1180 two local Ollama models (in which Q/R/M/C localizes a completion-limited failure, §4.5). ReAct,

Reflexion, Voyager, Generative Agents, AutoGen, MemGPT, and Mem0 are not reimplemented as full systems here; they define the current memory-agent design space against which Q/R/M/C is positioned as a diagnostic protocol.

Positioning against SOTA memory designs. The memory-agent design space can be organised on two axes: single- vs. multi-agent, and how memory is represented and shared (Table 17). Vector-embedding memory (MemGPT, Voyager single-agent; Generative Agents multi-agent) retrieves by opaque nearest-neighbour lookup, so a failure is not localisable to a stage. Shared-context designs (MetaGPT, AutoGen) pool everything into one window, dissolving per-agent state. Learned-communication MARL (CommNet, and stronger CTDE methods QMIX (Rashid et al., 2018), MAPPO (Yu et al., 2022), MADDPG (Lowe et al., 2017)) optimises a message channel end-to-end and, as we verify for CommNet PPO (§4.5), emits no separable query/lock/completion events. The cell our system occupies—per-agent *symbolic* memory with an explicit query/retrieve/materialise interface, exchanged at a tunable cadence—is empty in this matrix, and it is exactly the interface that makes the four Q/R/M/C stages observable. Our claim against SOTA is therefore not higher reward but higher *diagnosability at matched task success*: the same interface that a strong learned method trades away for end-to-end optimisation is what lets the framework localise where memory helps.

Table 17: Where the framework sits in the memory-agent design space. Q/R/M/C stage-observability requires an explicit query/lock interface, which the symbolic-per-agent cell provides.

Memory representation	Single-agent	Multi-agent
Vector-embedding memory	MemGPT, Voyager (Packer et al., 2023; Wang et al., 2023)	Generative Agents (Park et al., 2023)
Shared context window	—	AutoGen (Wu et al., 2023)
Learned-comm. MARL	—	CommNet, QMIX, MAPPO, MADDPG (Sukhbaatar et al., 2016; Rashid et al., 2018; Yu et al., 2022; Lowe et al., 2017)
Symbolic per-agent + explicit interface	Q/R/M/C single-agent (§3)	this work (peer memory, cadence K)

F.12 Multi-agent oracle-intervention numerics

F.13 Oracle C: definition and reading

Oracle C is the completion-stage analogue of oracle R and oracle M, added to close the R/M/C triple symmetrically. The real Q/R/M pipeline runs unchanged (real memory query, real materialisation), and once an agent has locked a target that is a genuine `water_source`, the oracle guarantees arrival by placing the agent on that cell; scarcity is preserved because a cell already claimed or occupied is not available, so oracle C cannot make more than T agents succeed. The runner is `experiments/big_experiment/exp_oracle_interventions.py` (mode `oracle_C`).

Two readings follow from Table 19. (i) Oracle C produces the *lowest* mean t_{succ} of the three interventions (~ 6.1 – 6.4 vs. ~ 7.2 for R and M), because it removes the entire post-lock travel time; this confirms that a real part of the excess time is controller travel to a correctly materialised target. (ii) Oracle C yields the *smallest* success-rate gain ($+0.01$ – 0.02 over baseline, below both oracle R at 0.614 and oracle M at 0.602), because in scarcity the constraint on *whether* an agent succeeds is upstream—retrieving and locking a still-available target—not reaching it. The two facts together are exactly the framework’s thesis in intervention form: completion is where time is spent, retrieval/materialisation is where success is won.

F.14 Stage-agnostic cadence hypotheses and what the framework changes

A first version of the multi-agent cadence sweep (12,960 runs at 20 seeds; subsumed by the 35,640-run sweep in this paper) was originally planned with four operational hypotheses formulated *before* Empirical Claim 1. They are retained here because they motivated the eventual framework. The framework was formulated after the sweep was complete; it is not claimed to have predicted the sweep numbers a priori (see “Chronology” in §4.3).

1219 Original hypotheses (pre-framework).

1220 **H1.** An interior $K^* \in \{2, 4, 8\}$ minimising mean time-to-success exists for most (N, T, layout)
1221 cells.

1222 **H2.** K^* grows monotonically with scarcity $\rho = N/T$.

1223 **H3.** When $\rho \leq 1$, $K^* = 1$.

1224 **H4.** Peer-broadcast architectures Pareto-dominate centralized aggregation on (mean time-to-success,
1225 success rate).

1226 **Stand-alone result on the original 20-seed sweep.** The practical claim was directionally supported:
1227 peer at the empirically best K was faster than centralized in the scarcity regime (a modest effect;
1228 Mann–Whitney $p=0.044$ on this small standalone sweep, reported as exploratory rather than as
1229 primary evidence, consistent with the effect-size-first line of Appendix F.3). Among the four
1230 structural hypotheses, only H4 was clearly supported in the random-layout subset. H1 held in 7 of
1231 18 cells; H2 had a non-significant negative correlation; H3 was supported in 4 of 9 admissible cells.
1232 The 35,640-run extended sweep used elsewhere in the paper produces matching stand-alone numbers
1233 within bootstrap uncertainty.

1234 **What the framework changes.** The hypotheses H1–H4 are stage-agnostic: each ties K^* to scarcity
1235 ρ without specifying mechanism. Empirical Claim 1, formulated after the sweep, ties the slopes of
1236 $P(M^*|R^*)$ and $P(C^*|M^*)$ to cadence; the slope conditions are testable on logs. The data verify
1237 both signs as moderate, cluster-robust effect sizes (Appendix F.3). We make no formal claim about
1238 the shape of the conditional product Φ , on which the data are monotone. The interior optimum holds
1239 on the efficiency metric mean time-to-success.

1240 G Assist on/off analysis

1241 Across all four single-agent task families, top-line success is essentially unchanged under assist on/off
1242 ($\Delta \leq 0.10$). The only non-trivial effect is on rest, where assists improve post-retrieval completion
1243 by +0.05 without changing query/retrieval rates, consistent with assists acting on materialization and
1244 post-retrieval execution.

1245 H Retrained retriever and CommNet baseline

1246 H.1 Retrained candidate-generation MLP (§3.5)

1247 **Dataset extraction.** From the 10-seed hazard-recovery benchmark (MiniGrid-LavaGapS7-v0,
1248 method `milestone_state_conditioned_hazard_recovery_v7`), an 18-dimensional feature
1249 vector was extracted for every step of every episode of every seed — 3 numeric agent-state scalars
1250 (energy, thirst, risk_budget), 9 semantic-tag confidences, 6 intent one-hot — and a binary label: *was*
1251 *this symbolic node ever selected as a candidate during the run?* Dataset totals: 814 examples, 113
1252 positives (13.9%).

1253 **Split.** By seed: train $\{2, 3, 4, 5, 6, 7\}$ (505 examples, 9.7% positive), val $\{8, 9\}$ (170, 22.4% positive),
1254 test $\{10, 11\}$ (139, 18.7% positive).

1255 **Architecture.** $18 \rightarrow 32 \rightarrow 32 \rightarrow 1$, ReLU activations, BCE with `pos_weight` ≈ 9.3 for class
1256 imbalance, Adam `lr` $= 2 \cdot 10^{-3}$ with `weight_decay` $= 10^{-4}$, batch 64, 200 epochs, best by val loss.

1257 Results.

1258 The MLP underperforms majority-class on accuracy across all splits and achieves only moderate
1259 precision at fixed recall. Instantaneous state features are insufficient to predict candidate eligibility
1260 downstream. This is the negative-but-informative finding referenced in §3.5: the framework’s
1261 “retrieval-limited” attribution refines to “memory-accumulation-limited” rather than “ranking-limited”.

1262 H.2 CommNet-style learned communication baseline (§4.5)

1263 **Policy architecture.** Per-agent observation is 4-dim direction one-hot + $5 \times 5 \times 5$ local cell-type window
1264 (125) + 2-dim normalised position = 131. Shared-weights encoder $131 \rightarrow 64 \rightarrow 64$ (ReLU), 16-dim




Figure 10: Refined attribution diagnostic. *Left*: 349 retrieval calls aggregated across 10 seeds; 317 (90.8%) returned empty candidate sets, 31 (8.9%) returned at least one satisfied candidate, only 1 (0.3%) returned candidates that were ranked but unsatisfied. The bottleneck is candidate generation, not candidate ranking. *Right*: the trained MLP underperforms a majority-class baseline on all three seed-disjoint splits, confirming that instantaneous state features alone are insufficient to predict candidate eligibility.

1265 comm projection, peer mean-pool (excluding self), combine $80 \rightarrow 64$ (ReLU), actor head $64 \rightarrow 4$
 1266 (TURN_LEFT, TURN_RIGHT, FORWARD, NOOP), critic head $64 \rightarrow 1$. This is a CommNet-style
 1267 learned-communication architecture (Sukhbaatar et al., 2016).

1268 **Training.** REINFORCE (Williams, 1992) with value-baseline; loss = $-\text{mean}(\log \pi(a) \cdot (R - V)) +$
 1269 $0.5 \cdot \text{MSE}(V, R) - 0.01 \cdot H[\pi]$; Adam lr = $3 \cdot 10^{-4}$, gradient clipping 1.0. Reward shaping: +1.0
 1270 on first success per agent, -0.01 step cost, -0.1 per hazard contact. 2000 episodes cycling 200
 1271 environment seeds, wall-clock 100 s on CPU. Held-out eval: 50 environment seeds 250–299.

1272 **Training curve.** Episode-level success (last-100 window) stays in $[0.13, 0.23]$ throughout training;
 1273 no monotonic improvement after ~ 200 episodes. The policy converges to a local optimum where
 1274 one of three agents finds water by chance.




Figure 11: Learned-communication baseline diagnostics. *Left*: CommNet training curve (smoothed success rate over a 100-episode window) over 2000 training episodes; success stays below 0.25 throughout. The dashed line is the symbolic peer architecture at $K=8$ on the same scenario for reference. *Right*: Q/R/M/C profile of the trained CommNet on 50 held-out seeds vs symbolic peer at $K=8$. Q^* is saturated at 1.0 on CommNet (degenerate by construction); R/M collapse together; only on the symbolic architecture are the conditional rates separated.

1275 **Q/R/M/C evaluation (50 held-out seeds, stochastic action sampling).**

1276 Compare to the symbolic peer architecture at $K=8$: $P(M^*|R^*)=0.69$, $P(C^*|M^*)=0.87$, success
1277 ~ 0.65 .

1278 **PPO variant: structural degeneracy verified at competitive success.** A PPO-trained (Schul-
1279 man et al., 2017) variant of the same CommNet architecture (clip $\epsilon=0.2$, generalized ad-
1280 vantage estimation (Schulman et al., 2016) with $\lambda=0.95$, $\gamma=0.99$, 4 PPO epochs per roll-
1281 out, 64 rollouts per update, 200 updates, ≈ 96 minutes on a single CPU core; full setup in
1282 `experiments/commnet_ppo_baseline/train_ppo_commnet.py`) reaches 0.667 success on 100
1283 held-out seeds — competitive with the symbolic peer architecture at $K=8$ (0.65).

1284 The PPO result converts the structural argument from an assertion (“training budget would not change
1285 this”) to a measurement: Q^* remains saturated at 1.00, $P(M^*|R^*)$ remains 1.00 (R/M collapse), even
1286 at competitive success. The Q stage is degenerate by construction in any policy whose communication
1287 channel is a real-valued vector with a learned linear projection — the channel is never informatively
1288 empty, regardless of training budget — and without an explicit symbolic target-lock interface the M
1289 event cannot separate from R. Only architectures with an explicit *query/lock* interface (the symbolic
1290 stack of §4.1) expose stage-separable Q/R/M/C events.

1291 I Reproducibility commands

1292 An anonymized public repository accompanies this submission. The main paper artifacts were
1293 generated from the following scripts; PYTHONPATH=. and a Python 3 virtual environment with numpy,
1294 scipy, pandas, matplotlib, imageio, gymnasium, and minigrid are assumed throughout.

1295 **Computing infrastructure.** All experiments ran on a single consumer machine: Apple Silicon
1296 (10 CPU cores used for the sweeps via ProcessPoolExecutor), macOS (Darwin 25.4), Python
1297 3.9.6. Key library versions: numpy 1.x, scipy, torch (CommNet-PPO baseline and VMAS), vmas
1298 1.5.2, minigrid/gymnasium. LLM experiments use local models via Ollama (llama3.1:latest
1299 8B, qwen3:4b) at temperature 0 with deterministic on-disk response caching; the external-framework
1300 study additionally uses openai 2.44, langgraph 0.6, langchain-ollama, and pyautogen 0.9
1301 (AutoGen classic). No GPU cluster or paid API was used; the full multi-agent block reproduces in
1302 ≈ 22 minutes on 8 cores, and each LLM study in minutes-to-hours of local inference.

1303 One-command reproduction (multi-agent block, ≈ 22 minutes on 8 cores)

1304 `bashscripts/reproduce_paper.sh`

1305 Clean Paper 1 suite

- 1306 • smoke suite (all local non-LLM blocks): `pythonscripts/run_paper1_clean_experiments.py`
1307 `--modesmoke--skip_llm--out_dir/tmp/paper1_clean_experiments_smoke`
- 1308 • full local suite (single-agent, oracle, multi-agent, CommNet, portability/noise):
1309 `pythonscripts/run_paper1_clean_experiments.py--modefull--skip_`
1310 `llm--workers8--out_dir/tmp/paper1_clean_experiments_full`
- 1311 • LLM TextNav step-budget sweep (requires local Ollama), run once per budget $K' \in \{6, 12, 25\}$:
1312 `python-mexperiments.llm_collective.run_model_sweep--modelsqwen3:4bllama3.`
1313 `1:latest--episodes8--step_limit<K'>--out_dir/tmp/llm_steplimit_<K'>`

1314 Single-agent (§3.2–§3.4)

- 1315 • article benchmark and learned-baseline suite: `scripts/benchmark_symbolic_memory_`
1316 `article.py`
- 1317 • article analysis and figure generation: `scripts/analyze_symbolic_memory_article.py`
- 1318 • synthetic factorization sanity check: `scripts/validate_qrmc_factorization.py`
- 1319 • oracle-stage intervention benchmark: `scripts/benchmark_oracle_stage_interventions.`
1320 `py`
- 1321 • BabyAI transfer benchmark: `scripts/compare_semnav_babyai.py`

1322 Multi-agent (§4.5)

- 1323 • Base batch — peer at $K \in \{1, 2, 4, 8, 16\}$ and the three cadence-free architectures (in-
1324 dependent, shared bus, central aggregator), 40 seeds (25,920 runs = 16,200 peer + 9,720
1325 cadence-free), ≈ 18 min: `pythonexperiments/big_experiment/exp_cadence_phase.
1326 py--modefull--workers8--out_dirtmp/big_experiment_qrmc_40`
- 1327 • Extended-cadence batch — peer at $K \in \{32, 48, 64\}$ only (9,720 runs), ≈ 4 min:
1328 `pythonexperiments/big_experiment/exp_extra_K.py--workers8--out_dirtmp/
1329 big_experiment_extraK`
- 1330 *Note.* These two batches partition the 35,640 total by *cadence range*; the architecture-wise
1331 partition of Appendix F.2 (25,920 peer over all eight cadences vs. 9,720 cadence-free) reuses
1332 the same two totals for *different* subsets (a coincidence of 3 extended cadences $\times$ 3,240 = 3
1333 cadence-free architectures $\times$ 3,240). The combined log is `tmp/paper1_clean_experiments_
1334 full/multiagent_cadence_full/runs.csv`.
- 1335 • Q/R/M/C stage analysis and bootstrap CIs: `pythonexperiments/big_experiment/
1336 analyze_qrmc.py--runs_csvtmp/big_experiment_qrmc_40/runs.csv--out_
1337 dirtmp/big_experiment_qrmc_40`
- 1338 • Multi-agent oracle interventions: `pythonexperiments/big_experiment/exp_oracle_
1339 interventions.py--workers8--out_dirtmp/big_experiment_oracle`
- 1340 • Phase 1–4 smoke tests: `pythonscripts/run_all_smokes.py--out_dirtmp/all_smokes`
- 1341 • PPO-trained CommNet baseline (Table 20, ≈ 96 min on a single CPU core):
1342 `python-mexperiments.commnet_ppo_baseline.train_ppo_commnet--total_
1343 updates200--rollouts_per_update64`
- 1344 • Q/R/M/C eval of the PPO policy on 100 held-out seeds: `pythonexperiments/commnet_
1345 baseline/eval_with_qrmc.py--policy_pathexperiments/commnet_ppo_baseline/
1346 ppo_policy.pt--out_direxperiments/commnet_ppo_baseline--n_episodes100`
- 1347 • MiniGrid-substrate portability sweep (§6, 100 runs, ≈ 2 min): `python-mexperiments.
1348 minigrid_multiagent_wrapper.run_portability_sweep--out_dirtmp/minigrid_
1349 portability`

1350 J Asset and license note

1351 The experiments use MiniGrid (Chevalier-Boisvert et al., 2023) and BabyAI (Chevalier-Boisvert
1352 et al., 2018), both released under the MIT License. The locally implemented baselines and benchmark
1353 runner scripts introduced in this paper will be released under the MIT License as part of the camera-
1354 ready artifact package.

Table 18: SOTA coverage for Paper 1. “Instrumented” means the run emits the operational Q^* , R^* , M^* , C^* events used by Definition 1.

Family		Representative method	Status in this paper	Q/R/M/C comparison point
Topological navigation memory		SPTM-like graph baseline (Savinov et al., 2018)	Direct MiniGrid baseline	Tests whether topological memory without semantic query/lock exposes non-trivial stage events. Strong success on easy tasks but no explicit query/lock interface, so diagnostics are structurally limited. Competitive success; Q^* saturates and R/M collapse, verifying observability loss.
Behavior cloning		Learned grid-observation BC	Direct MiniGrid baseline	
Learned communication		CommNet PPO (Sukhbaatar et al., 2016; Schulman et al., 2017)	Direct multi-agent baseline	
Cooperative MARL (CTDE / value decomposition)		QMIX, MAPPO, MADDPG (Rashid et al., 2018; Yu et al., 2022; Lowe et al., 2017)	Primary-source reference	
Reasoning/acting LLM agents		ReAct (Yao et al., 2023)	Primary-source reference	Stronger end-to-end baselines than CommNet; they optimise a joint policy/message channel without an explicit query/lock interface, so the same observability-collapse argument applies — the framework quantifies the diagnostic cost of implicit coordination rather than competing on reward. Q/R/M/C can instrument the resulting action traces if queries, retrieved evidence, target locks, and completions are logged. Reflection can be treated as upstream Q/R revision; the framework asks whether it improves R, M, or C. Skill/memory retrieval maps naturally onto R and M, but full reproduction is outside this paper. Their observation/memory/reflection loop motivates logging stage-specific memory use rather than only believability. Conversation traces provide candidate Q/R/M/C events; this paper supplies the measurement contract. Persistent memory systems can be evaluated by where they fail: retrieval, materialization, or completion. Q/R/M/C localizes a completion-limited failure ($C^* \rightarrow 0$ while $Q/R/M$ stay saturated) in real LLM agents; see Table 4.
Reflective LLM agents	LLM	Reflexion (Shinn et al., 2023)	Primary-source reference	
Embodied LLM agents	LLM	Voyager (Wang et al., 2023)	Primary-source reference	
Social/memory simulations		Generative Agents (Park et al., 2023)	Primary-source reference	
Multi-agent frameworks	LLM	AutoGen (Wu et al., 2023)	Primary-source reference	
Long-term memory systems	LLM	MemGPT / Mem0 (Packer et al., 2023; Chhikara et al., 2025)	Primary-source reference	
Local LLM sub-strate	LLM	qwen3:4b, llama3.1	Direct TextNav step-budget study	

Table 19: Multi-agent oracle interventions on the scarcity cases at three cadences. Mean time-to-success (and success rate) averaged across $(N, T) \in \{(3, 2), (5, 3), (8, 5)\}$, two layouts (asymmetric, random), 15 seeds. Oracle R returns ground-truth water positions; oracle M locks the nearest real water; **oracle C** guarantees completion of a real materialised lock (Appendix F.13). Oracle C gives the largest *time* reduction (it removes travel) but the smallest *success* gain — completion is a time cost, not the success bottleneck, which sits upstream at R/M.

K	Baseline	Oracle R	Oracle M	Oracle C
<i>mean t_{succ}</i>				
1	10.33	7.22	7.21	6.19
4	10.98	7.22	7.21	6.07
16	8.45	7.22	7.21	6.41
<i>success rate</i>				
1	0.591	0.614	0.602	0.591
4	0.575	0.614	0.602	0.600
16	0.578	0.614	0.602	0.598

Metric	Train	Val	Test
Accuracy	0.584	0.553	0.590
Majority-class baseline accuracy	0.903	0.776	0.813
Precision @ recall = 0.70	0.171	0.278	0.275

Quantity	Value
Q^* rate	1.00 (saturated)
R^* rate	0.50
M^* rate	0.50 (collapses onto R)
C^* rate	0.19
$P(R Q)$	0.50
$P(M R)$	1.00
$P(C M)$	0.42
Success rate	0.19
Mean steps	80 (budget exhausted)

Table 20: REINFORCE vs PPO CommNet vs symbolic peer ($K=8$) on the scarcity scenario ($N=3$, $T=2$, asymmetric, hazard 0.05). PPO is competitive on success rate, yet still exhibits Q-saturation and R/M collapse — the diagnostic claim verified, not asserted.

Metric	CommNet REINFORCE	CommNet PPO	Symbolic peer ($K=8$)
Success rate	0.19	0.667	0.65
Q^* rate	1.00	1.00	0.99
R^* rate	0.50	0.997	0.69
M^* rate	0.50	0.997	0.69
$P(M^* R^*)$	1.00	1.00	0.69
$P(C^* M^*)$	0.42	0.67	0.87
R/M separated?	no	no	yes

1355 NeurIPS Paper Checklist

1356 1. Claims

1357 Question: Do the main claims made in the abstract and introduction accurately reflect the
1358 paper’s contributions and scope?

1359 Answer: [Yes]

1360 Justification: The abstract and introduction present the paper primarily as a four-stage
1361 diagnostic framework with a reference symbolic-semantic instantiation, and they explicitly
1362 scope the evidence to benchmarked navigation environments with a limited portability check
1363 (Abstract; Sections 1 and 6).

1364 Guidelines:

- 1365 • The answer [N/A] means that the abstract and introduction do not include the claims
1366 made in the paper.
- 1367 • The abstract and/or introduction should clearly state the claims made, including the
1368 contributions made in the paper and important assumptions and limitations. A [No] or
1369 [N/A] answer to this question will not be perceived well by the reviewers.
- 1370 • The claims made should match theoretical and experimental results, and reflect how
1371 much the results can be expected to generalize to other settings.
- 1372 • It is fine to include aspirational goals as motivation as long as it is clear that these goals
1373 are not attained by the paper.

1374 2. Limitations

1375 Question: Does the paper discuss the limitations of the work performed by the authors?

1376 Answer: [Yes]

1377 Justification: The paper includes an explicit “Limitations and future work” section that dis-
1378 cusses environment difficulty, the observable-query and observable-occupancy assumptions,
1379 the slope-level scope of Empirical Claim 1, the learned-baseline scope, and agent-count
1380 scaling (Section 6).

1381 Guidelines:

- 1382 • The answer [N/A] means that the paper has no limitation while the answer [No] means
1383 that the paper has limitations, but those are not discussed in the paper.
- 1384 • The authors are encouraged to create a separate “Limitations” section in their paper.
- 1385 • The paper should point out any strong assumptions and how robust the results are to
1386 violations of these assumptions (e.g., independence assumptions, noiseless settings,
1387 model well-specification, asymptotic approximations only holding locally). The authors
1388 should reflect on how these assumptions might be violated in practice and what the
1389 implications would be.
- 1390 • The authors should reflect on the scope of the claims made, e.g., if the approach was
1391 only tested on a few datasets or with a few runs. In general, empirical results often
1392 depend on implicit assumptions, which should be articulated.
- 1393 • The authors should reflect on the factors that influence the performance of the approach.
1394 For example, a facial recognition algorithm may perform poorly when image resolution
1395 is low or images are taken in low lighting. Or a speech-to-text system might not be
1396 used reliably to provide closed captions for online lectures because it fails to handle
1397 technical jargon.
- 1398 • The authors should discuss the computational efficiency of the proposed algorithms
1399 and how they scale with dataset size.
- 1400 • If applicable, the authors should discuss possible limitations of their approach to
1401 address problems of privacy and fairness.
- 1402 • While the authors might fear that complete honesty about limitations might be used by
1403 reviewers as grounds for rejection, a worse outcome might be that reviewers discover
1404 limitations that aren’t acknowledged in the paper. The authors should use their best
1405 judgment and recognize that individual actions in favor of transparency play an impor-
1406 tant role in developing norms that preserve the integrity of the community. Reviewers
1407 will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: The paper states the conditional stage-Markov property and positivity assumptions in Section 2.2, gives a proof sketch in the main text, and includes a full formal argument for Definition 1 (single-agent stage estimability) in Appendix A together with the multi-agent extension (Definition 2) and the slope-level argument for Empirical Claim 1.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The paper specifies the task families, environments, seeds, episode counts, assist modes, baselines, statistical protocol, and benchmark semantics; the single-agent setup is in Section 3.1 and the multi-agent sweep protocol in Section 4.4; the operational stage definitions used to extract Q^* , R^* , M^* , C^* from runtime logs are in Appendix B.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

1463 (d) We recognize that reproducibility may be tricky in some cases, in which case
1464 authors are welcome to describe the particular way they provide for reproducibility.
1465 In the case of closed-source models, it may be that access to the model is limited in
1466 some way (e.g., to registered users), but it should be possible for other researchers
1467 to have some path to reproducing or verifying the results.

1468 5. Open access to data and code

1469 Question: Does the paper provide open access to the data and code, with sufficient instruc-
1470 tions to faithfully reproduce the main experimental results, as described in supplemental
1471 material?

1472 Answer: [Yes]

1473 Justification: The paper describes a unified benchmark runner and source-of-truth artifacts.
1474 An anonymized public release of the codebase — single-agent benchmark, multi-agent grid
1475 environment, 35,640-run sweep driver, Q/R/M/C analyser, oracle-intervention scripts, and
1476 exact commands to reproduce every table and figure — accompanies the submission as
1477 supplemental material (anonymized GitHub repository; single-command reproduction of
1478 the main multi-agent sweep on an 8-core machine in ≈ 22 minutes, matching Appendix H).

1479 Guidelines:

- 1480 • The answer [N/A] means that paper does not include experiments requiring code.
- 1481 • Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- 1482 • While we encourage the release of code and data, we understand that this might not
1483 be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not
1484 including code, unless this is central to the contribution (e.g., for a new open-source
1485 benchmark).
- 1486 • The instructions should contain the exact command and environment needed to run to
1487 reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- 1488 • The authors should provide instructions on data access and preparation, including how
1489 to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- 1490 • The authors should provide scripts to reproduce all experimental results for the new
1491 proposed method and baselines. If only a subset of experiments are reproducible, they
1492 should state which ones are omitted from the script and why.
- 1493 • At submission time, to preserve anonymity, the authors should release anonymized
1494 versions (if applicable).
- 1495 • Providing as much information as possible in supplemental material (appended to the
1496 paper) is recommended, but including URLs to data and code is permitted.

1499 6. Experimental setting/details

1500 Question: Does the paper specify all the training and test details (e.g., data splits, hyperpa-
1501 rameters, how they were chosen, type of optimizer) necessary to understand the results?

1502 Answer: [Yes]

1503 Justification: The benchmark protocol, task definitions, assist settings, seeds, episode counts,
1504 baseline families, and the learned baselines are described in the main text together with the
1505 measurement and statistical protocols: single-agent setup in Section 3.1, multi-agent sweep
1506 in Section 4.4, and per-baseline hyperparameters in Appendix G.

1507 Guidelines:

- 1508 • The answer [N/A] means that the paper does not include experiments.
- 1509 • The experimental setting should be presented in the core of the paper to a level of detail
1510 that is necessary to appreciate the results and make sense of them.
- 1511 • The full details can be provided either with the code, in appendix, or as supplemental
1512 material.

1513 7. Experiment statistical significance

1514 Question: Does the paper report error bars suitably and correctly defined or other appropriate
1515 information about the statistical significance of the experiments?

Answer: [Yes]

Justification: The paper reports 95% bootstrap confidence intervals over seed-level aggregates ($B=5,000$ resamples; Section 3.1, “Protocol” paragraph) and uses paired Wilcoxon signed-rank tests with Holm-Bonferroni correction for assist comparisons. CIs are shown in Tables 3 and 5 (Best semantic per-task rates; the 35,640-run MOP column) and in Figures 4–6 (bootstrap shaded bands).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: The paper discloses the hardware platform (10-core ARM CPU, 32 GB RAM; CPU-only), the scale of each benchmark, and wall-clock costs — single-agent oracle sweep < 1 minute (Section 3.1); main multi-agent sweep (35,640 runs) ≈ 22 minutes on 8 cores (Section 4.4); $N=12$ scale-up (8,640 runs) ≈ 12 minutes (Appendix D); PPO CommNet baseline (200 updates $\times$ 64 rollouts) ≈ 96 minutes on a single CPU core (Appendix G); MiniGrid portability sweep (100 runs) ≈ 2 minutes on a single core (§6).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn’t make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The work is a simulation-only study in standard benchmark environments, uses anonymized authorship for review, and does not involve human subjects, private data, or deployment claims that conflict with the NeurIPS Code of Ethics.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: The paper includes a dedicated “Broader impacts” section that discusses the primary methodological benefits of the work and the indirect risks of improving navigation systems without adequate safeguards (Section 7).

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper does not release high-risk models or scraped datasets; it studies navigation policies in simulated environments.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.

- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: The paper cites the upstream benchmark platforms (MiniGrid, BabyAI) in the main references. Appendix I states that both are MIT-licensed and that the benchmark scripts introduced here will be released under the MIT License at camera-ready.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, `paperswithcode.com/datasets` has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [No]

Justification: The submission introduces benchmark scripts and analysis artifacts. A formal anonymized release package with documentation is planned for camera-ready; Appendix H lists all scripts needed to reproduce each reported result, and Appendix I states the intended license for the new assets.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The paper does not involve crowdsourcing or human-subject experiments.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: The paper does not involve human subjects or human-participant data collection.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: LLMs are not an important, original, or non-standard component of the core methods or experiments reported in the paper.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.